

# Short-term Memory for Self-collecting Mutators

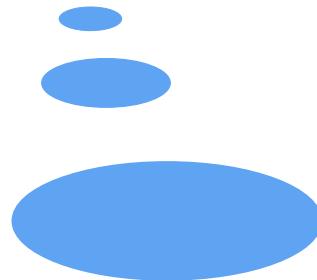
Martin Aigner, **Andreas Haas**, Christoph M. Kirsch, Michael Lippautz,  
Ana Sokolova, Stephanie Stroka, Andreas Unterweger

University of Salzburg

International Symposium on Memory Management 2011  
June 5, San Jose

# Overview

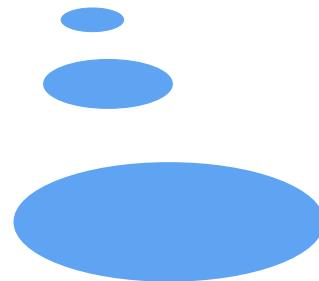
## Short-term Memory



new memory model  
for heap management

# Overview

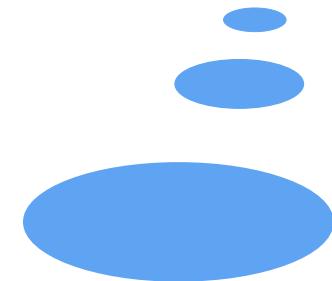
## Short-term Memory



new memory model  
for heap management

for

## Self-collecting Mutators



new explicit heap  
management system

# Idea

- ▶ Each memory object allocated in short-term memory has an **expiration date**
  - ◆ When an object **expires**, its memory may be reused

# Idea

- ▶ Each memory object allocated in short-term memory has an **expiration date**
  - ◆ When an object **expires**, its memory may be reused



# Idea

- ▶ Each memory object allocated in short-term memory has an **expiration date**
  - ◆ When an object **expires**, its memory may be reused



# Idea

- ▶ Each memory object allocated in short-term memory has an **expiration date**
  - ◆ When an object **expires**, its memory may be reused



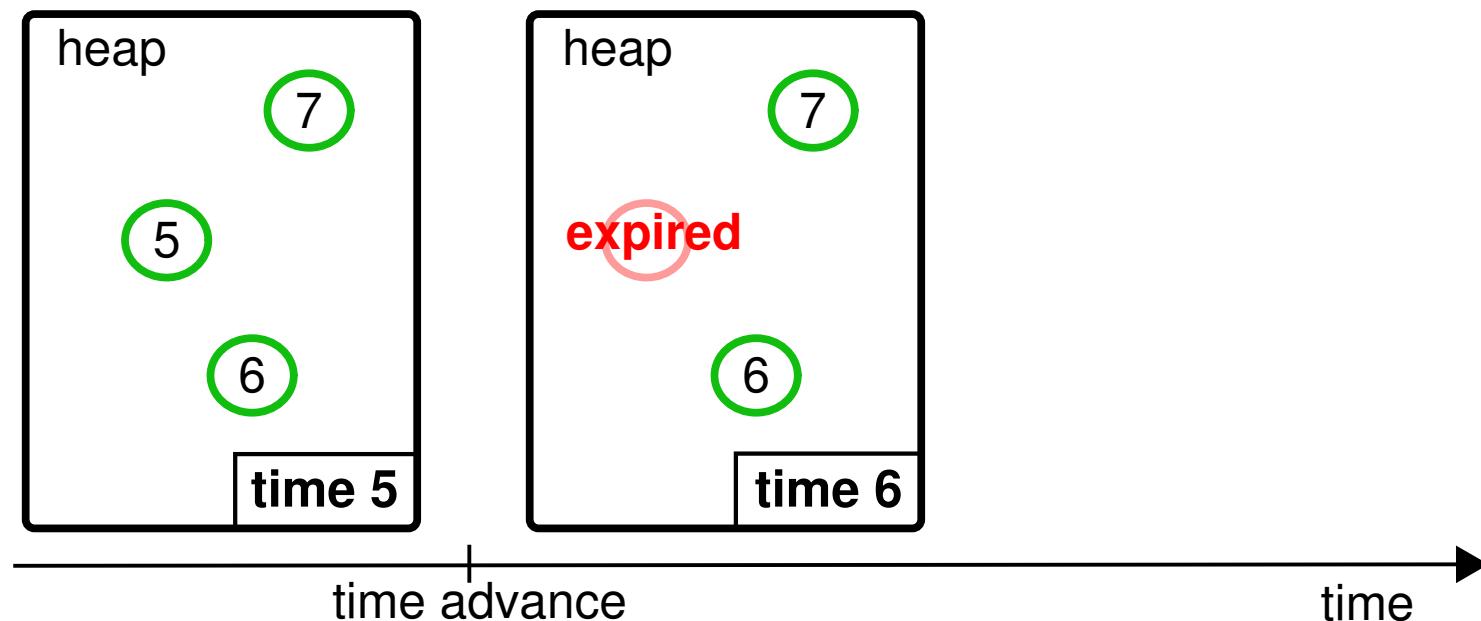
# Idea

- ▶ Each memory object allocated in short-term memory has an **expiration date**
  - ◆ When an object **expires**, its memory may be reused



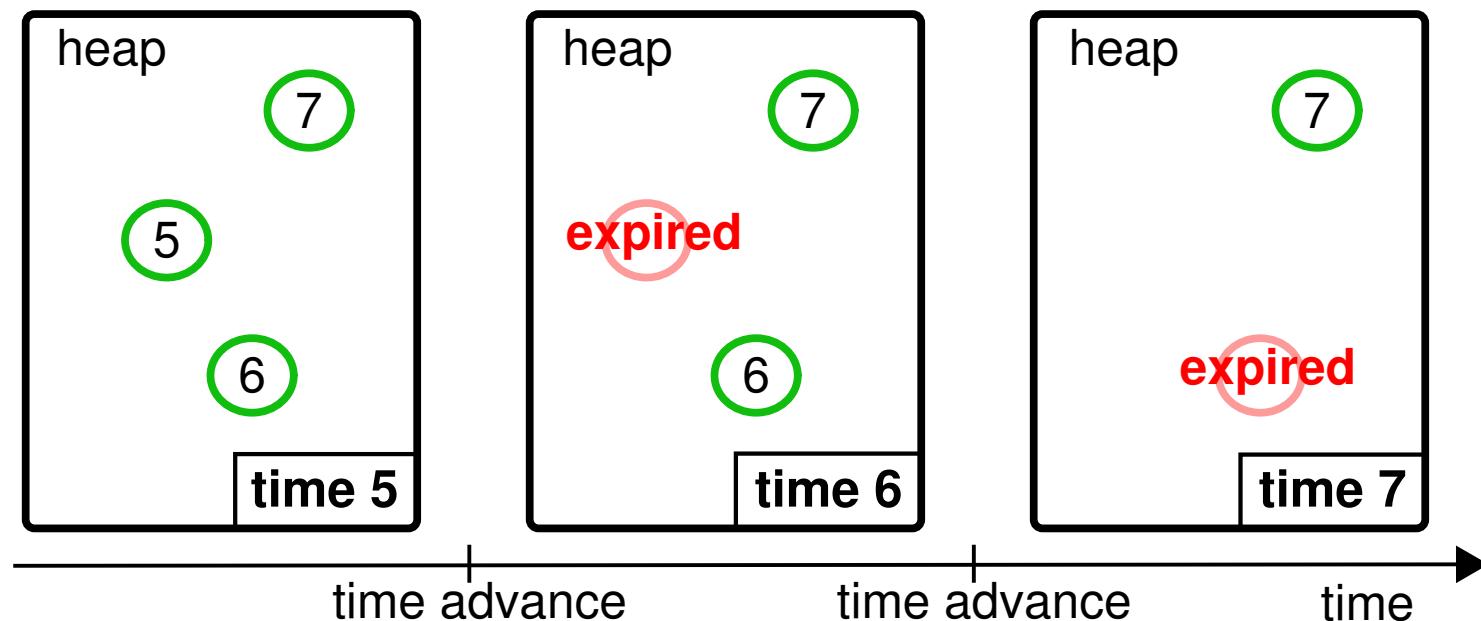
# Idea

- ▶ Each memory object allocated in short-term memory has an **expiration date**
  - ◆ When an object **expires**, its memory may be reused



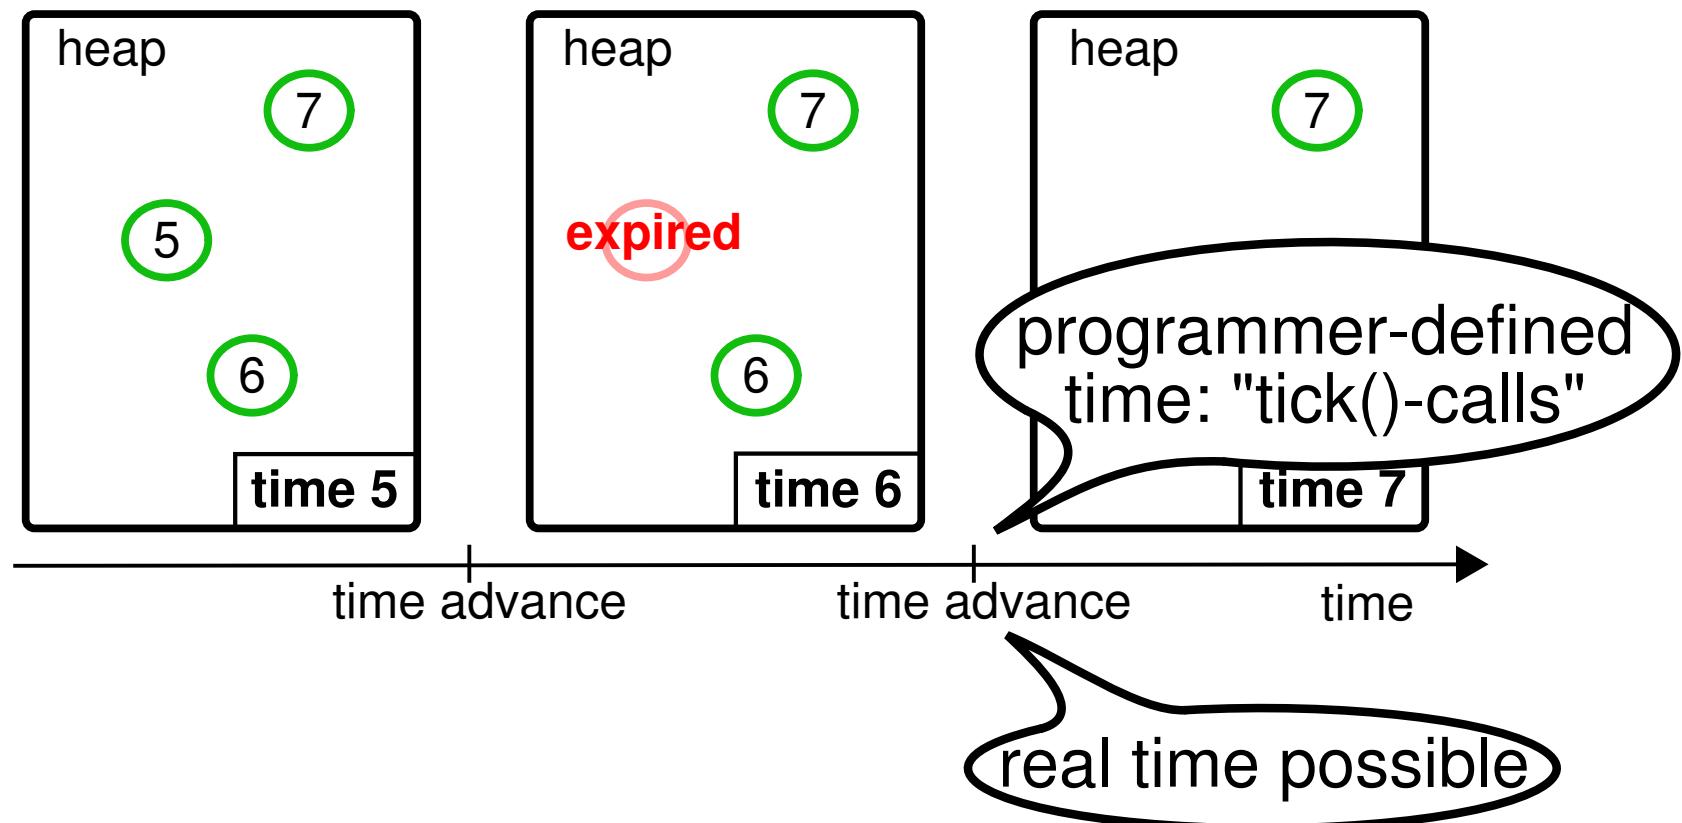
# Idea

- ▶ Each memory object allocated in short-term memory has an **expiration date**
  - ◆ When an object **expires**, its memory may be reused



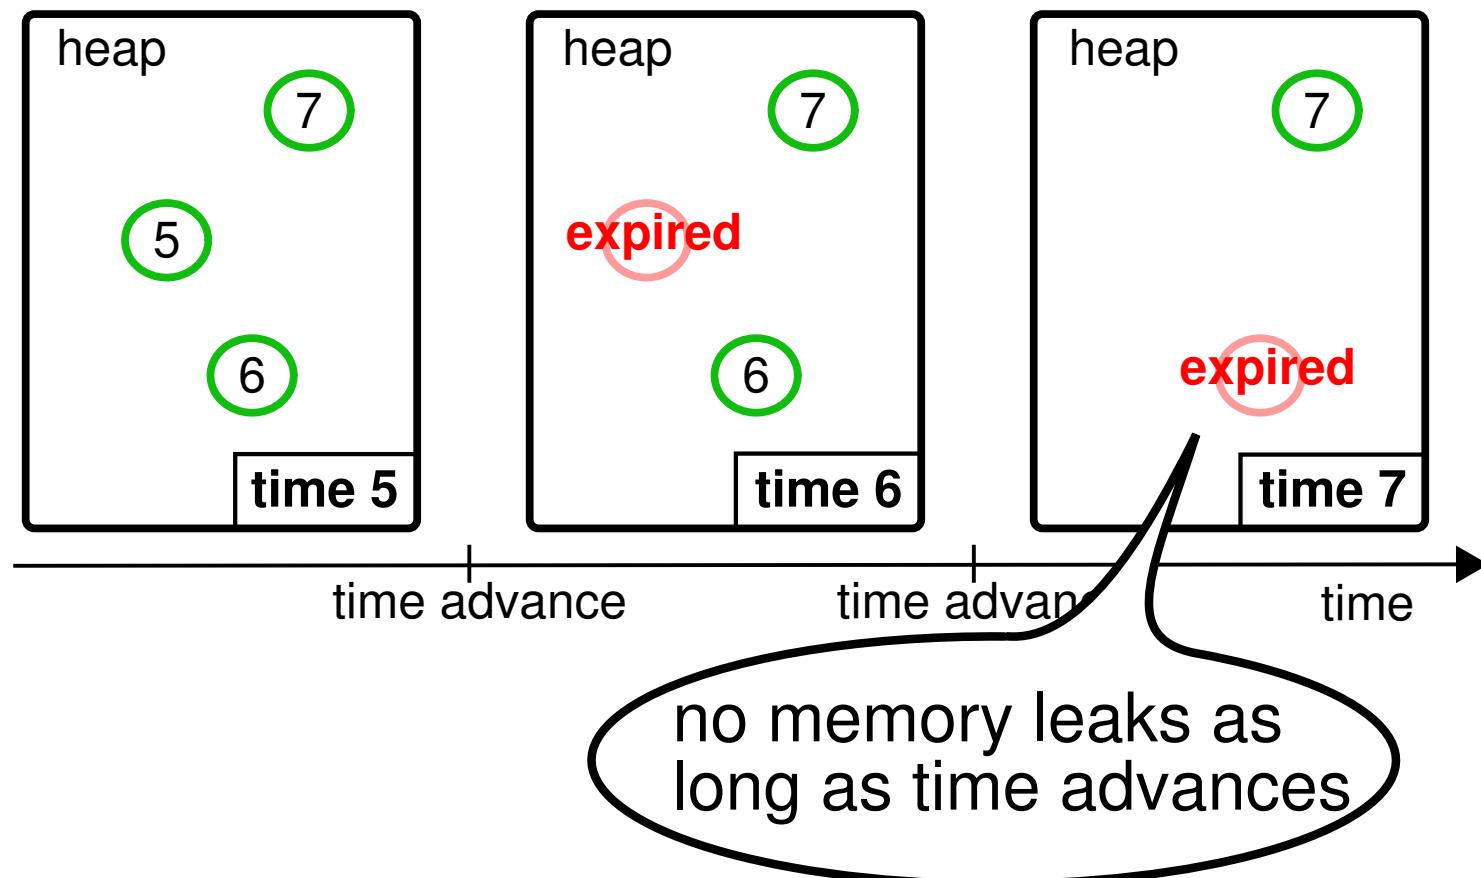
# Idea

- ▶ Each memory object allocated in short-term memory has an **expiration date**
  - ◆ When an object **expires**, its memory may be reused



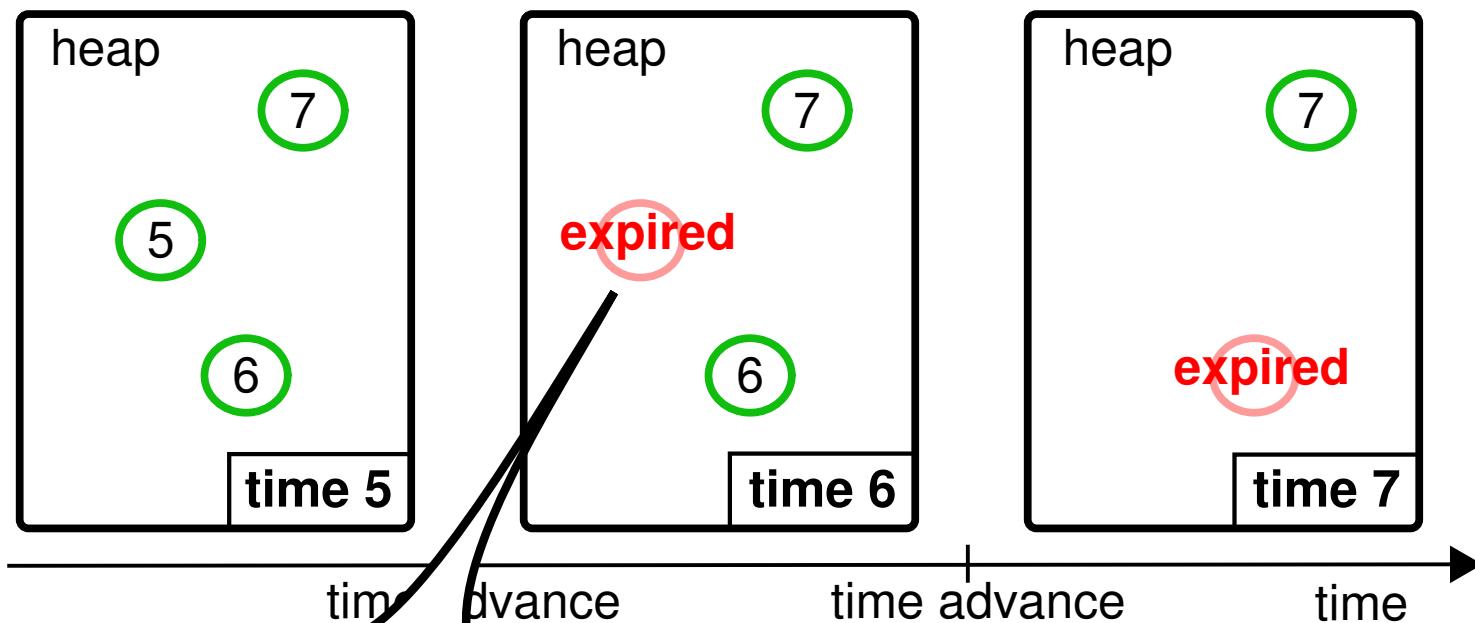
# Idea

- ▶ Each memory object allocated in short-term memory has an **expiration date**
  - ◆ When an object **expires**, its memory may be reused



# Idea

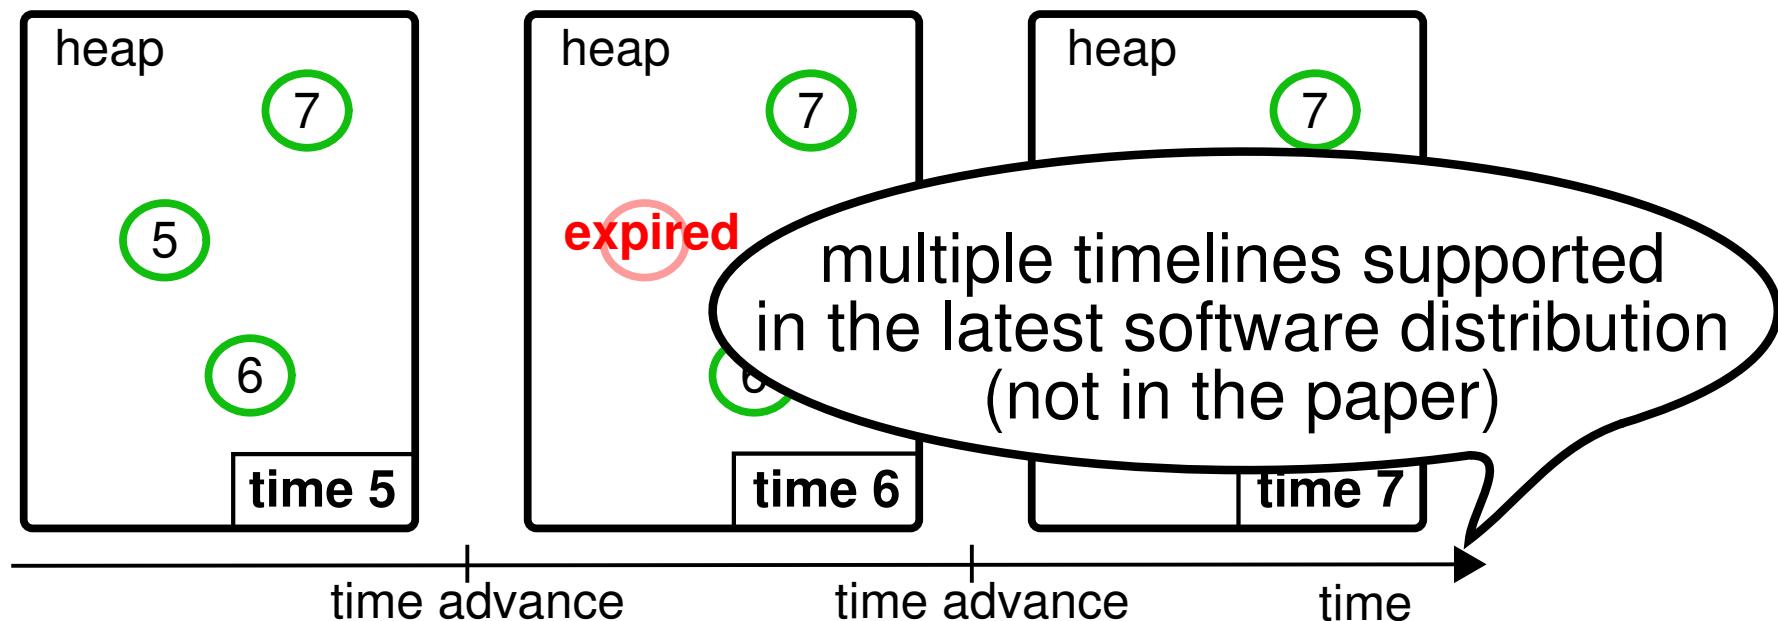
- ▶ Each memory object allocated in short-term memory has an **expiration date**
  - ◆ When an object **expires**, its memory may be reused



premature expiration  
may create dangling  
pointers

# Idea

- ▶ Each memory object allocated in short-term memory has an **expiration date**
  - ◆ When an object **expires**, its memory may be reused



# Refresh

- ▶ Extend the expiration date of an object by refreshing

# Refresh

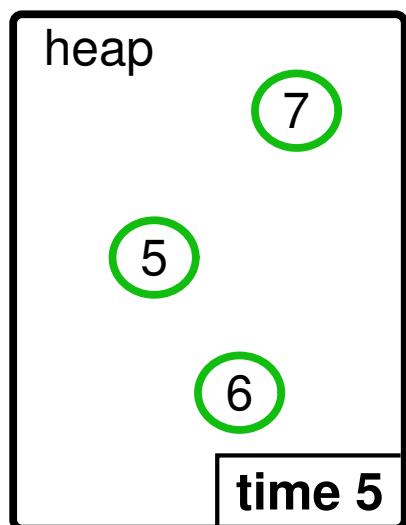
- ▶ Extend the expiration date of an object by refreshing

**new expiration date = max { current time + extension  
current expiration date}**

# Refresh

- ▶ Extend the expiration date of an object by refreshing

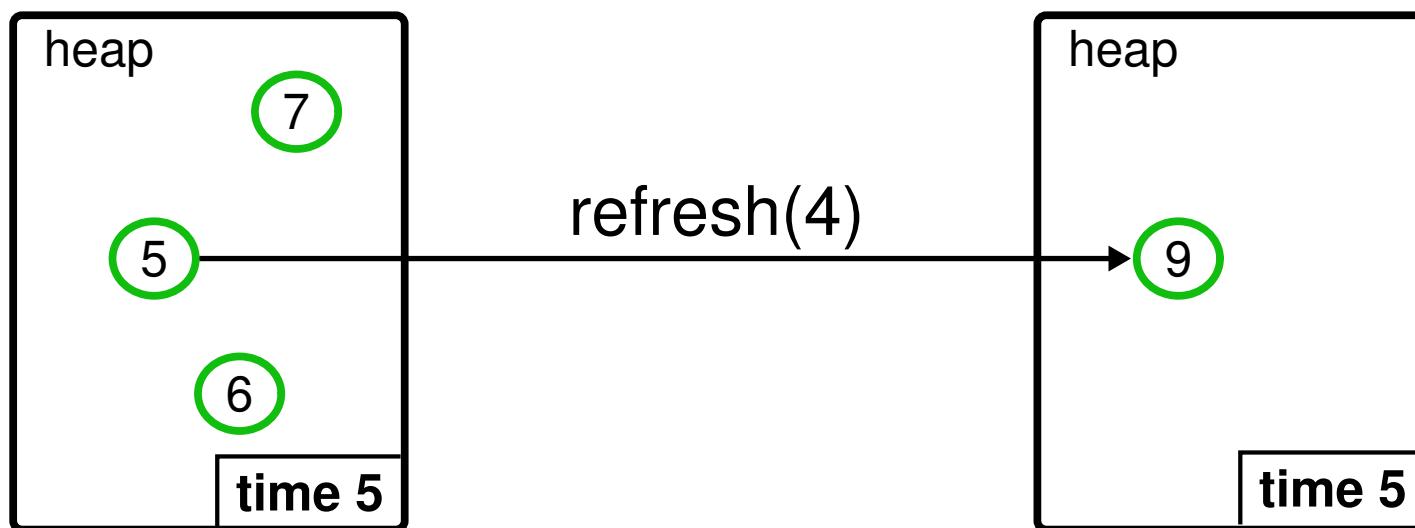
**new expiration date = max { current time + extension  
current expiration date }**



# Refresh

- Extend the expiration date of an object by refreshing

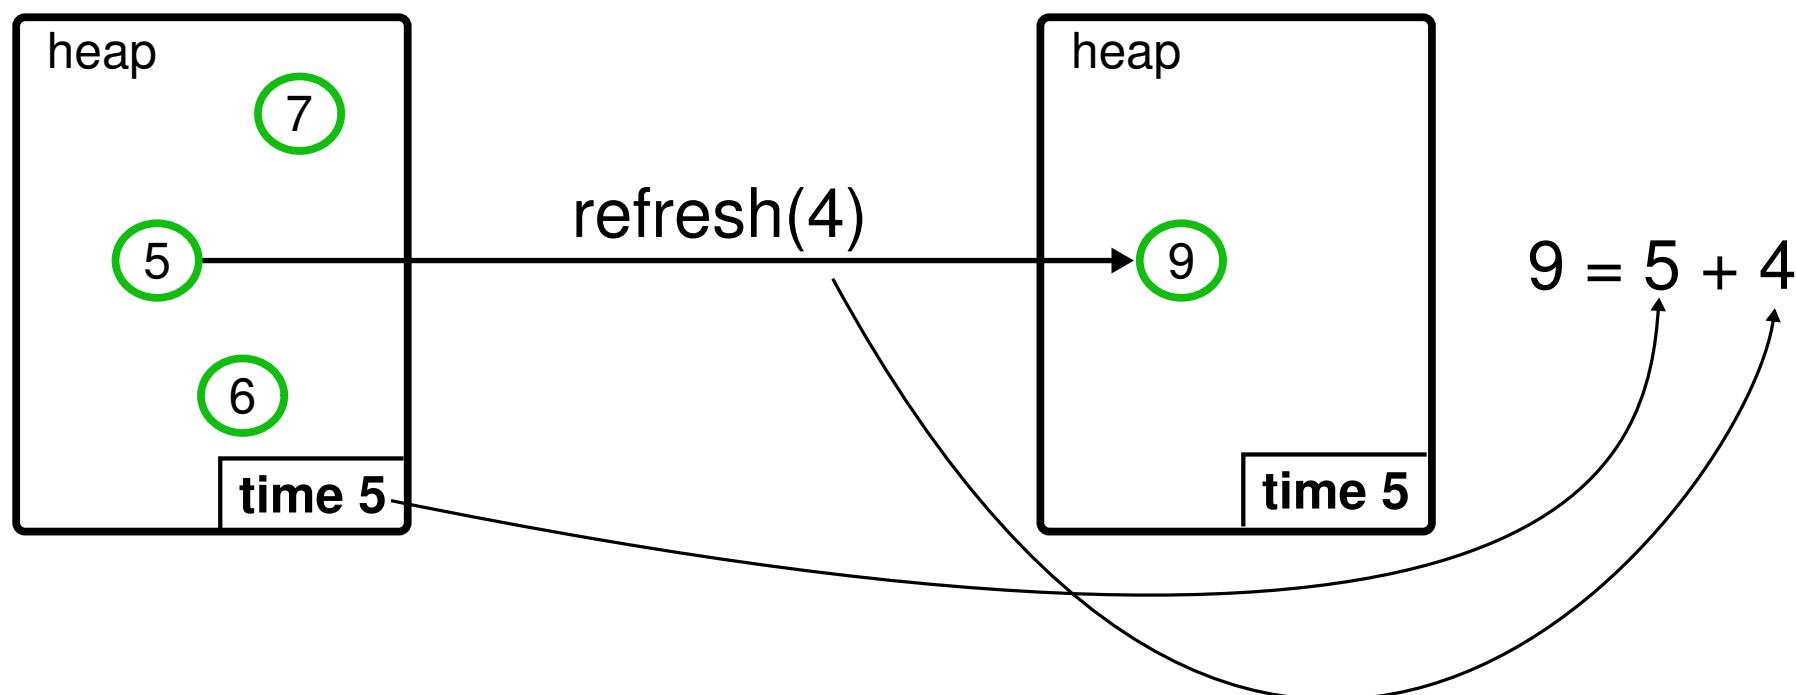
**new expiration date = max { current time + extension  
current expiration date }**



# Refresh

- Extend the expiration date of an object by refreshing

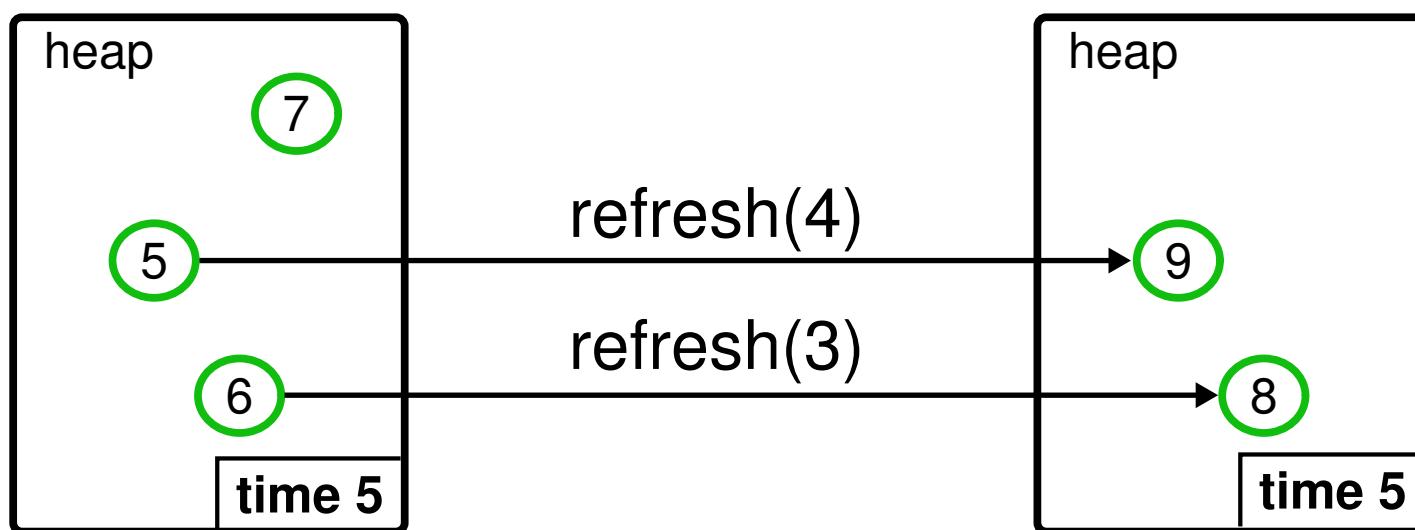
**new expiration date = max { current time + extension  
current expiration date }**



# Refresh

- Extend the expiration date of an object by refreshing

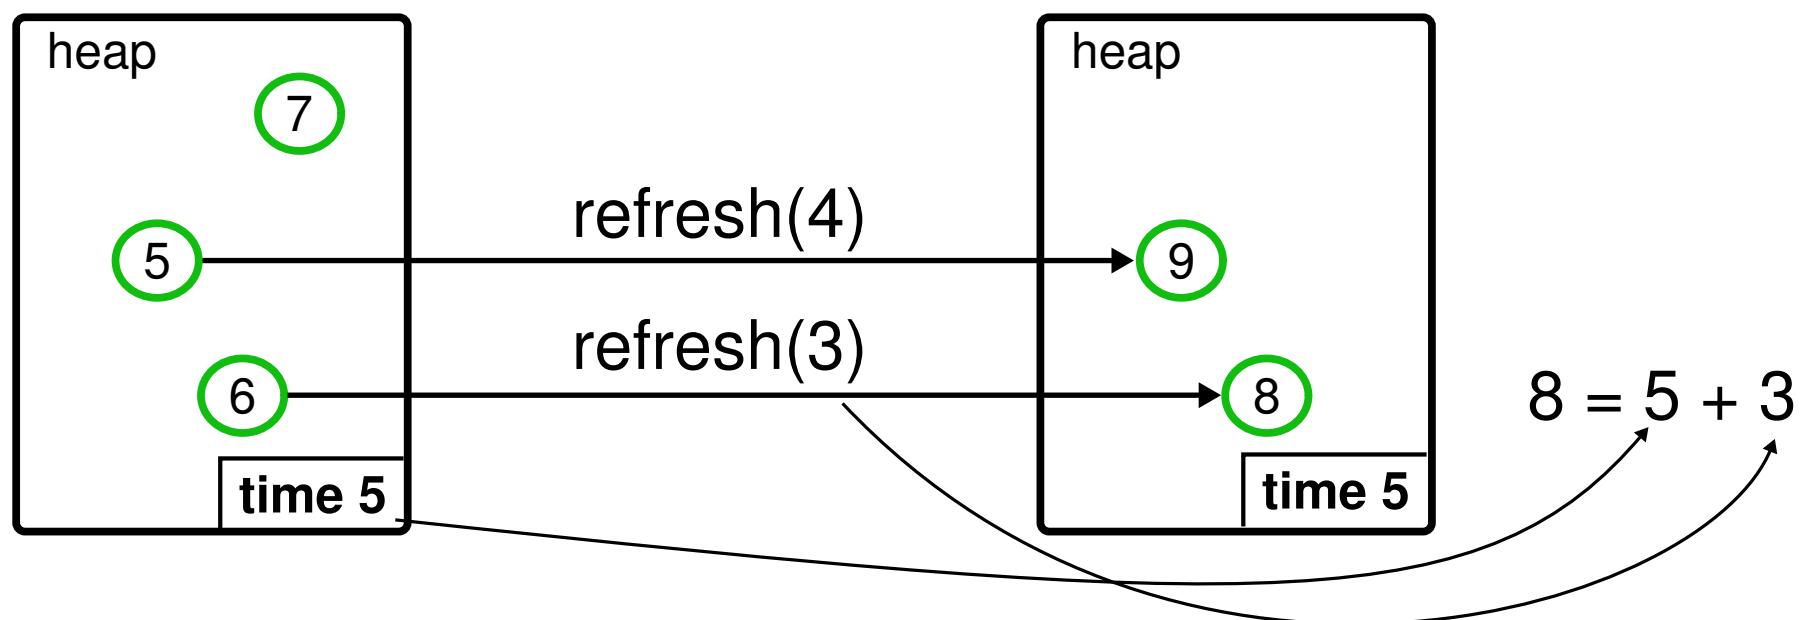
**new expiration date = max { current time + extension  
current expiration date }**



# Refresh

- Extend the expiration date of an object by refreshing

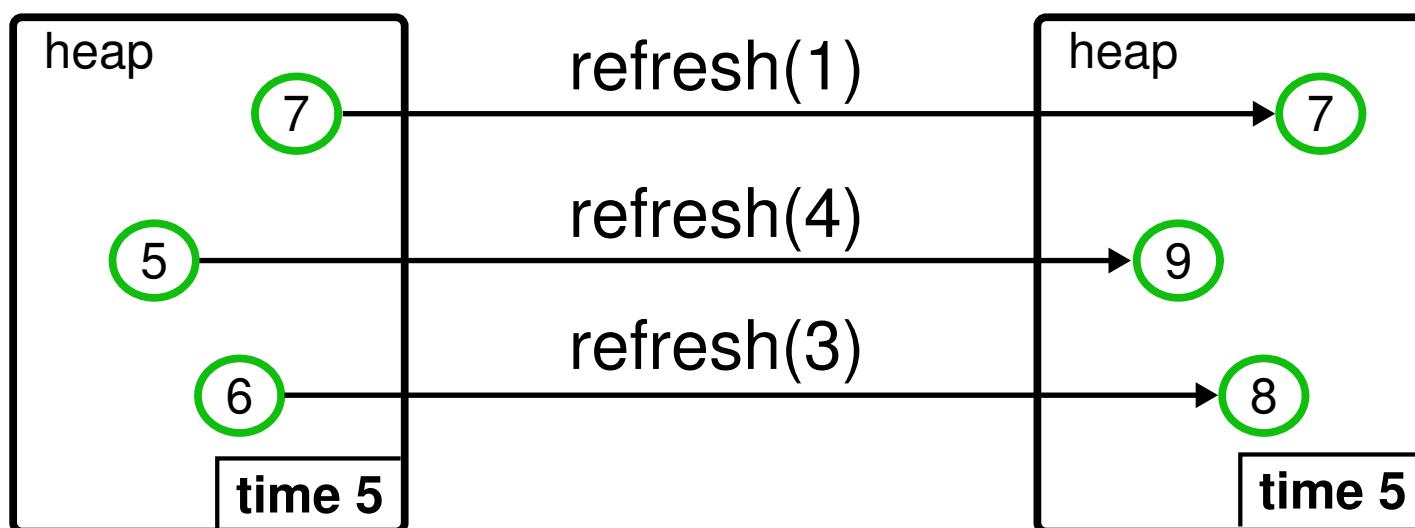
**new expiration date = max { current time + extension  
current expiration date }**



# Refresh

- Extend the expiration date of an object by refreshing

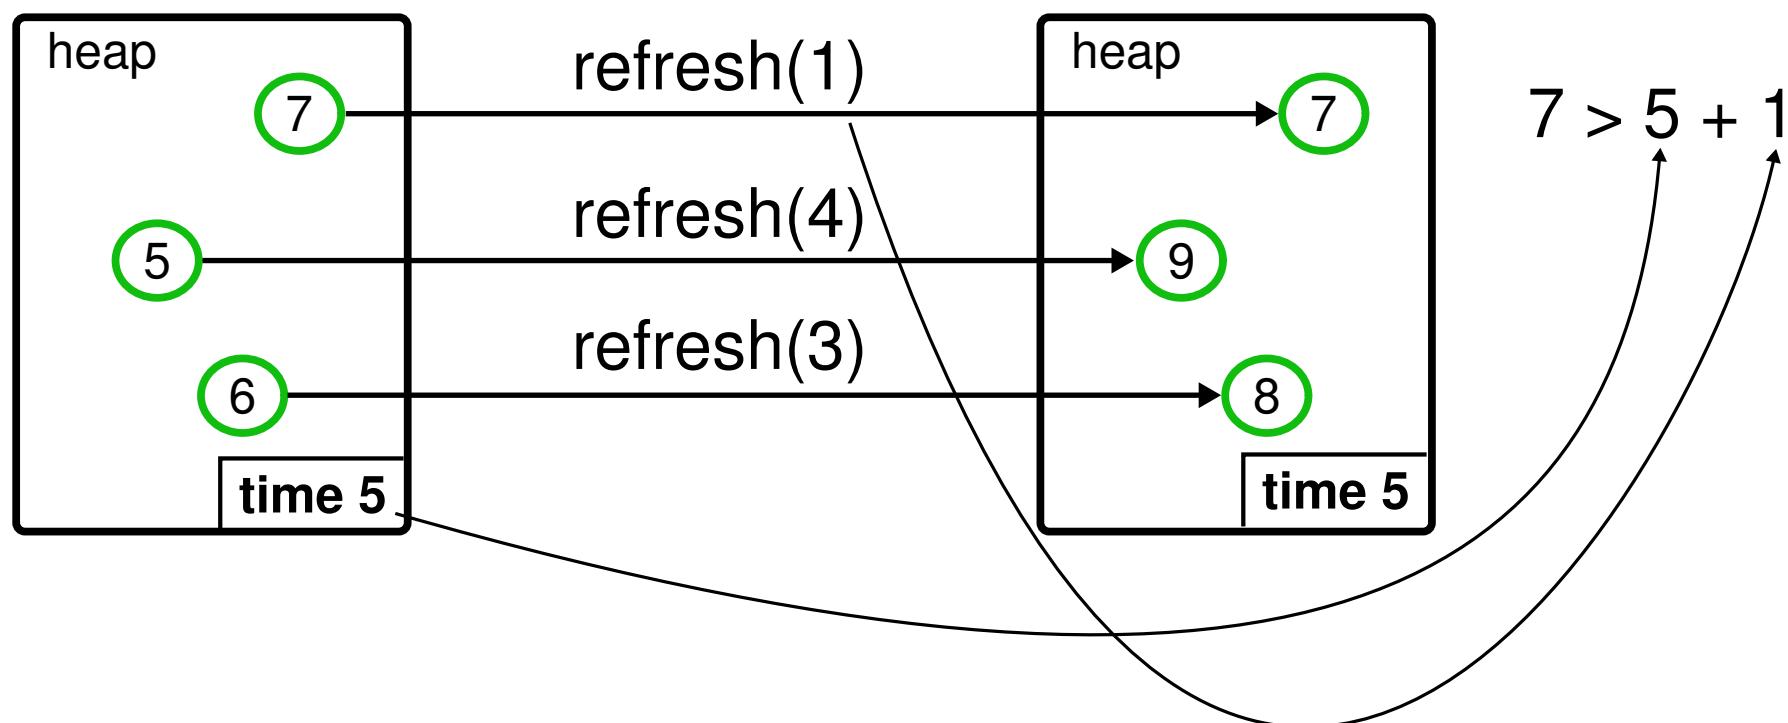
**new expiration date = max** { current time + extension  
current expiration date



# Refresh

- Extend the expiration date of an object by refreshing

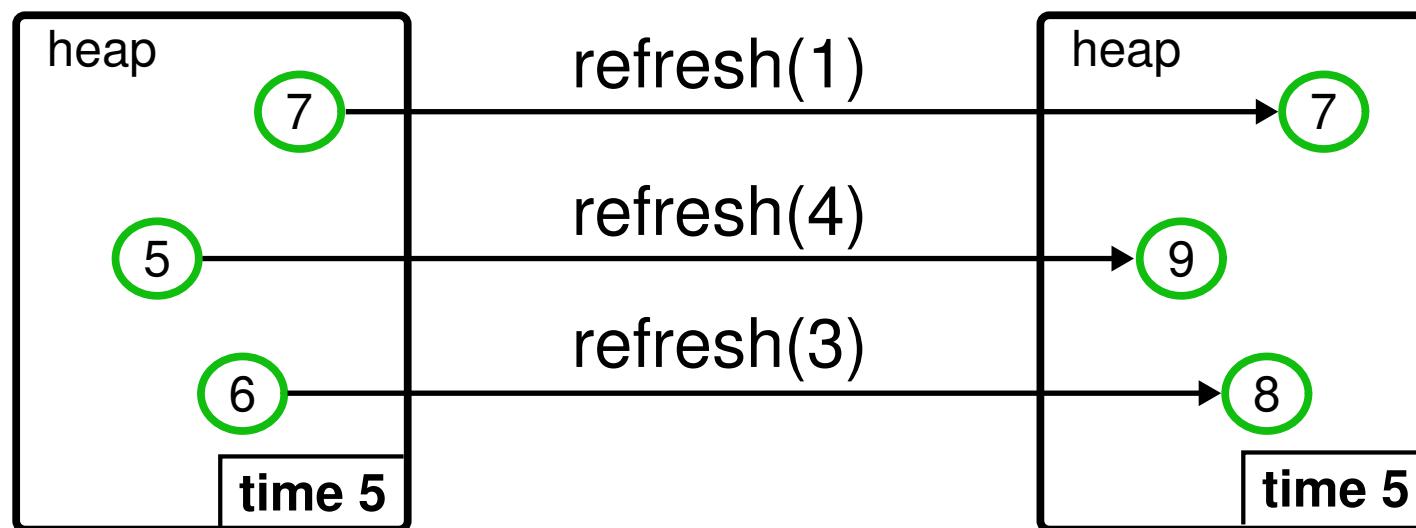
**new expiration date = max** { current time + extension  
current expiration date



# Refresh

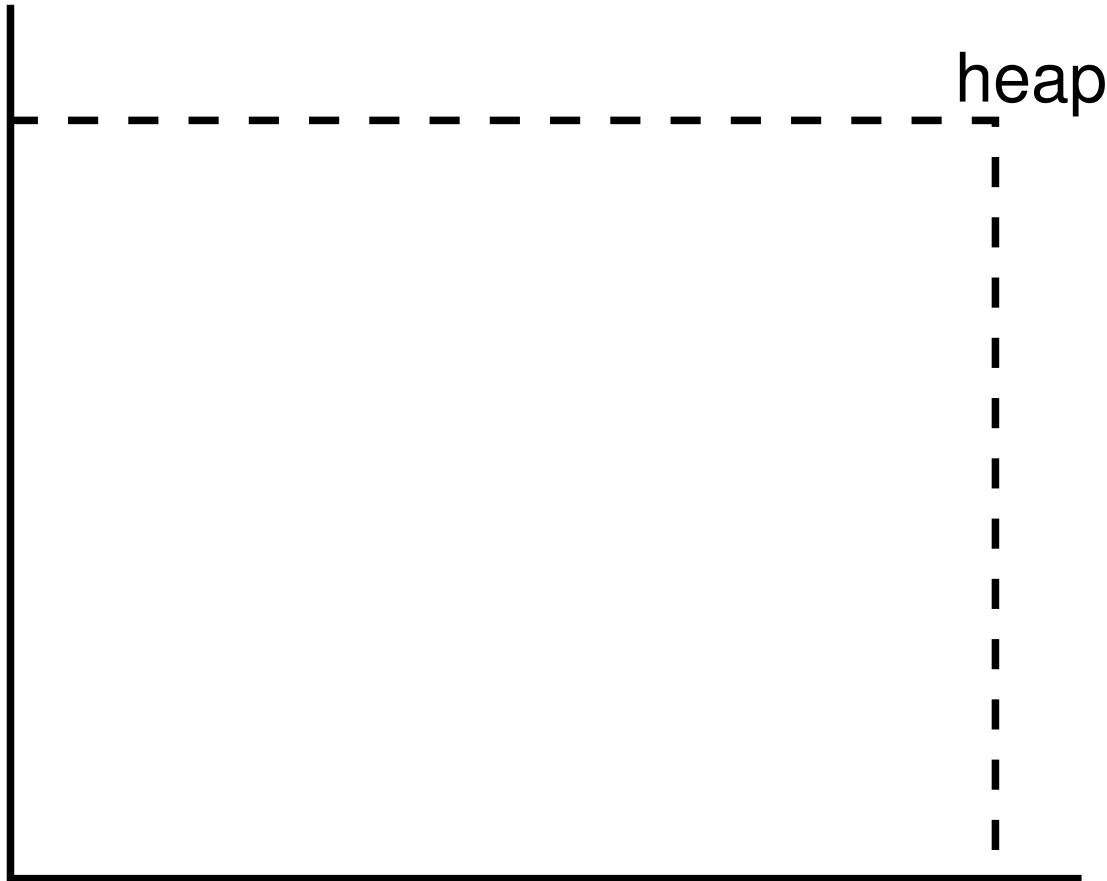
- Extend the expiration date of an object by refreshing

**new expiration date = max** { current time + extension  
current expiration date

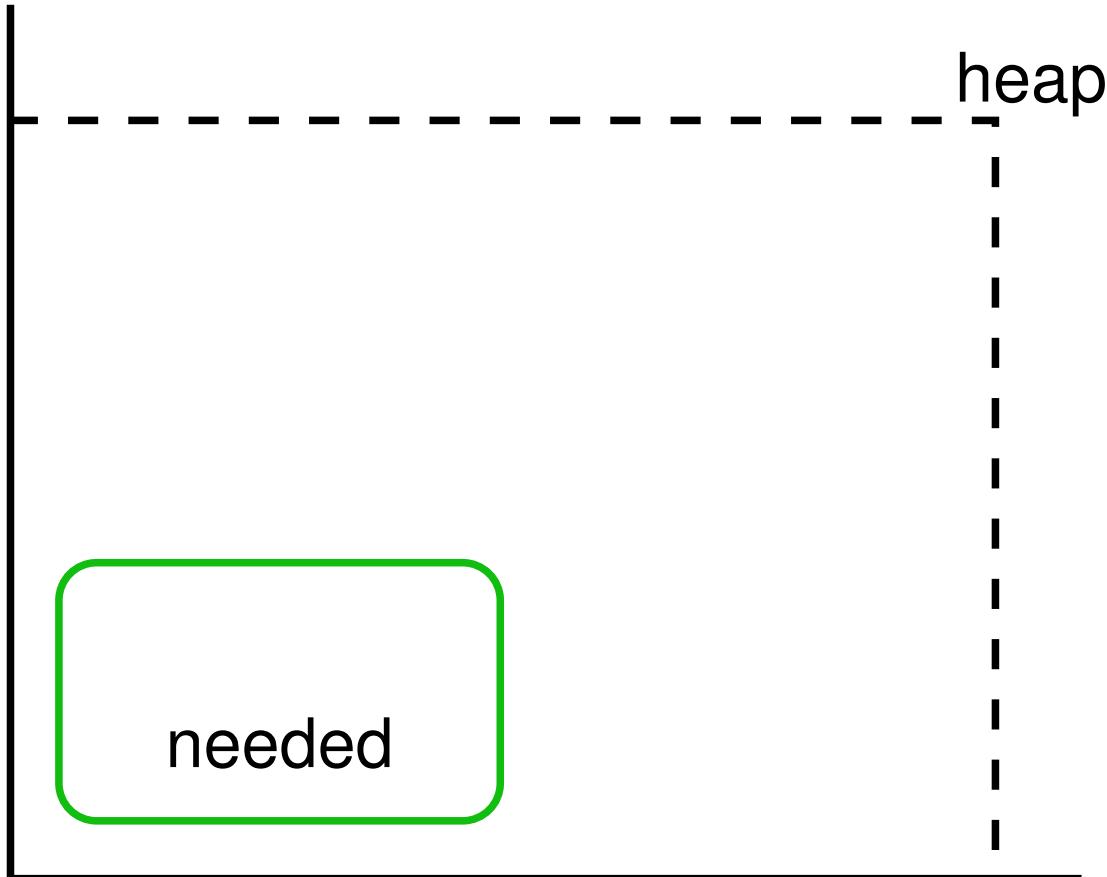


- Expiration extensions do not accumulate
- Objects may be refreshed multiple times as long as they have not expired

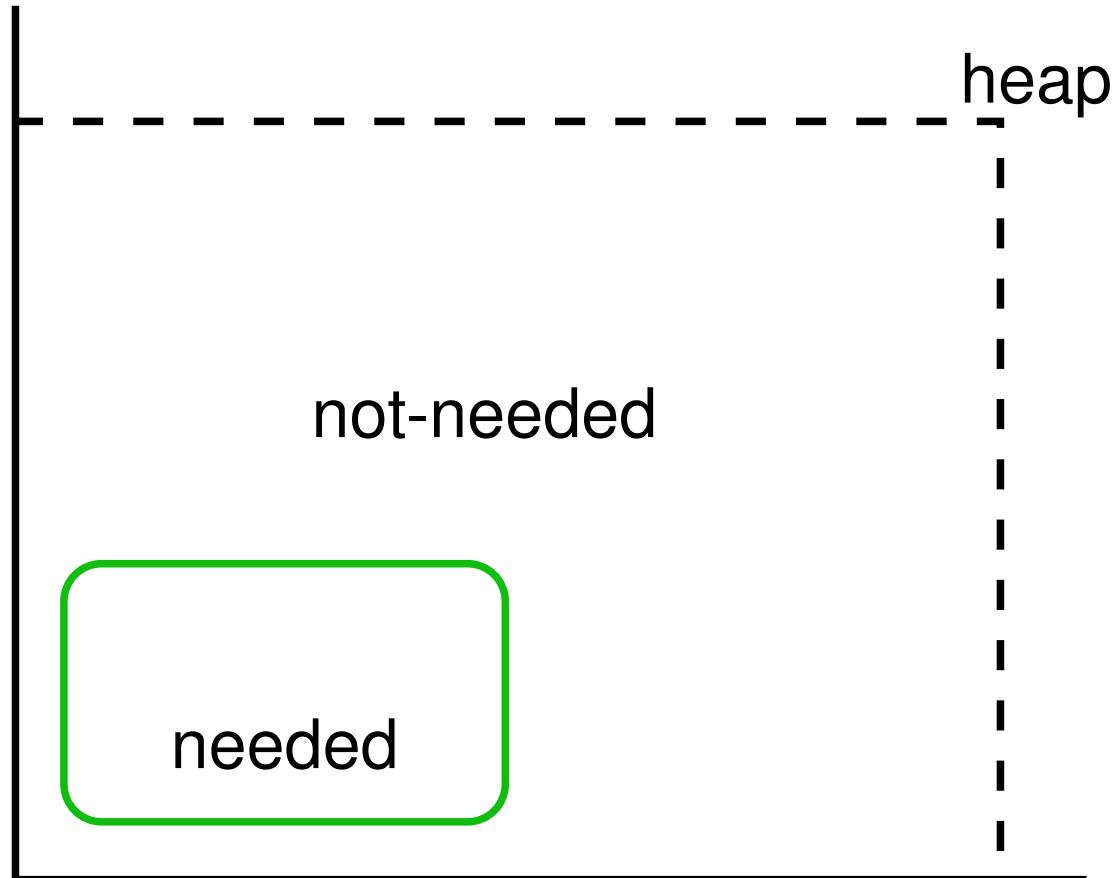
# Heap Management



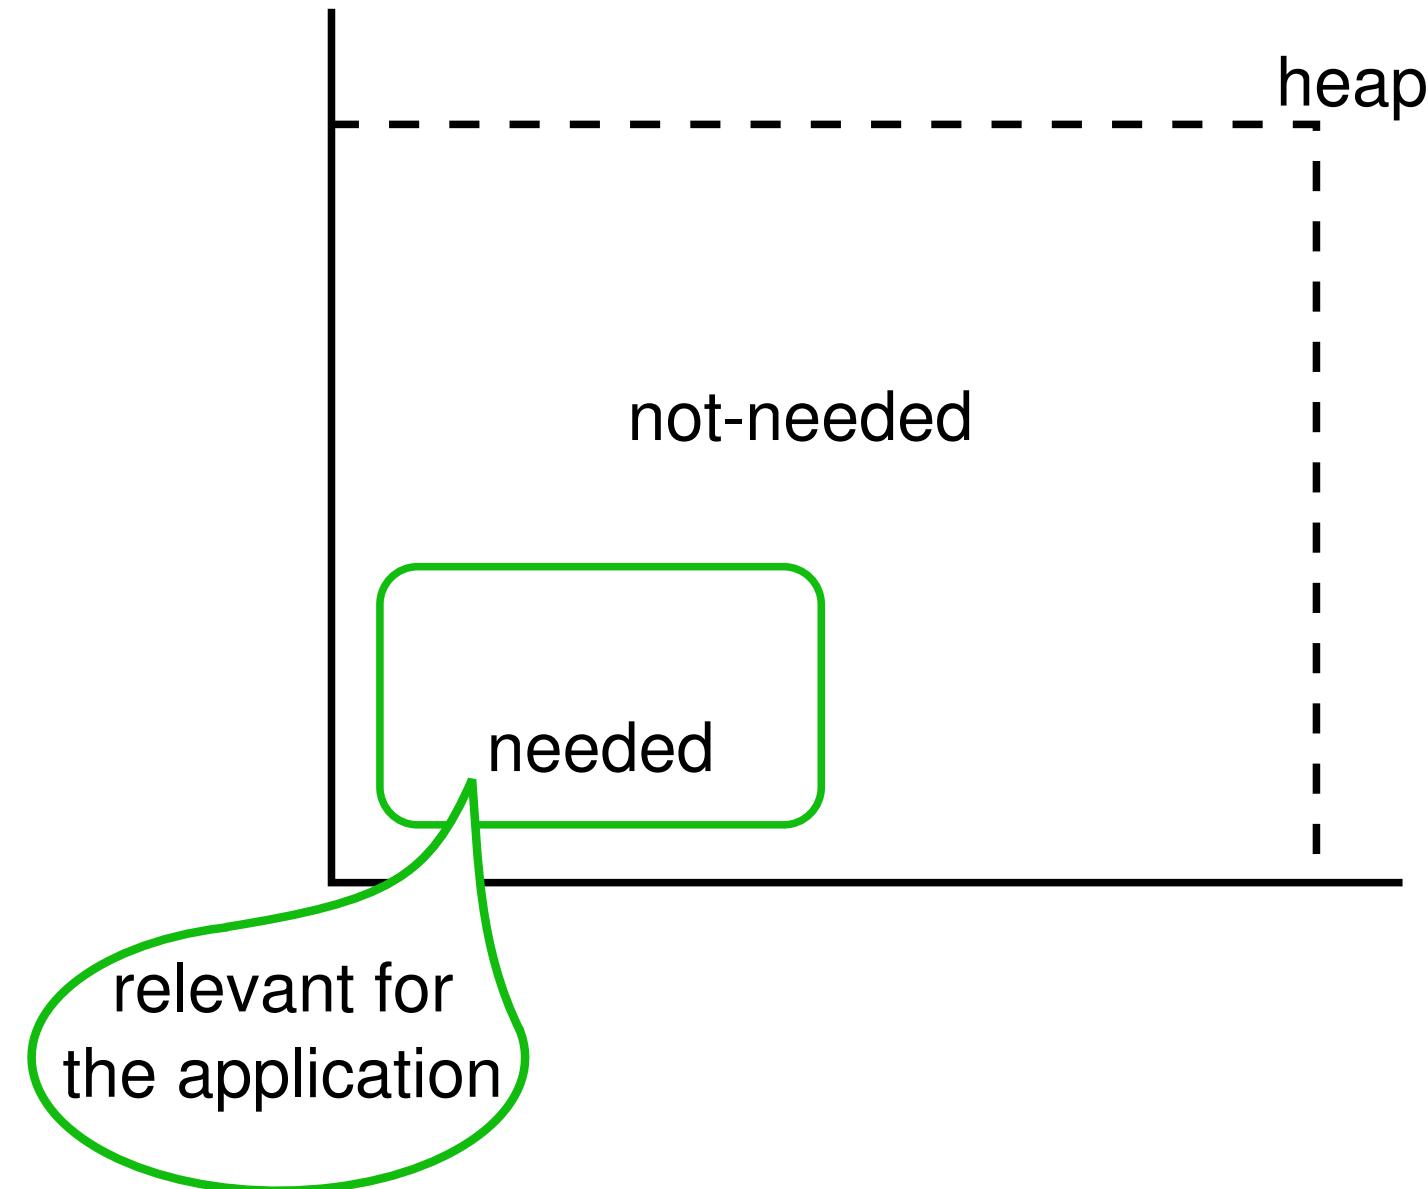
# Heap Management



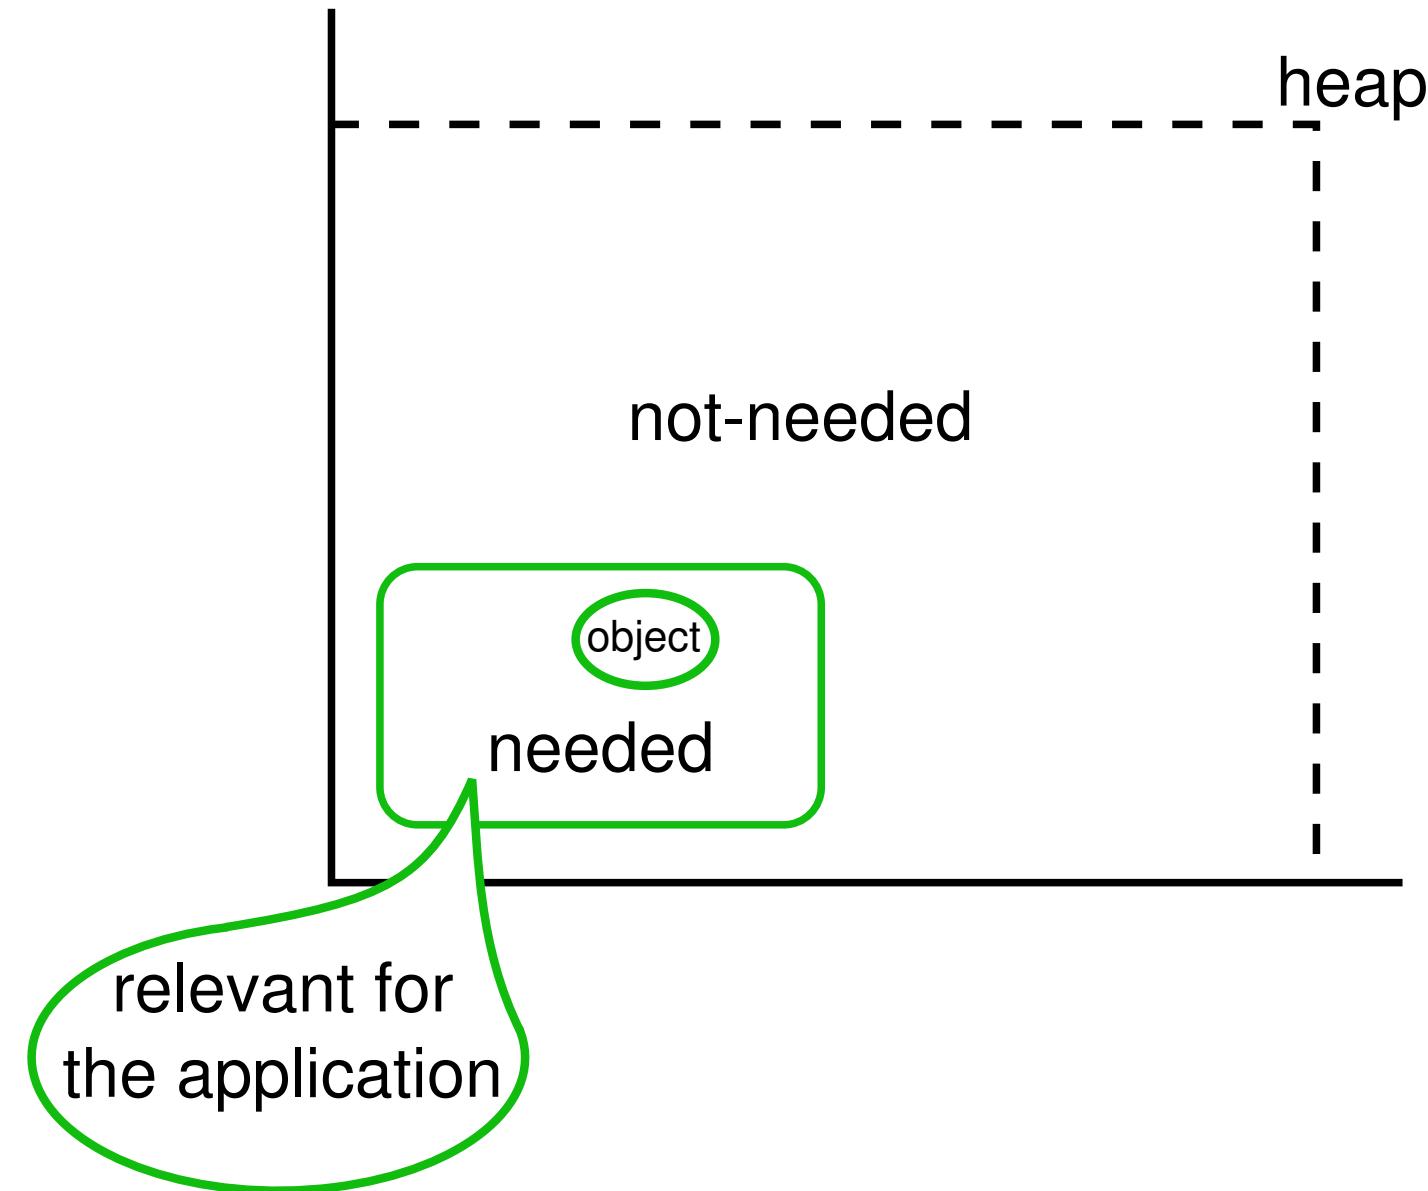
# Heap Management



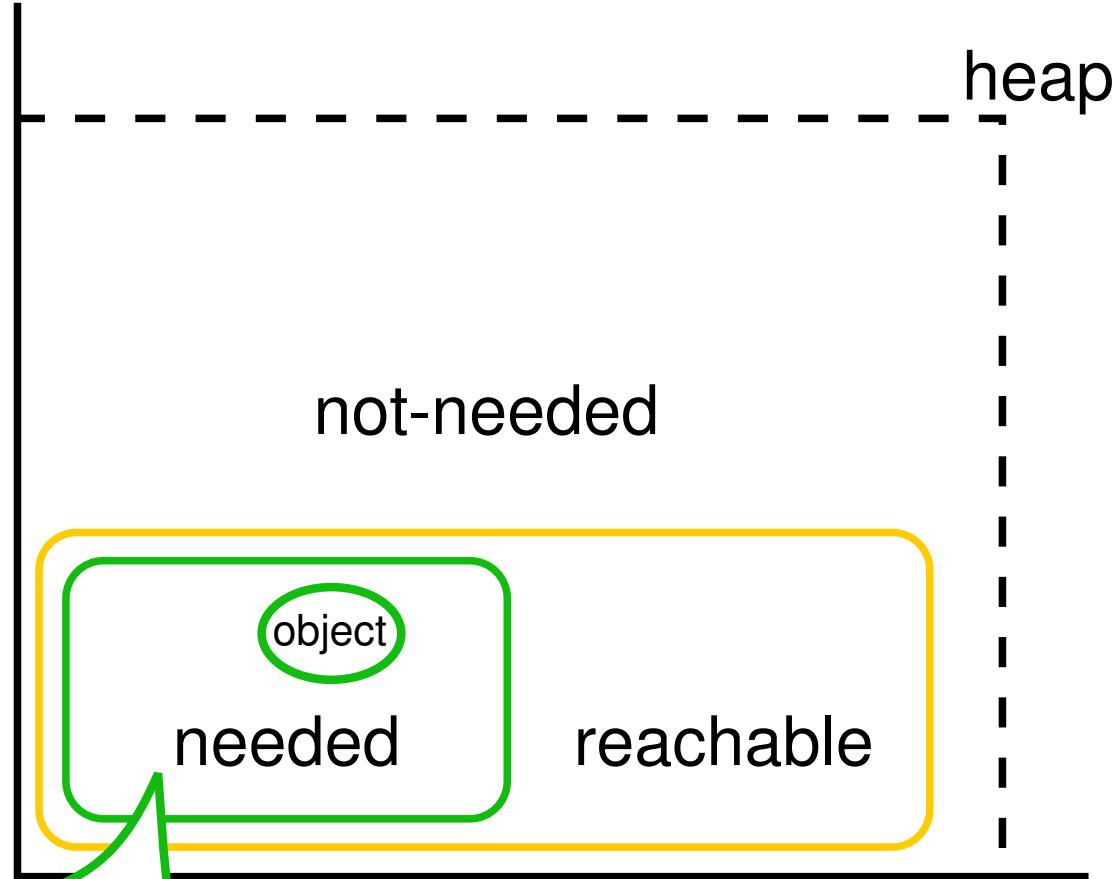
# Heap Management



# Heap Management

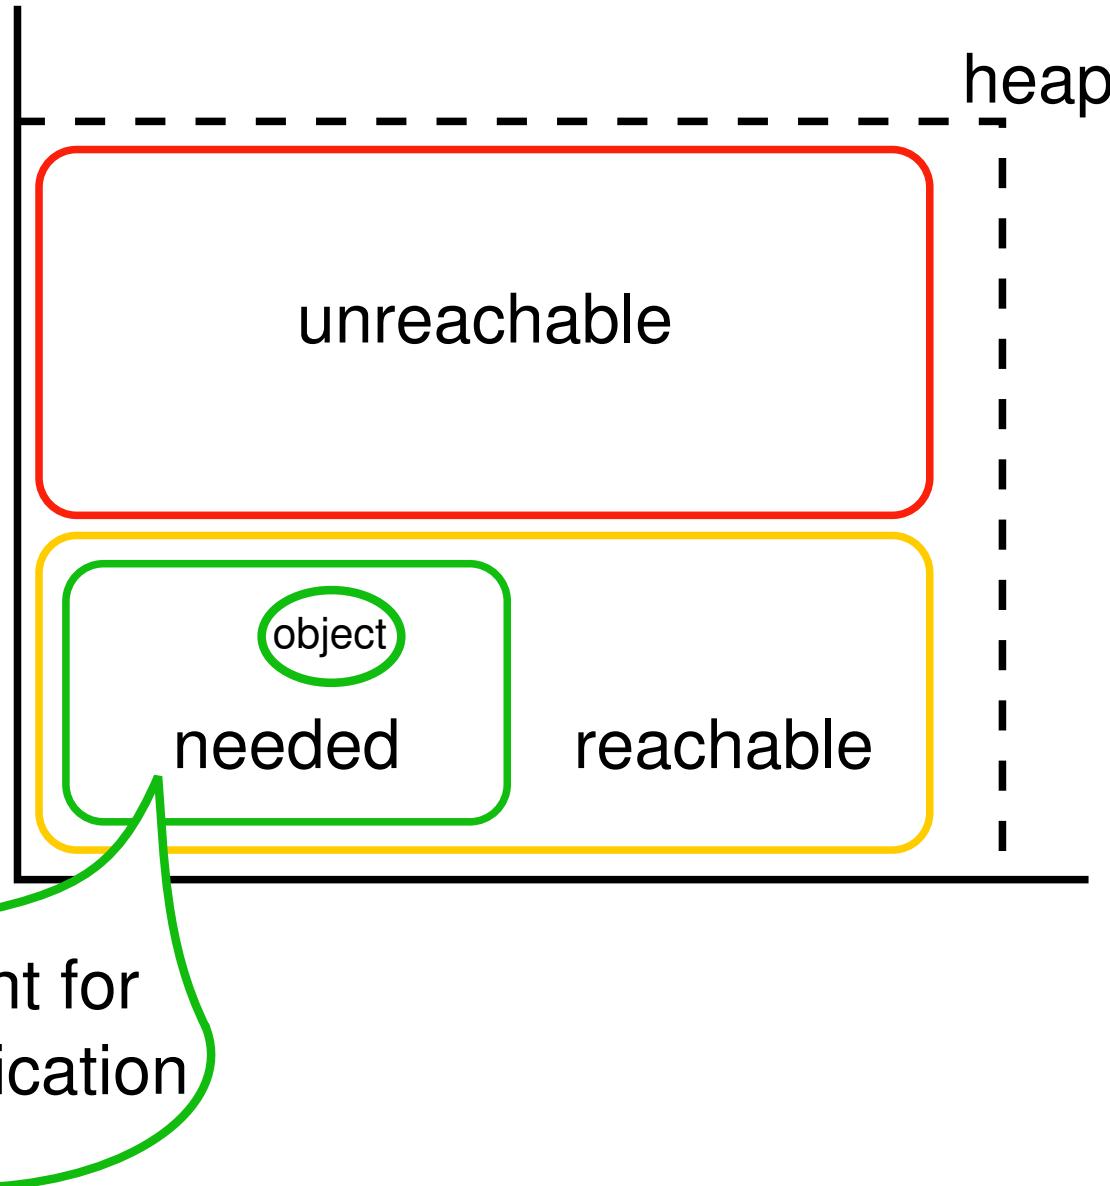


# Heap Management

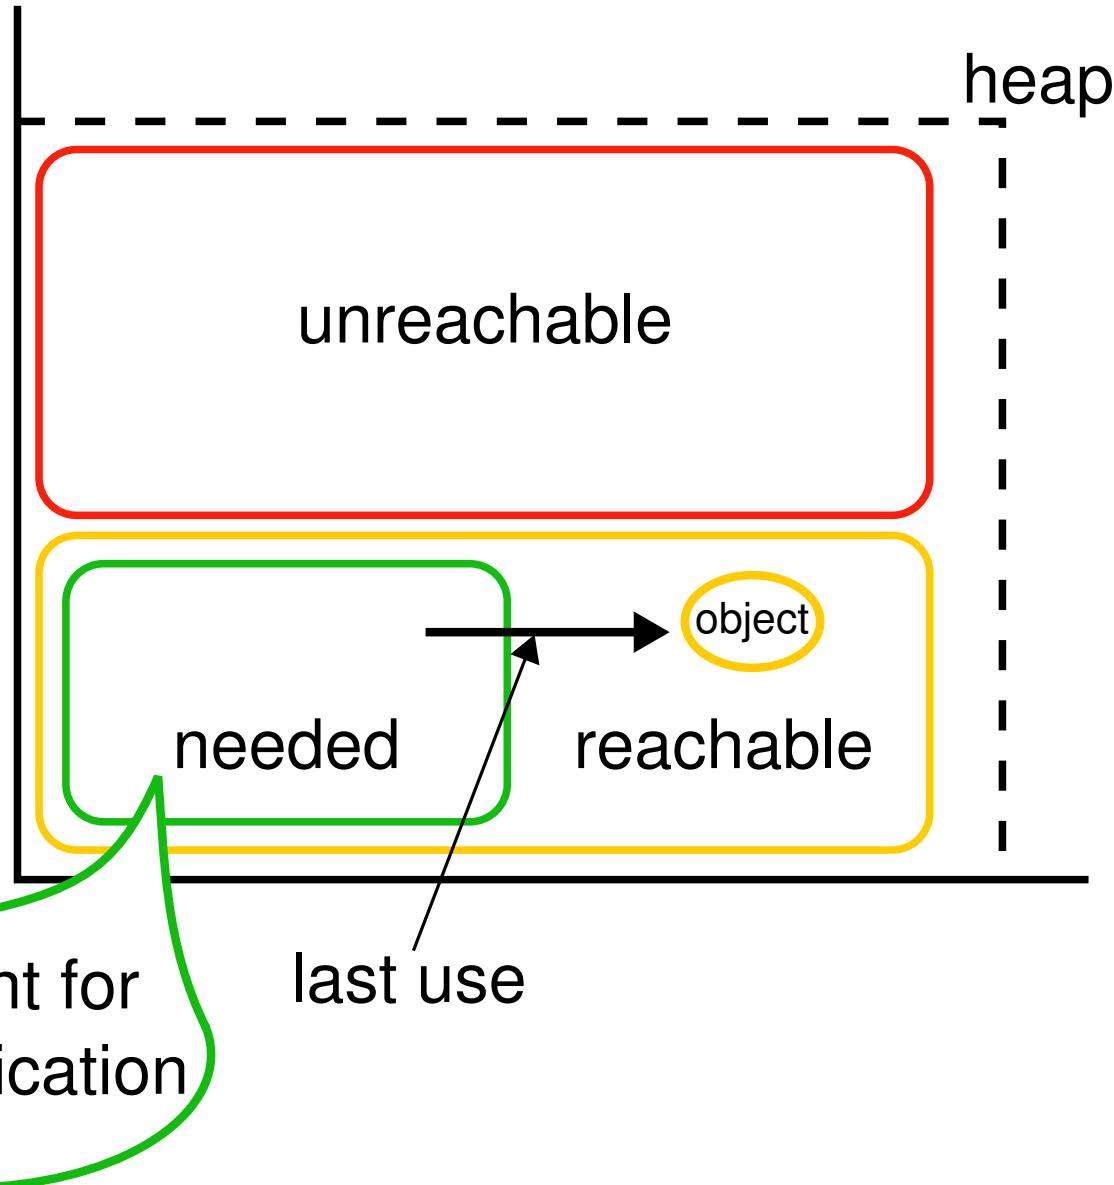


relevant for  
the application

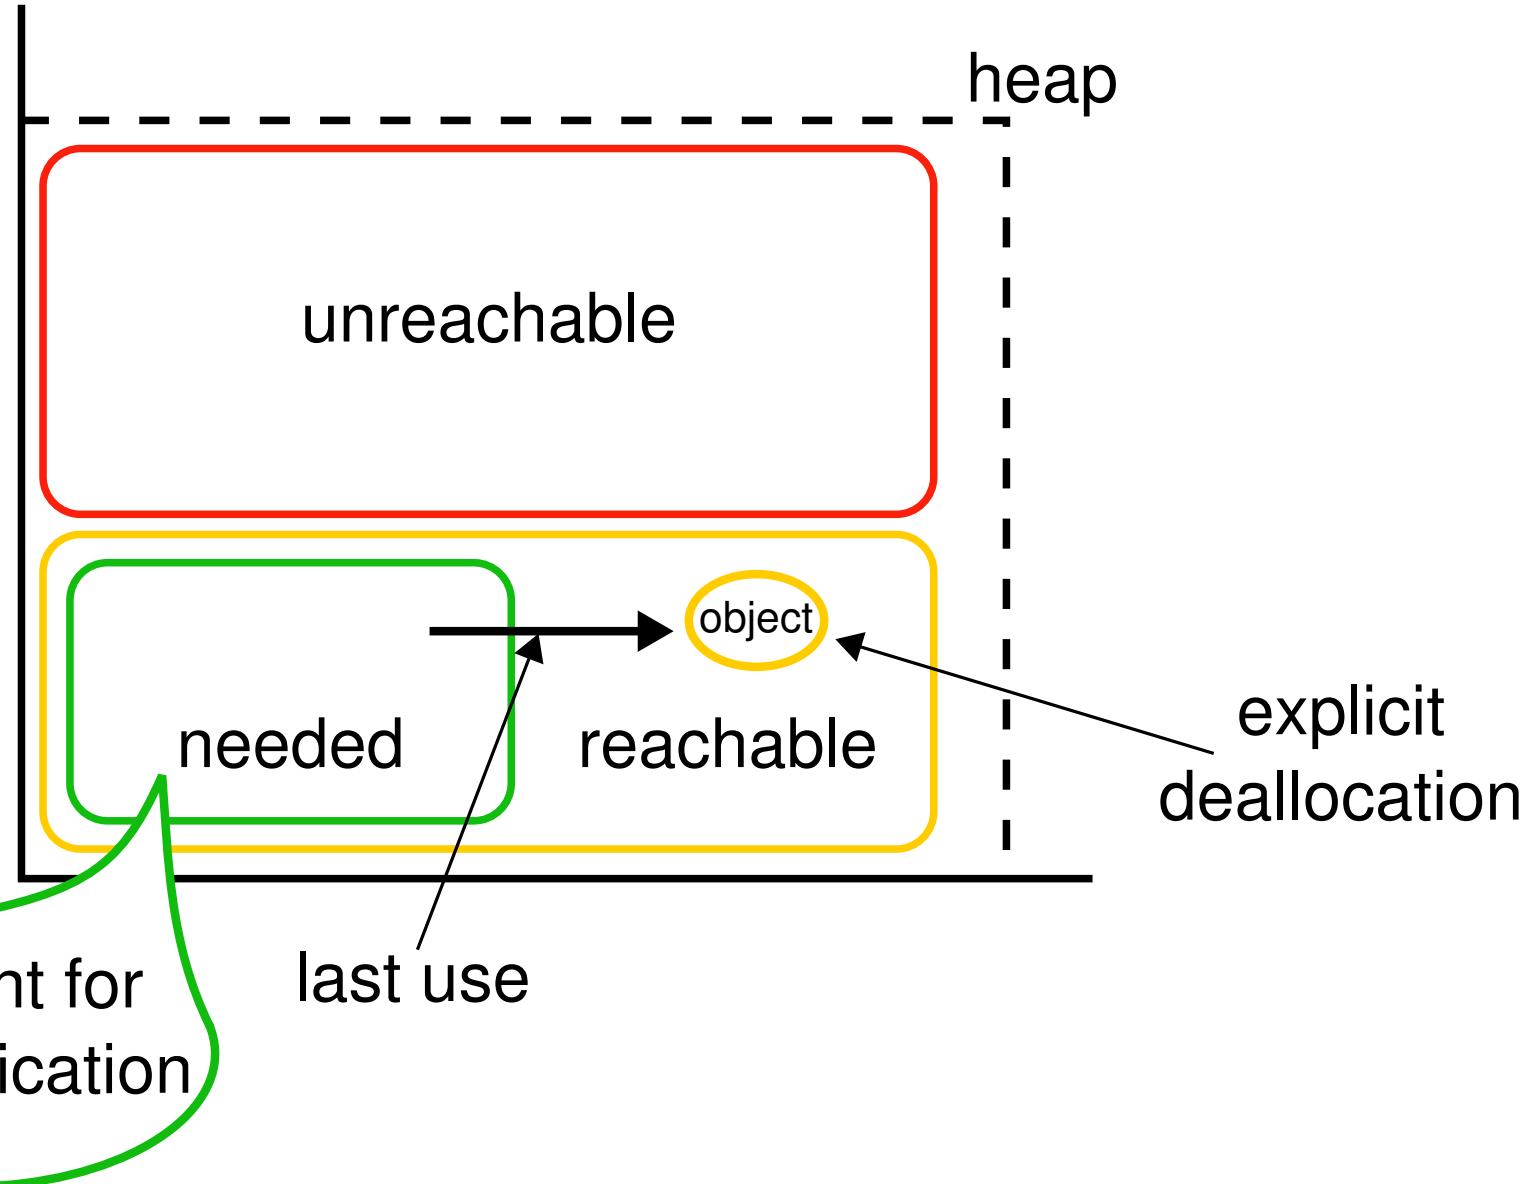
# Heap Management



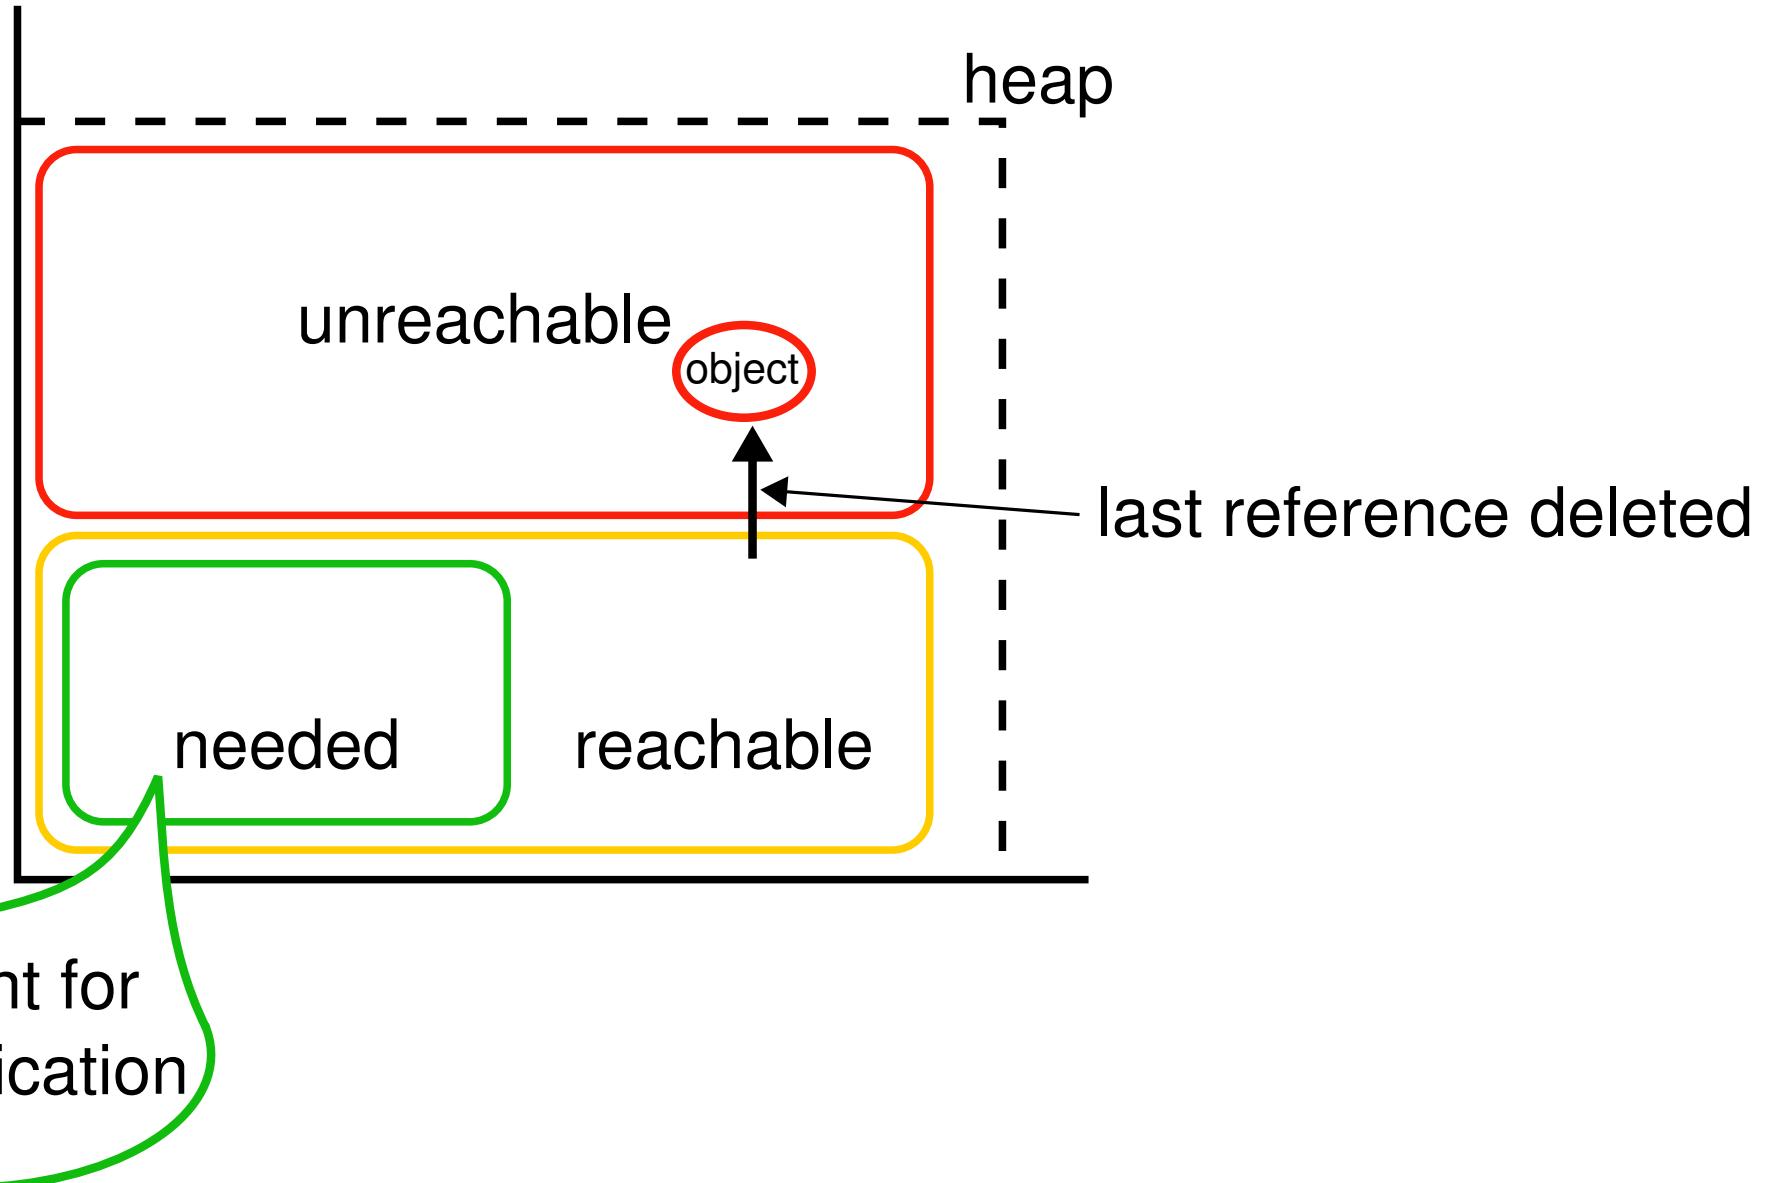
# Heap Management



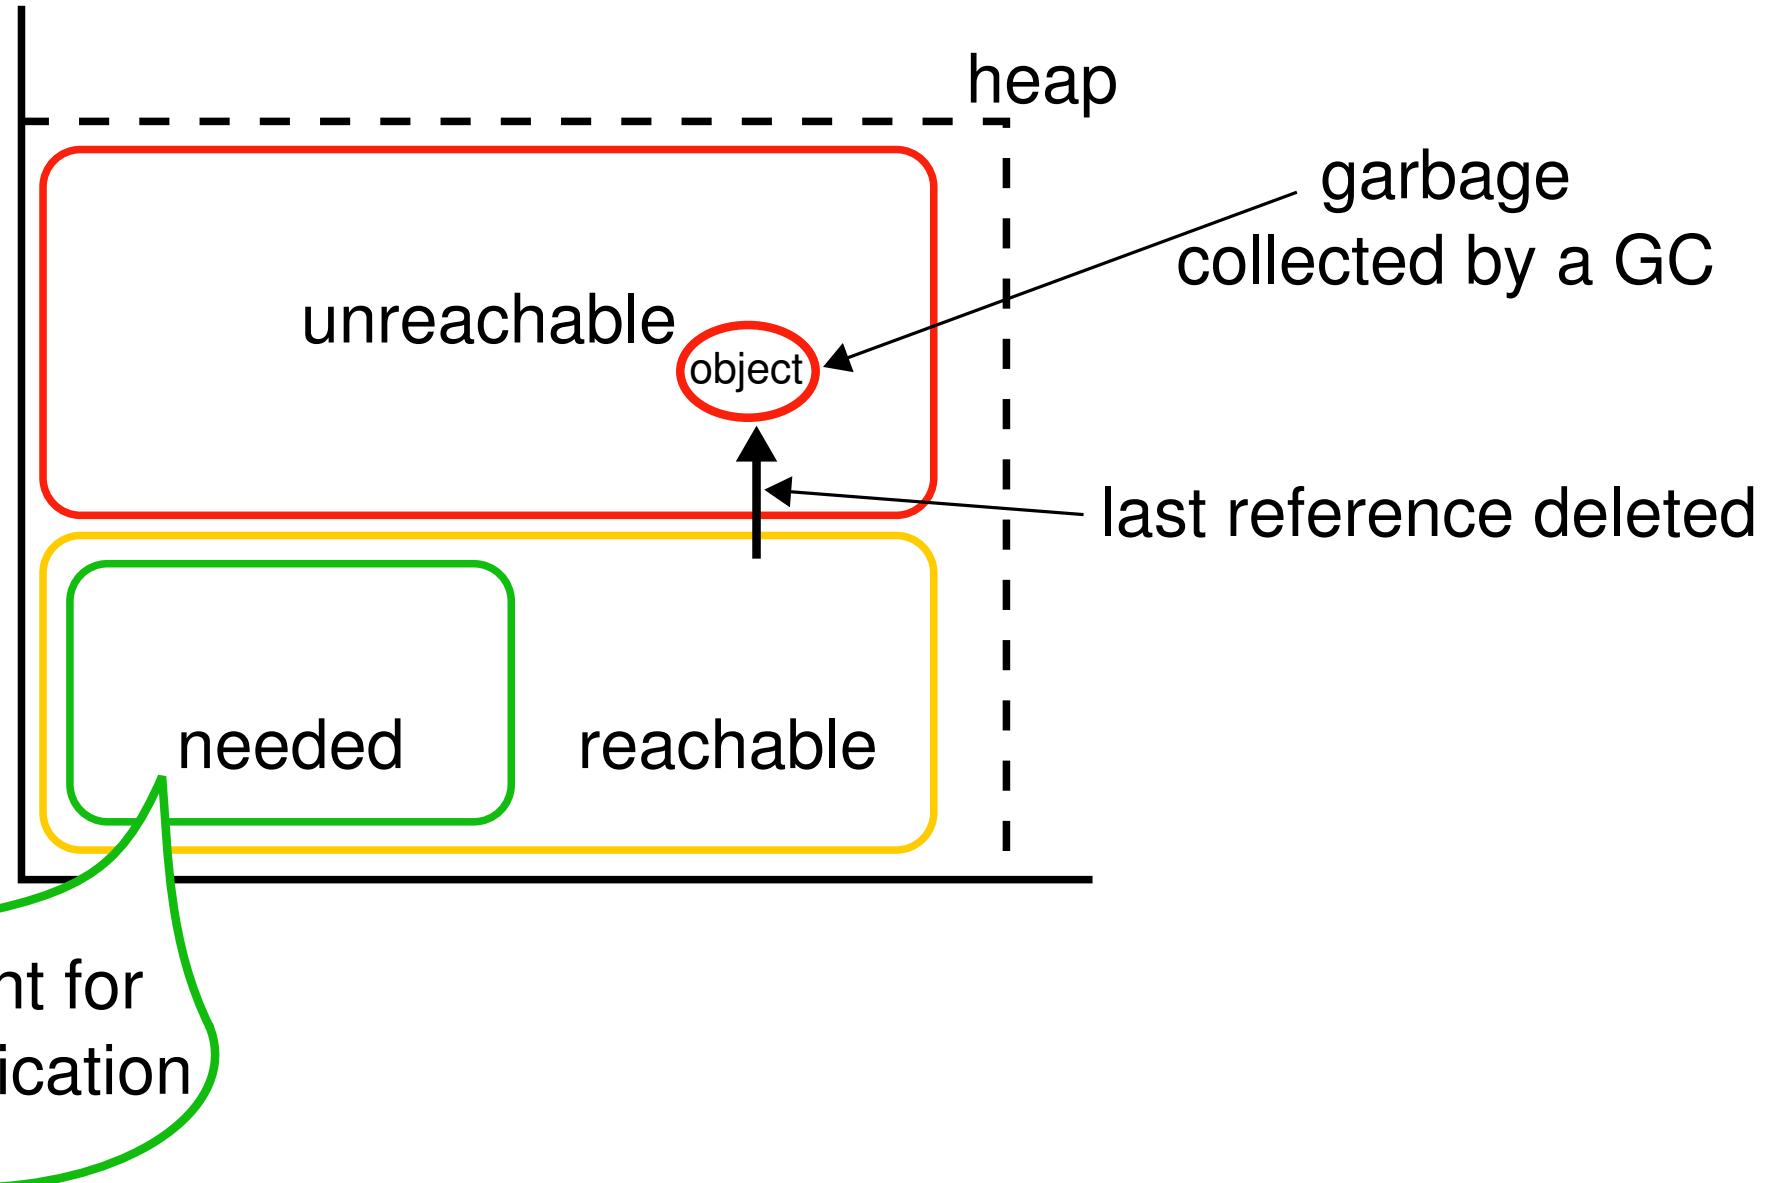
# Heap Management



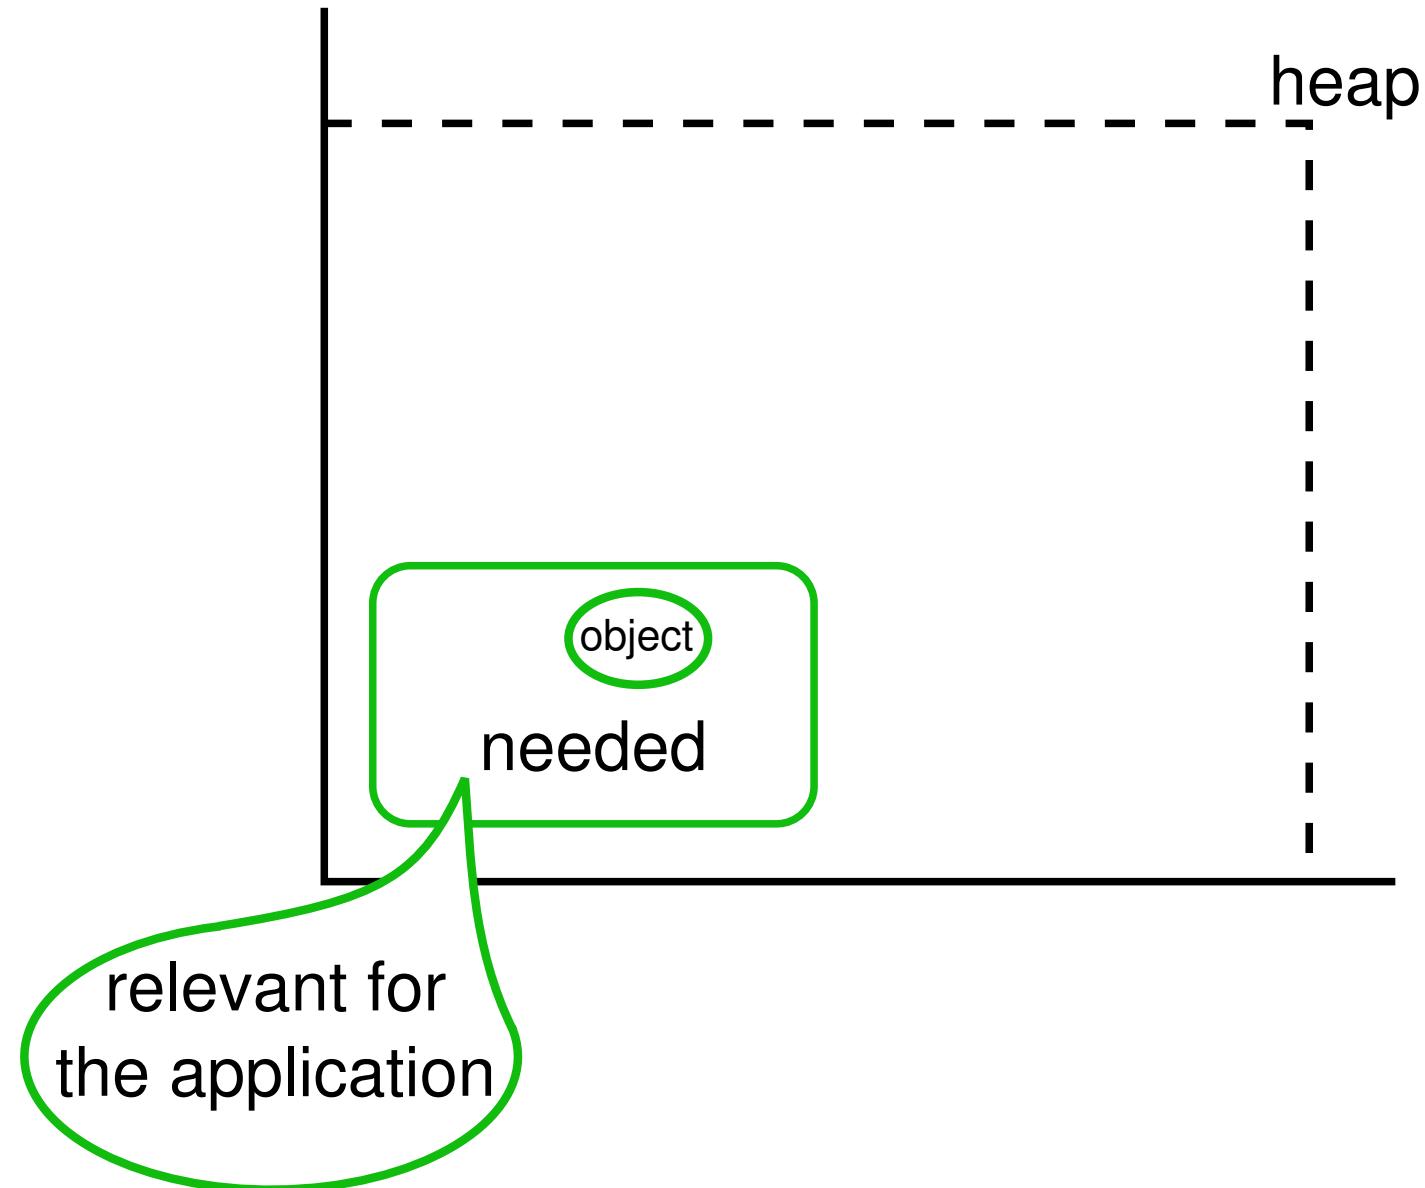
# Heap Management



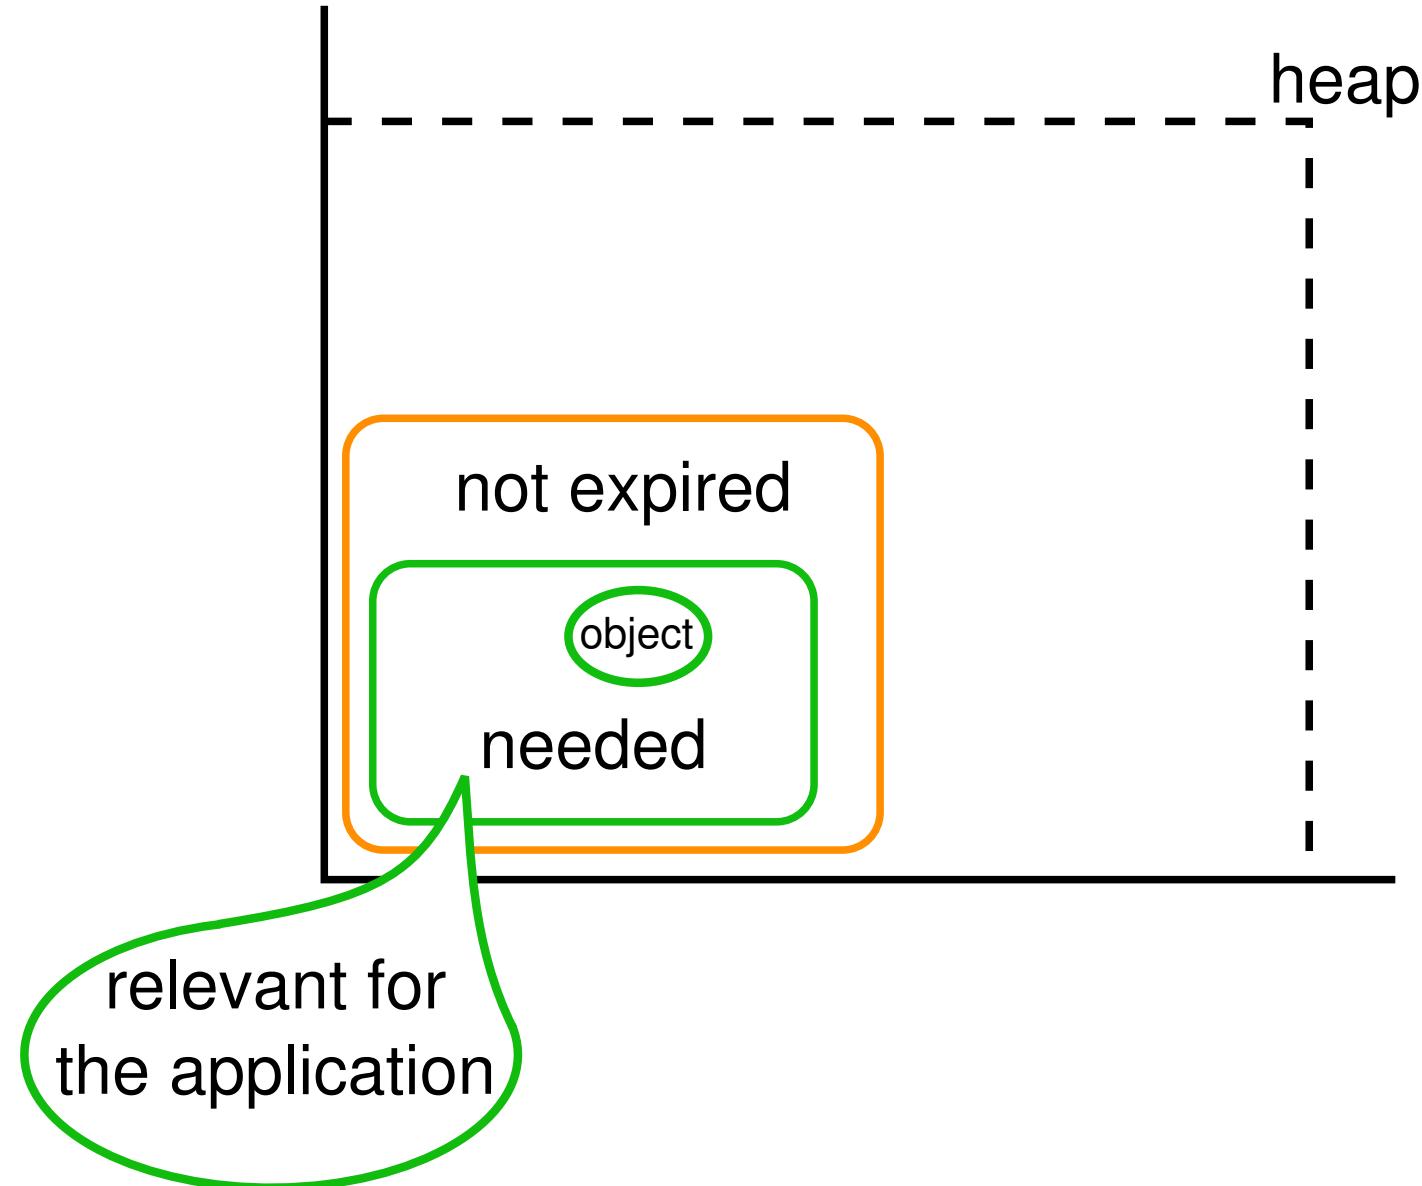
# Heap Management



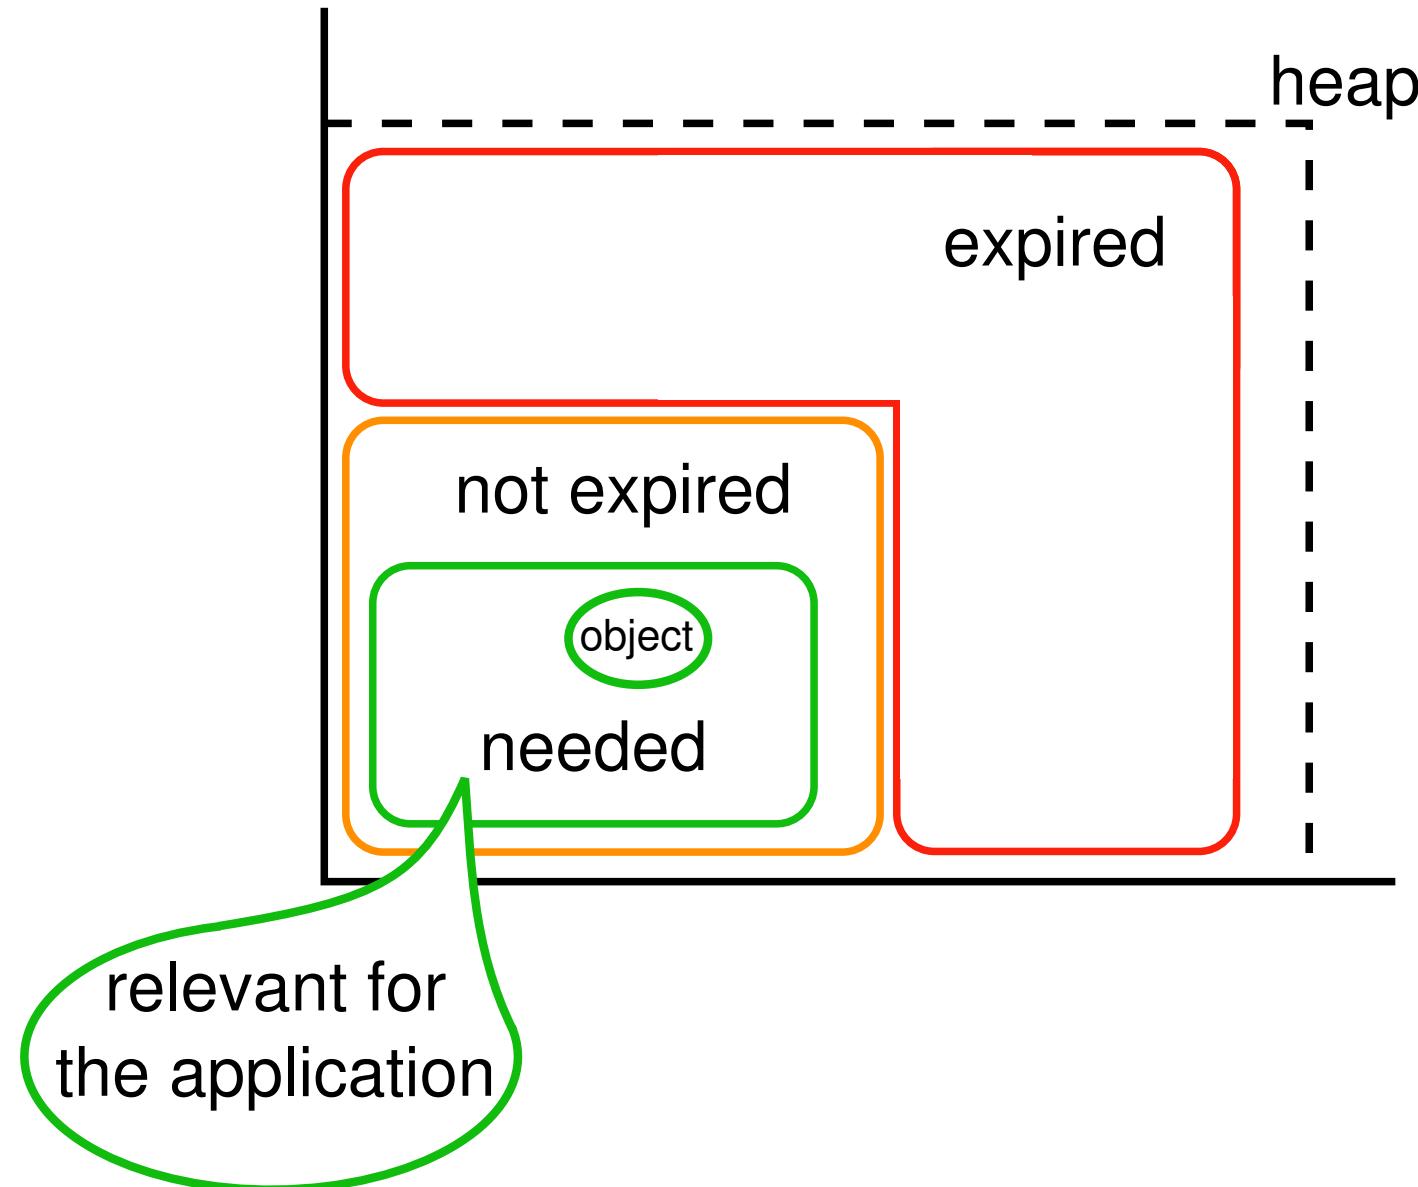
# Short-term memory



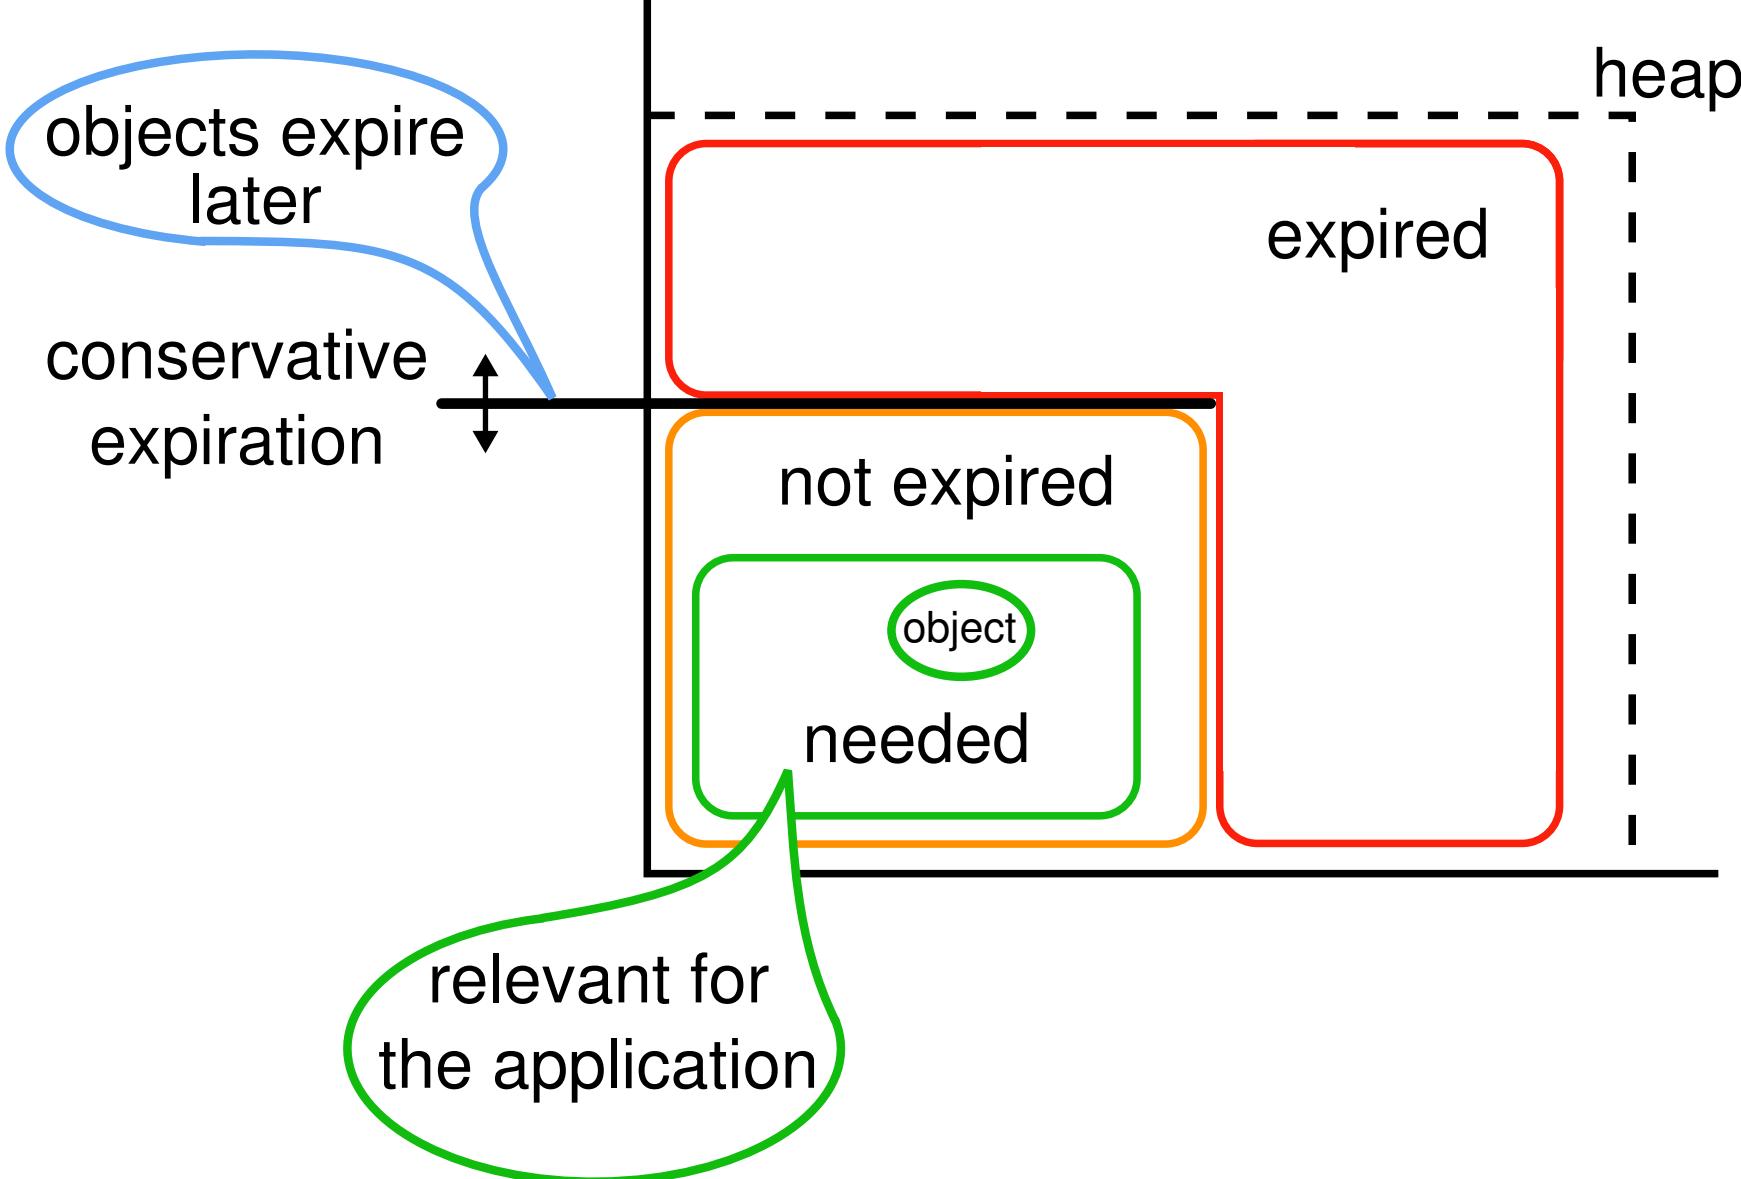
# Short-term memory



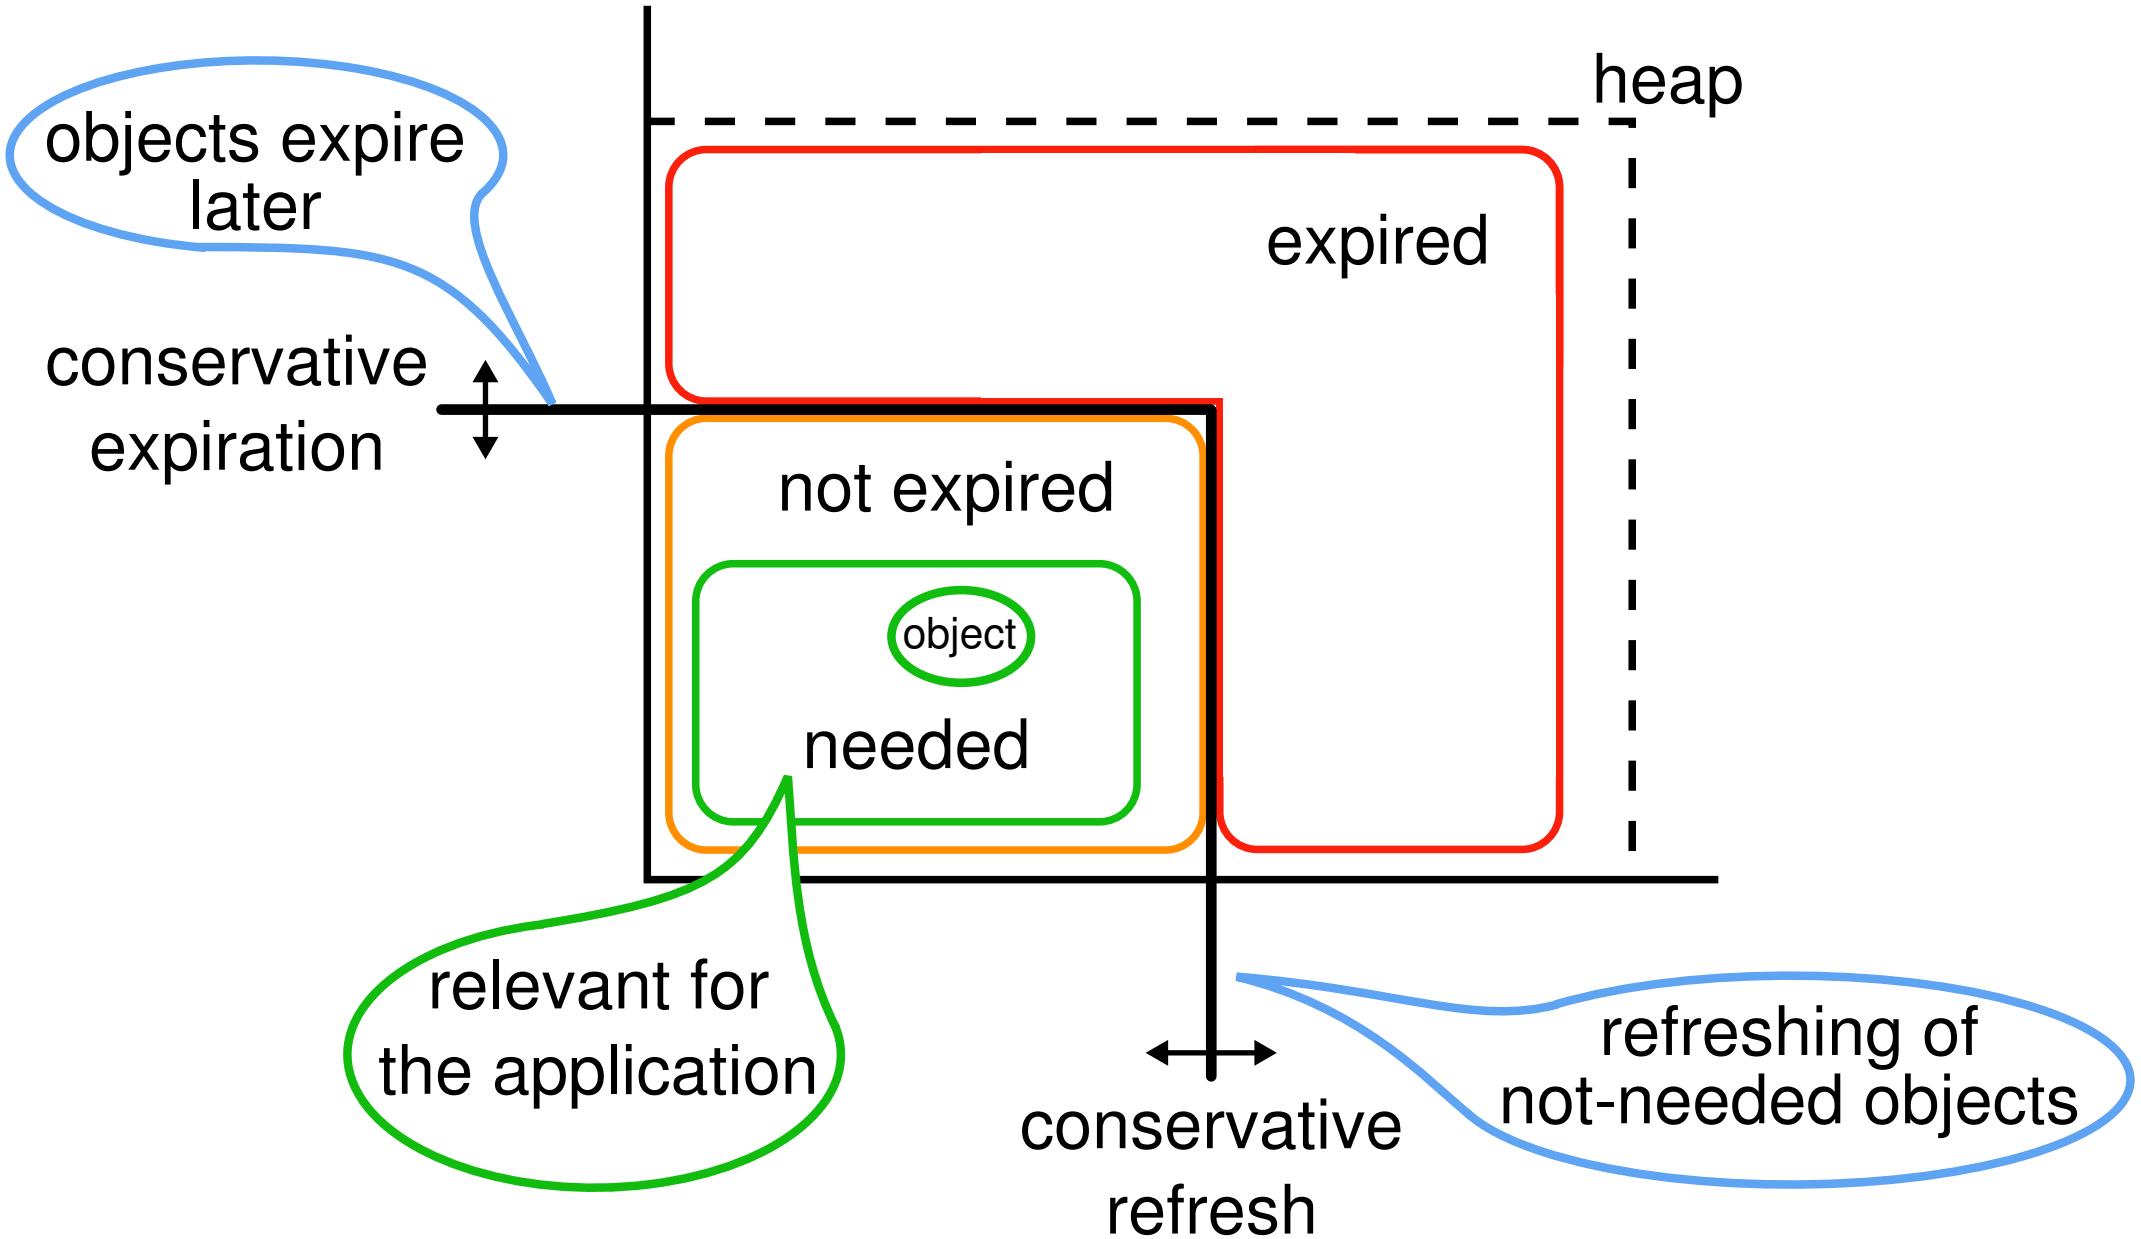
# Short-term memory



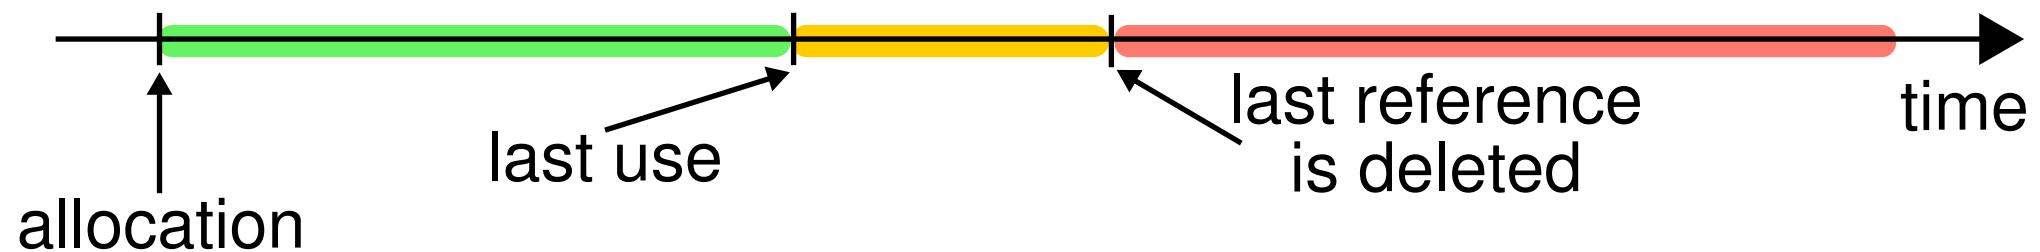
# Short-term memory



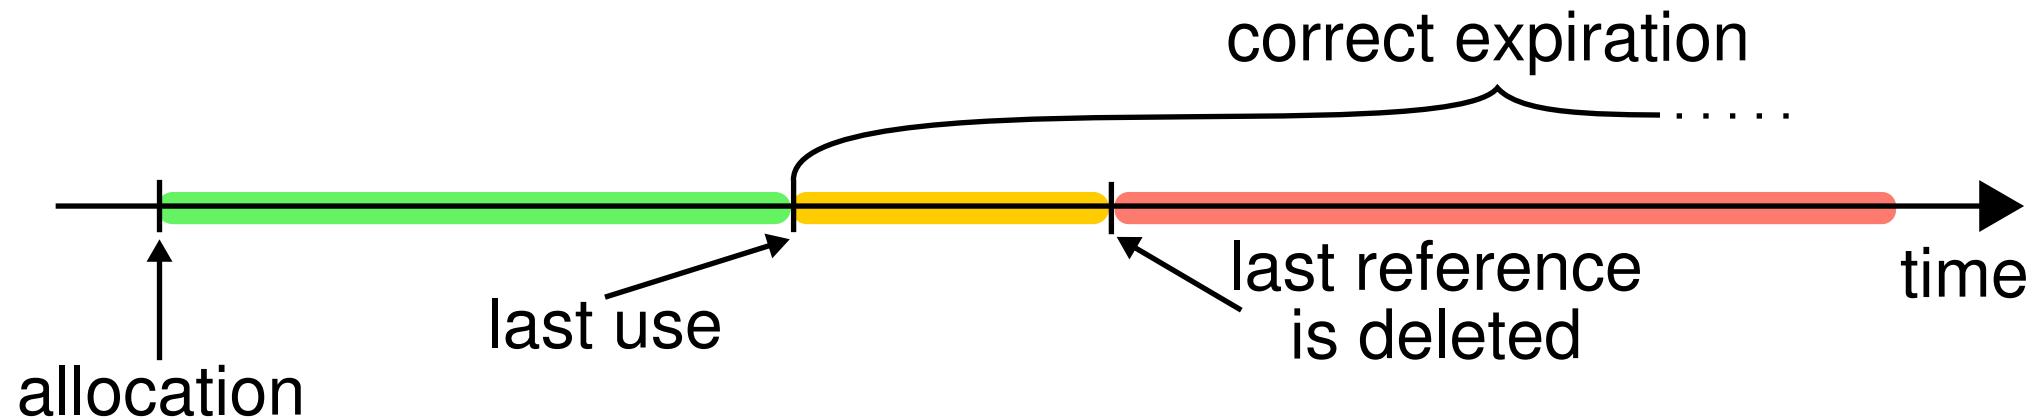
# Short-term memory



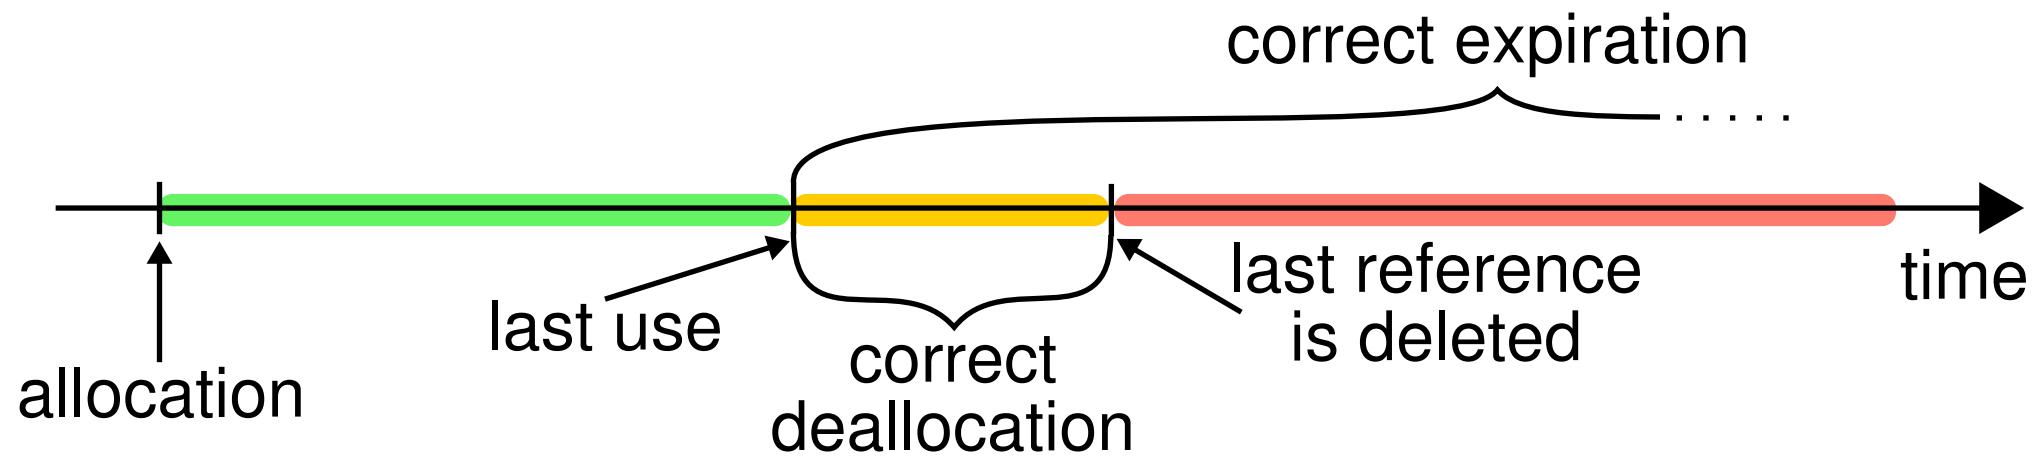
# Comparison



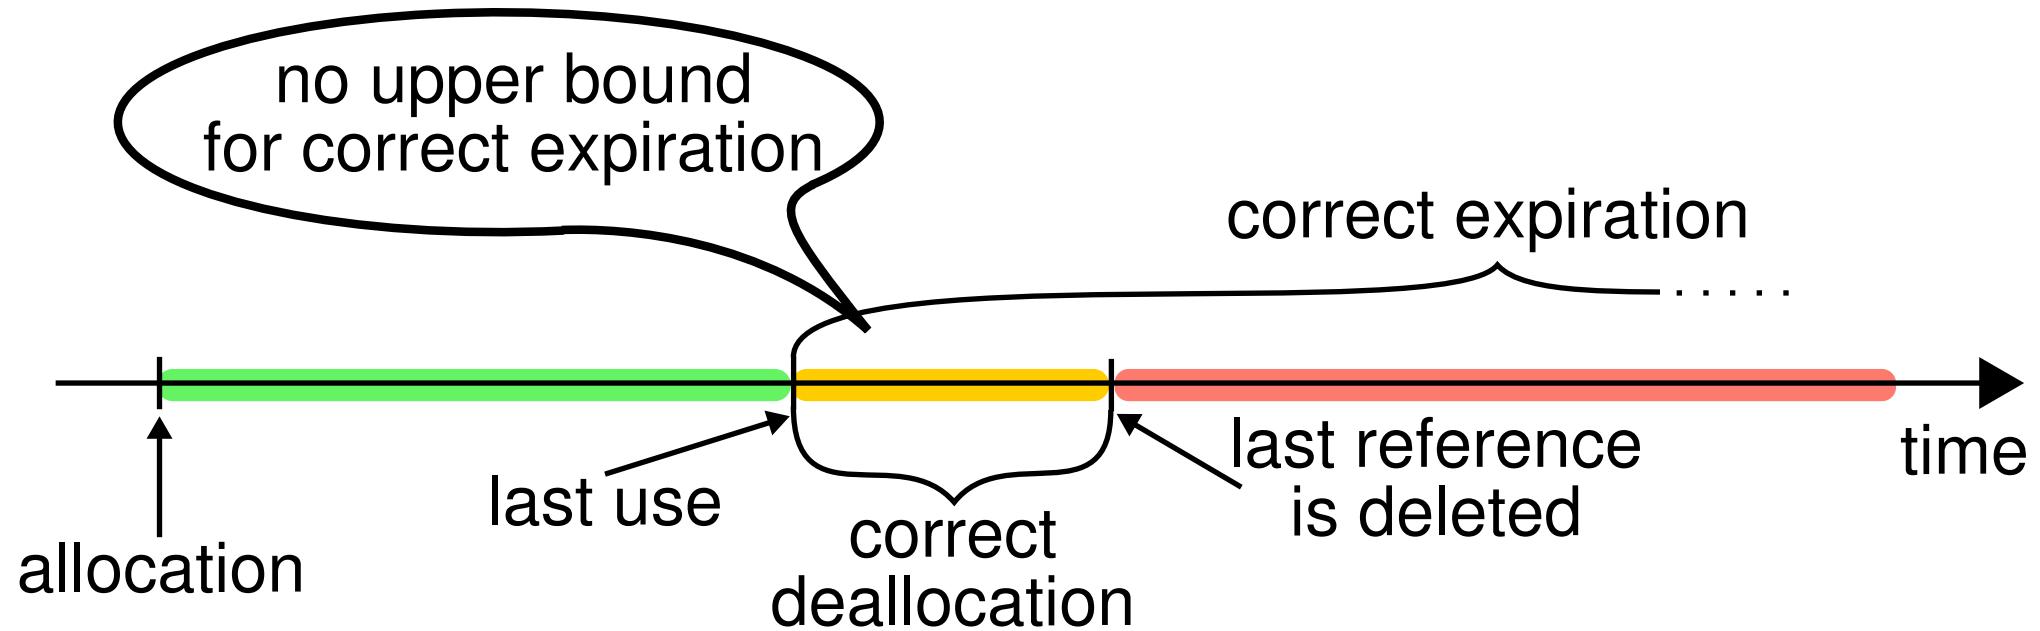
# Comparison



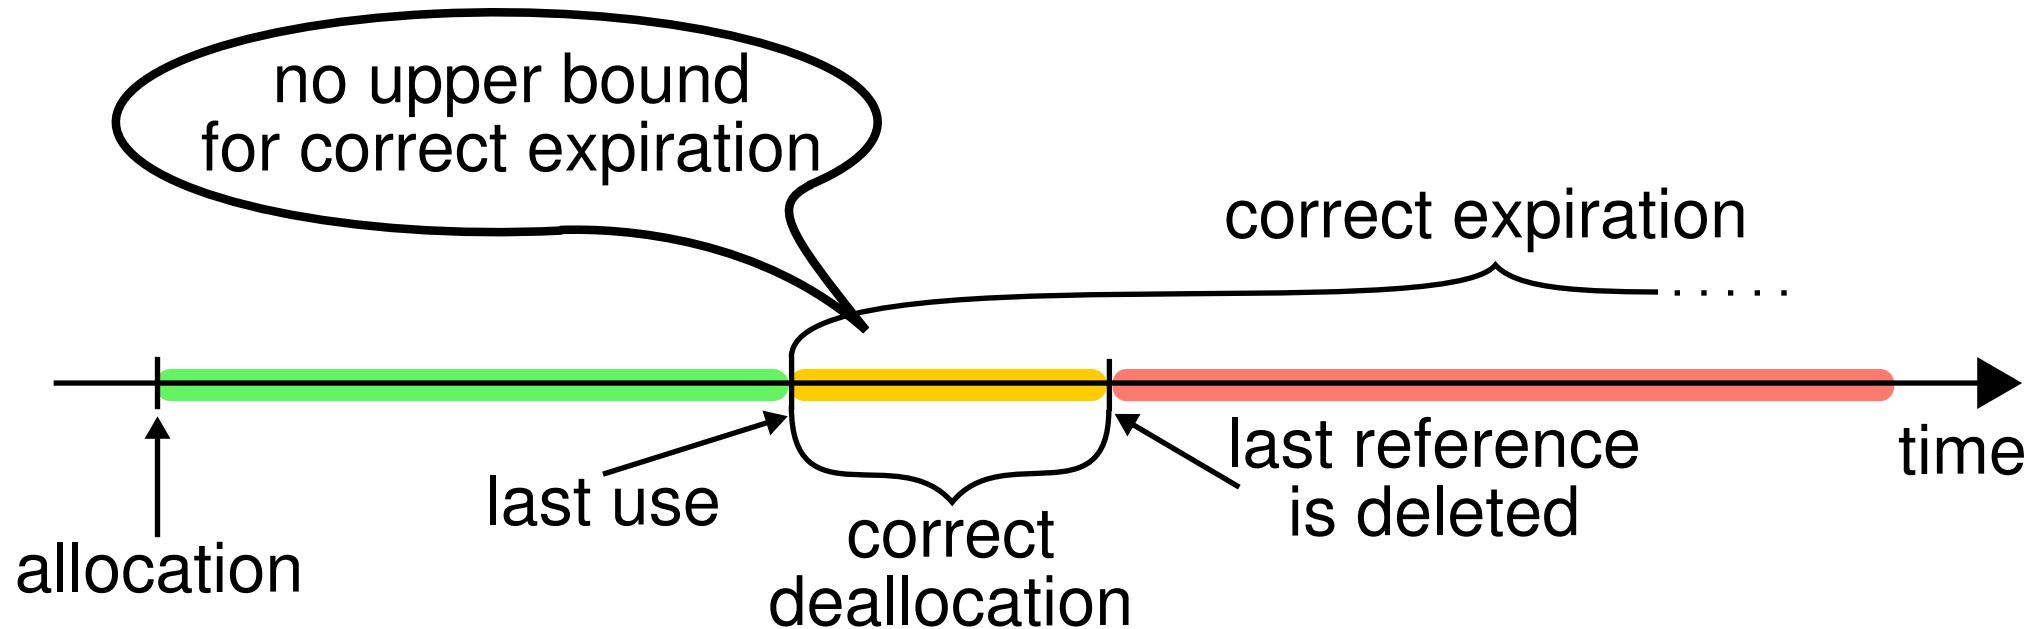
# Comparison



# Comparison



# Comparison



- ▶ Conservative lifetime approximations possible in short-term memory
  - ◆ May be easier to determine

# Self-collecting Mutators

explicit heap  
management system  
based on short-term  
memory

# Self-collecting Mutators

explicit heap  
management system  
based on short-term  
memory

explicit refresh

explicit time advance

# Self-collecting Mutators

explicit heap  
management system  
based on short-term  
memory

explicit refresh

explicit time advance

can be unsafe

# Self-collecting Mutators

explicit heap management system based on short-term memory

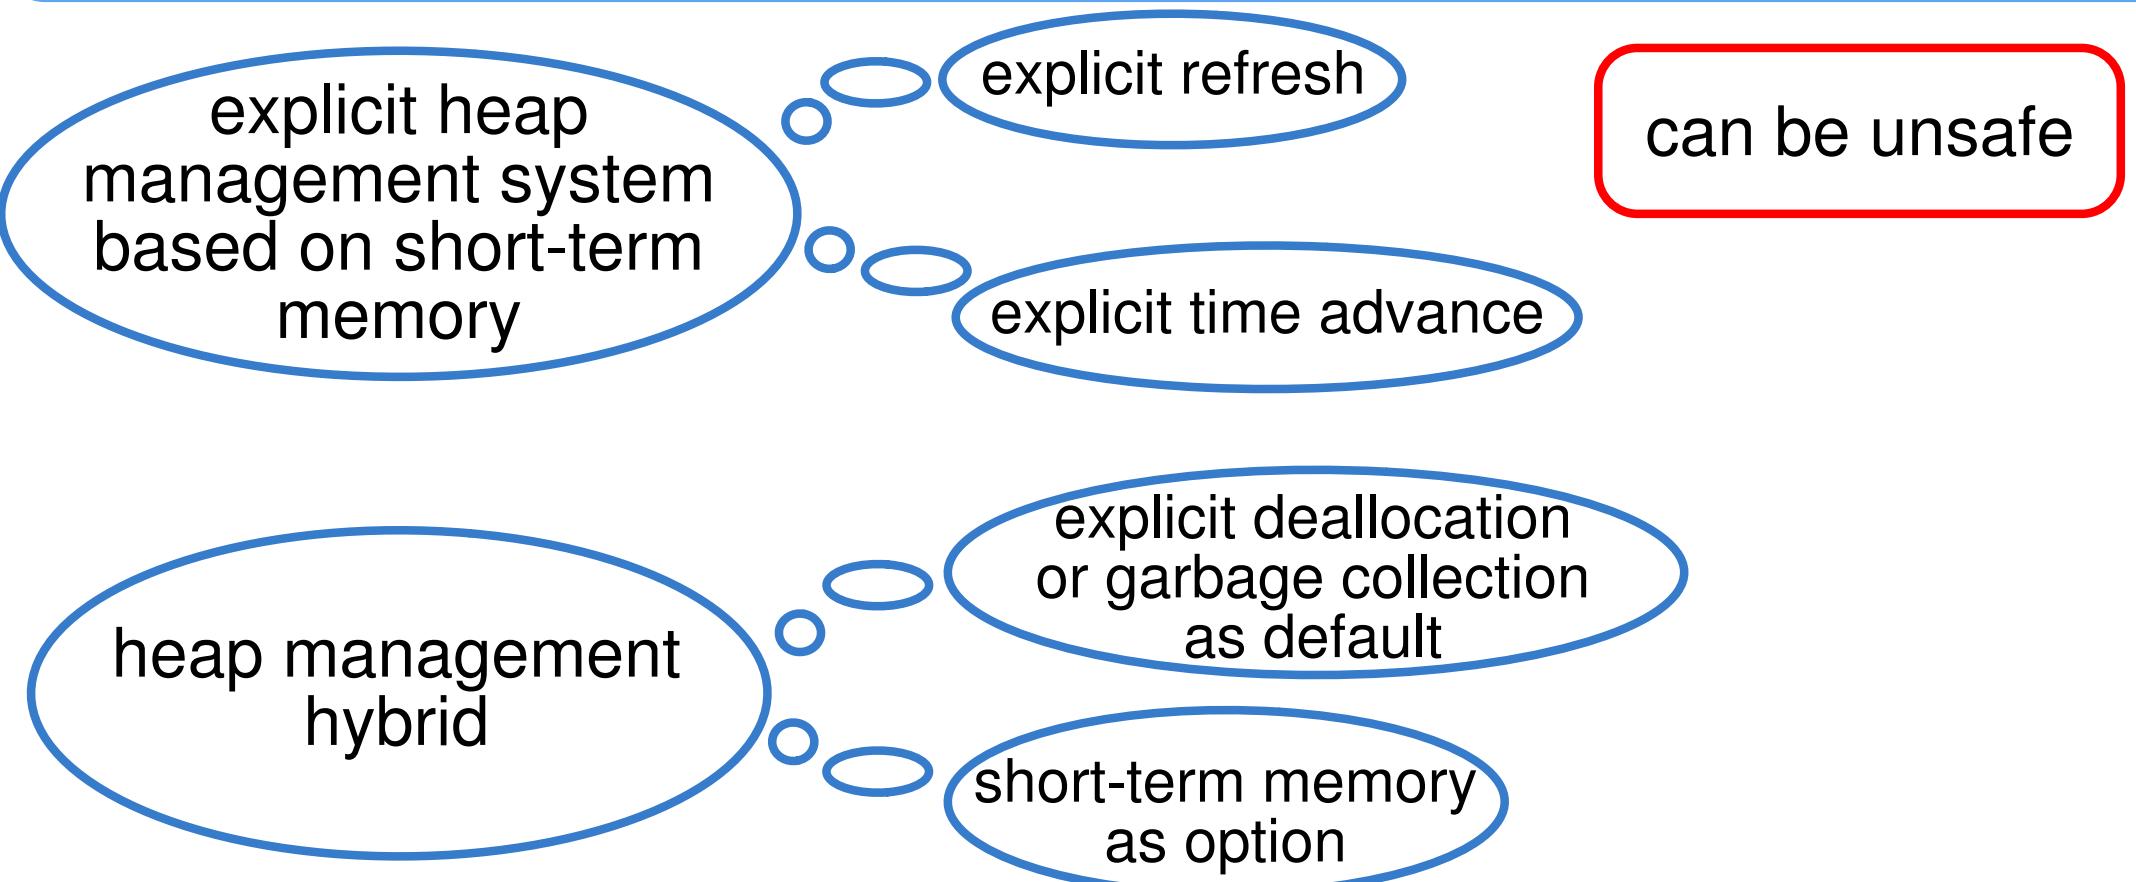
explicit refresh

explicit time advance

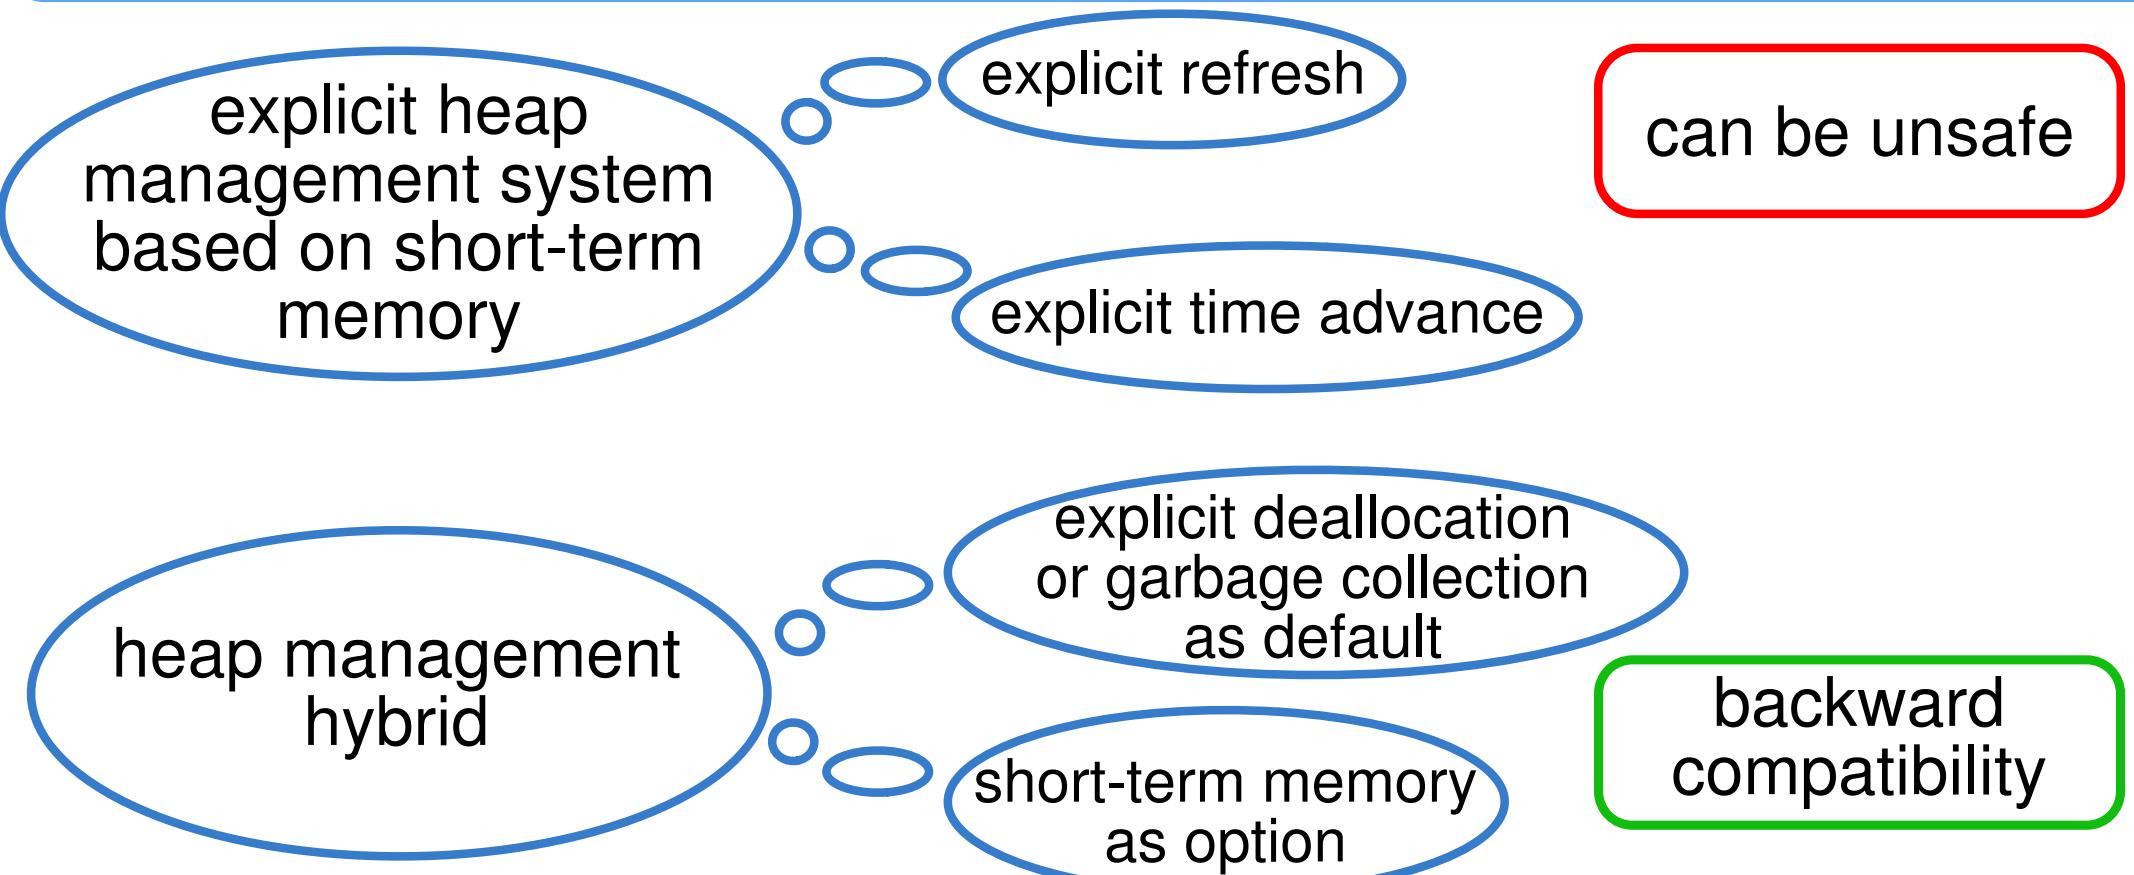
can be unsafe

heap management hybrid

# Self-collecting Mutators



# Self-collecting Mutators



# Self-collecting Mutators

explicit heap management system based on short-term memory

explicit refresh

can be unsafe

explicit time advance

heap management hybrid

explicit deallocation or garbage collection as default

backward compatibility

short-term memory as option

no additional threads for memory management

hence the name

# Self-collecting Mutators

explicit heap management system based on short-term memory

explicit refresh

can be unsafe

explicit time advance

heap management hybrid

explicit deallocation or garbage collection as default

backward compatibility

short-term memory as option

no additional threads for memory management

hence the name

fully scalable

# Programming Model

```
object = malloc(size);  
:  
refresh(object, 3);  
:  
tick();
```

time = 5

# Programming Model

```
object = malloc(size);
```

```
refresh(object, 3);
```

```
tick();
```

```
time = 5
```

allocates memory  
with one additional  
header word

an object expires  
when all its expiration  
dates have expired

0

object

counts  
expiration dates

# Programming Model

```
object = malloc(size);  
⋮  
refresh(object, 3);  
⋮  
tick();
```

hybrid heap  
management: the object  
can be deallocated here  
explicitly

time = 5

0 object

counts  
expiration dates

# Programming Model

```
object = malloc(size);
```

```
⋮
```

```
refresh(object, 3);
```

```
⋮
```

```
tick();
```

creates an expiration date for the object

1

object

8

time = 5

counts  
expiration dates

# Programming Model

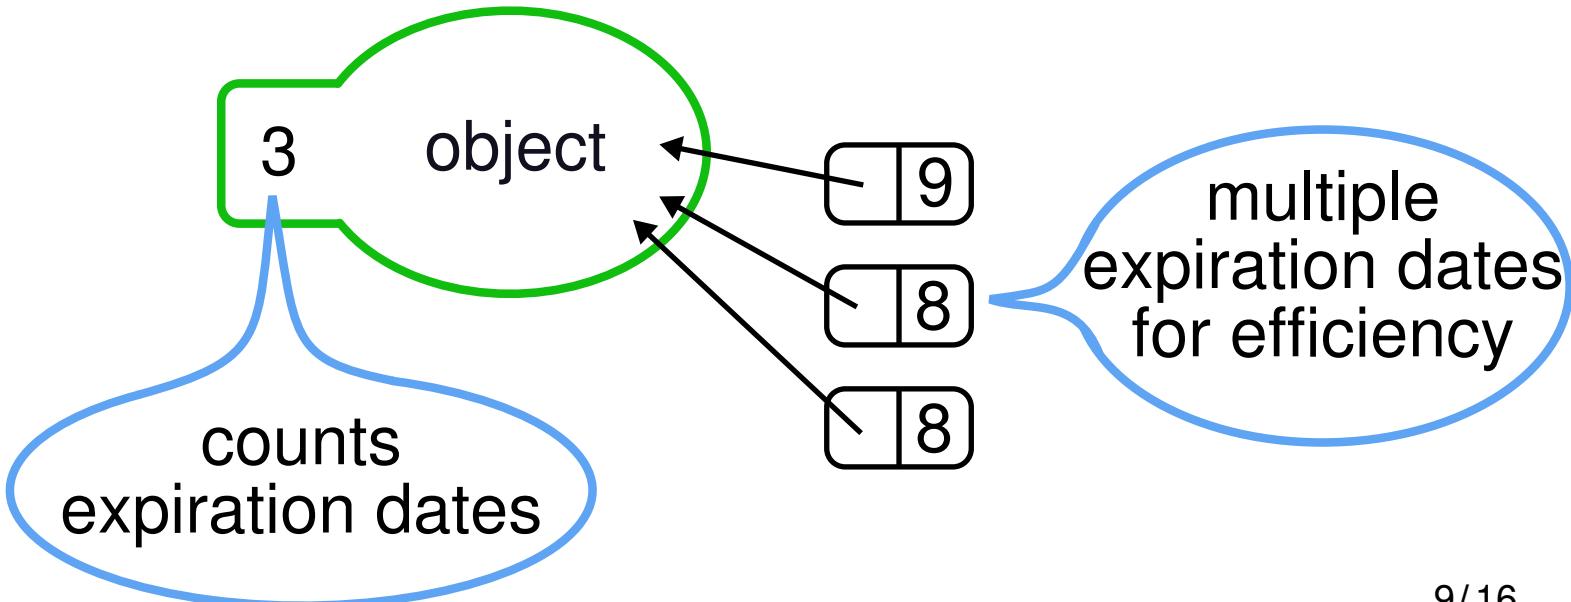
```
object = malloc(size);
```

```
refresh(object, 3);
```

```
tick();
```

```
time = 5
```

creates an expiration date for the object

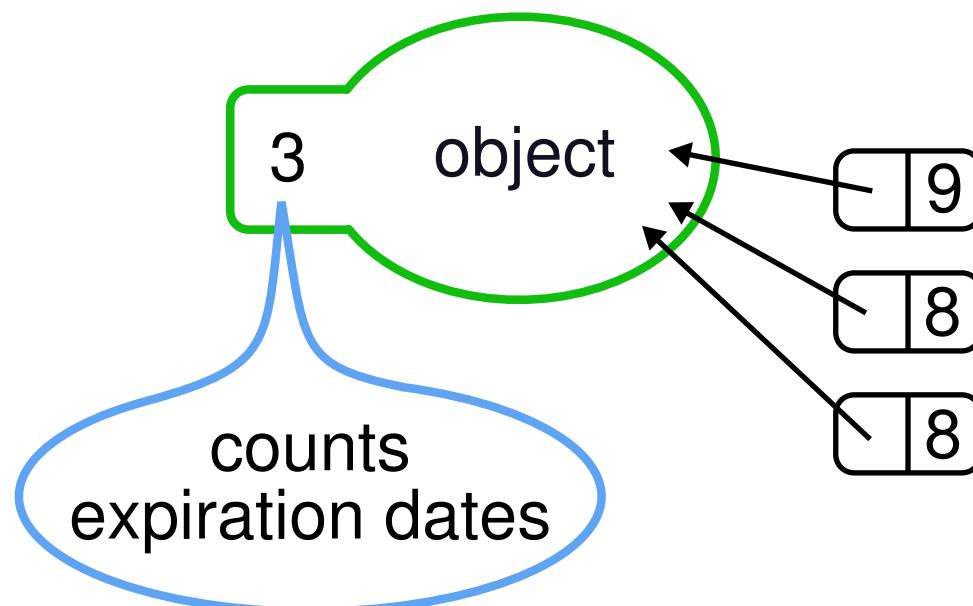


# Programming Model

```
object = malloc(size);  
:  
refresh(object, 3);  
:  
tick();
```

object is short-term:  
deallocation-calls  
are ignored

time = 5



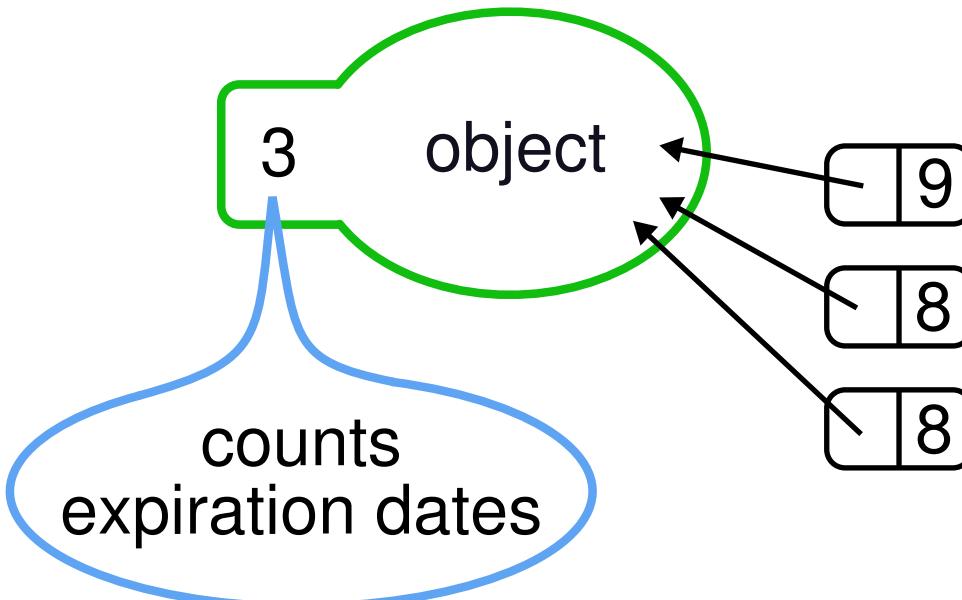
# Programming Model

```
object = malloc(size);  
:  
refresh(object, 3);  
:  
tick();
```

port an application  
without removing  
free-calls

object is short-term:  
deallocation-calls  
are ignored

time = 5



# Programming Model

```
object = malloc(size);
```

```
⋮
```

```
refresh(object, 3);
```

```
⋮
```

```
tick();
```

increments time

**time = 6**

3 object

counts  
expiration dates

9

8

8

# Programming Model

```
object = malloc(size);  
:  
refresh(object, 3);  
:  
tick();
```

expiration dates  
are managed in buffers

**time = 6**

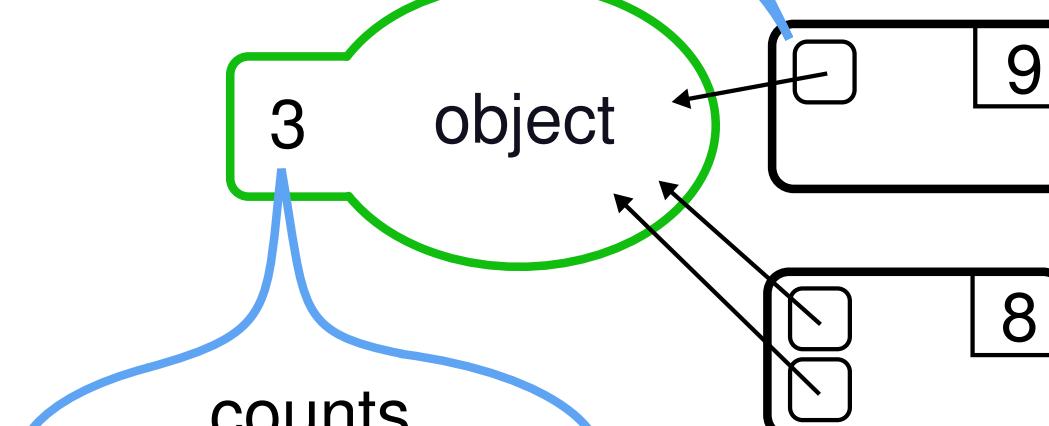
counts  
expiration dates

3

object

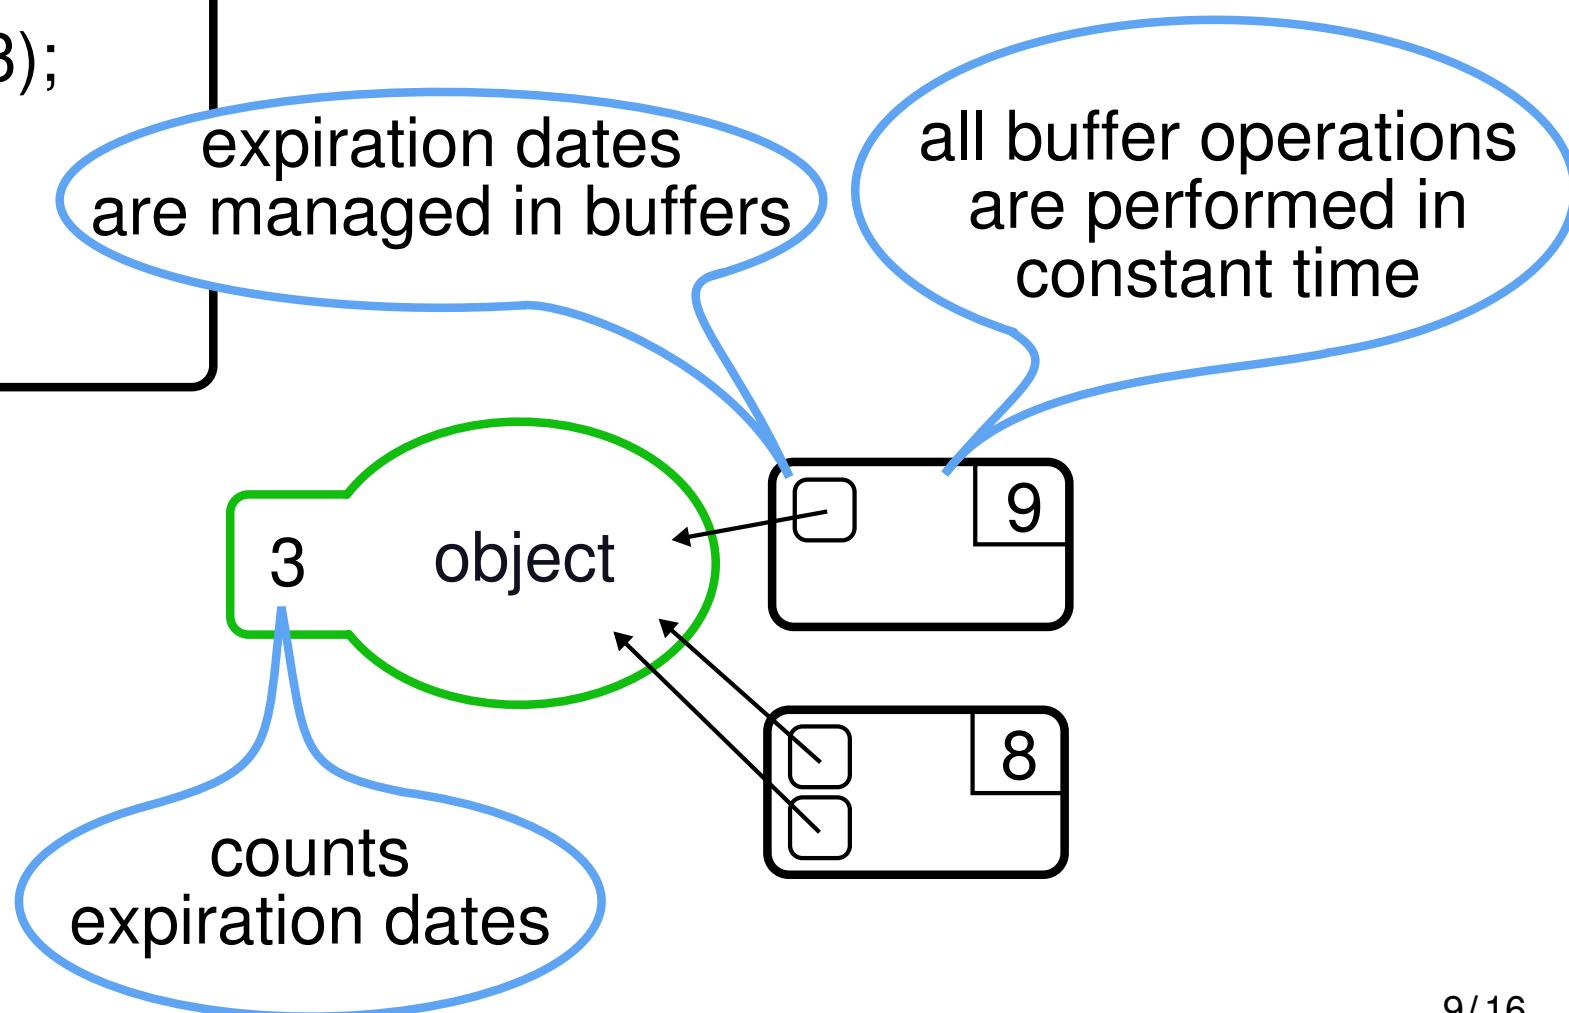
9

8



# Programming Model

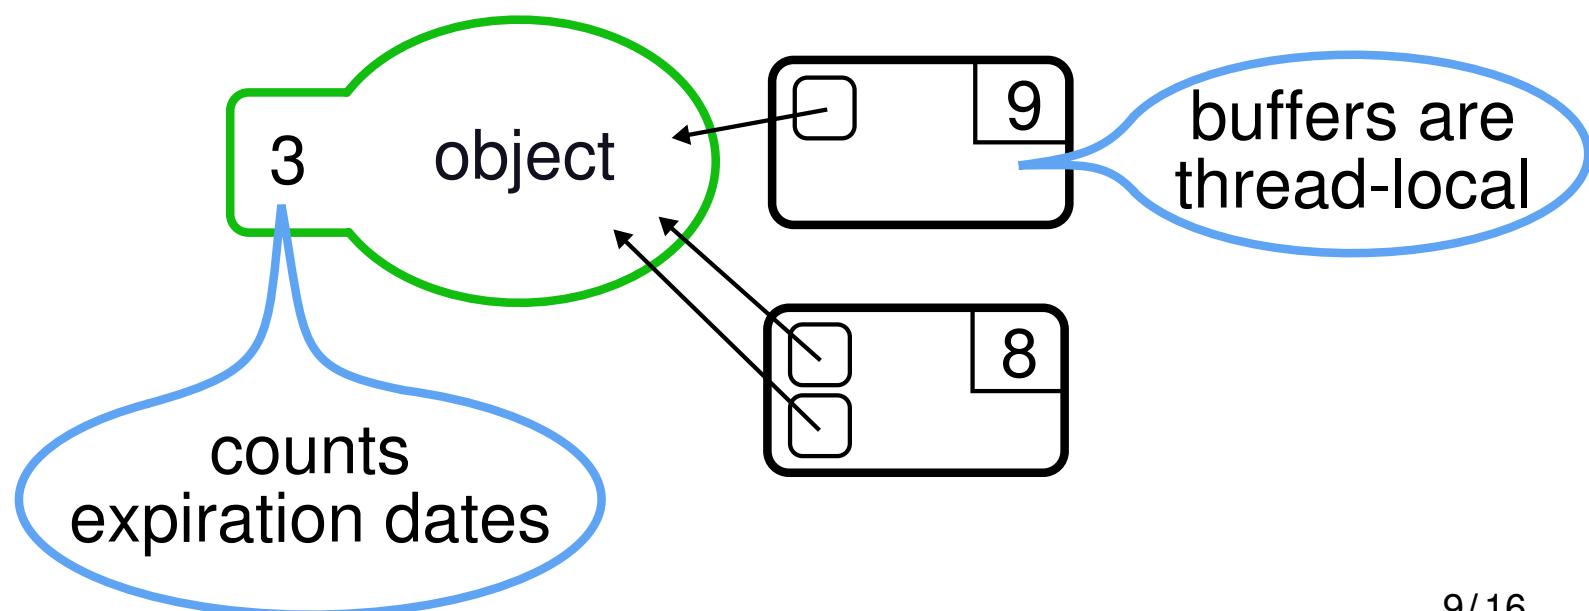
```
object = malloc(size);  
:  
refresh(object, 3);  
:  
tick();
```



# Programming Model

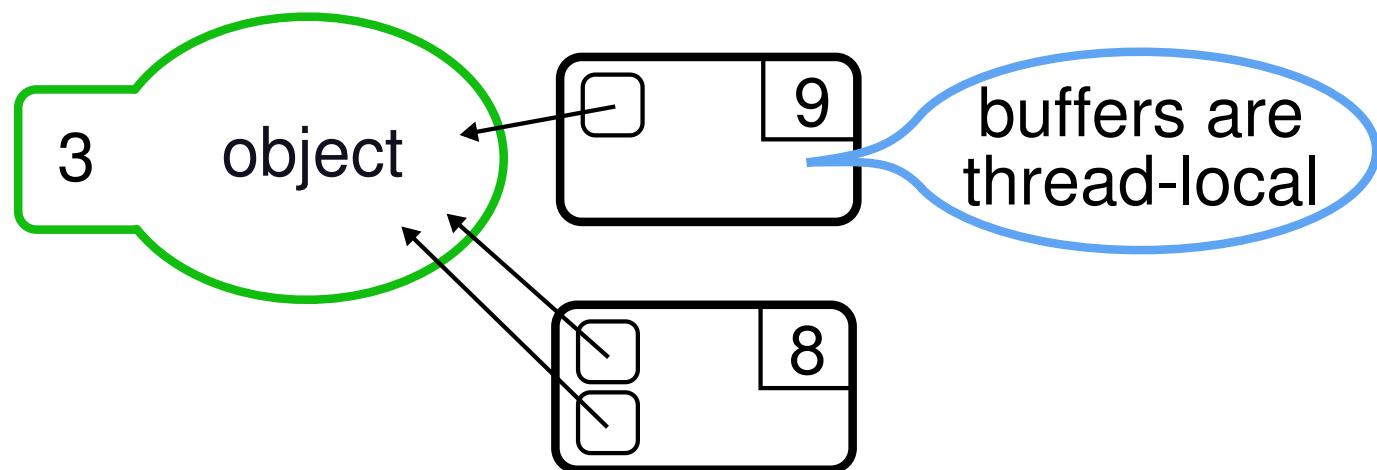
```
object = malloc(size);  
⋮  
refresh(object, 3);  
⋮  
tick();
```

**time = 6**



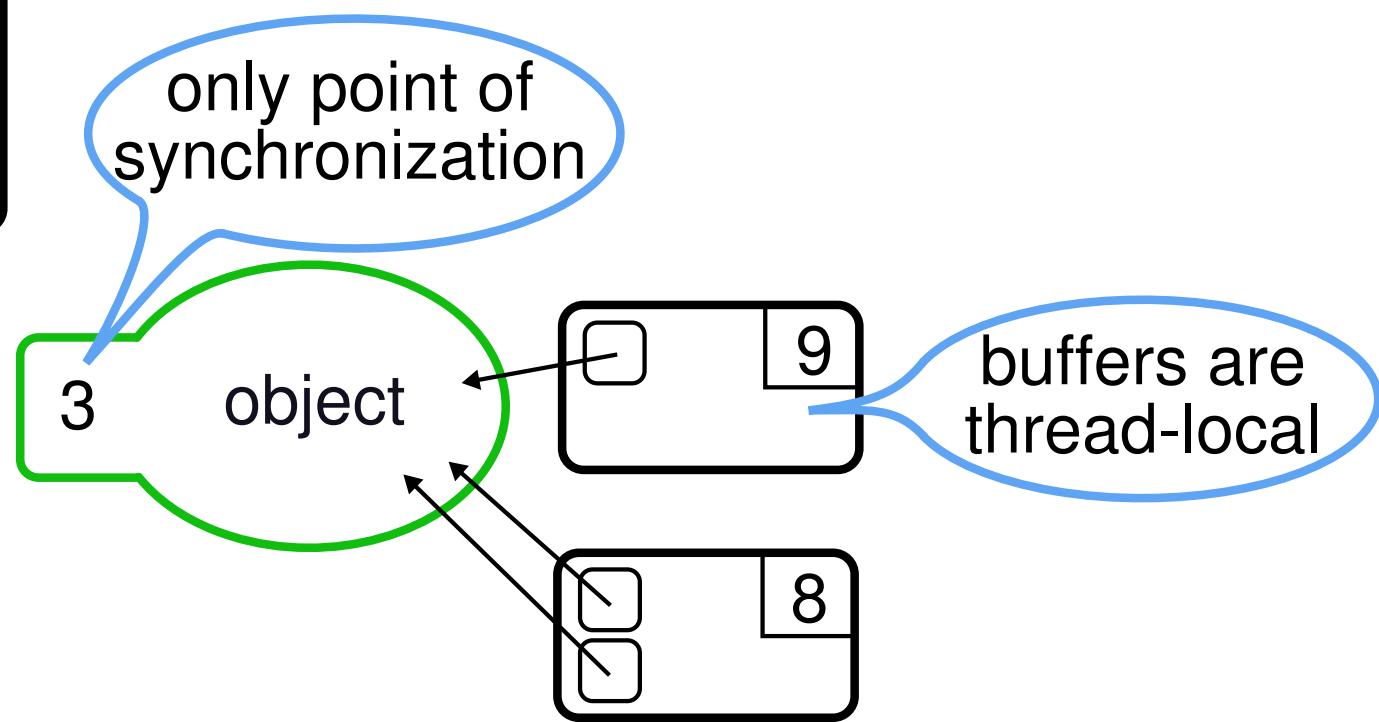
# Programming Model

```
object = malloc(size);  
:  
refresh(object, 3);  
:  
tick();
```

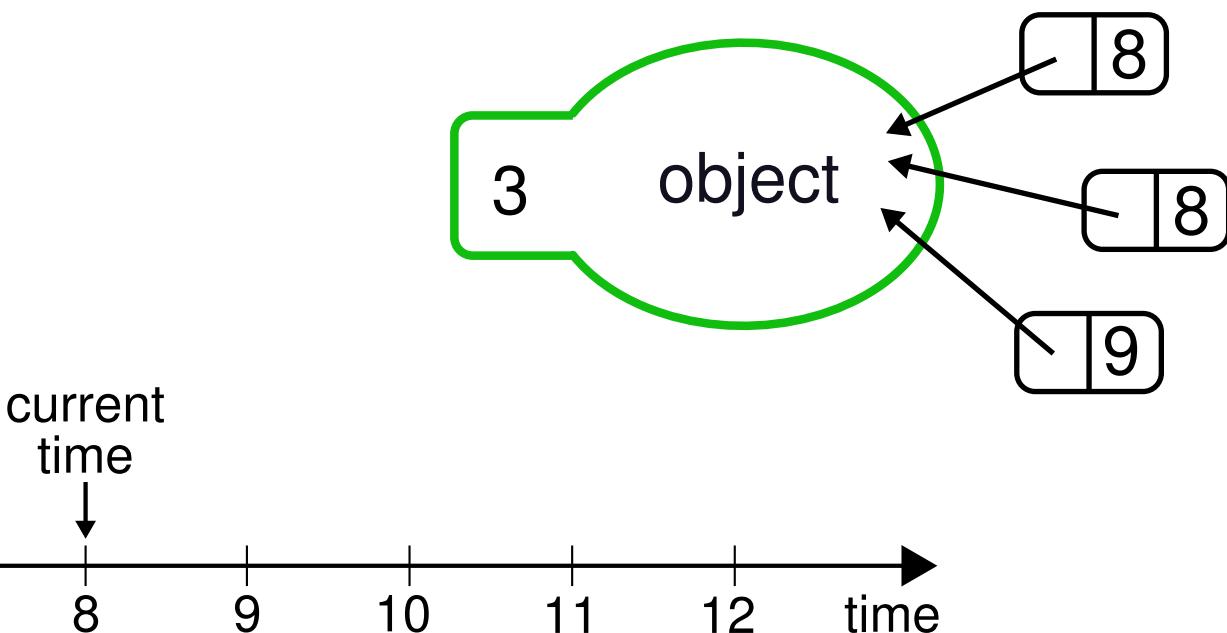


# Programming Model

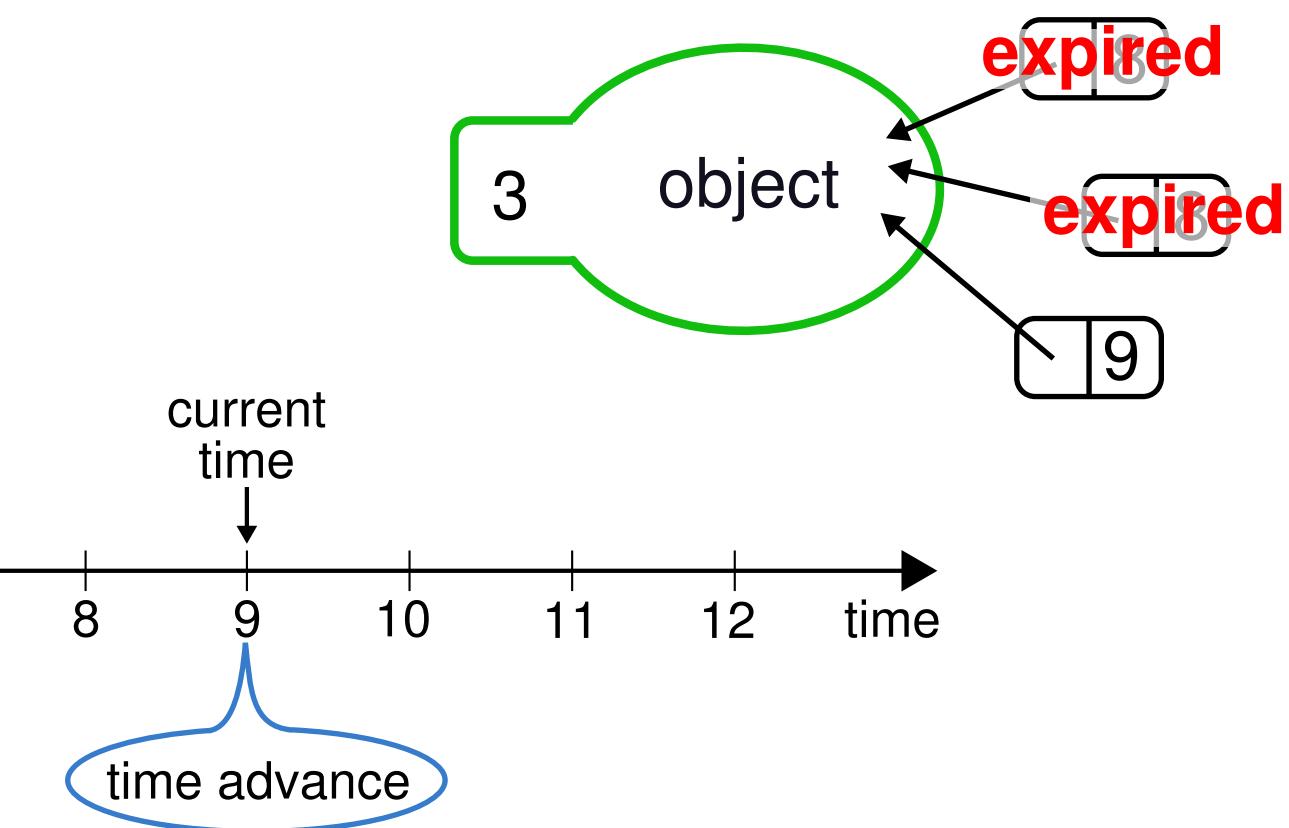
```
object = malloc(size);  
:  
refresh(object, 3);  
:  
tick();
```



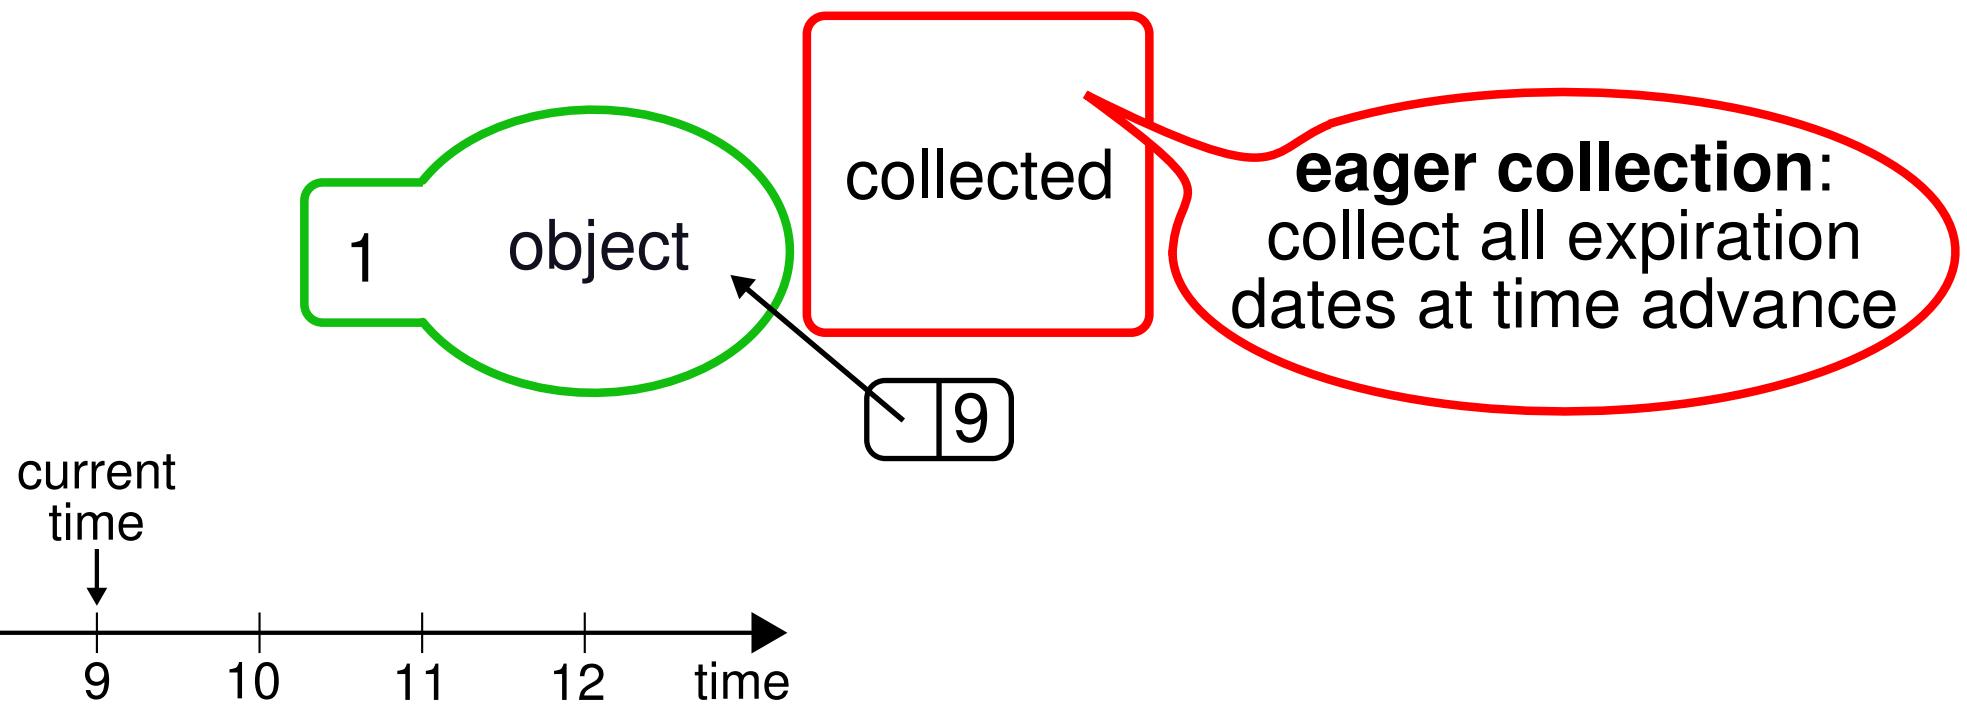
# Collection of Expired Objects



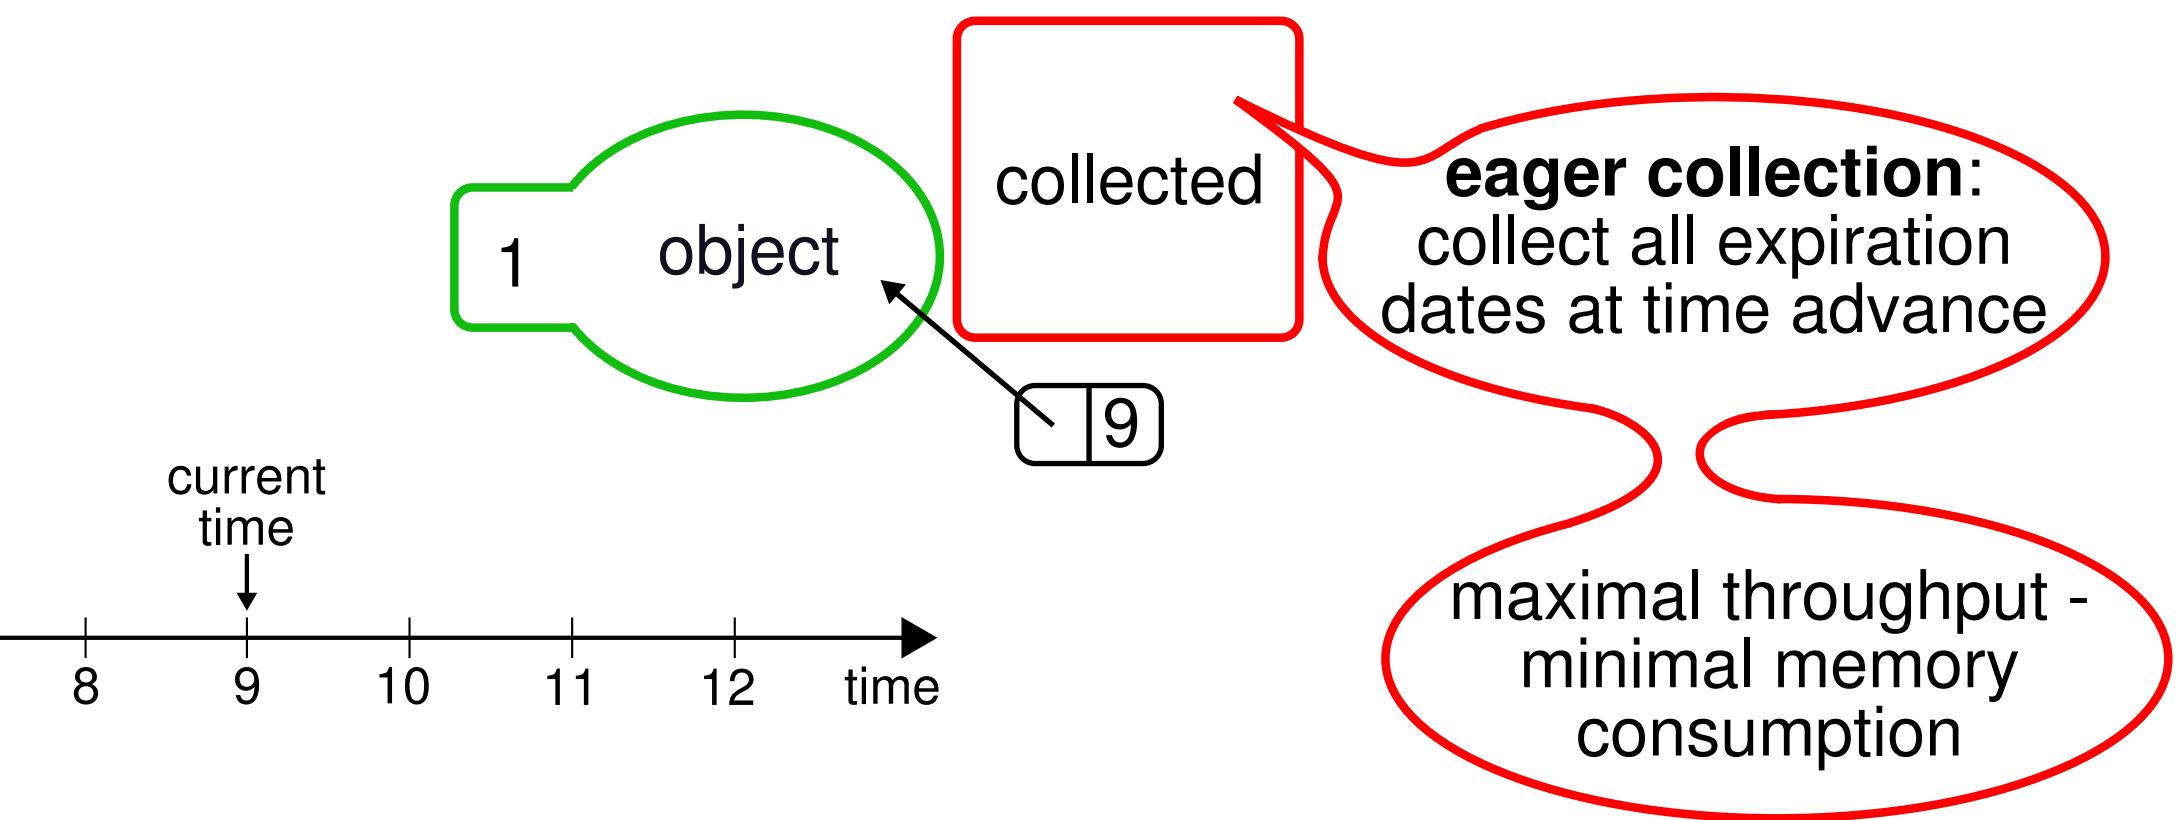
# Collection of Expired Objects



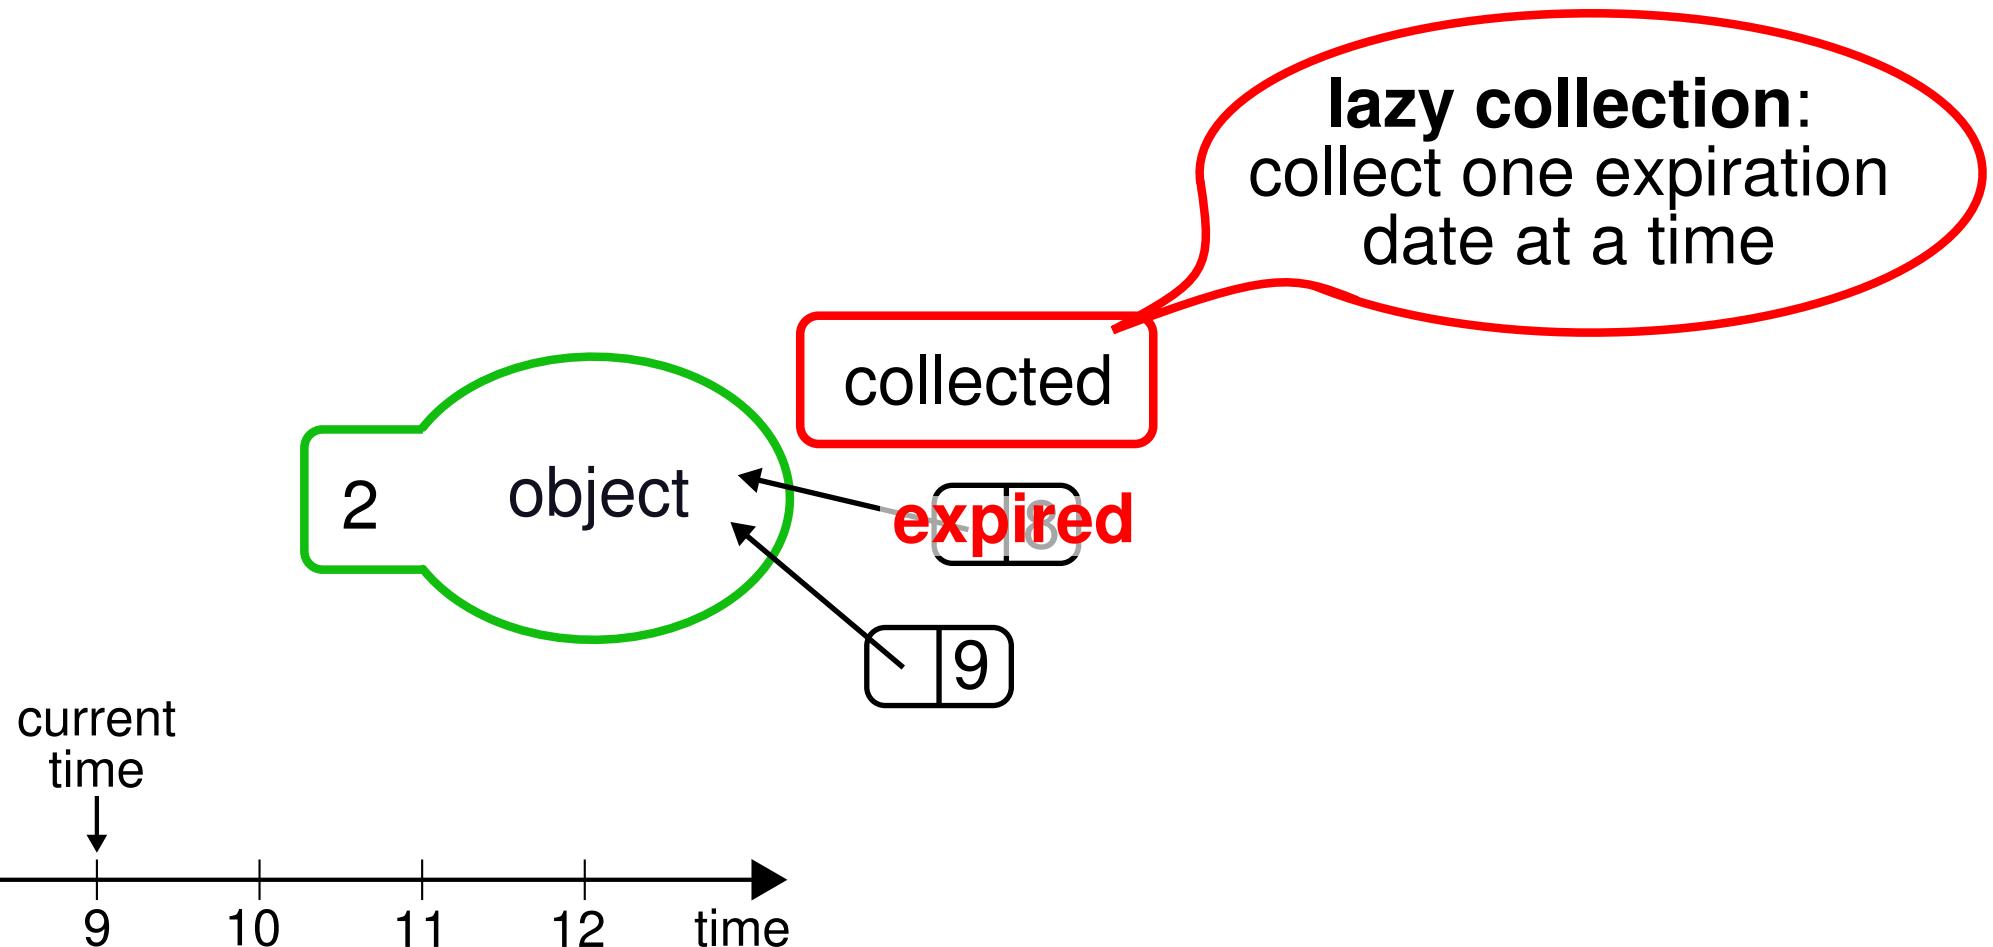
# Collection of Expired Objects



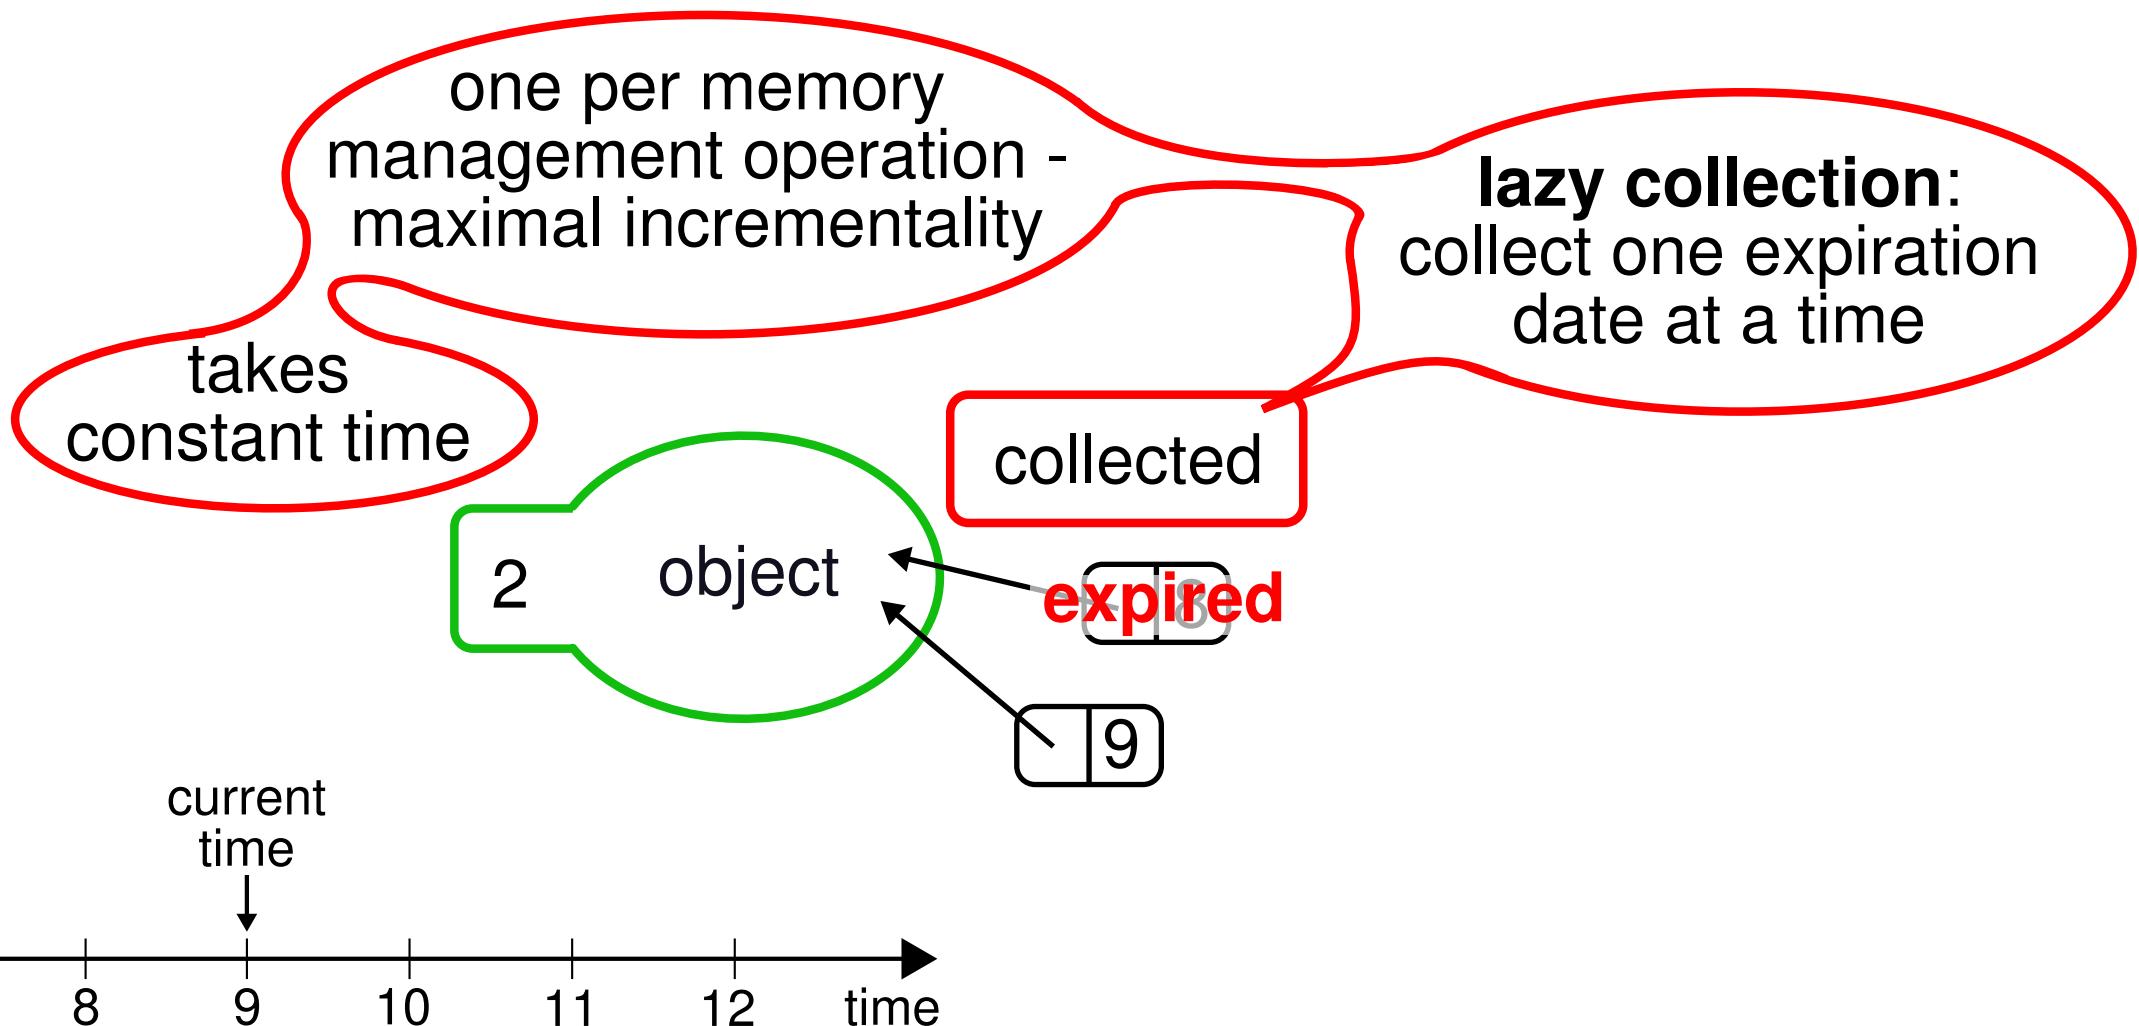
# Collection of Expired Objects



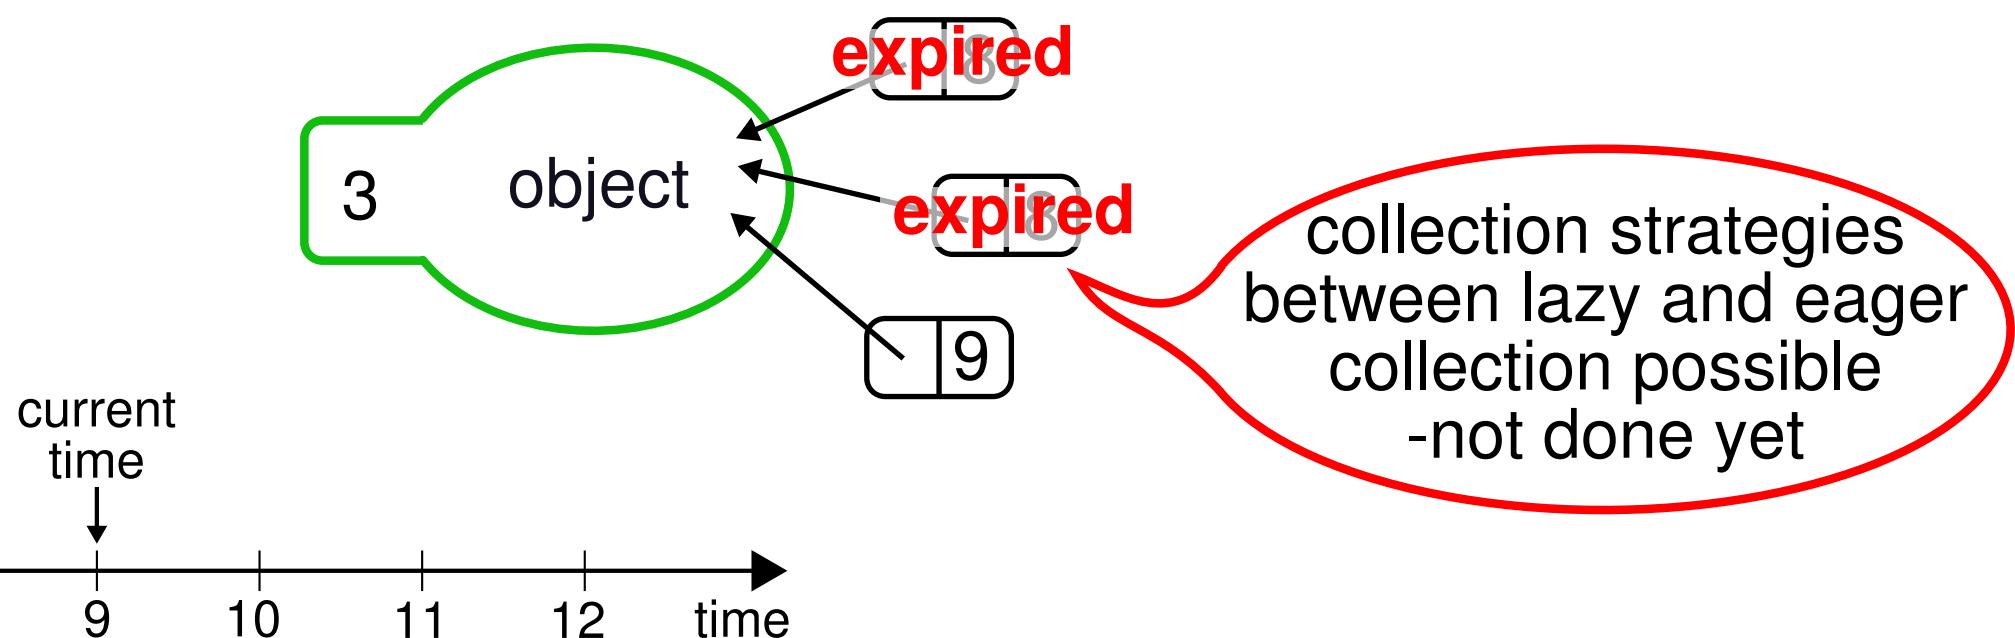
# Collection of Expired Objects



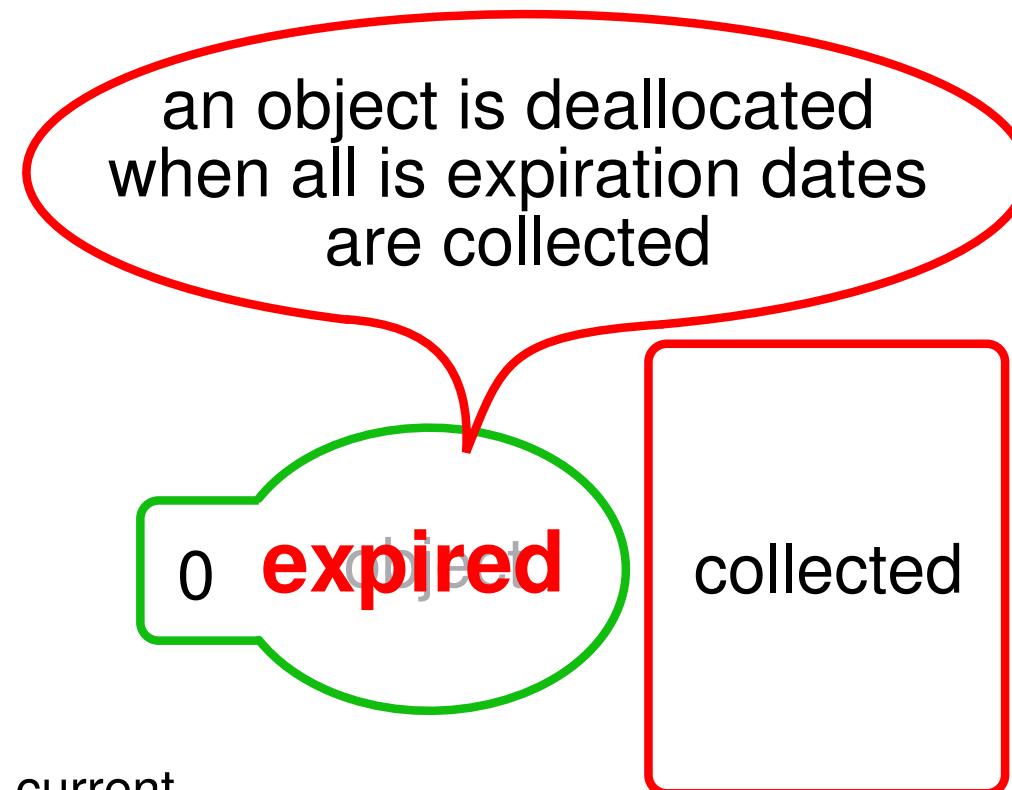
# Collection of Expired Objects



# Collection of Expired Objects



# Collection of Expired Objects



# Implementation and Evaluation

**Java**  
in Jikes RVM

**C**  
based on ptmalloc2

**Go**  
in the 6g runtime

# Implementation and Evaluation

**Java**  
in Jikes RVM  
Monte Carlo

**C**  
based on ptmalloc2  
mpg123  
x264

**Go**  
in the 6g runtime  
tree benchmark  
webserver

# Implementation and Evaluation

**Java**  
in Jikes RVM

Monte Carlo

**C**  
based on ptmalloc2

mpg123  
x264

**Go**  
in the 6g runtime

tree benchmark  
webserver

performance

ease of use

- ▶ total execution time  
(throughput)
- ▶ execution time per operation  
(latency)
- ▶ memory consumption

- ▶ lines of code
- ▶ difficulty of placing code

# Implementation and Evaluation

**Java**  
in Jikes RVM

Monte Carlo

**C**  
based on ptmalloc2

mpg123  
x264

**Go**  
in the 6g runtime  
tree benchmark  
webserver

- 
- ▶ 1450 LoC
  - ▶ 1 tick-call
  - ▶ refresh all short-living objects right after allocation

# Implementation and Evaluation

**Java**  
in Jikes RVM

Monte Carlo

**C**  
based on ptmalloc2

mpg123  
x264

**Go**  
in the 6g runtime

tree benchmark  
webserver

- ▶ 1450 LoC
- ▶ 1 tick-call
- ▶ refresh all short-living objects right after allocation

to make all short-living objects short-term

# Implementation and Evaluation

**Java**  
in Jikes RVM  
Monte Carlo

**C**  
based on ptmalloc2  
mpg123  
x264

**Go**  
in the 6g runtime  
tree benchmark  
webserver

- 
- ▶ 16043 LoC
  - ▶ 1 tick-call
  - ▶ refresh all objects right after allocation

# Implementation and Evaluation

**Java**  
in Jikes RVM  
Monte Carlo

**C**  
based on ptmalloc2  
  
mpg123  
x264

**Go**  
in the 6g runtime  
  
tree benchmark  
webserver

- ▶ 61722 LoC
- ▶ 1 tick-call

**or**

- ▶ **single refresh**: refresh all frames after allocation with a conservative expiration date
- ▶ **continuous refresh**: refresh all needed frames before time advance with extension '1'

# Implementation and Evaluation

**Java**  
in Jikes RVM  
Monte Carlo

**C**  
based on ptmalloc2  
mpg123  
x264

**Go**  
in the 6g runtime  
tree benchmark  
webserver

- ▶ 61722 LoC
- ▶ 1 tick-call

2 tick-calls in the  
multi-threaded version:  
in the paper

**or**

- ▶ **single refresh:** refresh all frames after allocation  
with a conservative expiration date
- ▶ **continuous refresh:** refresh all needed frames  
before time advance with extension '1'

# Implementation and Evaluation

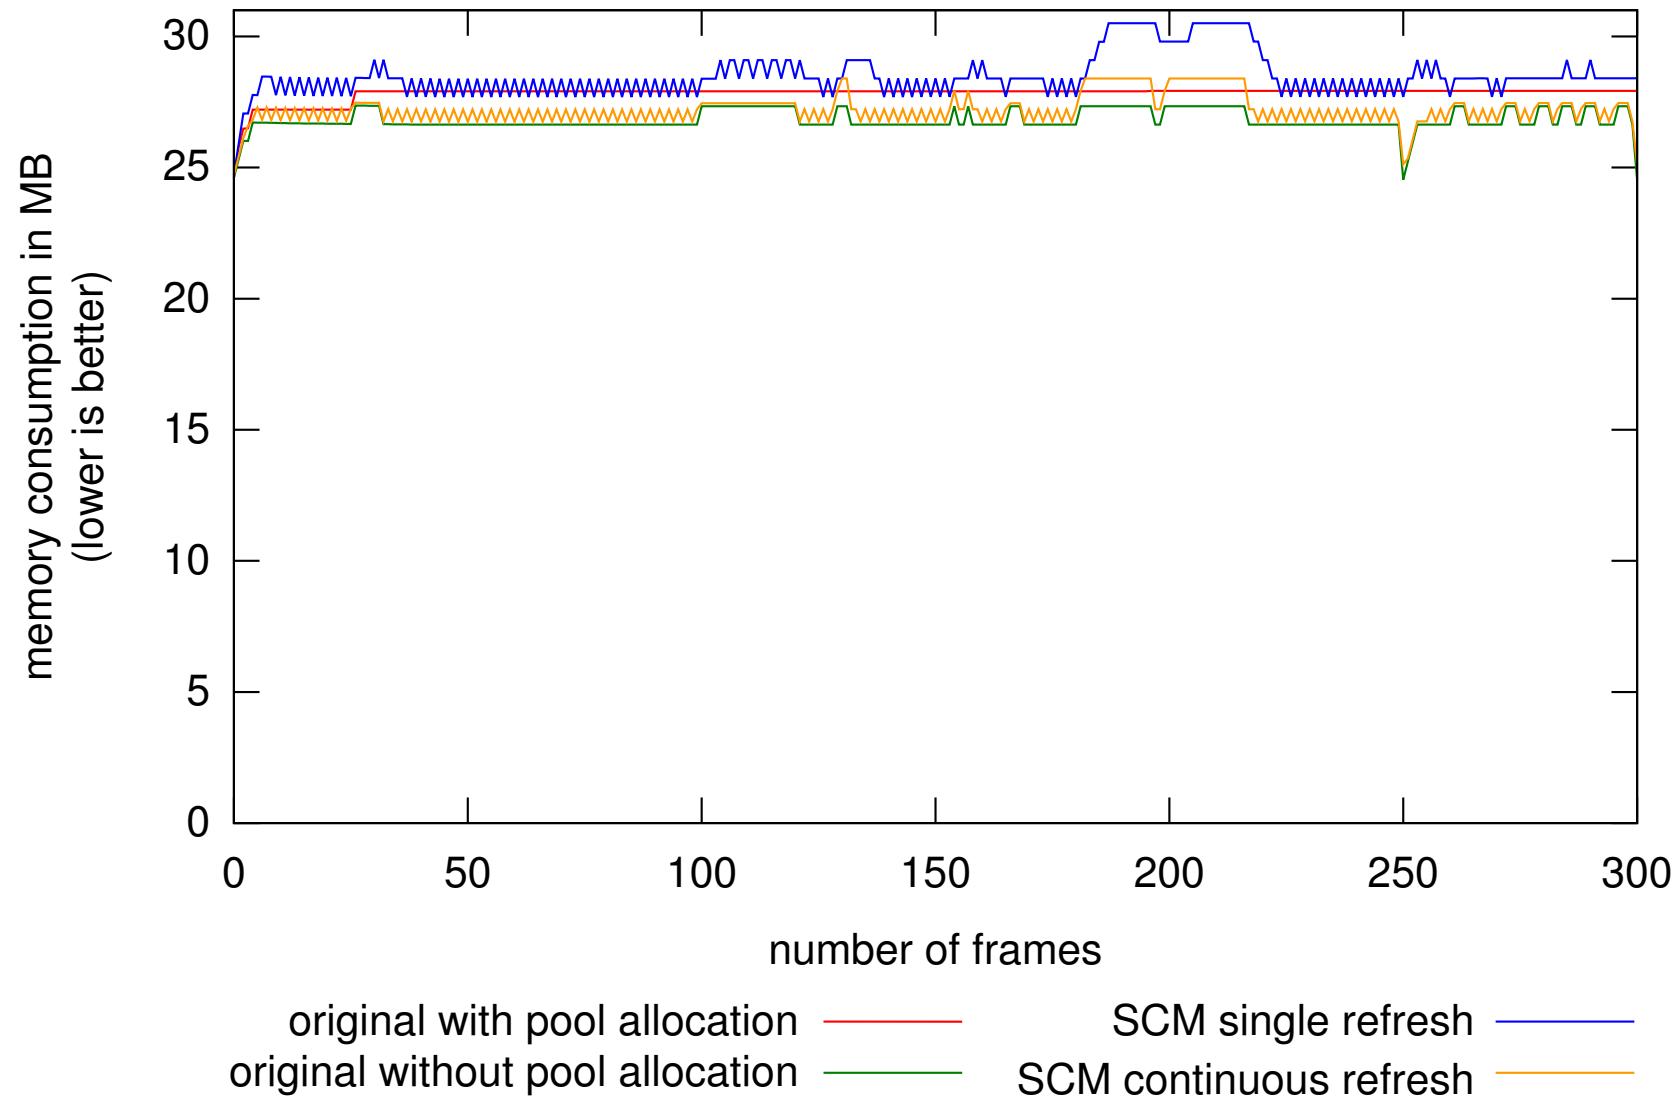
**Java**  
in Jikes RVM  
Monte Carlo

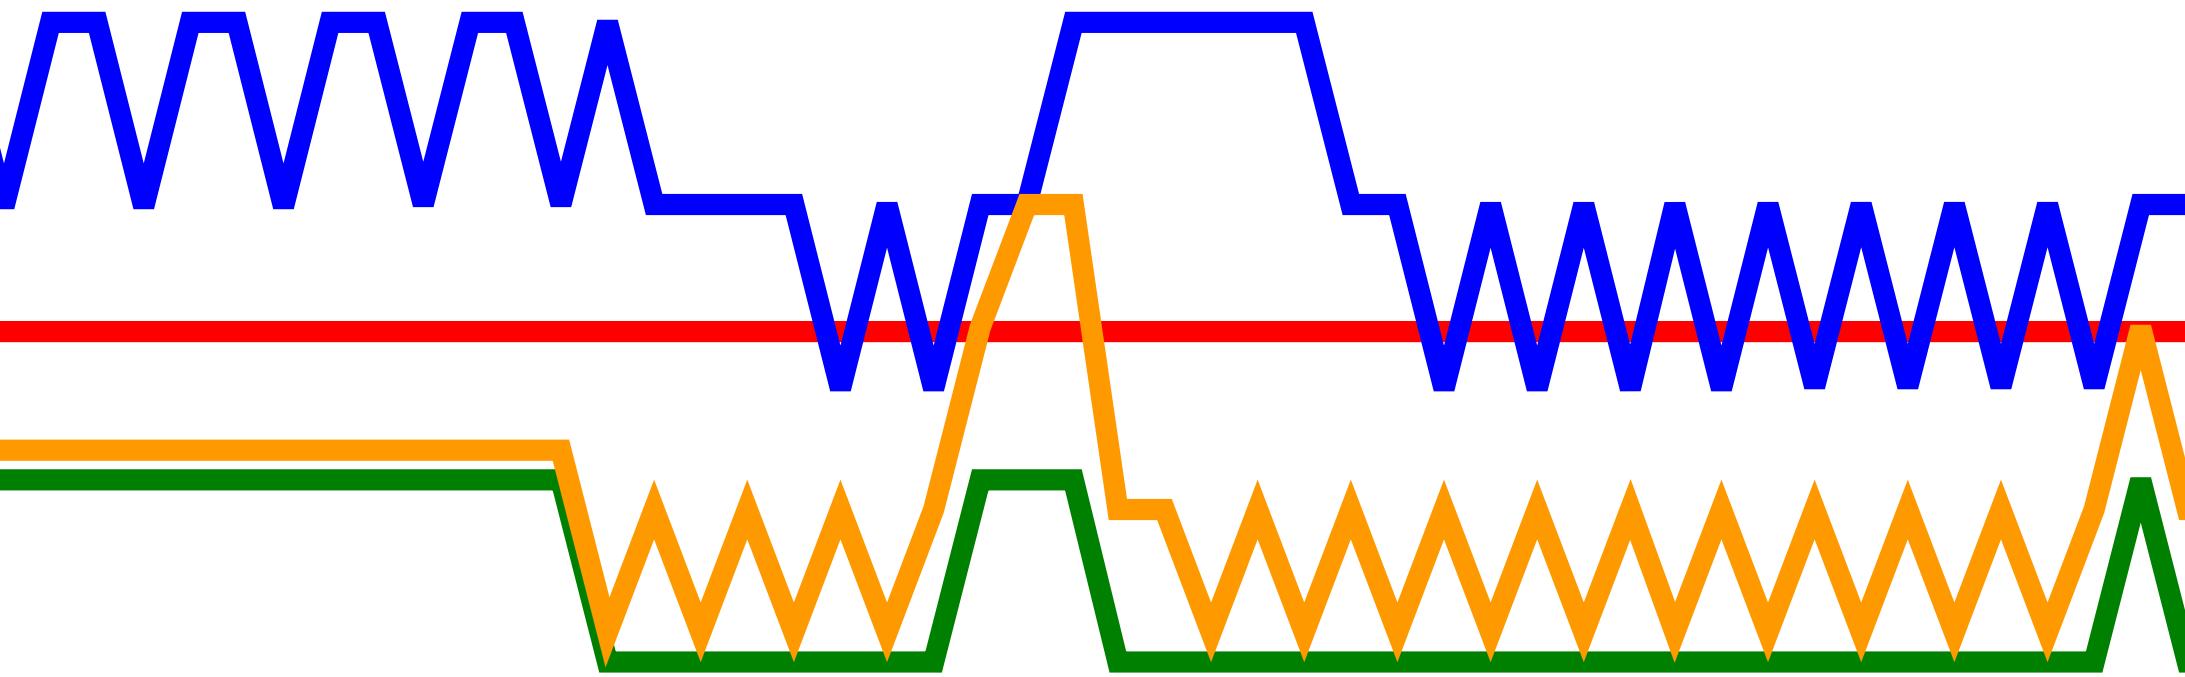
**C**  
based on ptmalloc2  
mpg123  
x264

**Go**  
in the 6g runtime  
tree benchmark  
webserver

in the paper

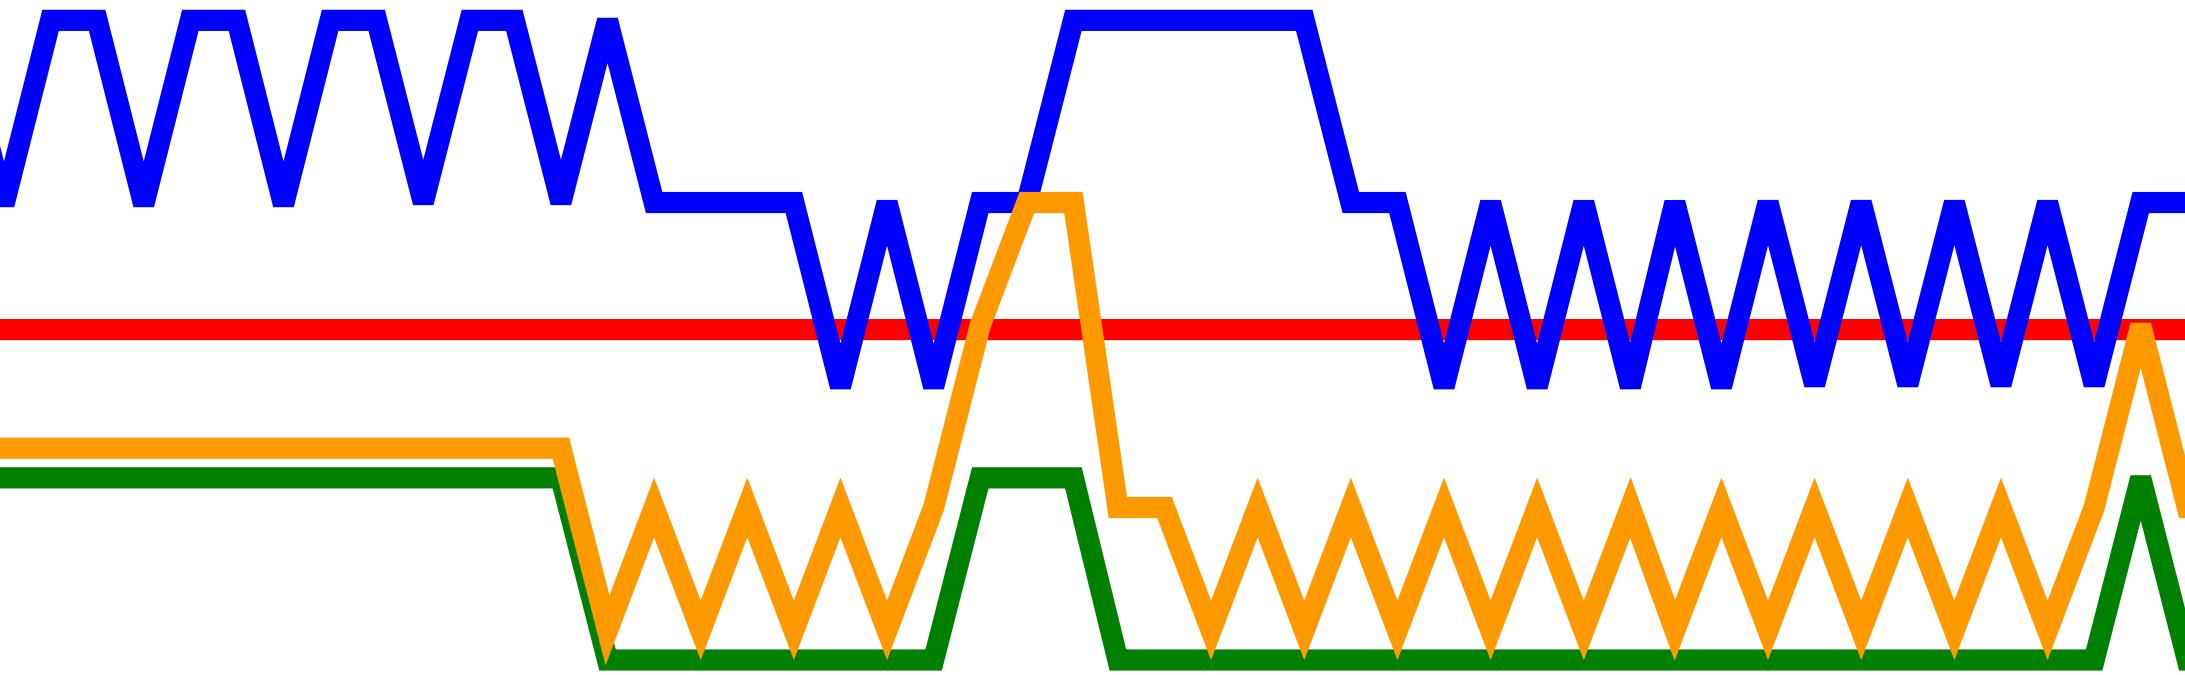
# Memory Consumption: x264





original with pool allocation  
original without pool allocation

SCM single refresh  
SCM continuous refresh

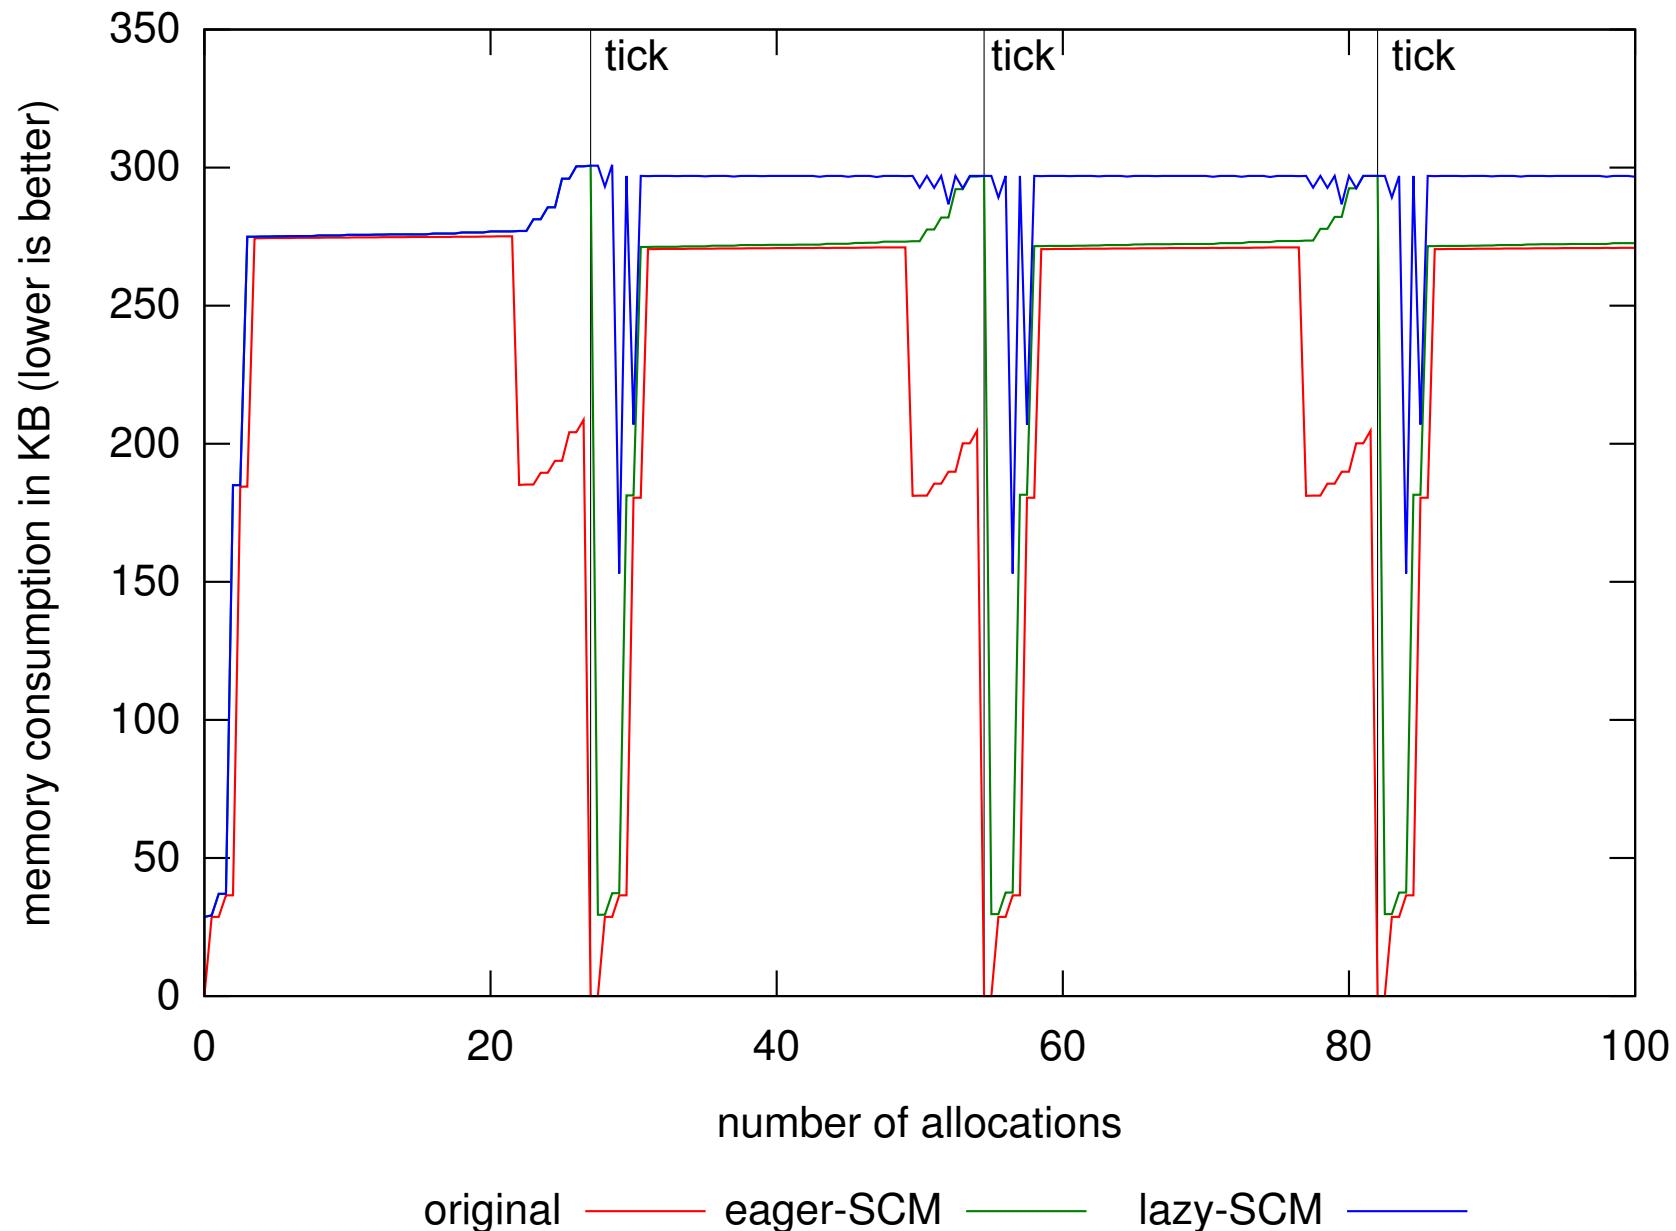


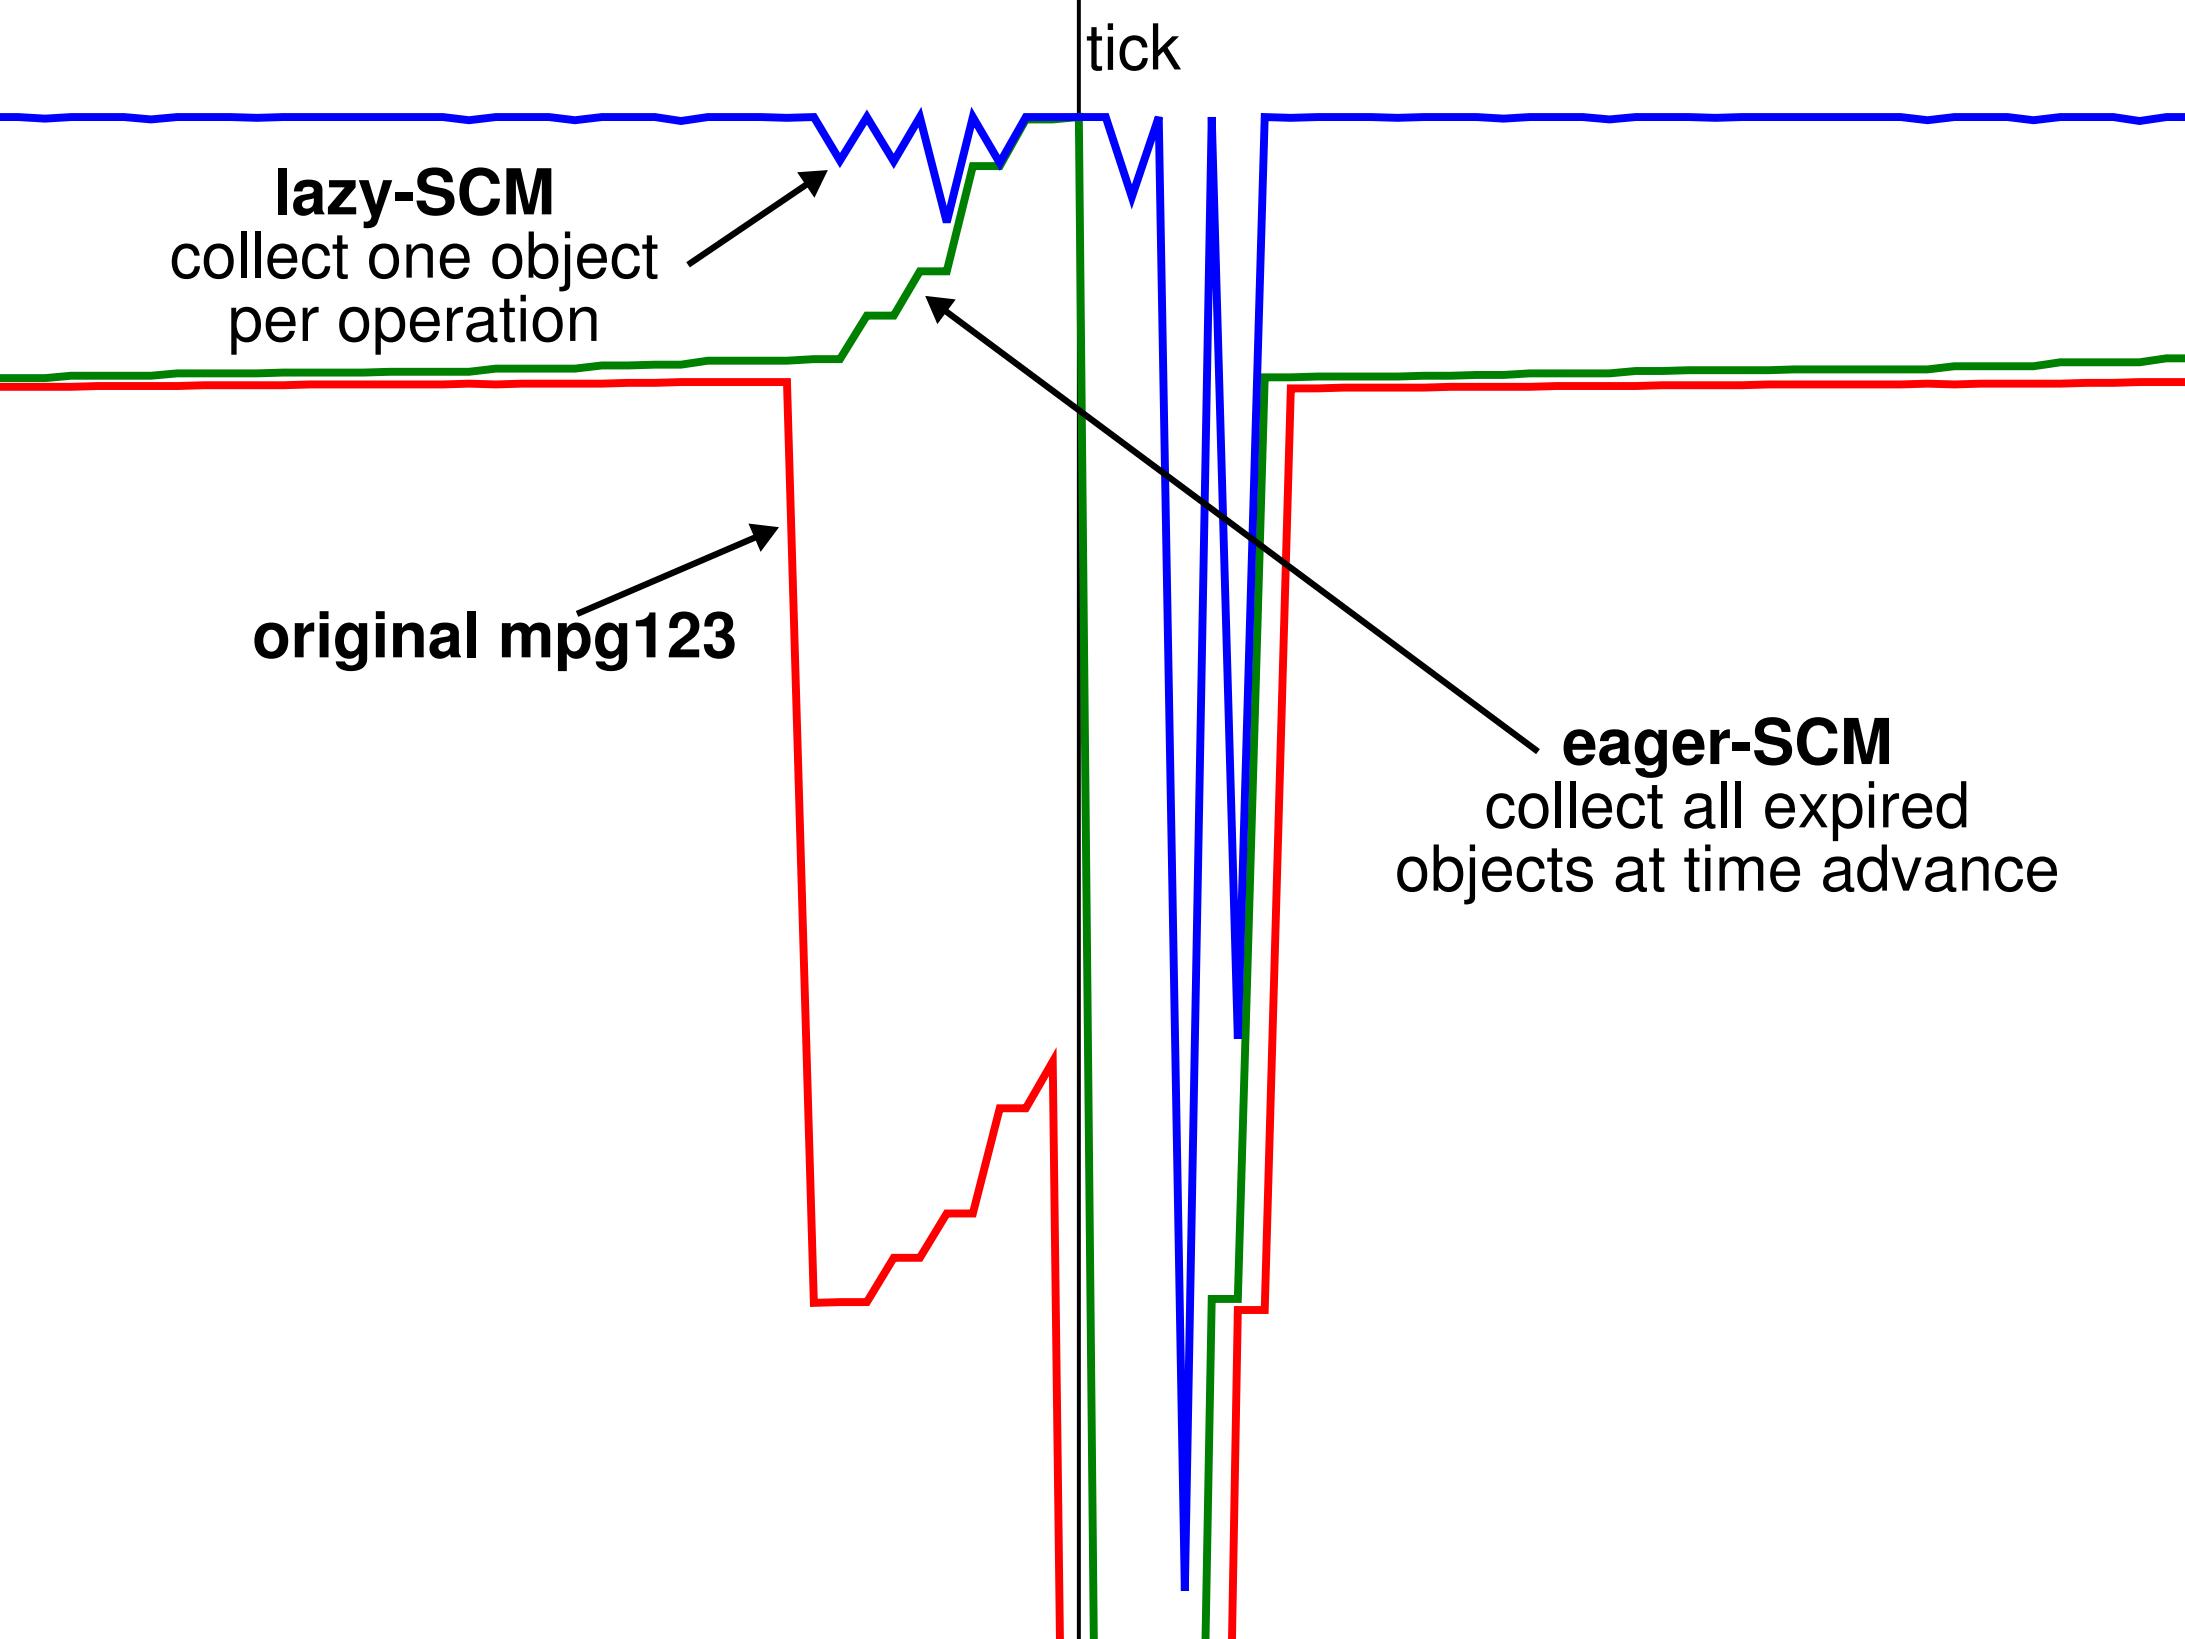
- ▶ Bounded memory overhead
- ▶ Easier to use than explicit deallocation (no reference counting)

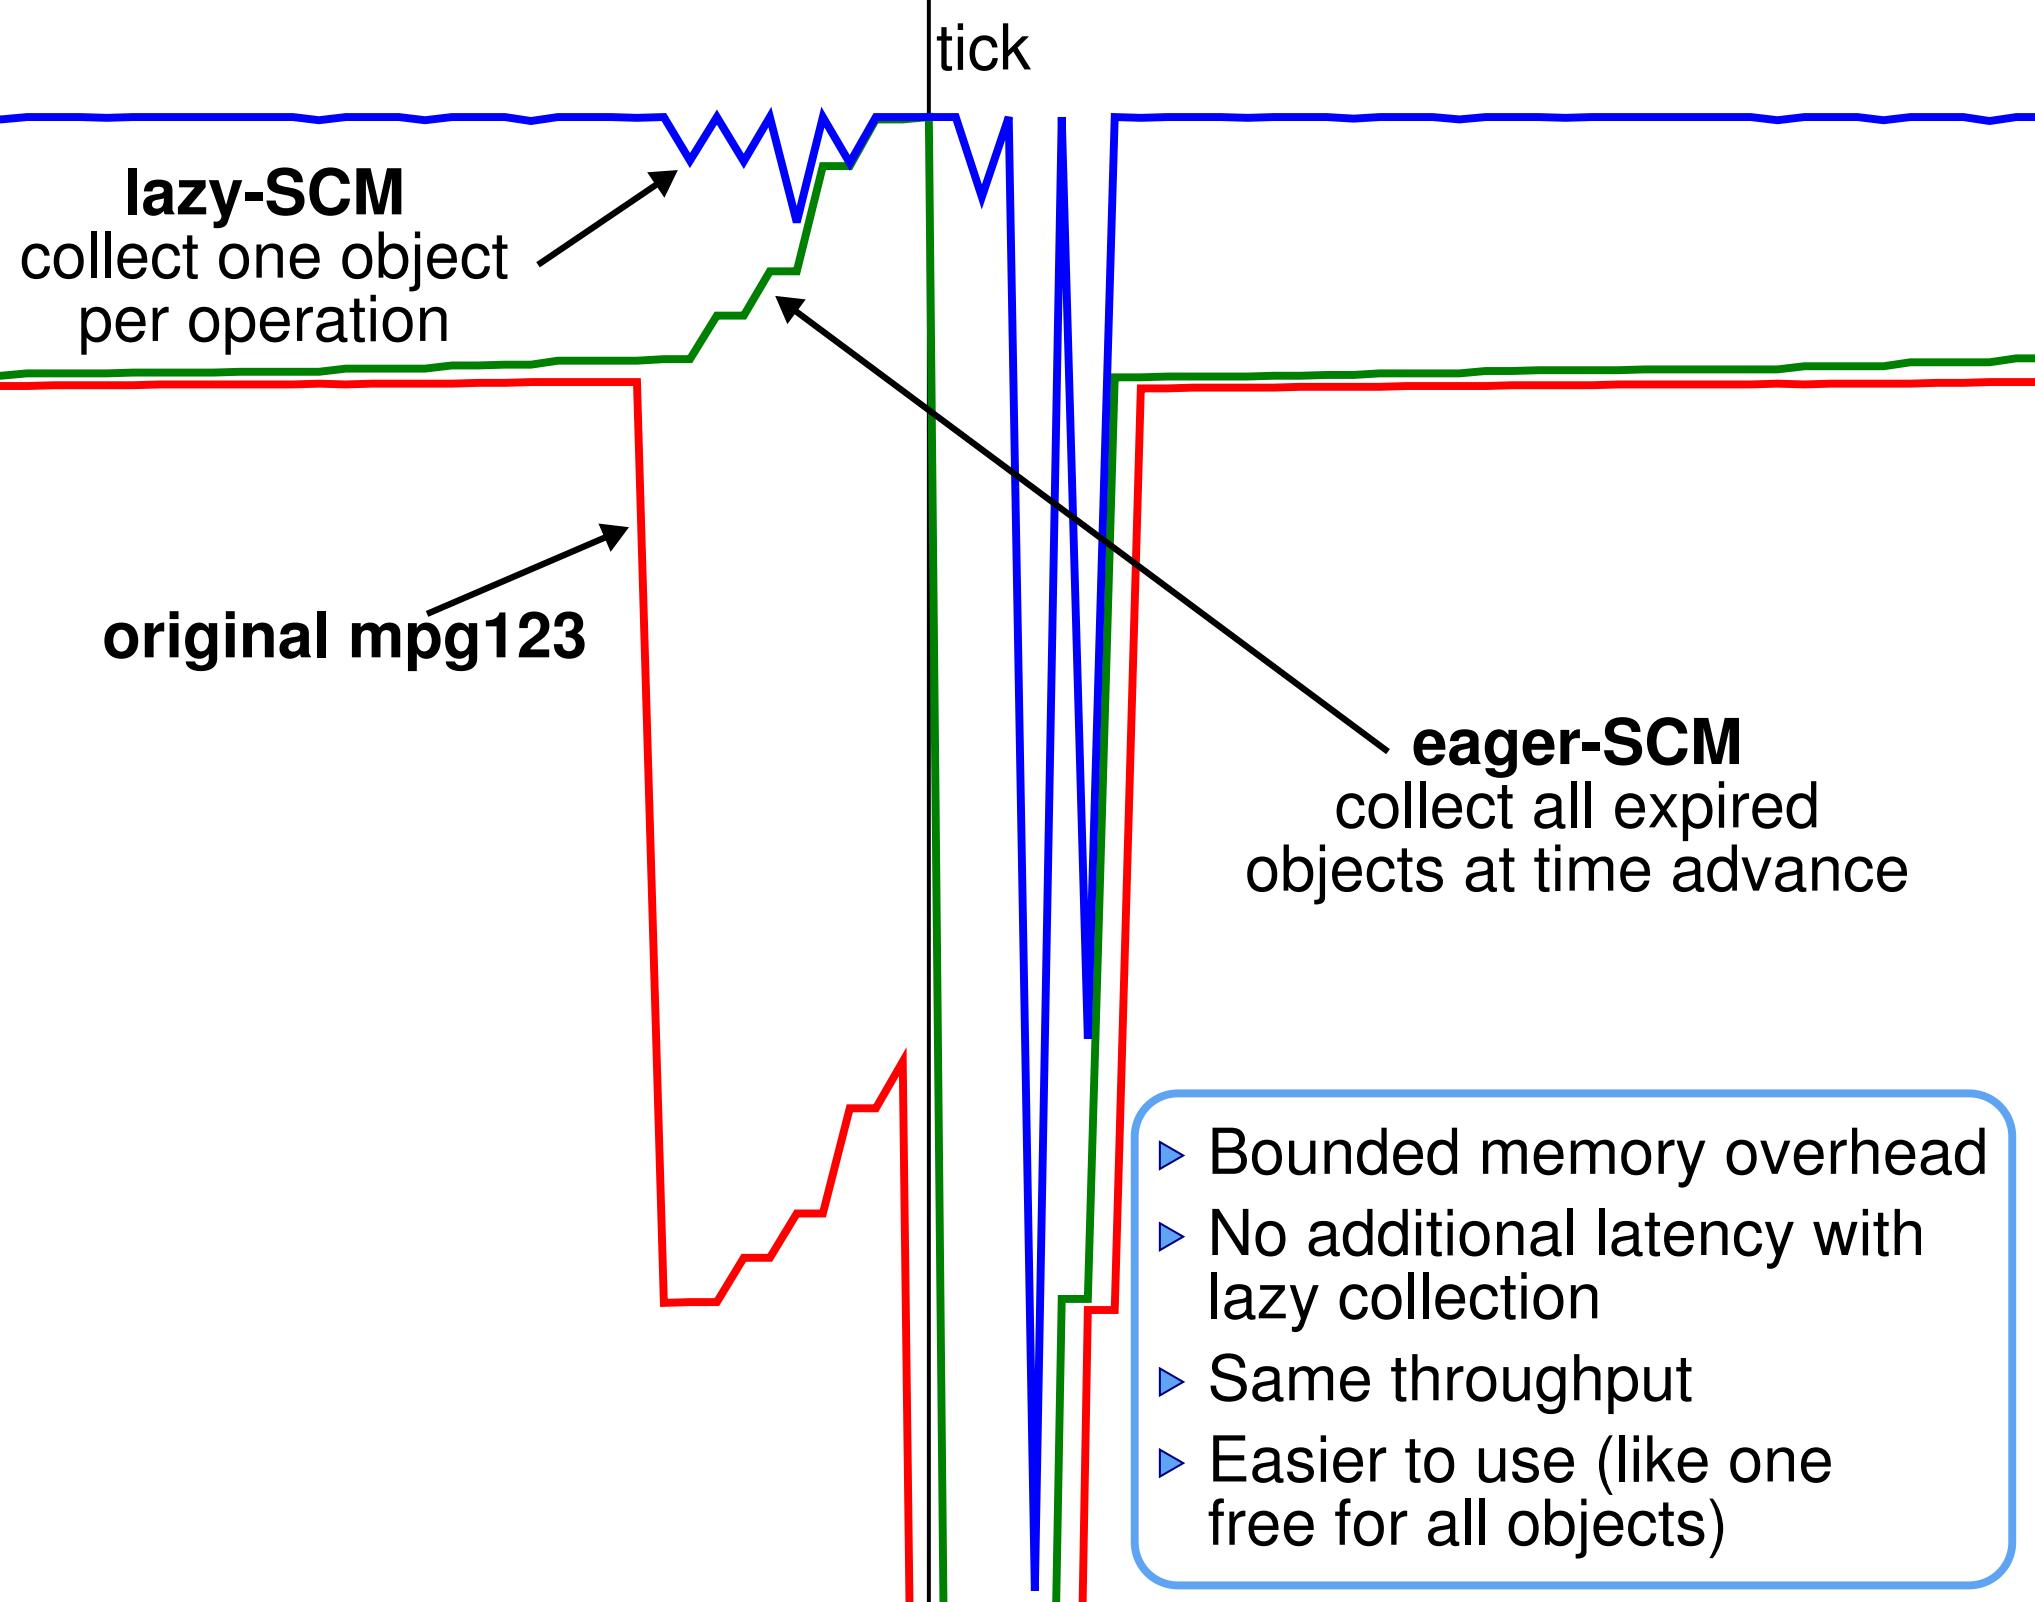
original with pool allocation —————  
original without pool allocation —————

SCM single refresh —————  
SCM continuous refresh —————

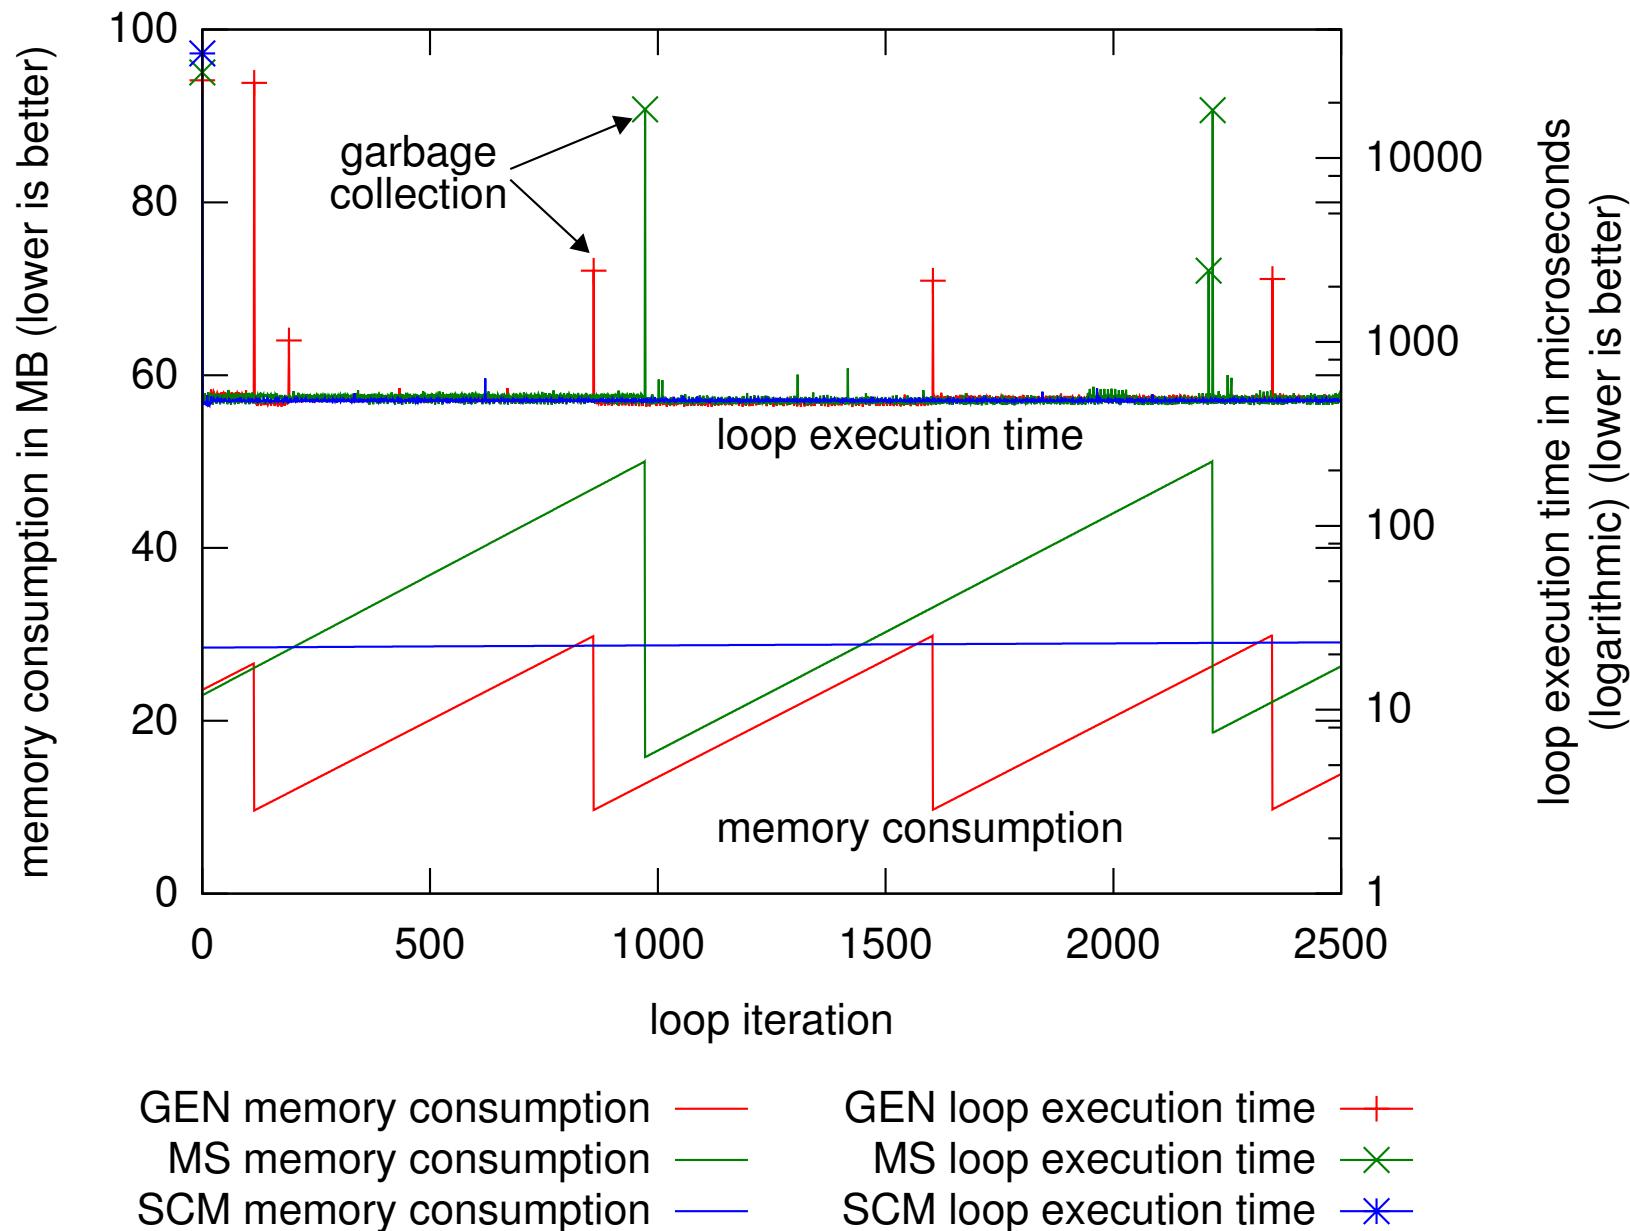
# Lazy vs Eager Collection: mpg123







# Latency: Monte Carlo



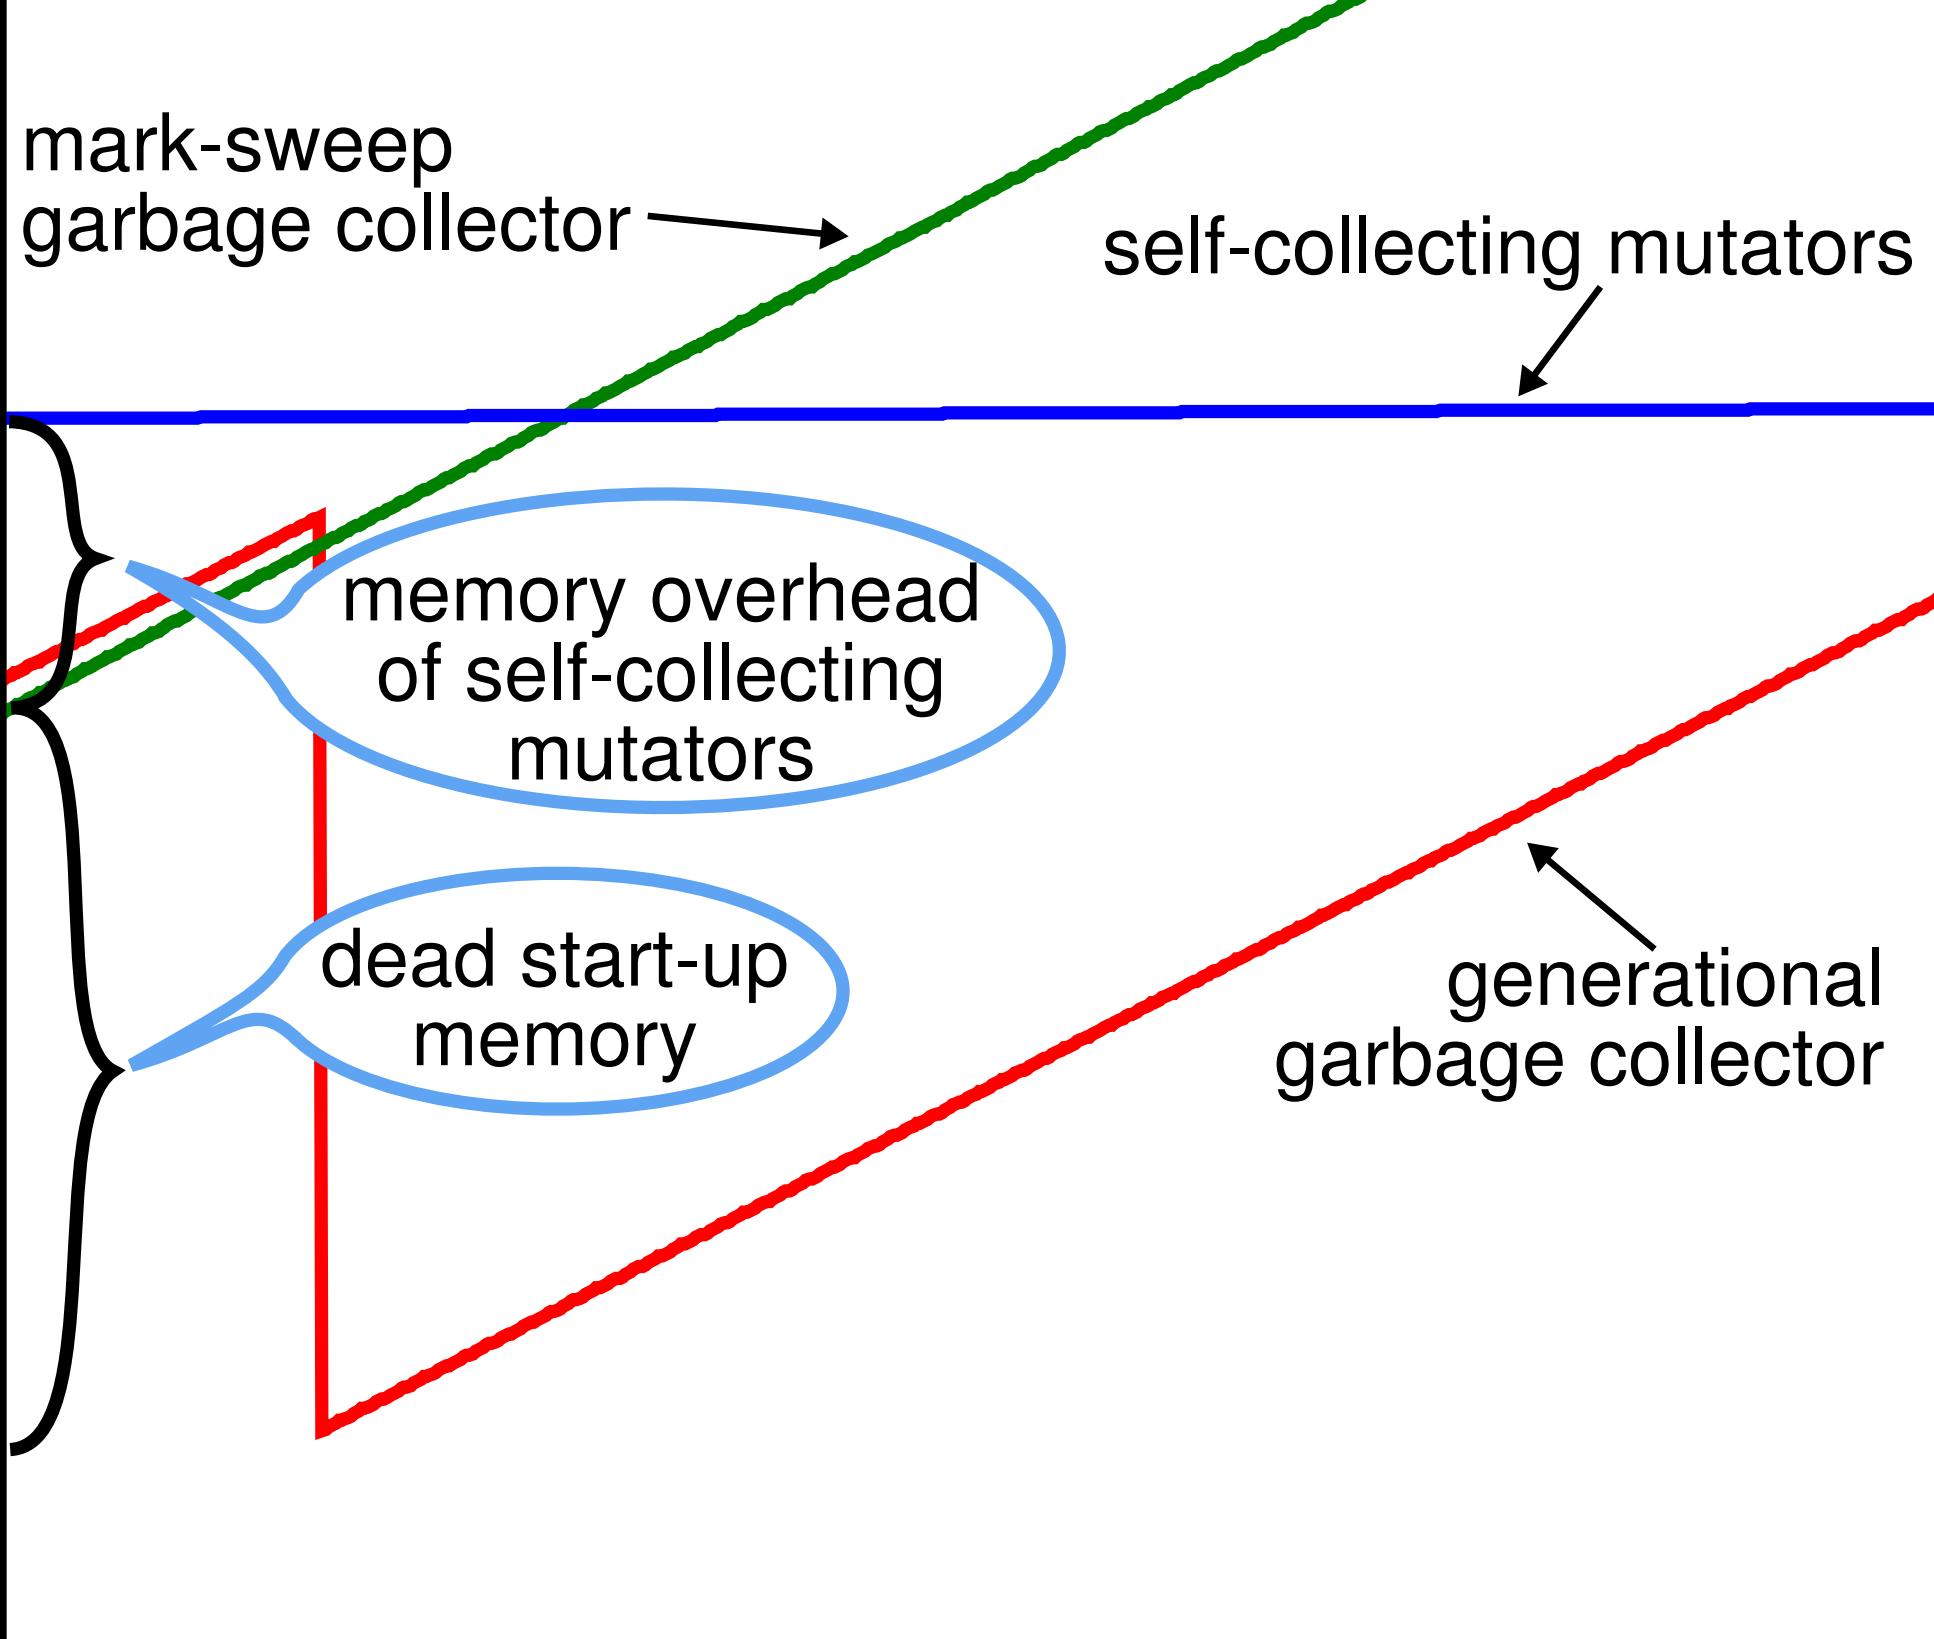
mark-sweep  
garbage collector

self-collecting mutators

memory overhead  
of self-collecting  
mutators

dead start-up  
memory

generational  
garbage collector



mark-sweep  
garbage collector

self-collecting mutators

memory overhead  
of self-collecting  
mutators

dead start-up  
memory

generational  
garbage collector

► Low latency at the expense of safety

# Conclusion

- ▶ Get the code from here:
  - ◆ <http://tiptoe.cs.uni-salzburg.at/short-term-memory/>
- ▶ Future work (memory management community)
  - ◆ **Efficiency:** integrate short-term memory into other allocators
  - ◆ **Correctness:** make short-term memory safe
- ▶ Future work (embedded software community)
  - ◆ **Time as first class citizen:** capture execution progress relevant for memory management

# Thank You

<http://tiptoe.cs.uni-salzburg.at/short-term-memory/>