



Selfie and the Basics

Christoph M. Kirsch, University of Salzburg, Austria

Onward! 2017 @ SPLASH, Vancouver, British Columbia, Canada

What are the absolute basics of
computer science that
everyone
should know about and
understand?

1. Identify a **concept** that you feel everyone should know about and understand
2. Write a **program** that exemplifies that concept in different ways
3. List the **basics** that you need to know about and understand to understand that program

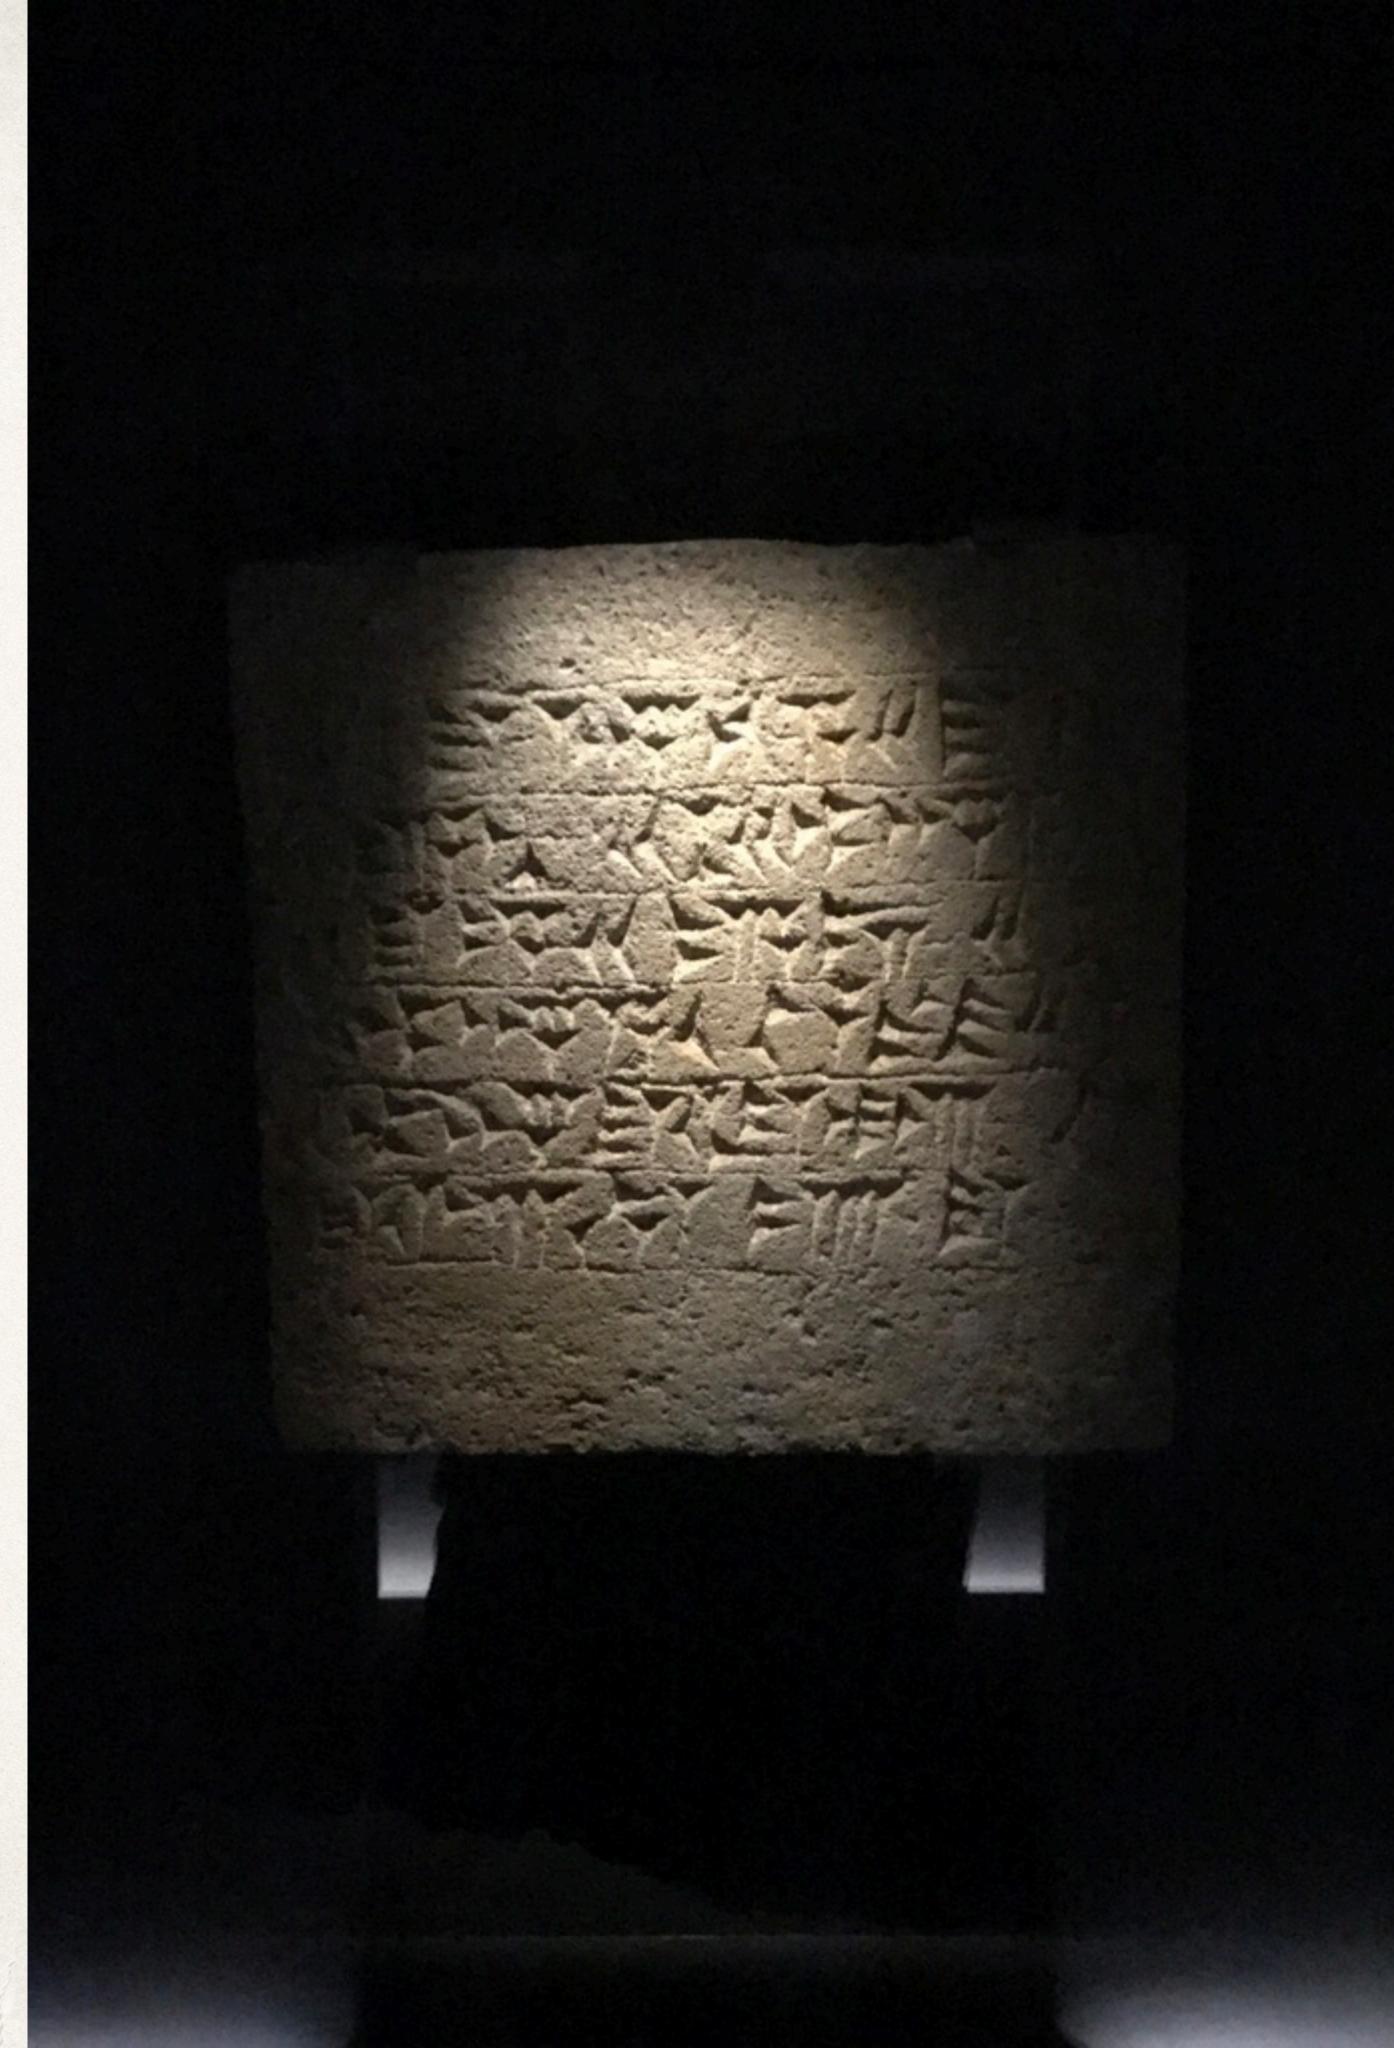
selfie.cs.uni-salzburg.at

...and the Basics:

12 (!) basic principles
essential (!) for understanding
selfie and (?) computer science

What is the meaning
of this sentence?

Selfie as in
self-referentiality



Do people need to understand self-referentiality?

Programming languages
resemble languages but are
really just formalisms with
(hopefully) precise semantics



Interpretation

Compilation

Teaching the Construction of Semantics of Formalisms

Virtualization

Verification

Joint Work

- ❖ Alireza Abyaneh
- ❖ Martin Aigner
- ❖ Sebastian Arming
- ❖ Christian Barthel
- ❖ Simon Bauer
- ❖ Thomas Hütter
- ❖ Alexander Kollert
- ❖ Michael Lippautz
- ❖ Cornelia Mayer
- ❖ Philipp Mayer
- ❖ Christian Moesl
- ❖ Simone Oblasser
- ❖ Clement Poncelet
- ❖ Sara Seidl
- ❖ Ana Sokolova
- ❖ Manuel Widmoser

Inspiration

- ❖ Armin Biere: SAT/SMT Solvers
- ❖ Donald Knuth: Art
- ❖ Jochen Liedtke: Microkernels
- ❖ David Patterson: RISC
- ❖ Niklaus Wirth: Compilers



Selfie: Teaching Computer Science

[selfie.cs.uni-salzburg.at]

- ✿ *Selfie* is a self-referential 7k-line C implementation (in a single file) of:
 1. a self-compiling compiler called *starc* that compiles a tiny subset of C called C Star (C*) to a tiny subset of MIPS64/RISC-V called MIPSter,
 2. a self-executing emulator called *mipster* that executes MIPSter code including itself when compiled with starc,
 3. a self-hosting hypervisor called *hypster* that virtualizes mipster and can host all of selfie including itself,
 4. a tiny C* library called *libcstar* utilized by all of selfie, and
 5. a tiny, experimental SAT solver called *babysat*.

Also, there is a...

- ✿ linker (in-memory only)
- ✿ disassembler (w/ source code line numbers)
- ✿ debugger (tracks full machine state)
- ✿ profiler (#proc-calls, #loop-iterations, #loads, #stores)

Discussion of Selfie recently reached
3rd place on Hacker News

news.ycombinator.com

Website

selfie.cs.uni-salzburg.at

Book (Draft)

[leanpub.com / selfie](http://leanpub.com/selfie)

Code

[github.com / ckstystemsteaching / selfie](https://github.com/cksystemsteaching/selfie)

nsf.gov / csforall

code.org

computingatschool.org.uk

programbydesign.org

k12cs.org

bootstrapworld.org

csfieldguide.org.nz

5 statements:
assignment
while
if
return
procedure()

```
int atoi(int *s) {  
    int i;  
    int n;  
    int c;  
  
    i = 0;  
    n = 0;  
    c = *(s+i);  
  
    while (c != 0) {  
        n = n * 10 + c - '0';  
        if (n < 0)  
            return -1;  
  
        i = i + 1;  
        c = *(s+i);  
    }  
  
    return n;  
}
```

integer arithmetics
pointer arithmetics

no data types other
than int and int*
and dereferencing:
the * operator

character literals
string literals

no bitwise operators
no Boolean operators

library: exit, malloc, open, read, write

Minimally complex, maximally self- contained system

Programming languages vs systems engineering?



```
> make  
cc -w -m32 -D'main(a,b)=main(a,char**argv)' selfie.c -o selfie
```

*bootstrapping **selfie.c** into x86 **selfie** executable
using standard C compiler*

(also available for RISC-V machines)

```
> ./selfie
./selfie: usage: selfie { -c { source } | -o binary | -s assembly
| -l binary } [ ( -m | -d | -y | -min | -mob ) size ... ]
```

selfie usage

```
> ./selfie -c selfie.c
```

./selfie: this is selfie's starc compiling **selfie.c**

./selfie: 176408 characters read in 7083 lines and 969 comments
./selfie: with 97779(55.55%) characters in 28914 actual symbols
./selfie: 261 global variables, 289 procedures, 450 string literals
./selfie: 1958 calls, 723 assignments, 57 while, 572 if, 243 return
./selfie: 121660 bytes generated with 28779 instructions and 6544 bytes of data

compiling selfie.c with x86 selfie executable

(takes seconds)

```
> ./selfie -c selfie.c -m 2 -c selfie.c
```

./selfie: this is selfie's starc compiling **selfie.c**

./selfie: this is selfie's mipster executing **selfie.c** with **2MB** of physical memory

selfie.c: this is selfie's starc compiling **selfie.c**

selfie.c: exiting with exit code **0** and **1.05MB** of mallocated memory

./selfie: this is selfie's mipster terminating **selfie.c** with exit code **0** and **1.16MB** of mapped memory

*compiling **selfie.c** with x86 **selfie** executable into a MIPster executable
and*

*then running that MIPster executable to compile **selfie.c** again
(takes ~6 minutes)*

```
> ./selfie -c selfie.c -o selfie1.m -m 2 -c selfie.c -o selfie2.m
```

./selfie: this is selfie's starc compiling **selfie.c**

./selfie: 121660 bytes with 28779 instructions and 6544 bytes of data
written into **selfie1.m**

./selfie: this is selfie's mipster executing **selfie1.m** with 2MB of
physical memory

selfie1.m: this is selfie's starc compiling **selfie.c**

selfie1.m: 121660 bytes with 28779 instructions and 6544 bytes of data
written into **selfie2.m**

selfie1.m: exiting with exit code 0 and 1.05MB of mallocated memory

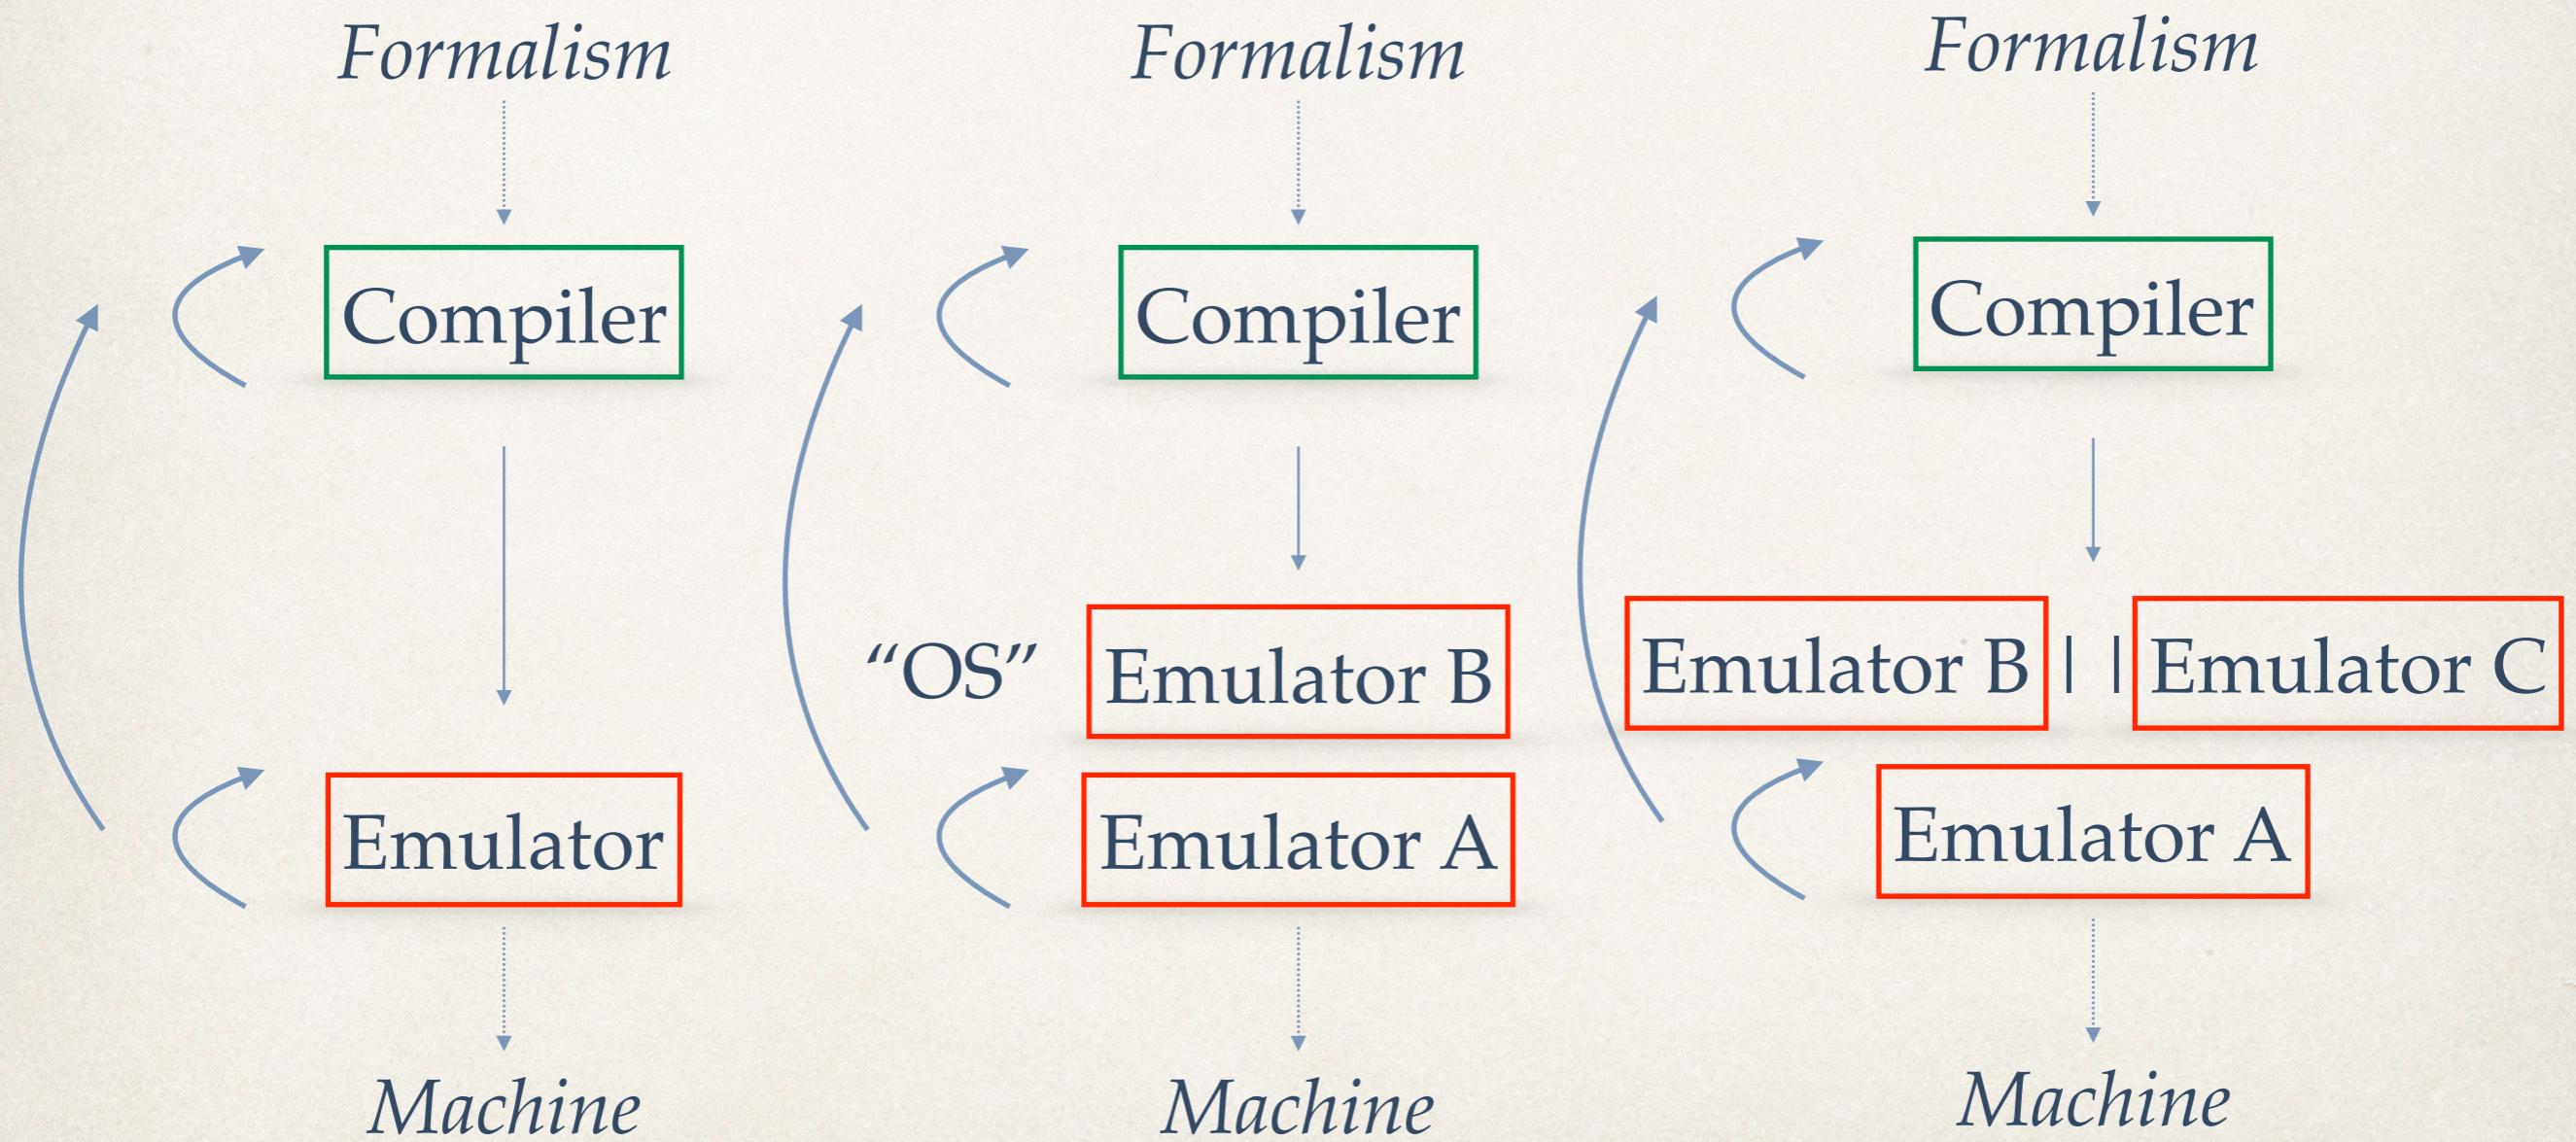
./selfie: this is selfie's mipster terminating **selfie1.m** with exit
code 0 and 1.16MB of mapped memory

*compiling **selfie.c** into a MIPSter executable **selfie1.m***

and

*then running **selfie1.m** to compile **selfie.c**
into another MIPSter executable **selfie2.m**
(takes ~6 minutes)*

Implementing an OS Kernel: 1-Week Homework Assignment



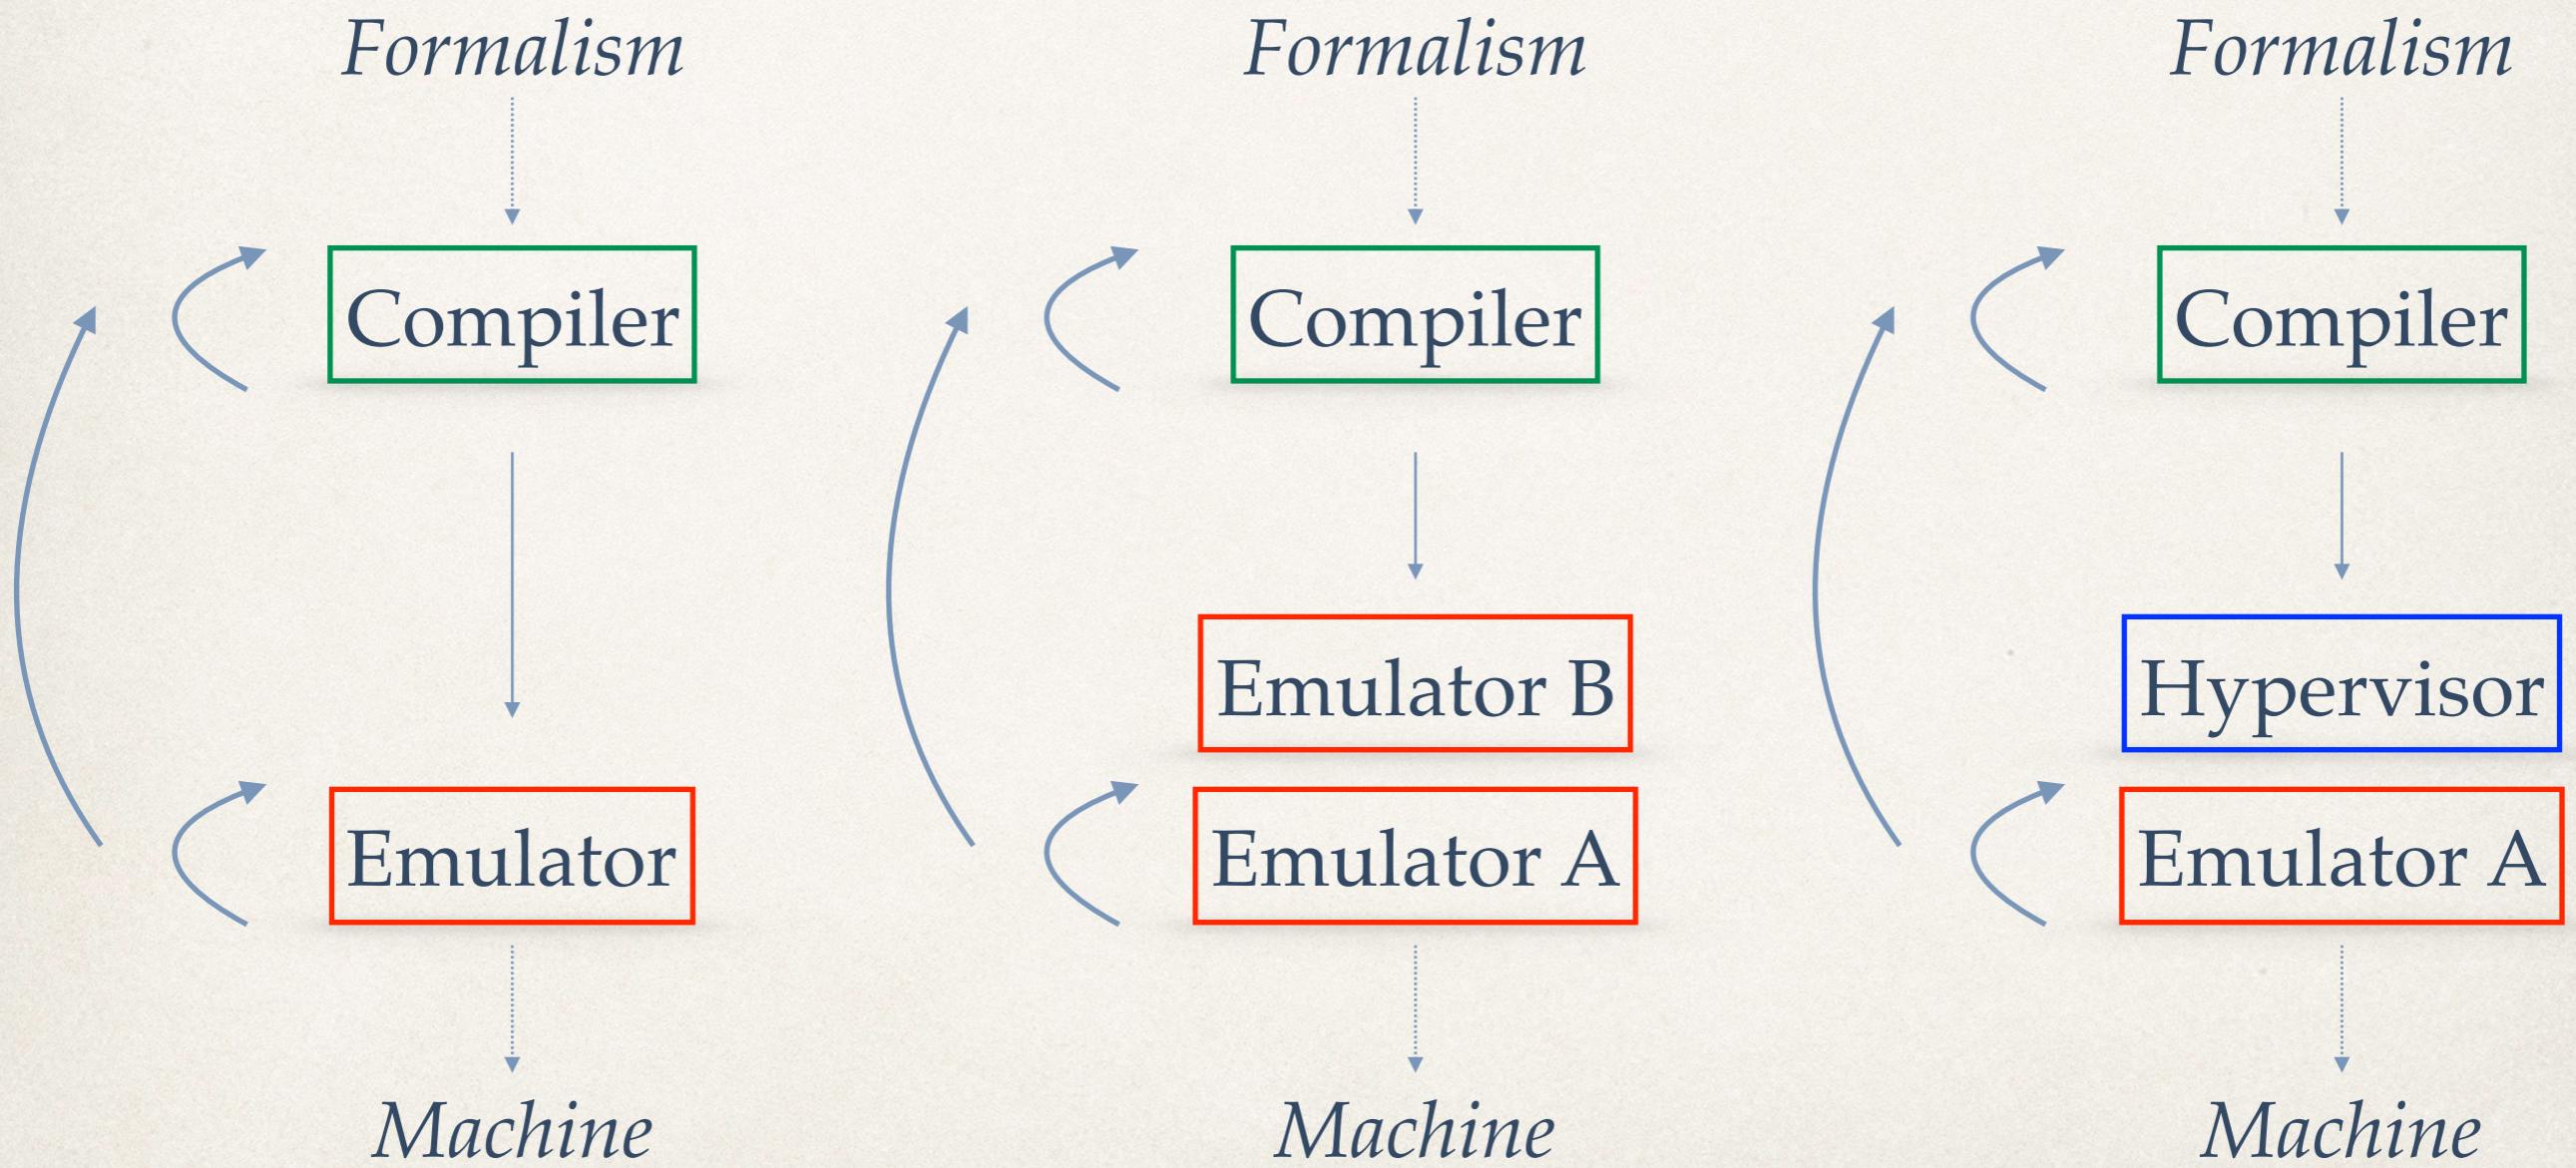
```
> ./selfie -c selfie.c -m 2 -c selfie.c -m 2 -c selfie.c
```

*compiling **selfie.c** with x86 **selfie** executable
and*

*then running that executable to compile **selfie.c** again
and*

*then running that executable to compile **selfie.c** again
(takes ~24 hours)*

Emulation versus Virtualization



```
> ./selfie -c selfie.c -m 2 -c selfie.c -y 2 -c selfie.c
```

*compiling **selfie.c** with x86 **selfie** executable
and*

*then running that executable to compile **selfie.c** again
and*

*then **hosting** that executable in a virtual machine to compile **selfie.c** again
(takes ~12 minutes)*



Ongoing Work

Verification

- ✿ SAT/SMT Solvers (microsat/boolector)
- ✿ Symbolic Execution Engine (KLEE/SAGE)
- ✿ Inductive Theorem Prover (ACL2)

-> microsat in C* is as fast as in C (forget structs, arrays, &&, ||, goto)

ISAs

1. Large memory and multicore support
2. x86 support through binary translation
3. ARM support?



Selfie and the Basics

Library

Compiler

Emulator

Hypervisor

SAT Solver

`selfie.c`

- ❖ Building and Using Selfie:
- ❖ Handling C* Literals:
- ❖ Program/Machine State:
- ❖ C*/Command Line Scanners:
- ❖ C* Parser and Procedures:
- ❖ Symbol Table and the Heap:
- ❖ MIPSter Code Generator:
- ❖ Address Spaces and Storage:
- ❖ (Composite) Data Types:
- ❖ MIPSter Boot Loader:
- ❖ MIPSter Emulator:
- ❖ MIPSter Hypervisor:
- 1. Semantics
- 2. Encoding
- 3. State
- 4. Regularity
- 5. Stack
- 6. Name
- 7. Time
- 8. Memory
- 9. Type
- 10. Bootstrapping
- 11. Interpretation
- 12. Virtualization

Thank you!



AUSTRIAN COMPUTER SCIENCE DAY 2018



15.06.2018 / SALZBURG

acsd2018.cs.uni-salzburg.at