

Distributed Queues: Faster Pools and Better Queues

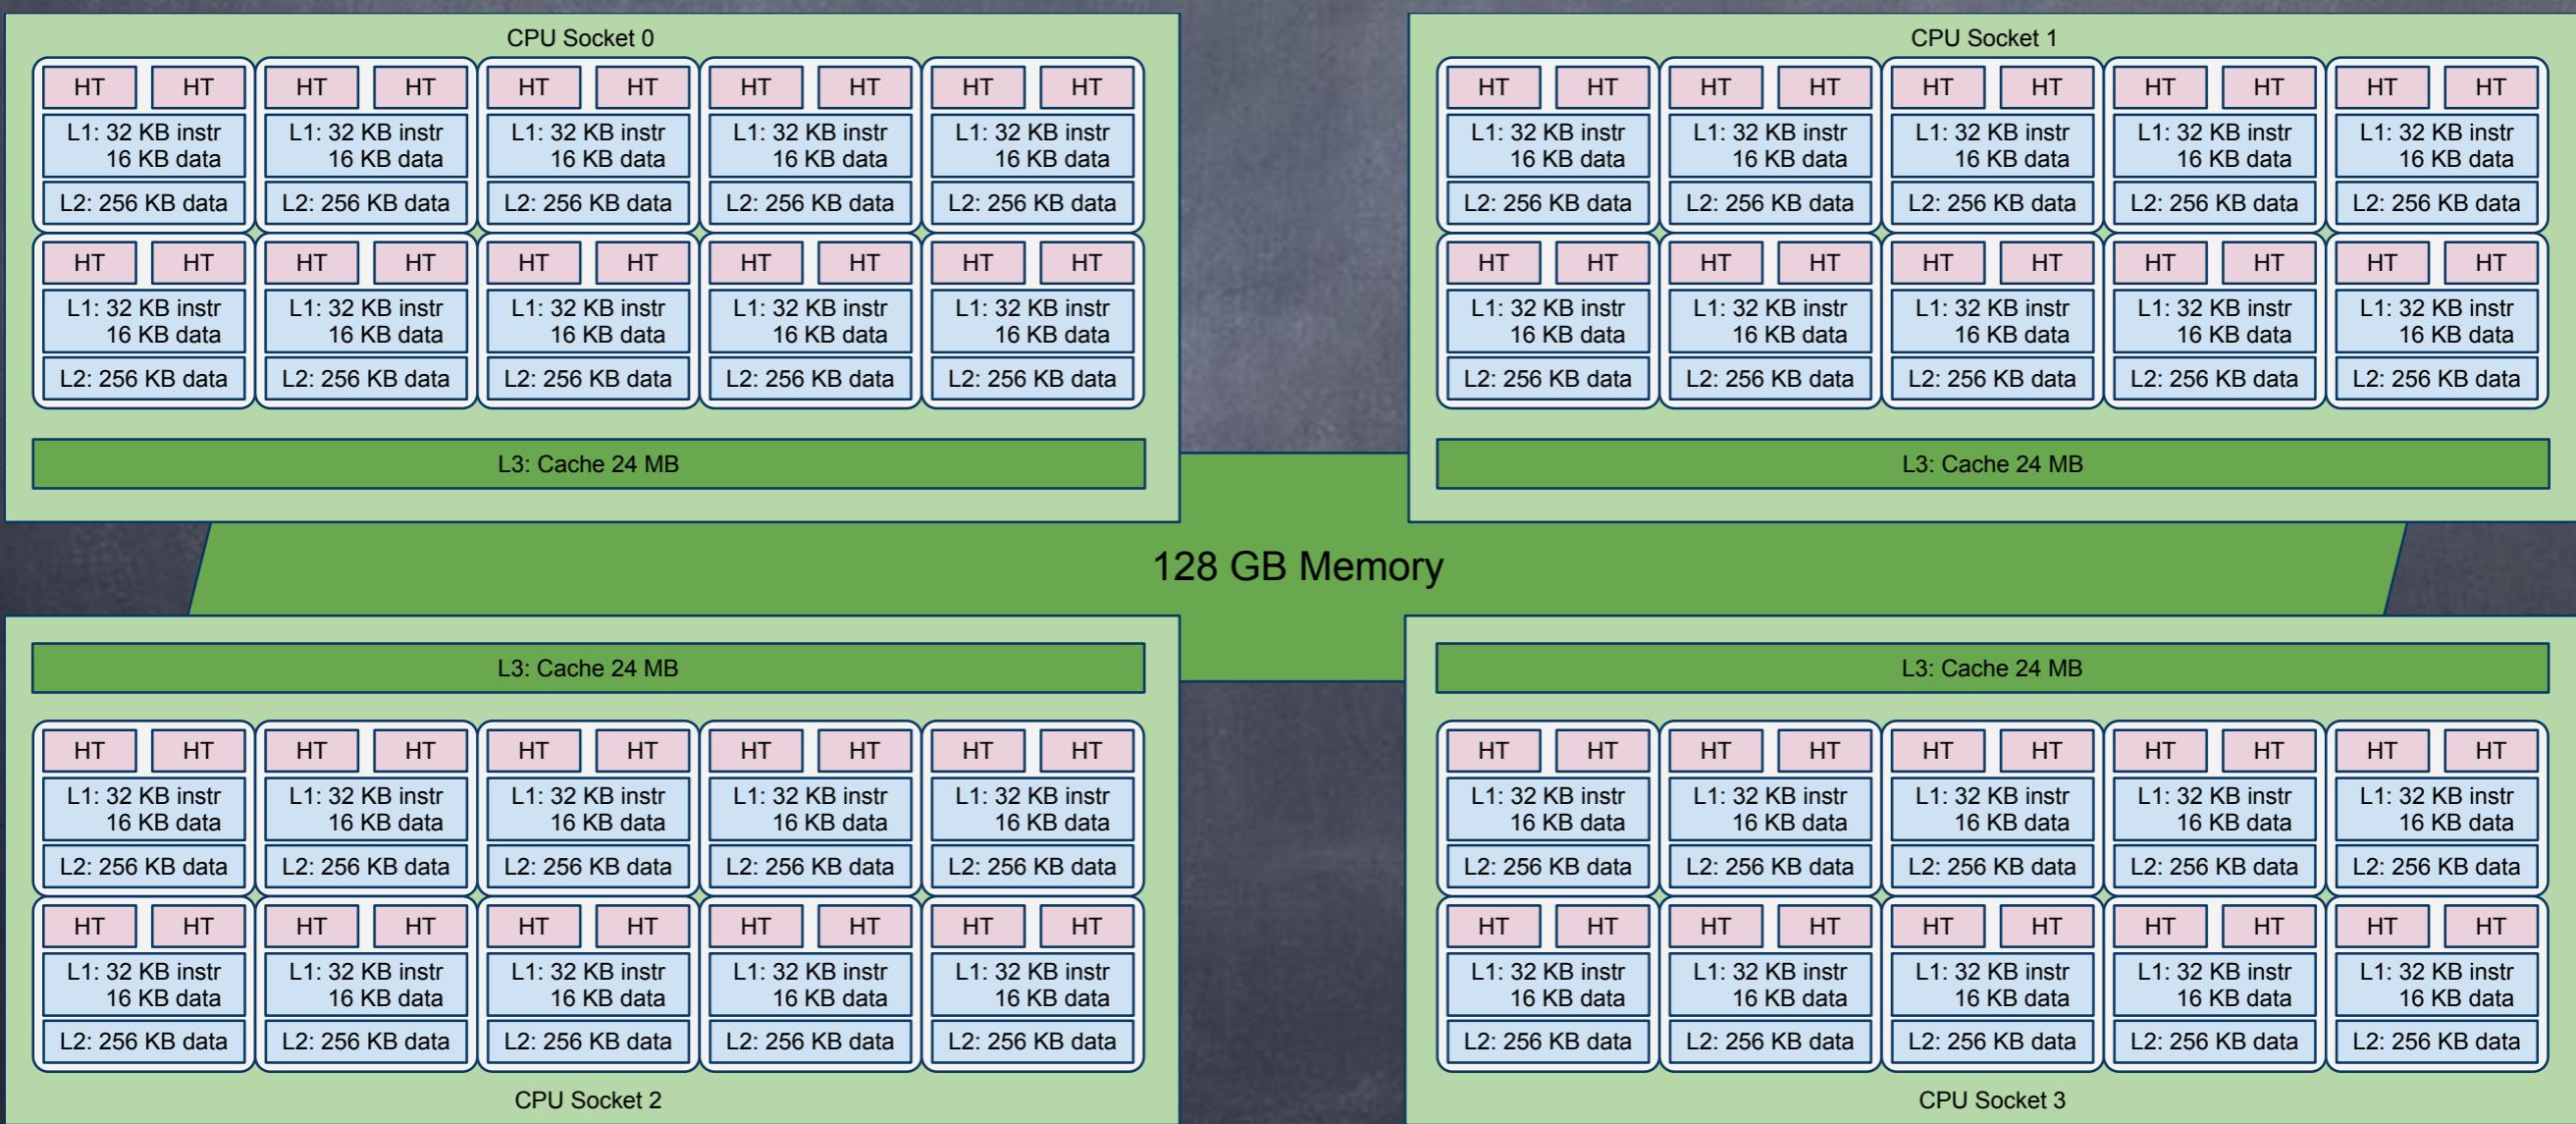
Christoph Kirsch
Universität Salzburg



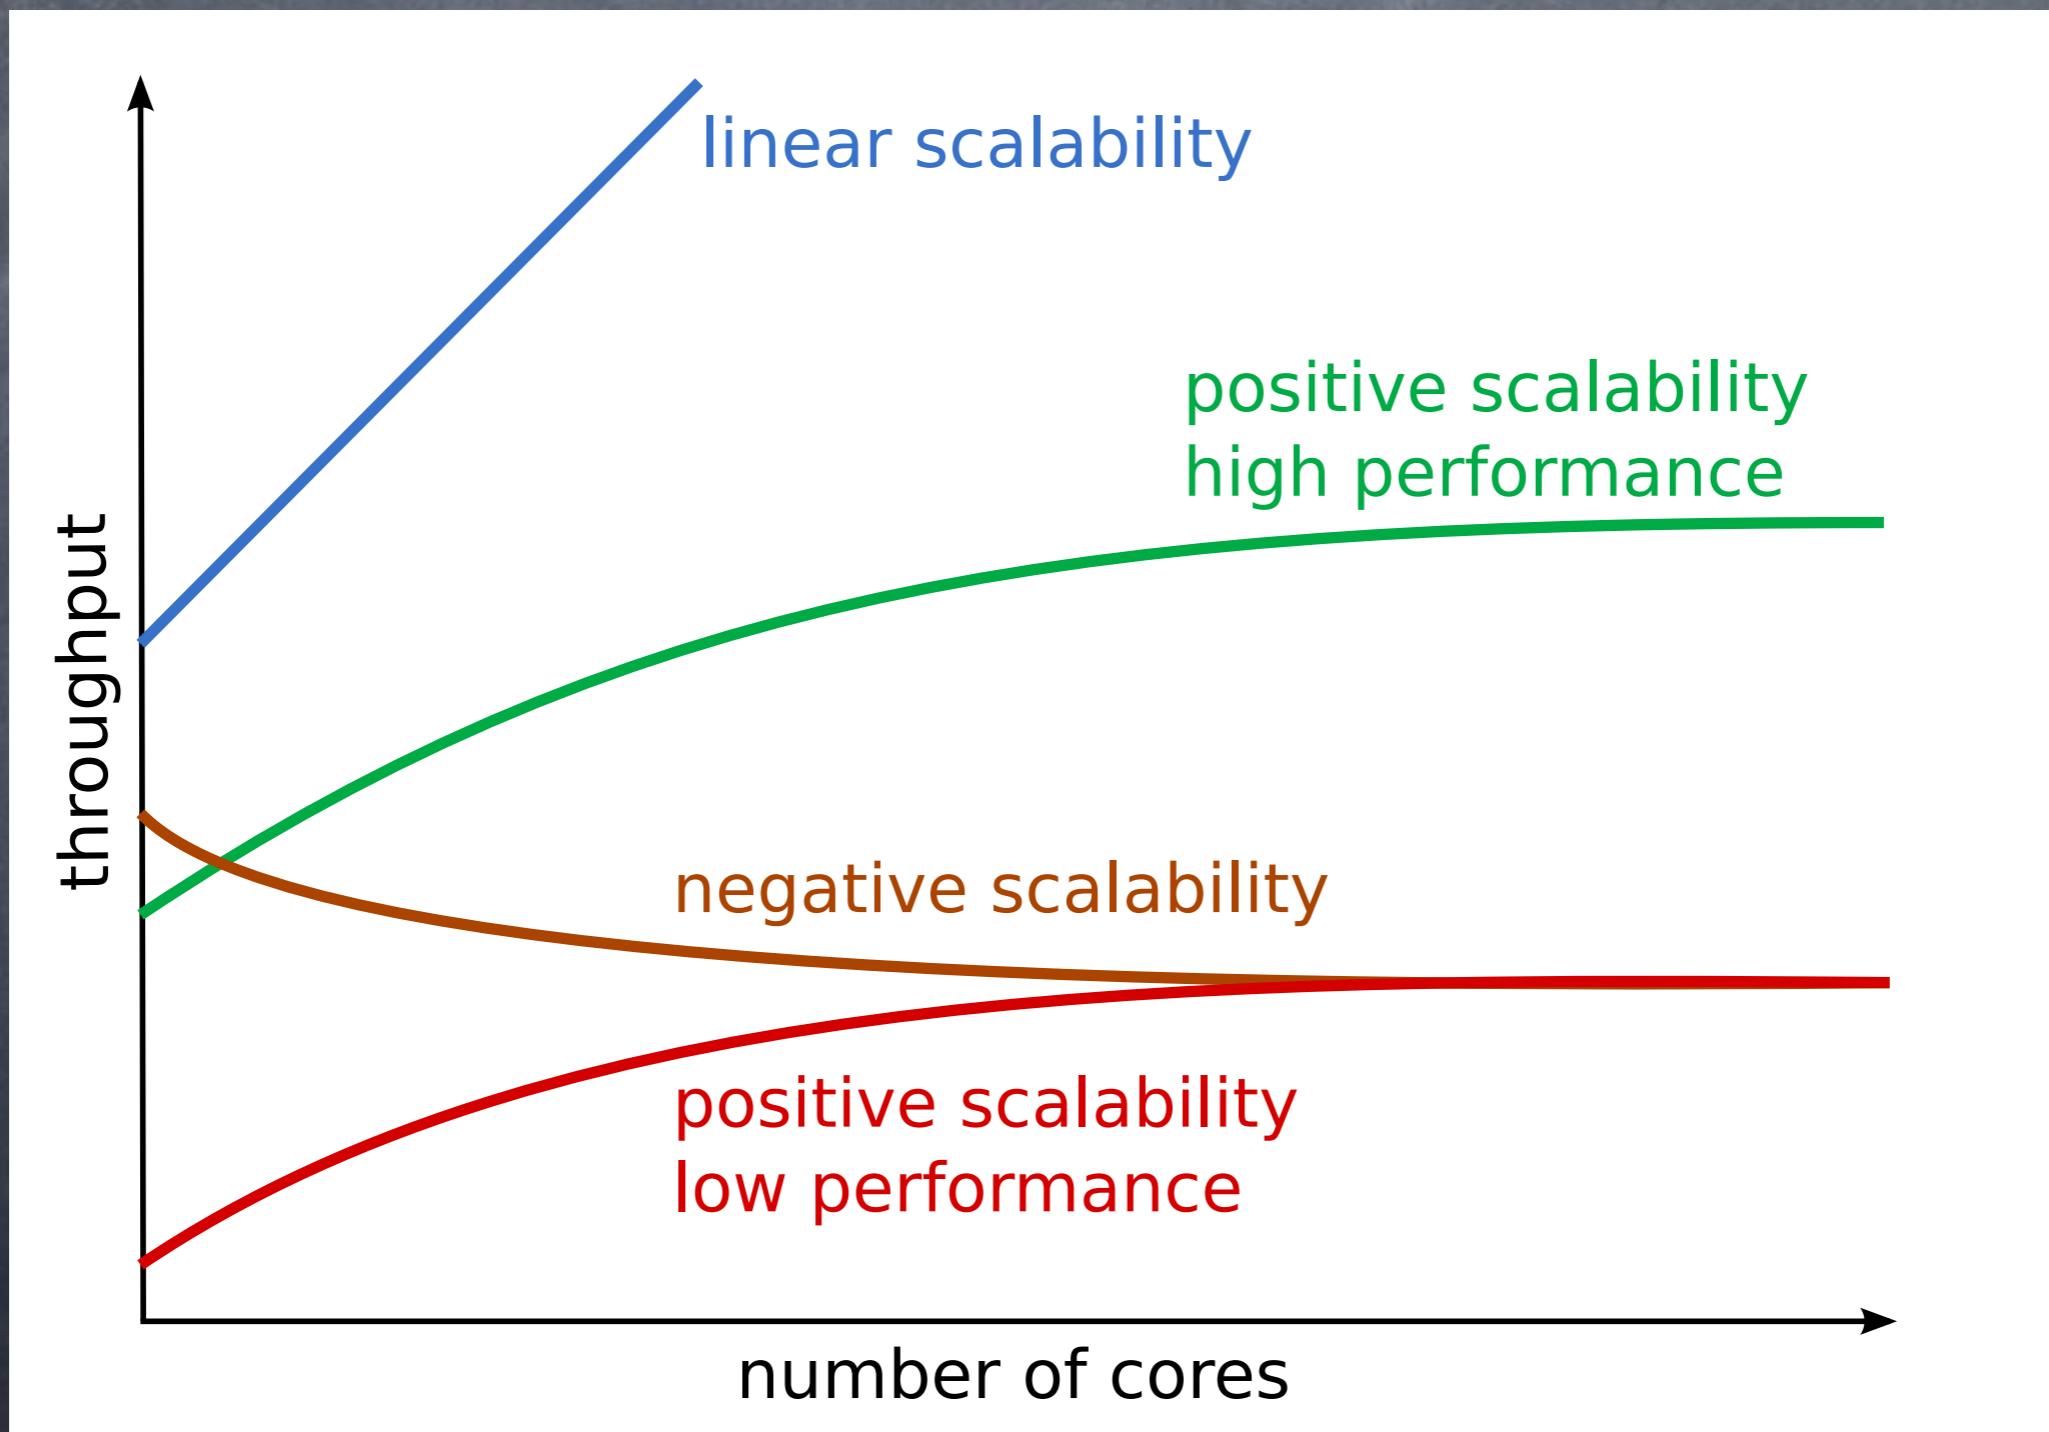
Oracle, Belmont, California, December 2012

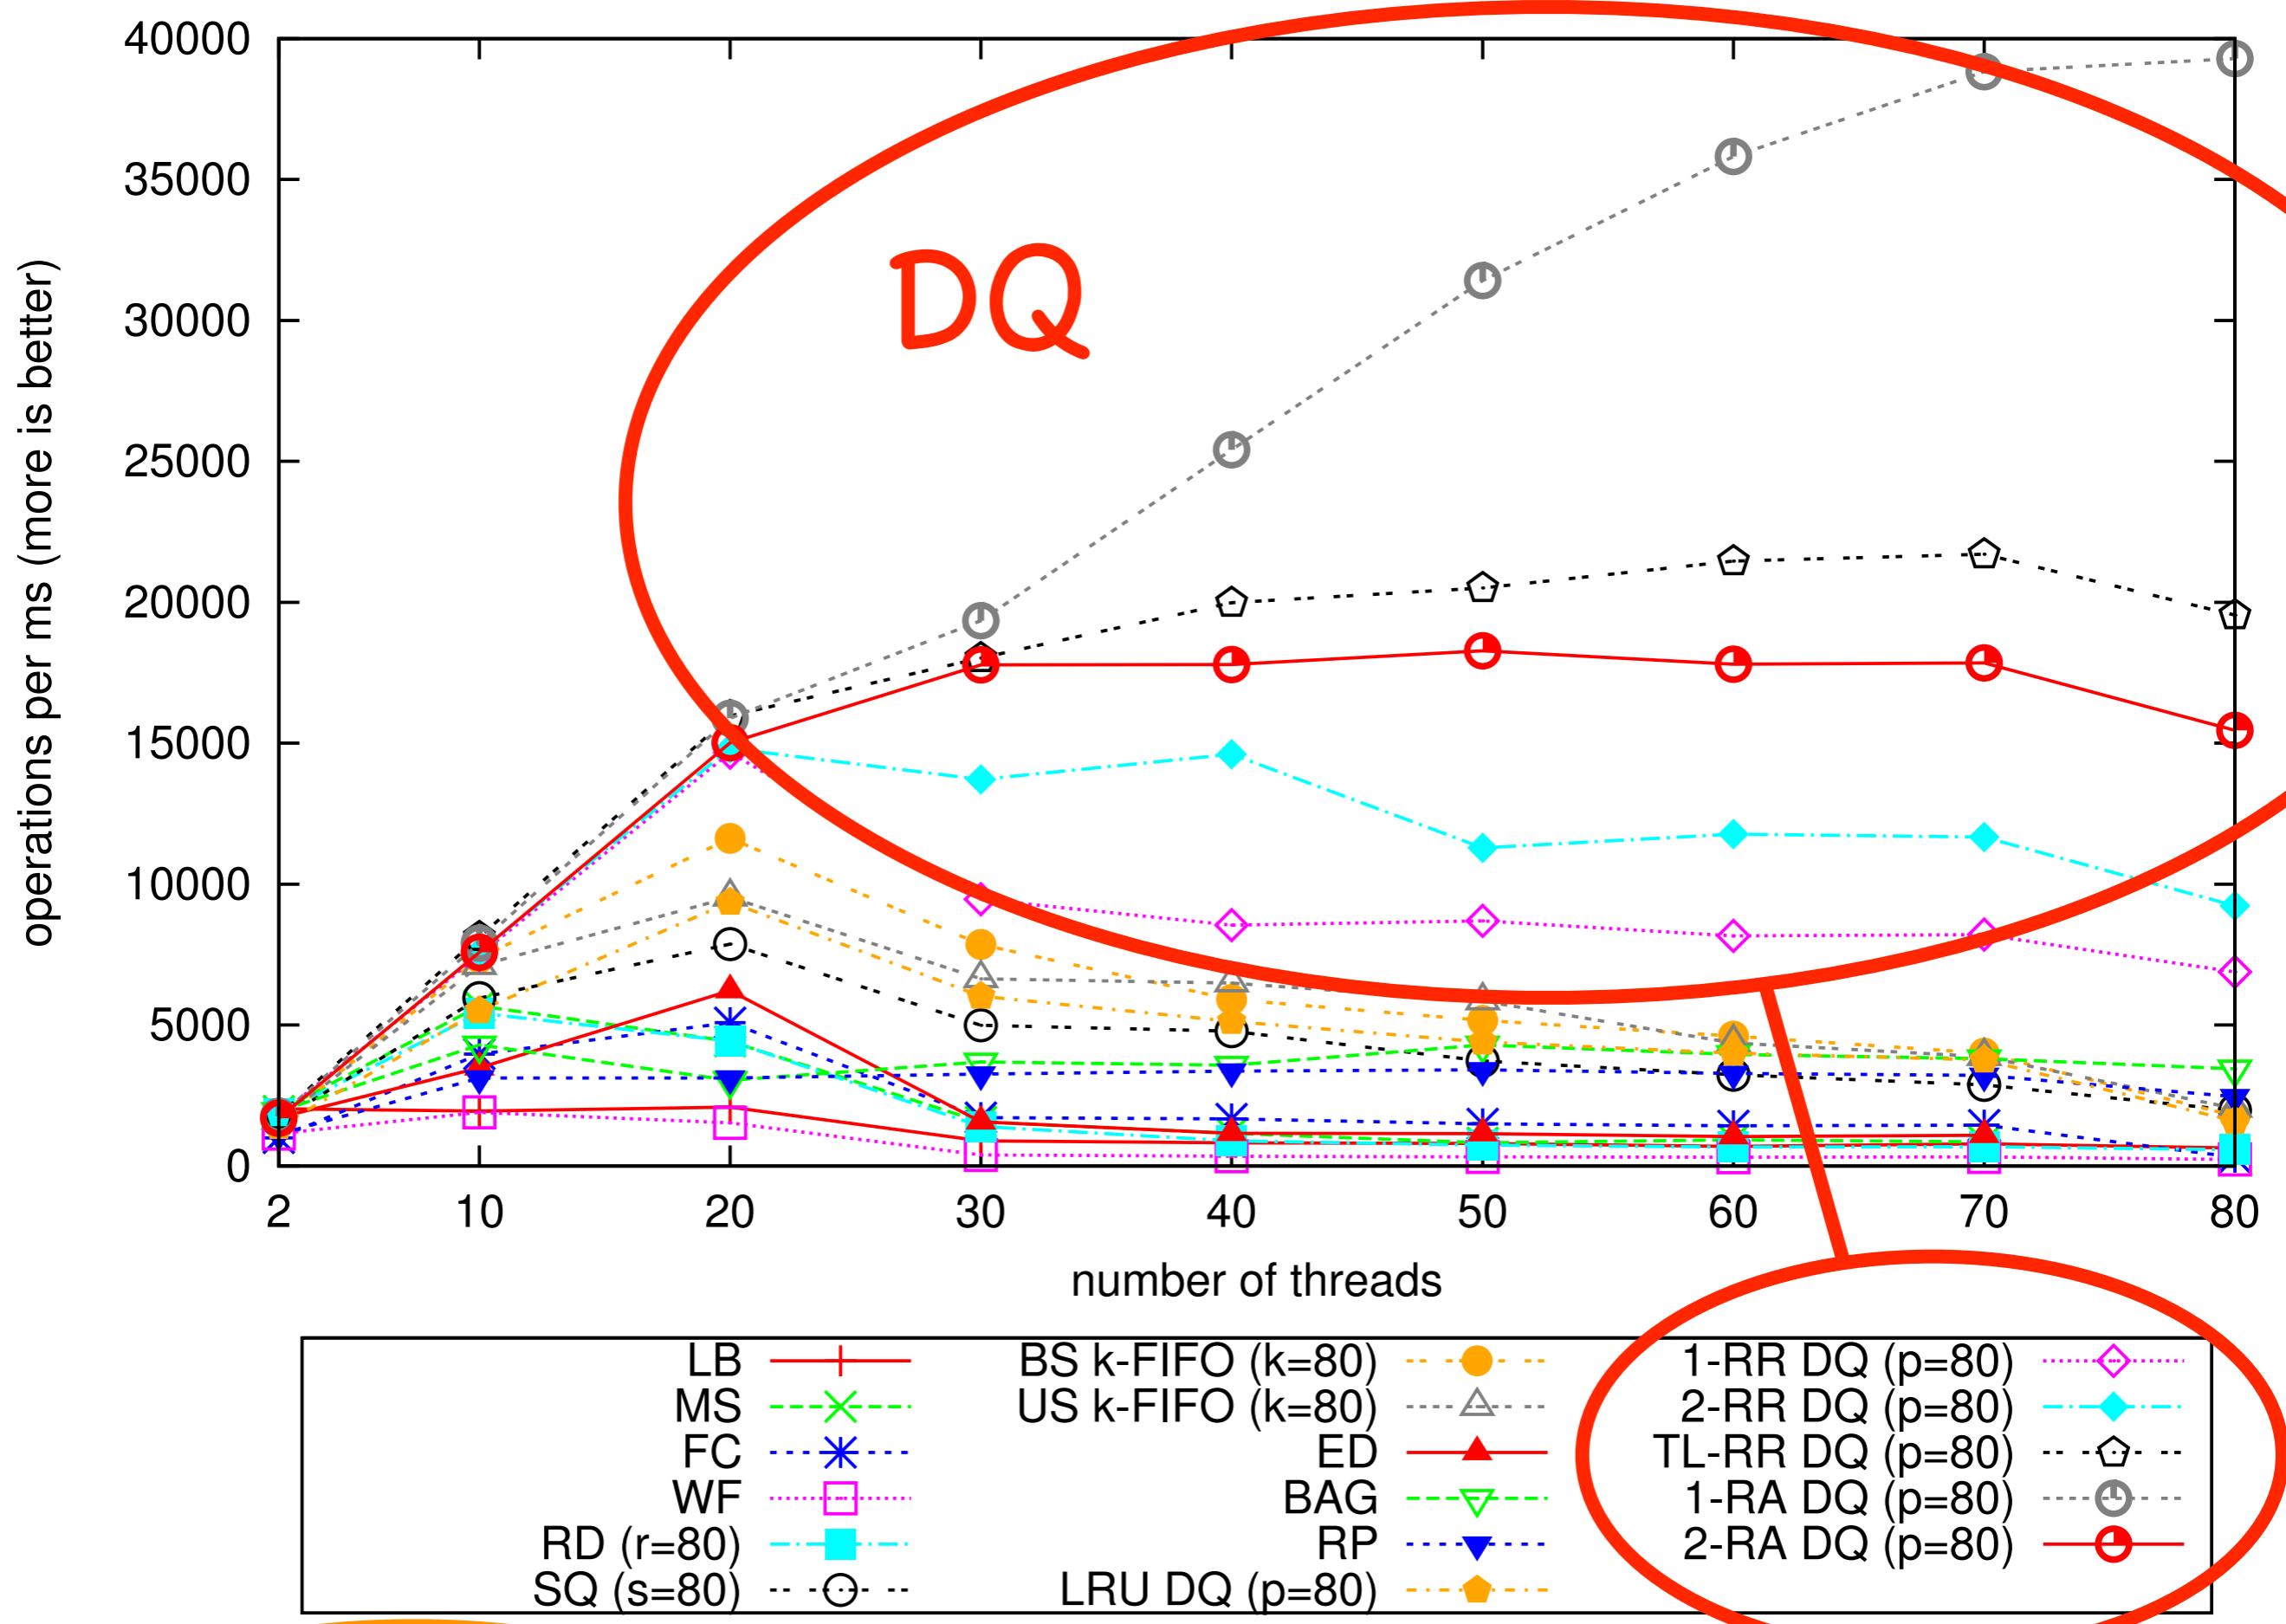
Joint work w/ A. Haas,
M. Lippautz, H. Payer,
A. Sokolova and our
collaborators at IST Austria
T. Henzinger, A. Sezgin

4 processors × 10 cores ×
 2 hardware threads =
 80 hardware threads



Performance & Scalability

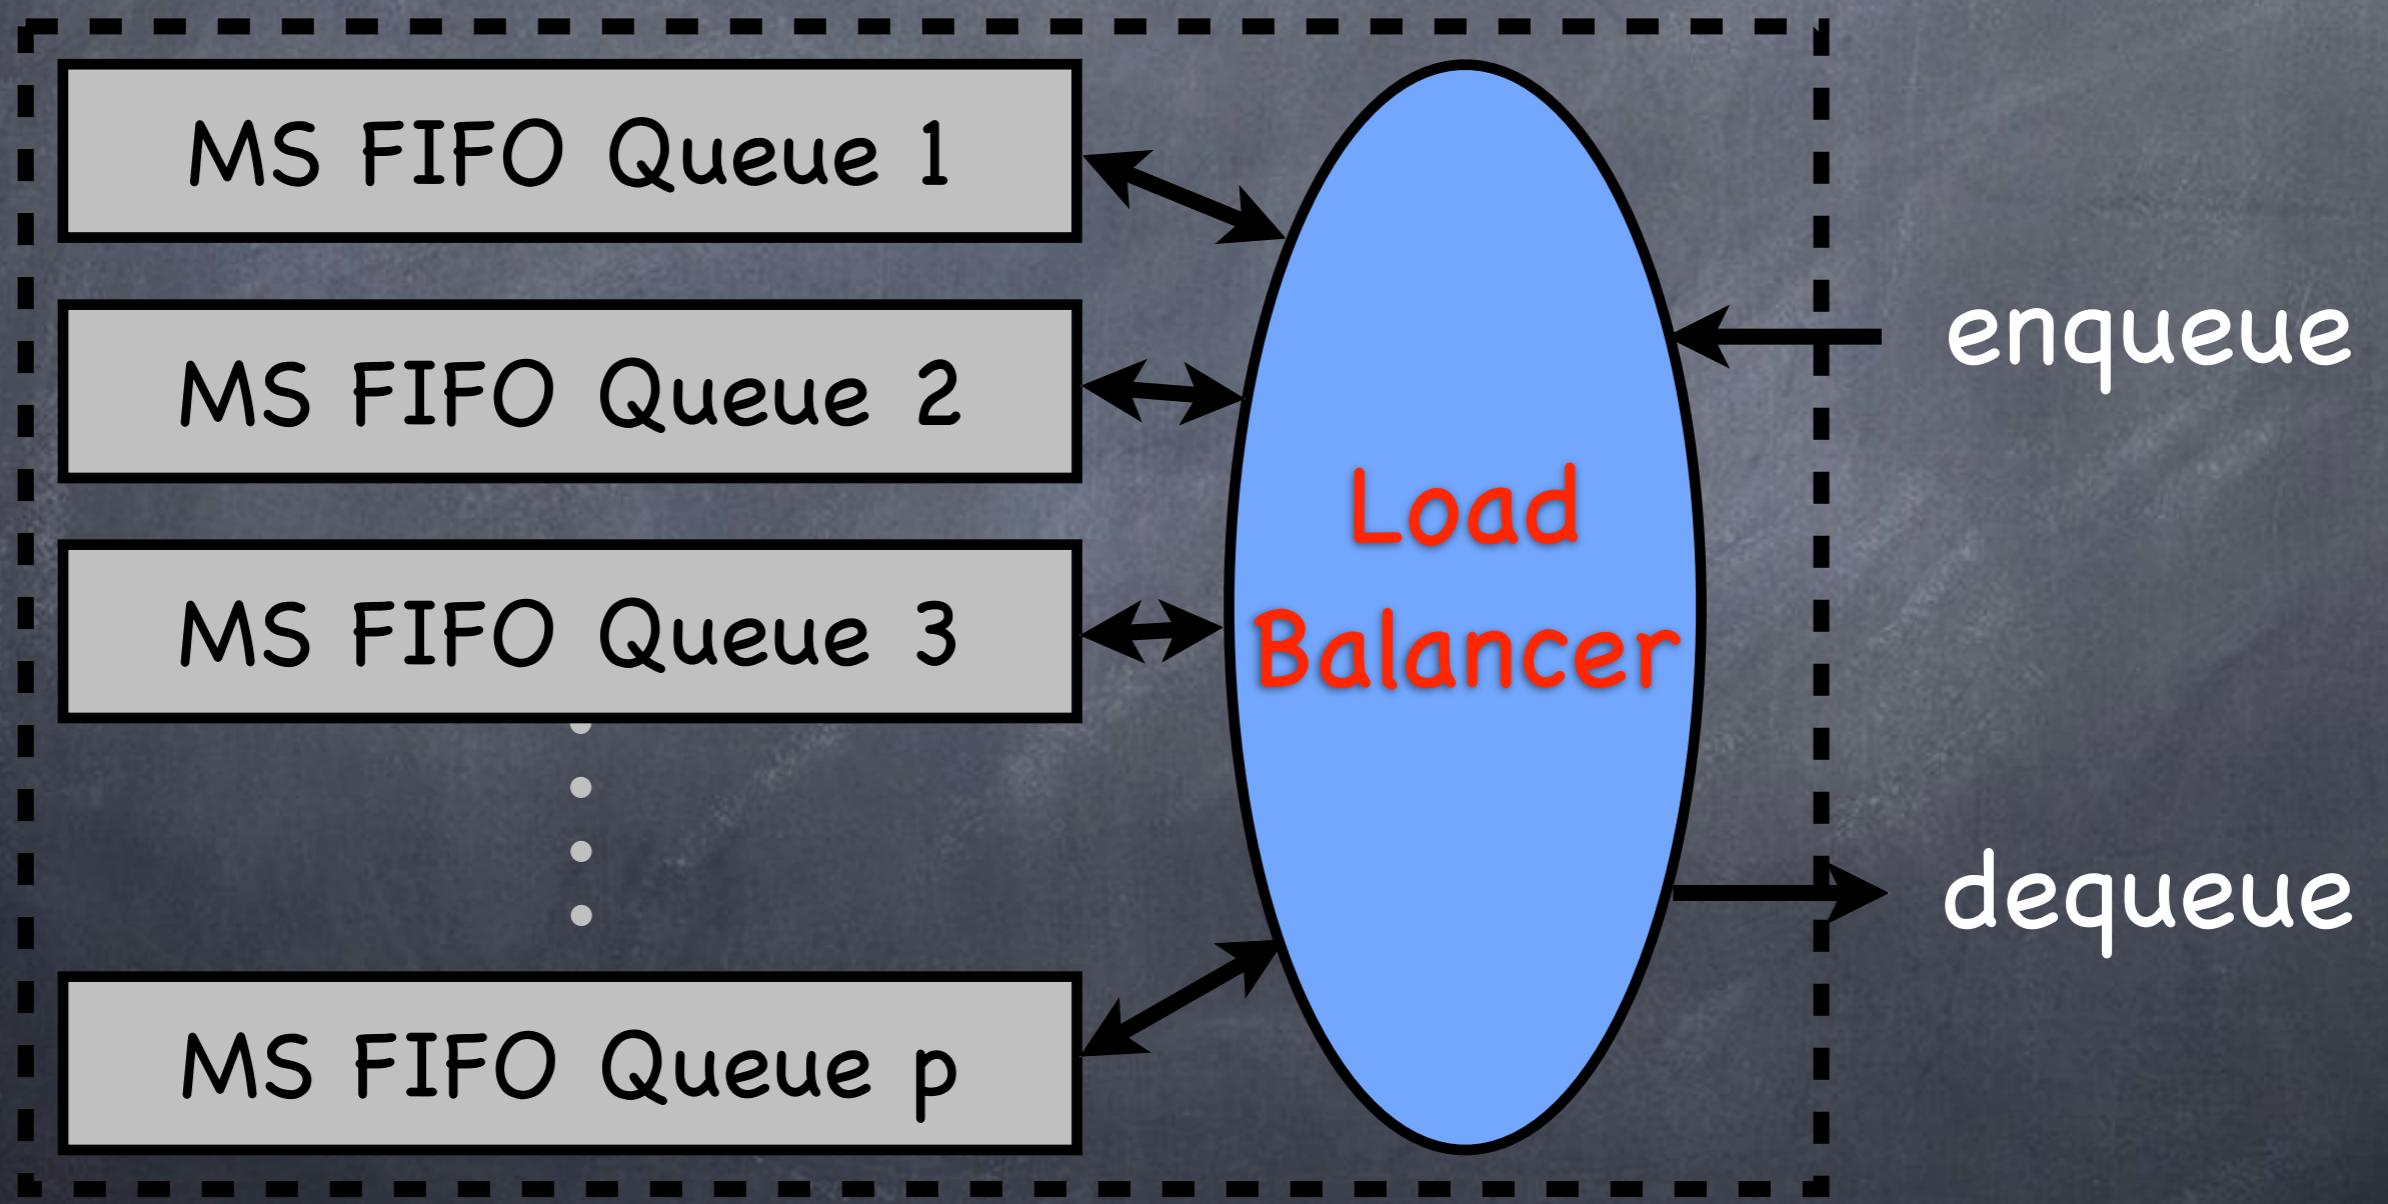




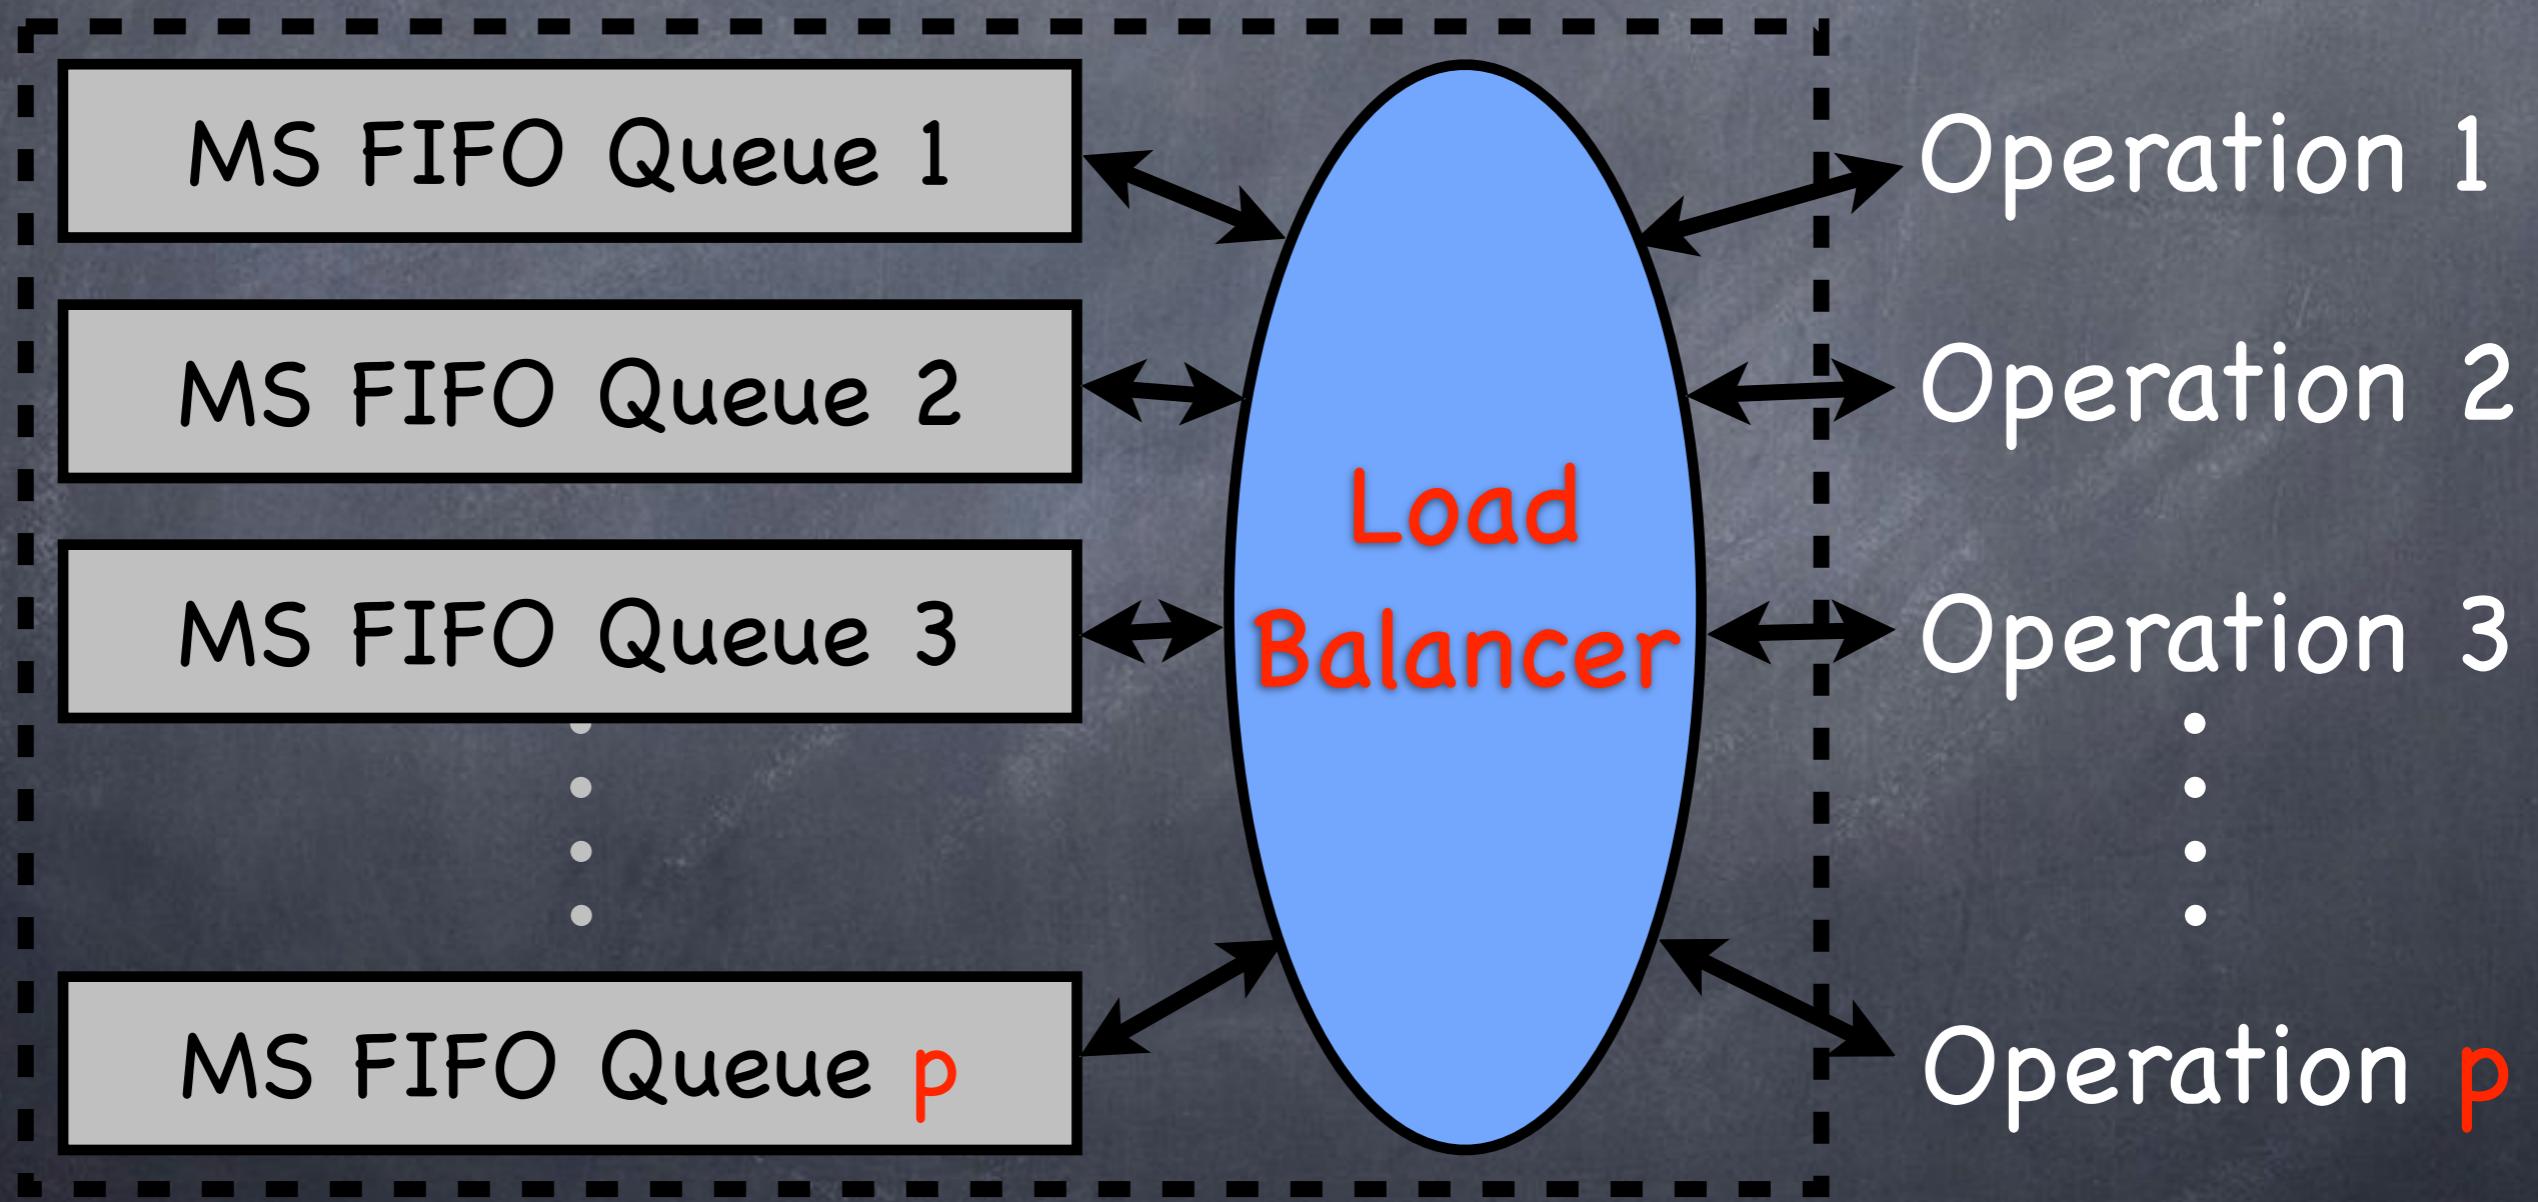
(a) High contention producer-consumer microbenchmark ($c = 250$)

Distributed Queues

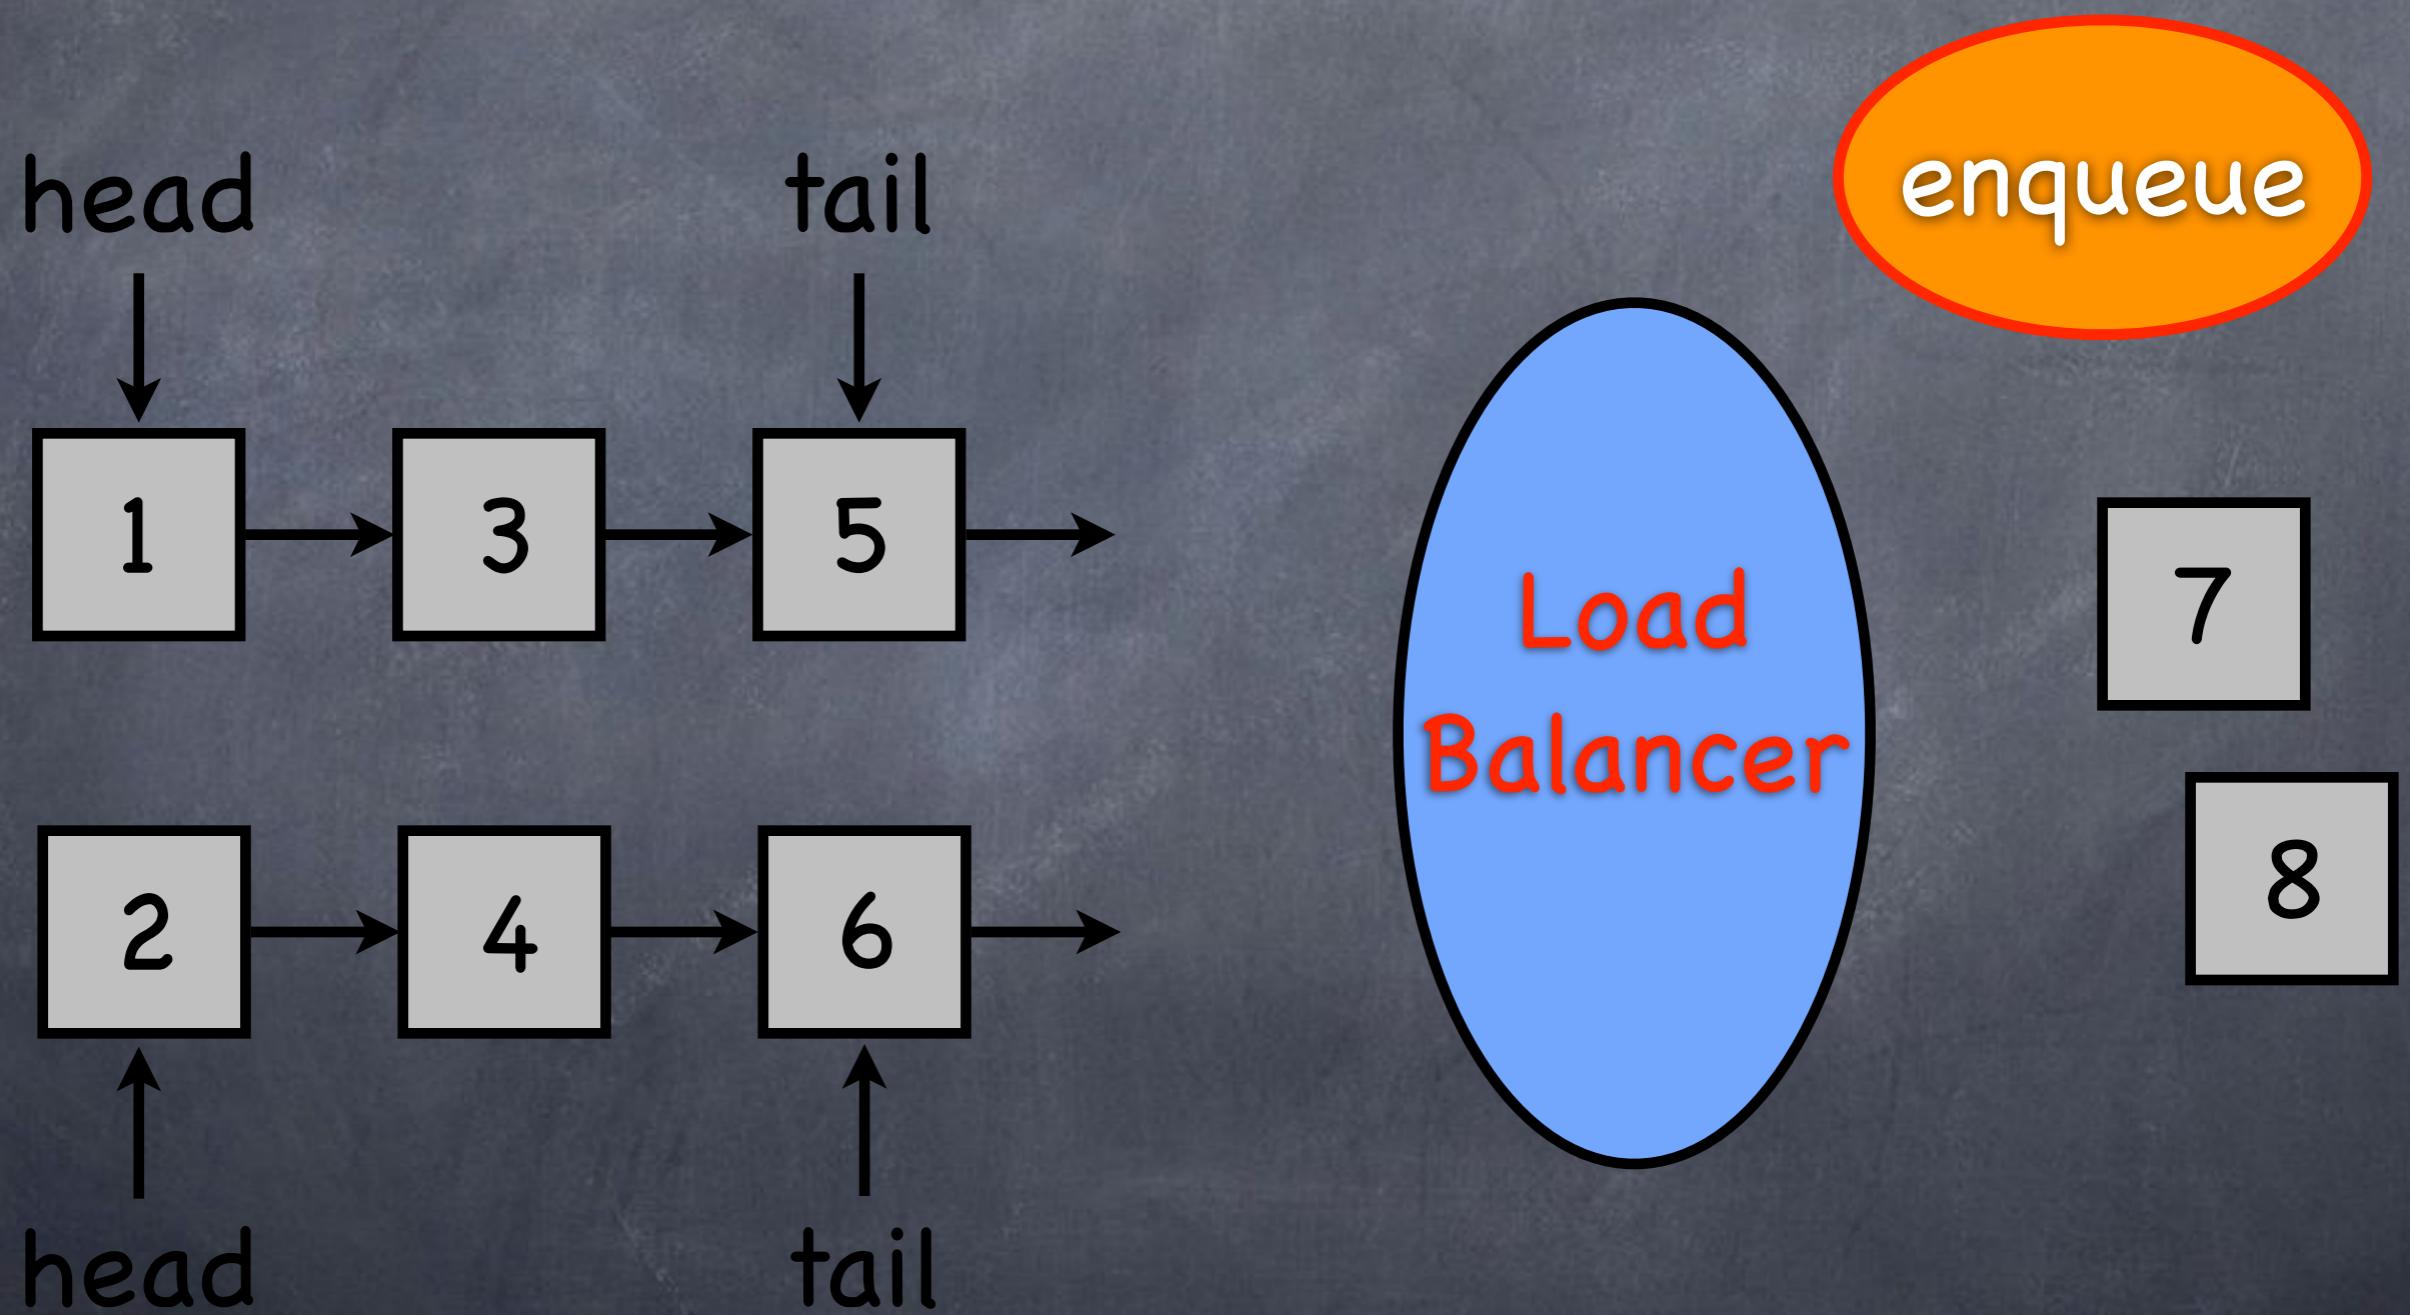
[PODC BA 2011, ICA3PP 2012, Submitted 2013]



Up to p Parallel Enqueues and p Parallel Dequeues



Parallel Access



Emptiness Check?

->

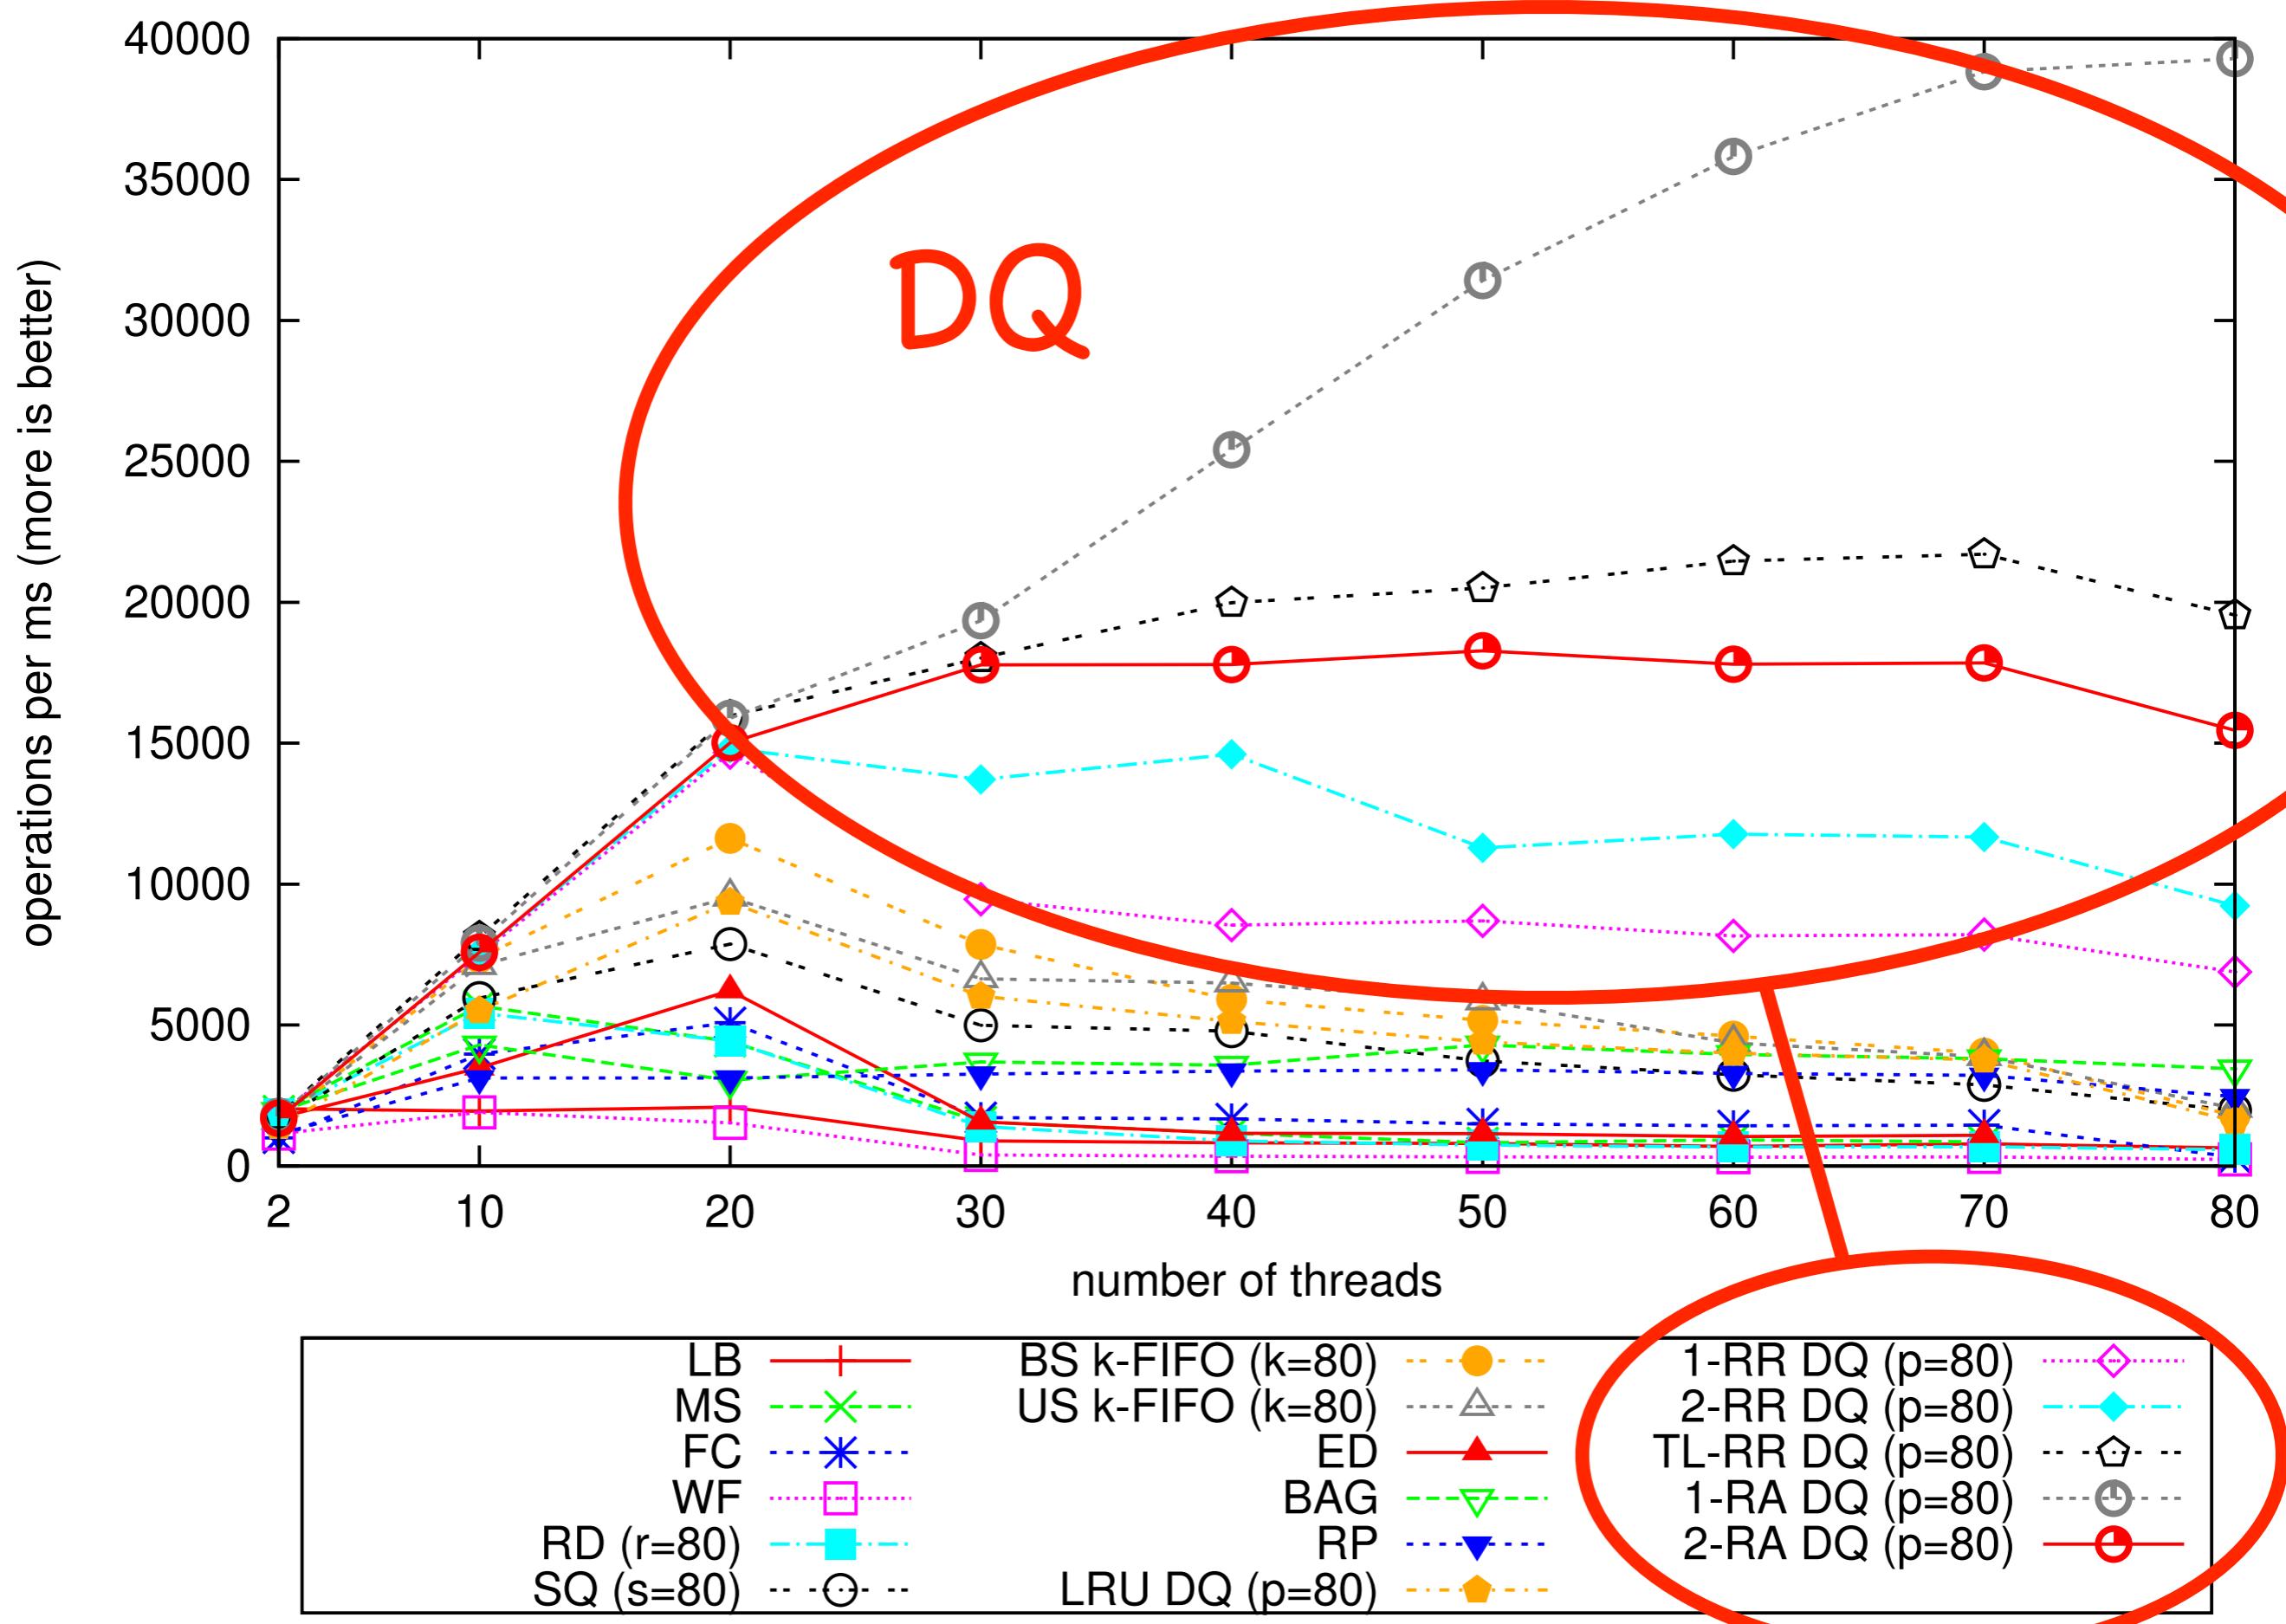
Not Relaxed!

Listing 1: Lock-free load-balanced distributed queue algorithm

```
1 enqueue (element) :  
2     index = load_balancer ();  
3     DQ [index] .MS_enqueue (element);  
4  
5 dequeue () :  
6     start = load_balancer ();  
7     while true:  
8         for i in 0 to p-1:  
9             index = (start + i) % p;  
10            element, current_tail = DQ [index] .MS_dequeue ();  
11            if element != null:  
12                return element;  
13            else:  
14                tail_old [index] = current_tail;  
15            for i in 0 to p-1:  
16                if get_tail (DQ [i]) != tail_old [i]:  
17                    start = i;  
18                    break;  
19                if i == p-1:  
20                    return null;
```

DQ[p]: array of MS queues

tail_old[p]: array of MS tails



(a) High contention producer-consumer microbenchmark ($c = 250$)

Semantics

[Related Work]

Our Stuff

Pools

k-FIFO ($k \geq 0$)

1-RA DQ
2-RA DQ

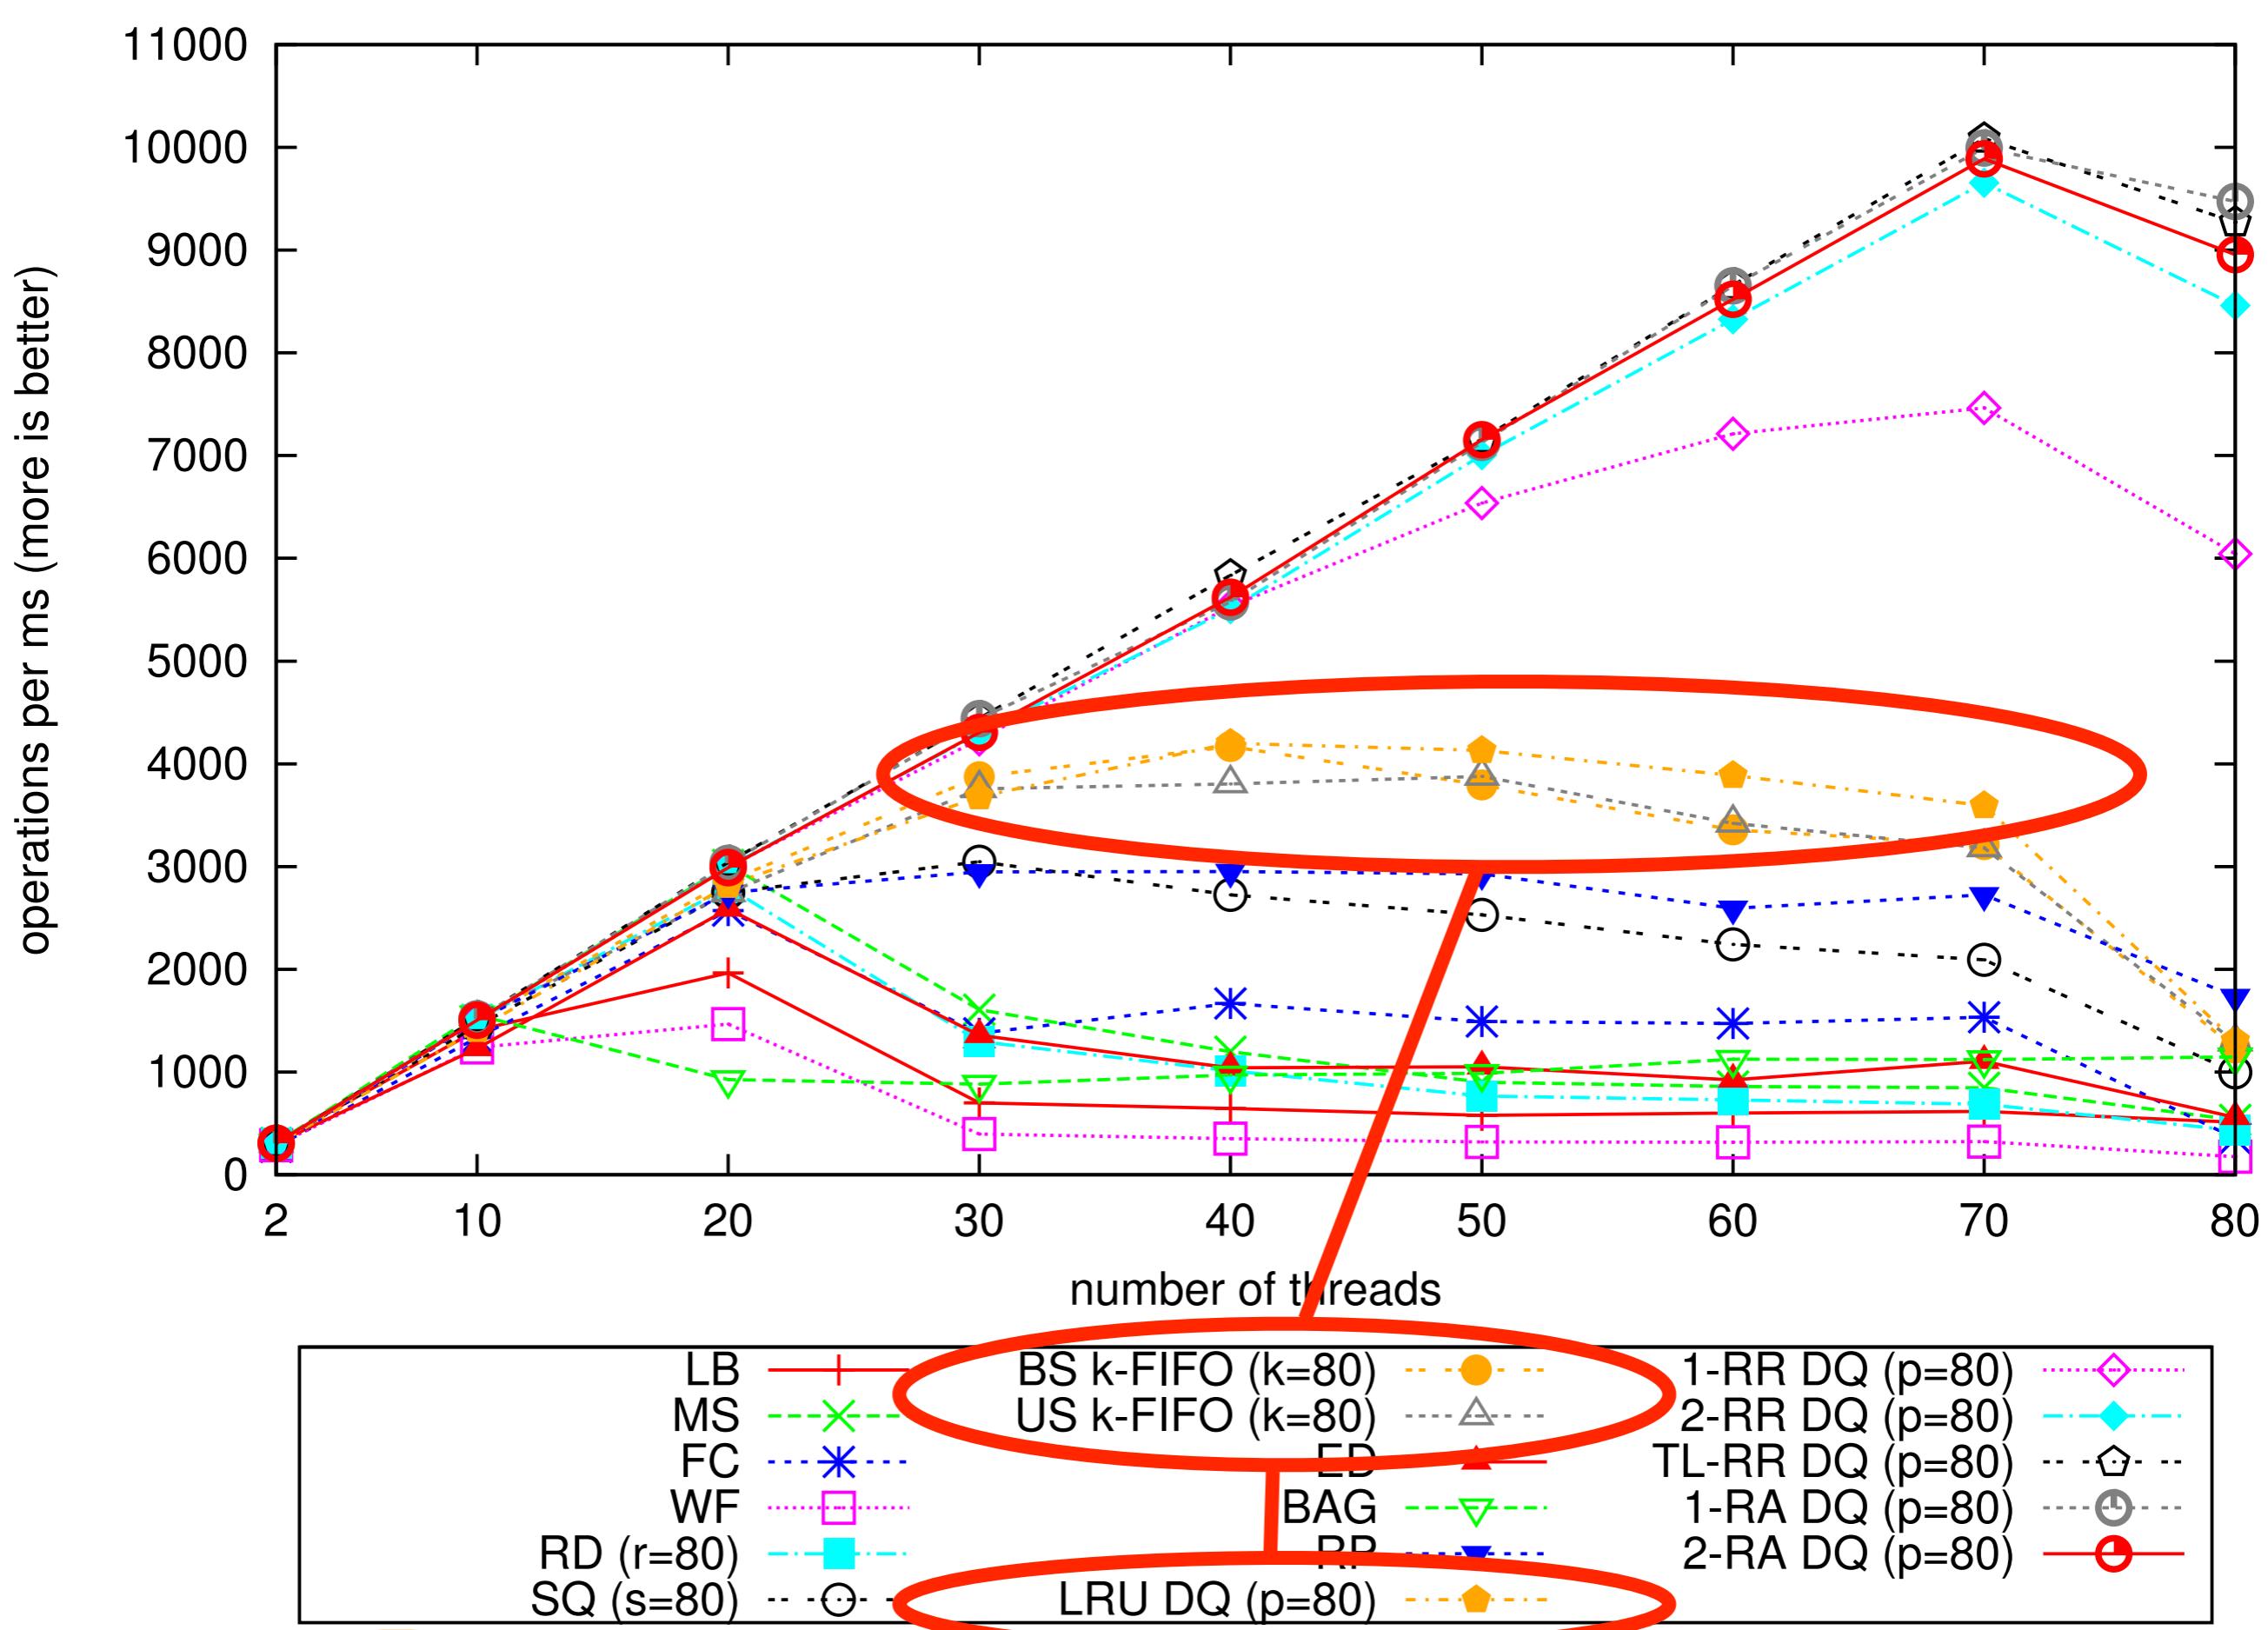
TL-RR DQ
2-RR DQ
1-RR DQ

ED

BAG

RP

[Sundell et al.'11]
[Afek et al.'11, '10]



(b) Low contention producer-consumer microbenchmark ($c = 2000$)

Listing 2: Lock-free LRU distributed queue algorithm

```
1 enqueue(element):
2     start = random();
3     while true:
4         aba_index, aba_count = lowest_aba_tail(start);
5         for i in 0 to p-1:
6             index = (aba_index + i) % p;
7             current_tail = get_tail(DQ[index]);
8             if current_tail.aba == aba_count &&
9                 DQ[index].try_MS_enqueue(element, current_tail):
10                return;
11
12 dequeue():
13     start = random();
14     while true:
15         aba_index, aba_count = lowest_aba_head(start);
16         check_emptiness = true;
17         clear(empty_queue);
18         for i in 0 to p-1:
19             index = (aba_index + i) % p;
20             current_head = get_head(DQ[index]);
21             if current_head.aba == aba_count:
22                 element, current_tail =
23                     DQ[index].try_MS_dequeue(current_head);
24                 if element == FAILED:
25                     check_emptiness = false;
26                 else if element == null:
27                     tail_old[index] = current_tail;
28                     empty_queue[index] = true;
29                 else:
30                     return element;
31
32             if check_emptiness && there_is_any(empty_queue):
33                 for i in 0 to p-1:
34                     if empty_queue[i] &&
35                         (get_tail(DQ[i]) != tail_old[i]):
36                         start = i;
37                         break;
38                     if i == p-1:
39                         return null;
```

LRU DQ:

max difference of
tail/head
ABA counters
is one!

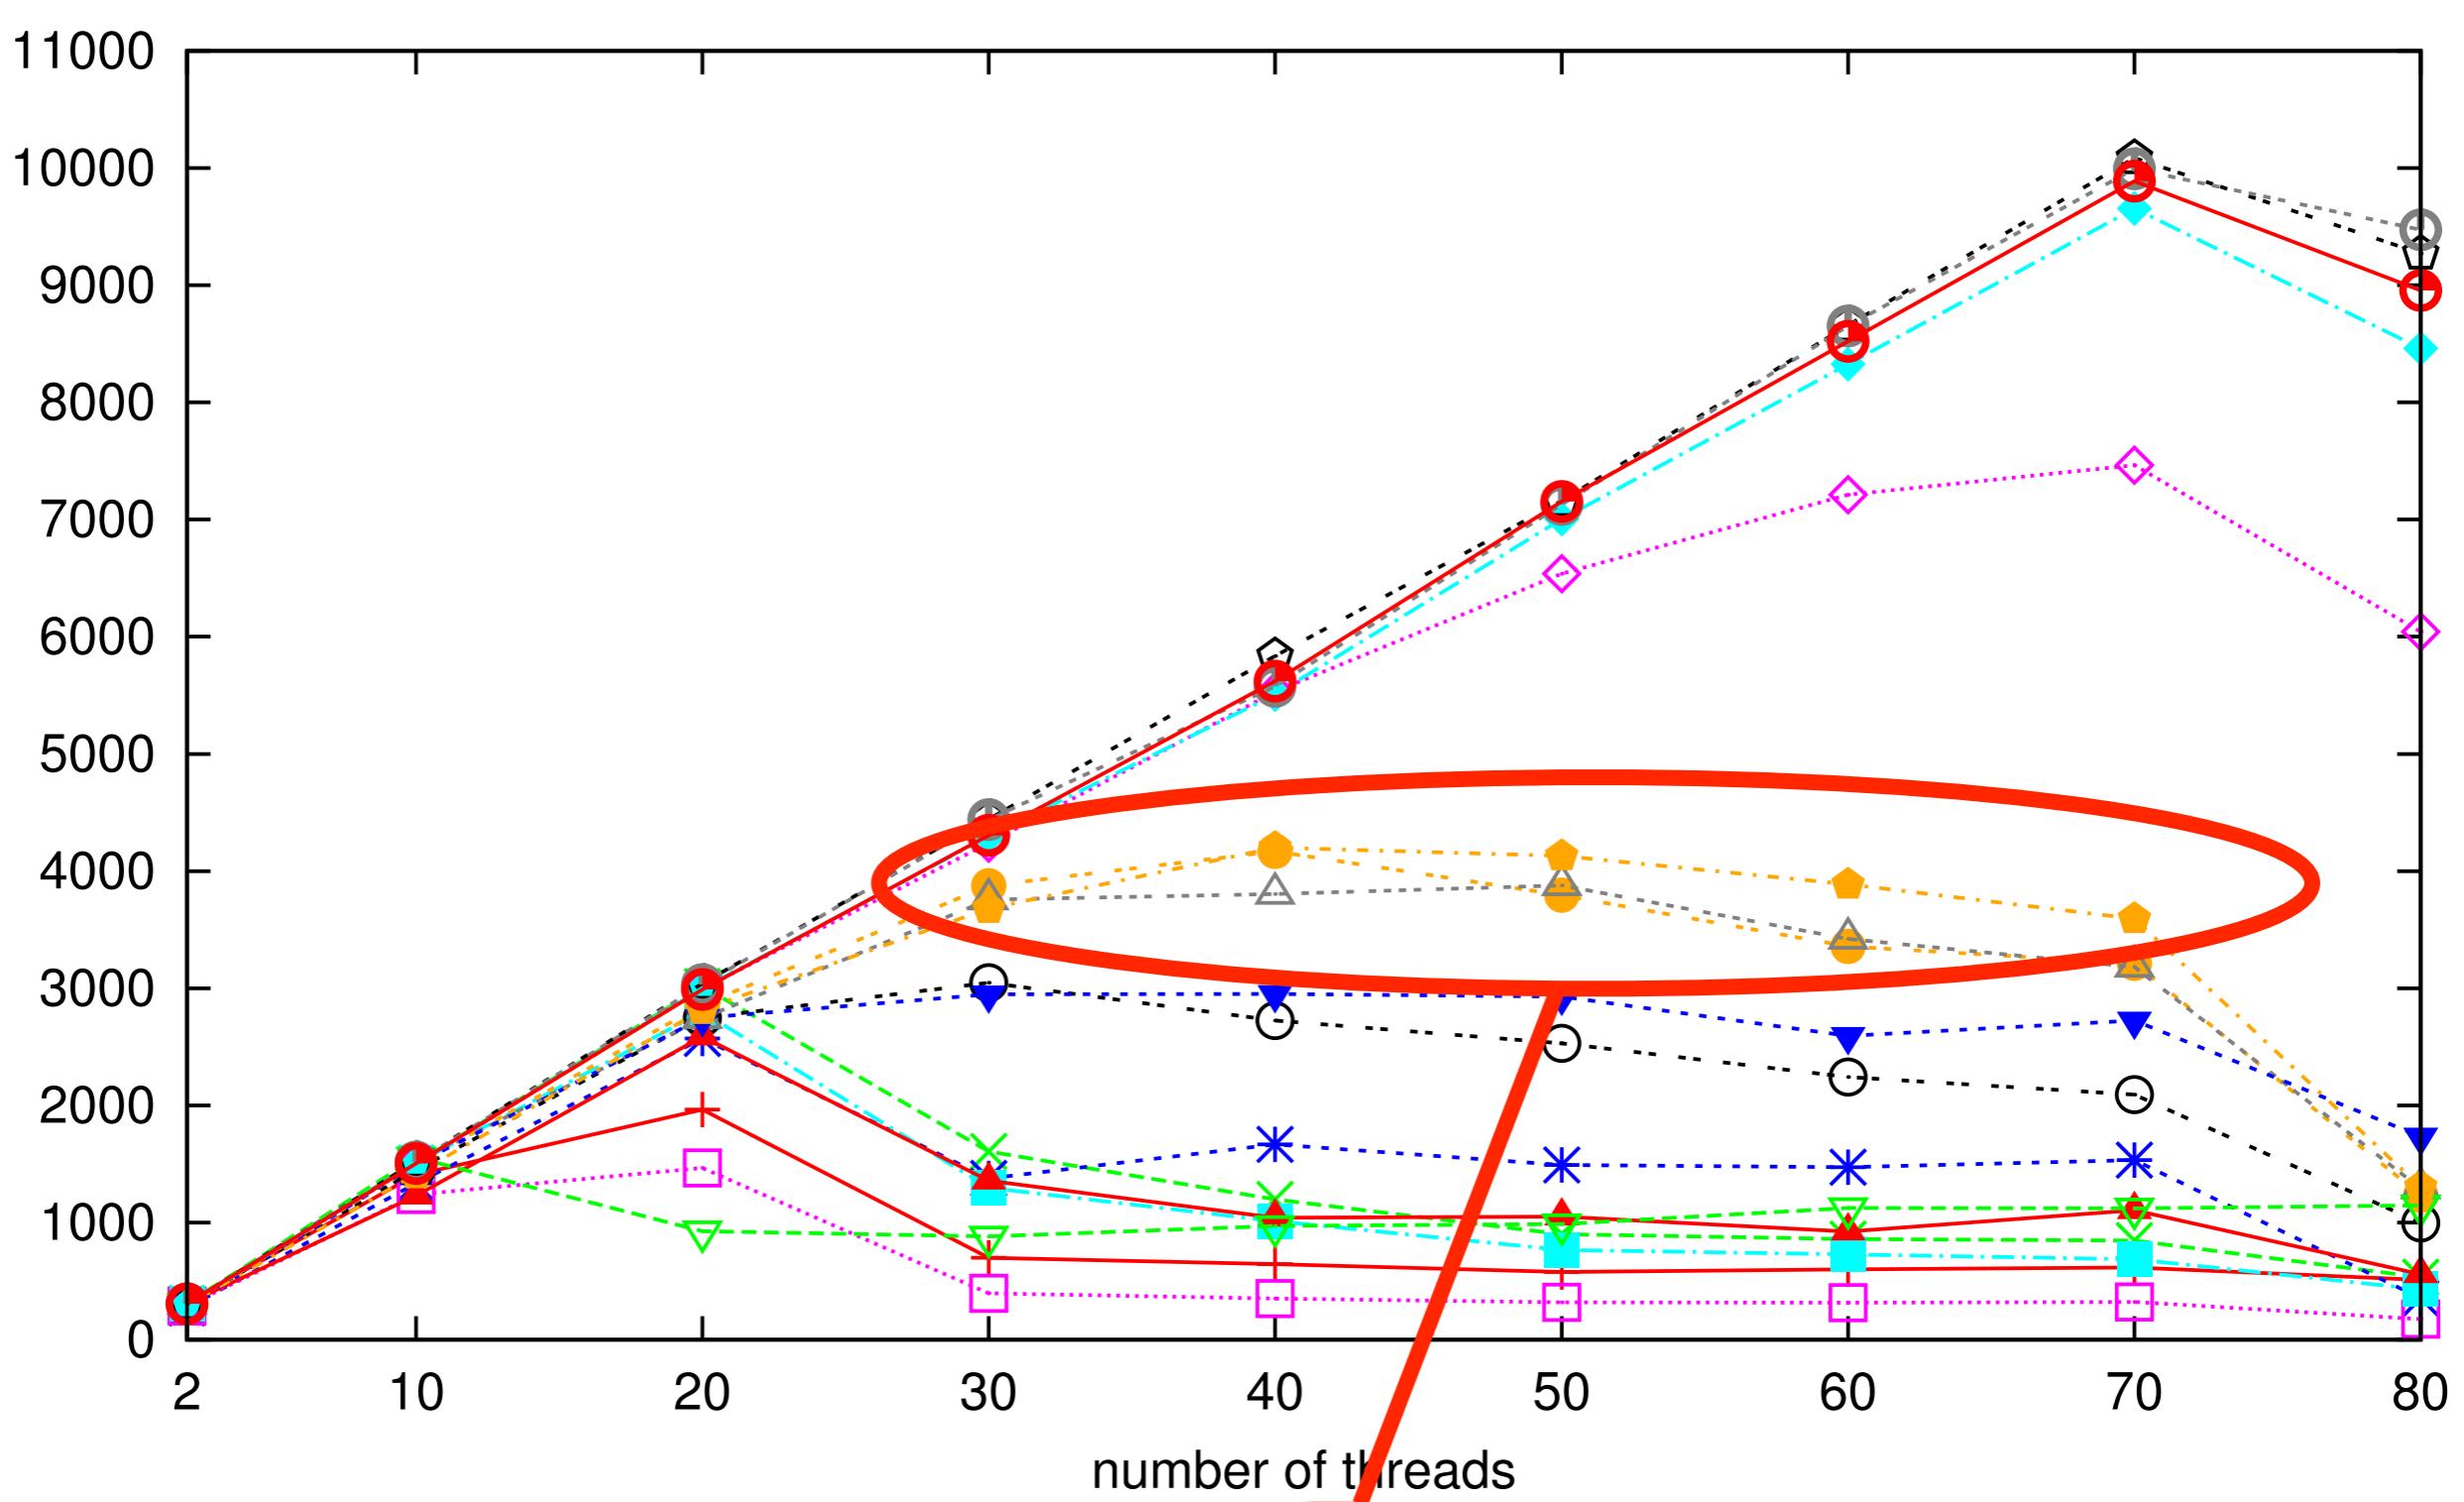


there are two
partitions of MS queues
with lowest/highest
ABA counters

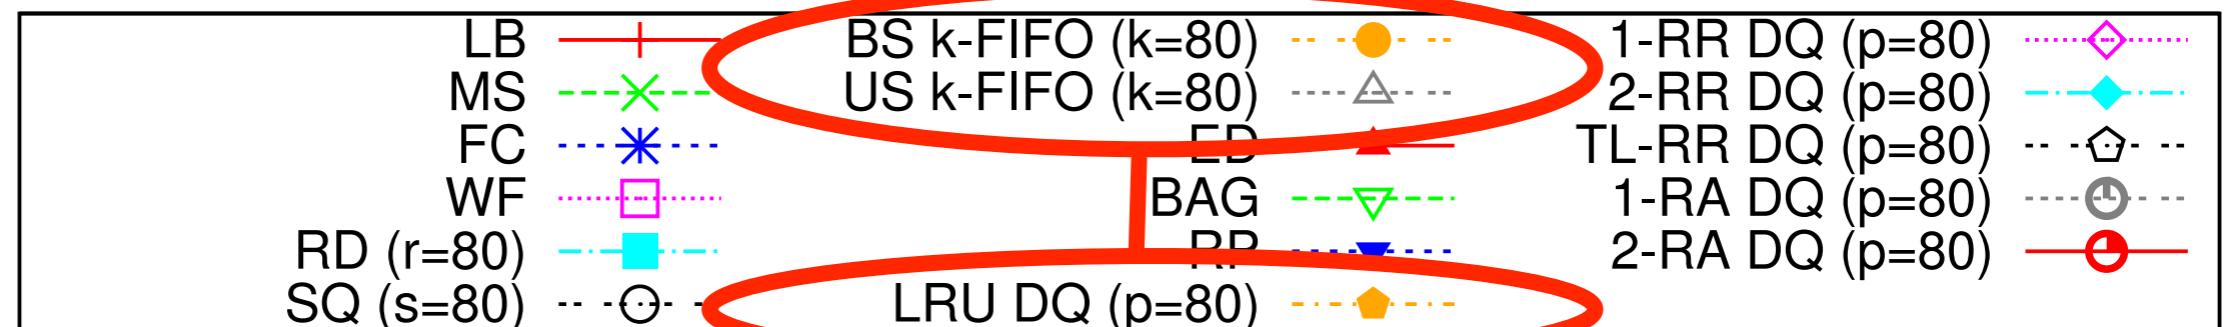


enqueue/dequeue
@one_of_lowest

operations per ms (more is better)



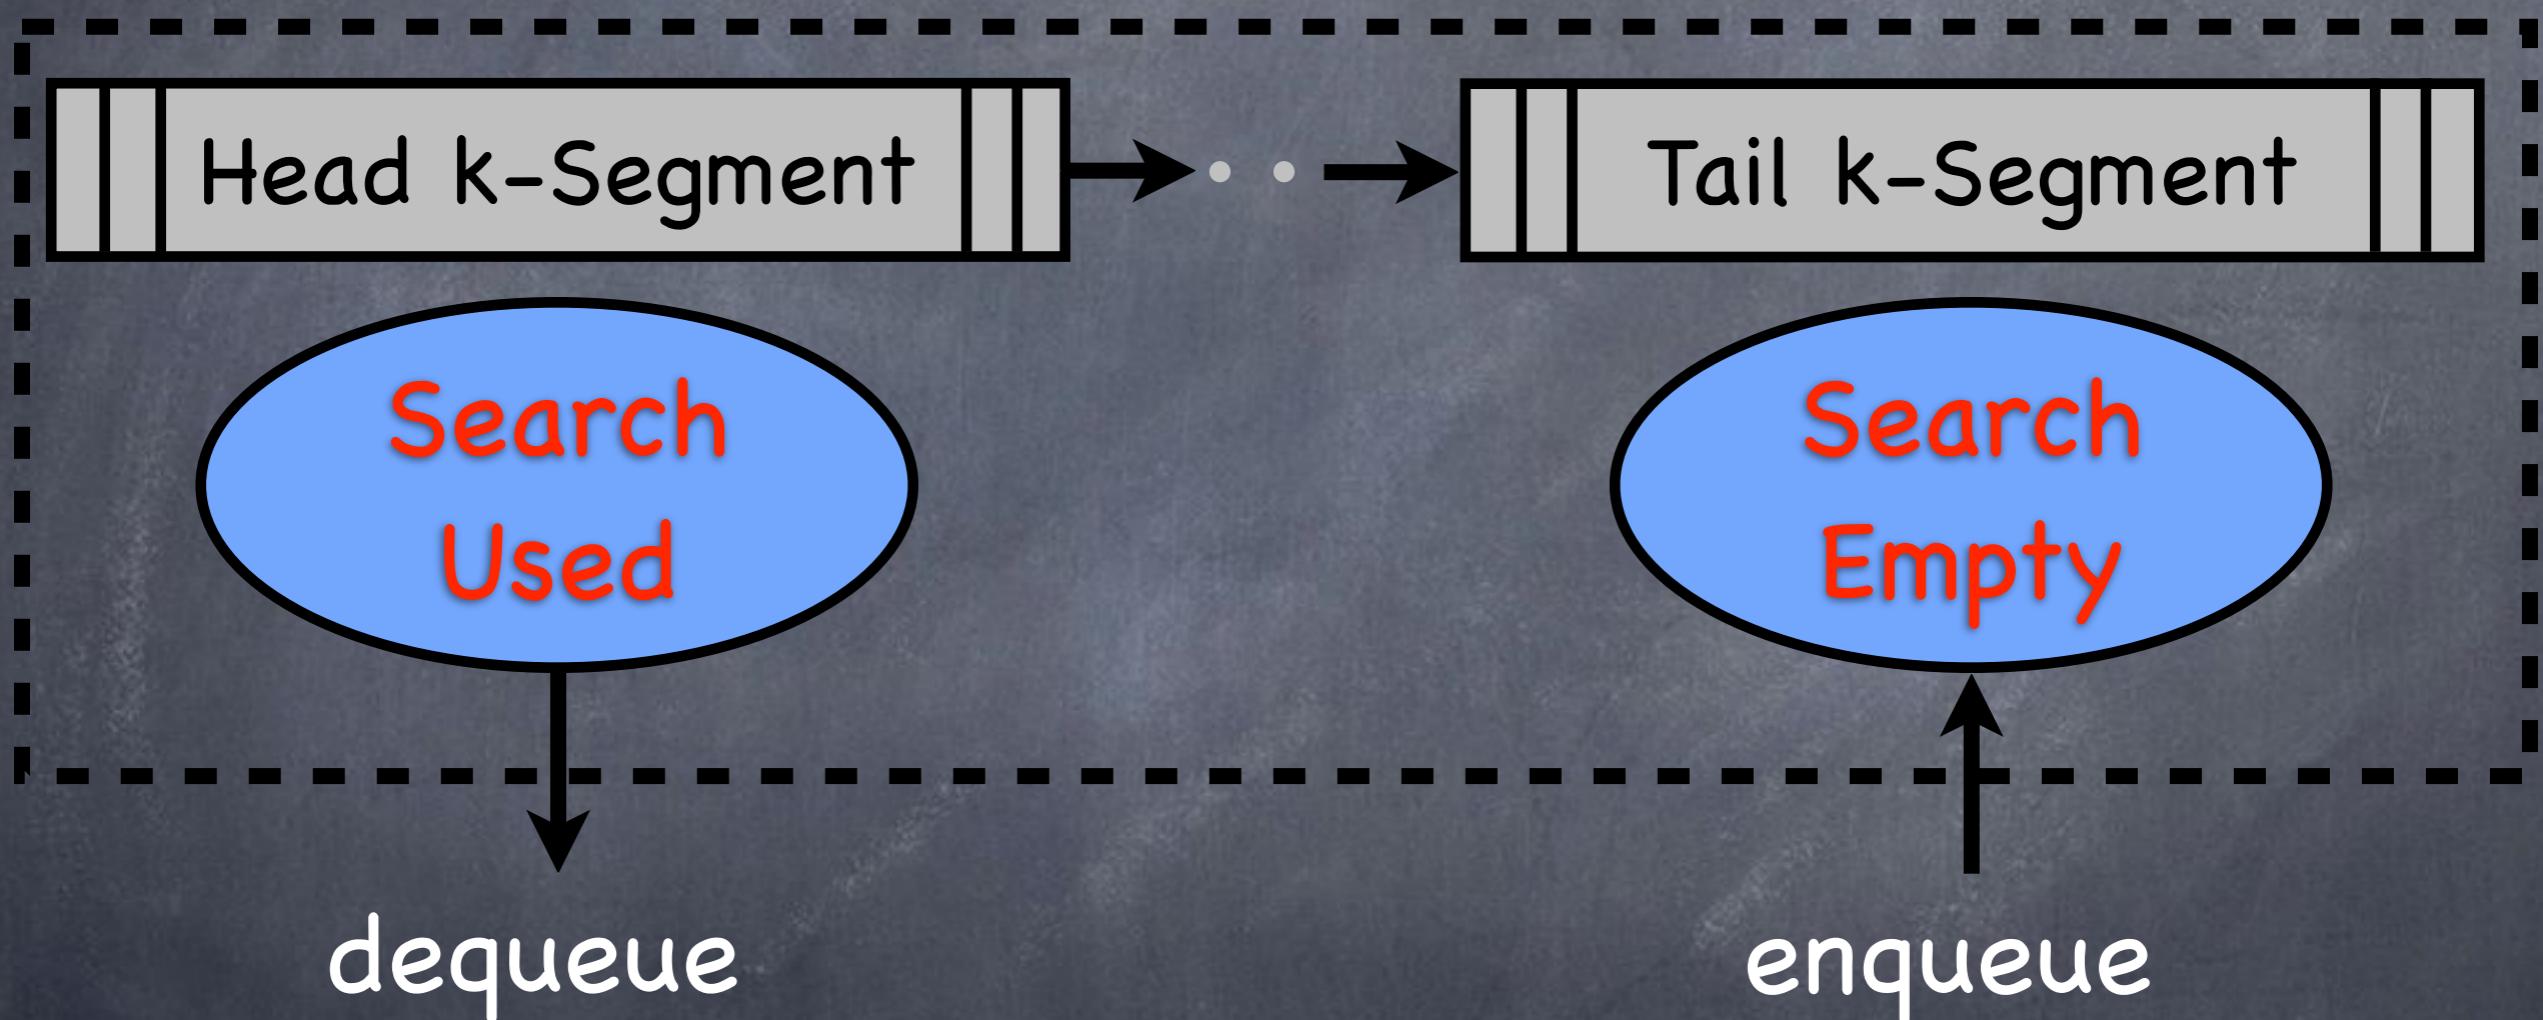
number of threads



(b) Low contention producer-consumer microbenchmark ($c = 2000$)

Segmented Queues (SQ)

[Afek,Korland,Yanovsky 2010]



→ BS, US *k*-FIFO Queues

[Lippautz,Payer 2012]

Emptiness Check?

->

Not Relaxed!

```

1 bool enqueue(item):
2     while true:
3         tail_old = get_tail();
4         head_old = get_head();
5         item_old, index = find_empty_slot(tail_old, k, TESTS);
6         if tail_old == get_tail():
7             if item_old.value == EMPTY:
8                 item_new = atomic_value(item, item_old.counter + 1);
9                 if CAS(&tail_old[index], item_old, item_new):
10                     if committed(tail_old, item_new, index):
11                         return true;
12             else:
13                 if queue_full(head_old, tail_old):
14                     if segment_not_empty(head_old, k) && head == get_head():
15                         return false;
16                     advance_head(head_old, k);
17                     advance_tail(tail_old, k);
18
19 bool committed(tail_old, item_new, index):
20     if tail_old[index] != item_new:
21         return true;
22     head_current = get_head();
23     tail_current = get_tail();
24     item_empty = atomic_value(EMPTY, item_new.counter + 1);
25     if in_queue_after_head(tail_old, tail_current, head_current):
26         return true;
27     else if not_in_queue(tail_old, tail_current, head_current):
28         if !CAS(&tail_old[index], item_new, item_empty):
29             return true;
30     else: //in queue at head
31         head_new = atomic_value(head_current.value, head_current.counter + 1);
32         if CAS(&head, head_current, head_new):
33             return true;
34         if !CAS(&tail_old[index], item_new, item_empty):
35             return true;
36     return false;

```

enqueue

dequeue

```
38 item dequeue():
39     while true:
40         tail_old = get_tail();
41         head_old = get_head();
42         item_old, index = find_item(head_old, k);
43         if head_old == head:
44             if item_old.value != EMPTY:
45                 if head_old.value == tail_old.value:
46                     advance_tail(tail_old, k);
47                     item_empty = atomic_value(EMPTY, item_old.counter + 1);
48                     if CAS(&head_old[index], item_old, item_empty):
49                         return item_old.value;
50             else:
51                 if head_old.value == tail_old.value && tail_old == get_tail():
52                     return null;
53                     advance_head(head_old, k);
```

Semantics

[Related Work]

Our Stuff

[Afek et al.'10]

Pools

k-FIFO ($k \geq 0$)

TL-RR DQ
2-RR DQ
1-RR DQ

(SQ)
RD

FIFO

LB MS
WF FC

LRU DQ
BS, US

configurable k

1-RA DQ
2-RA DQ

ED
BAG
RP

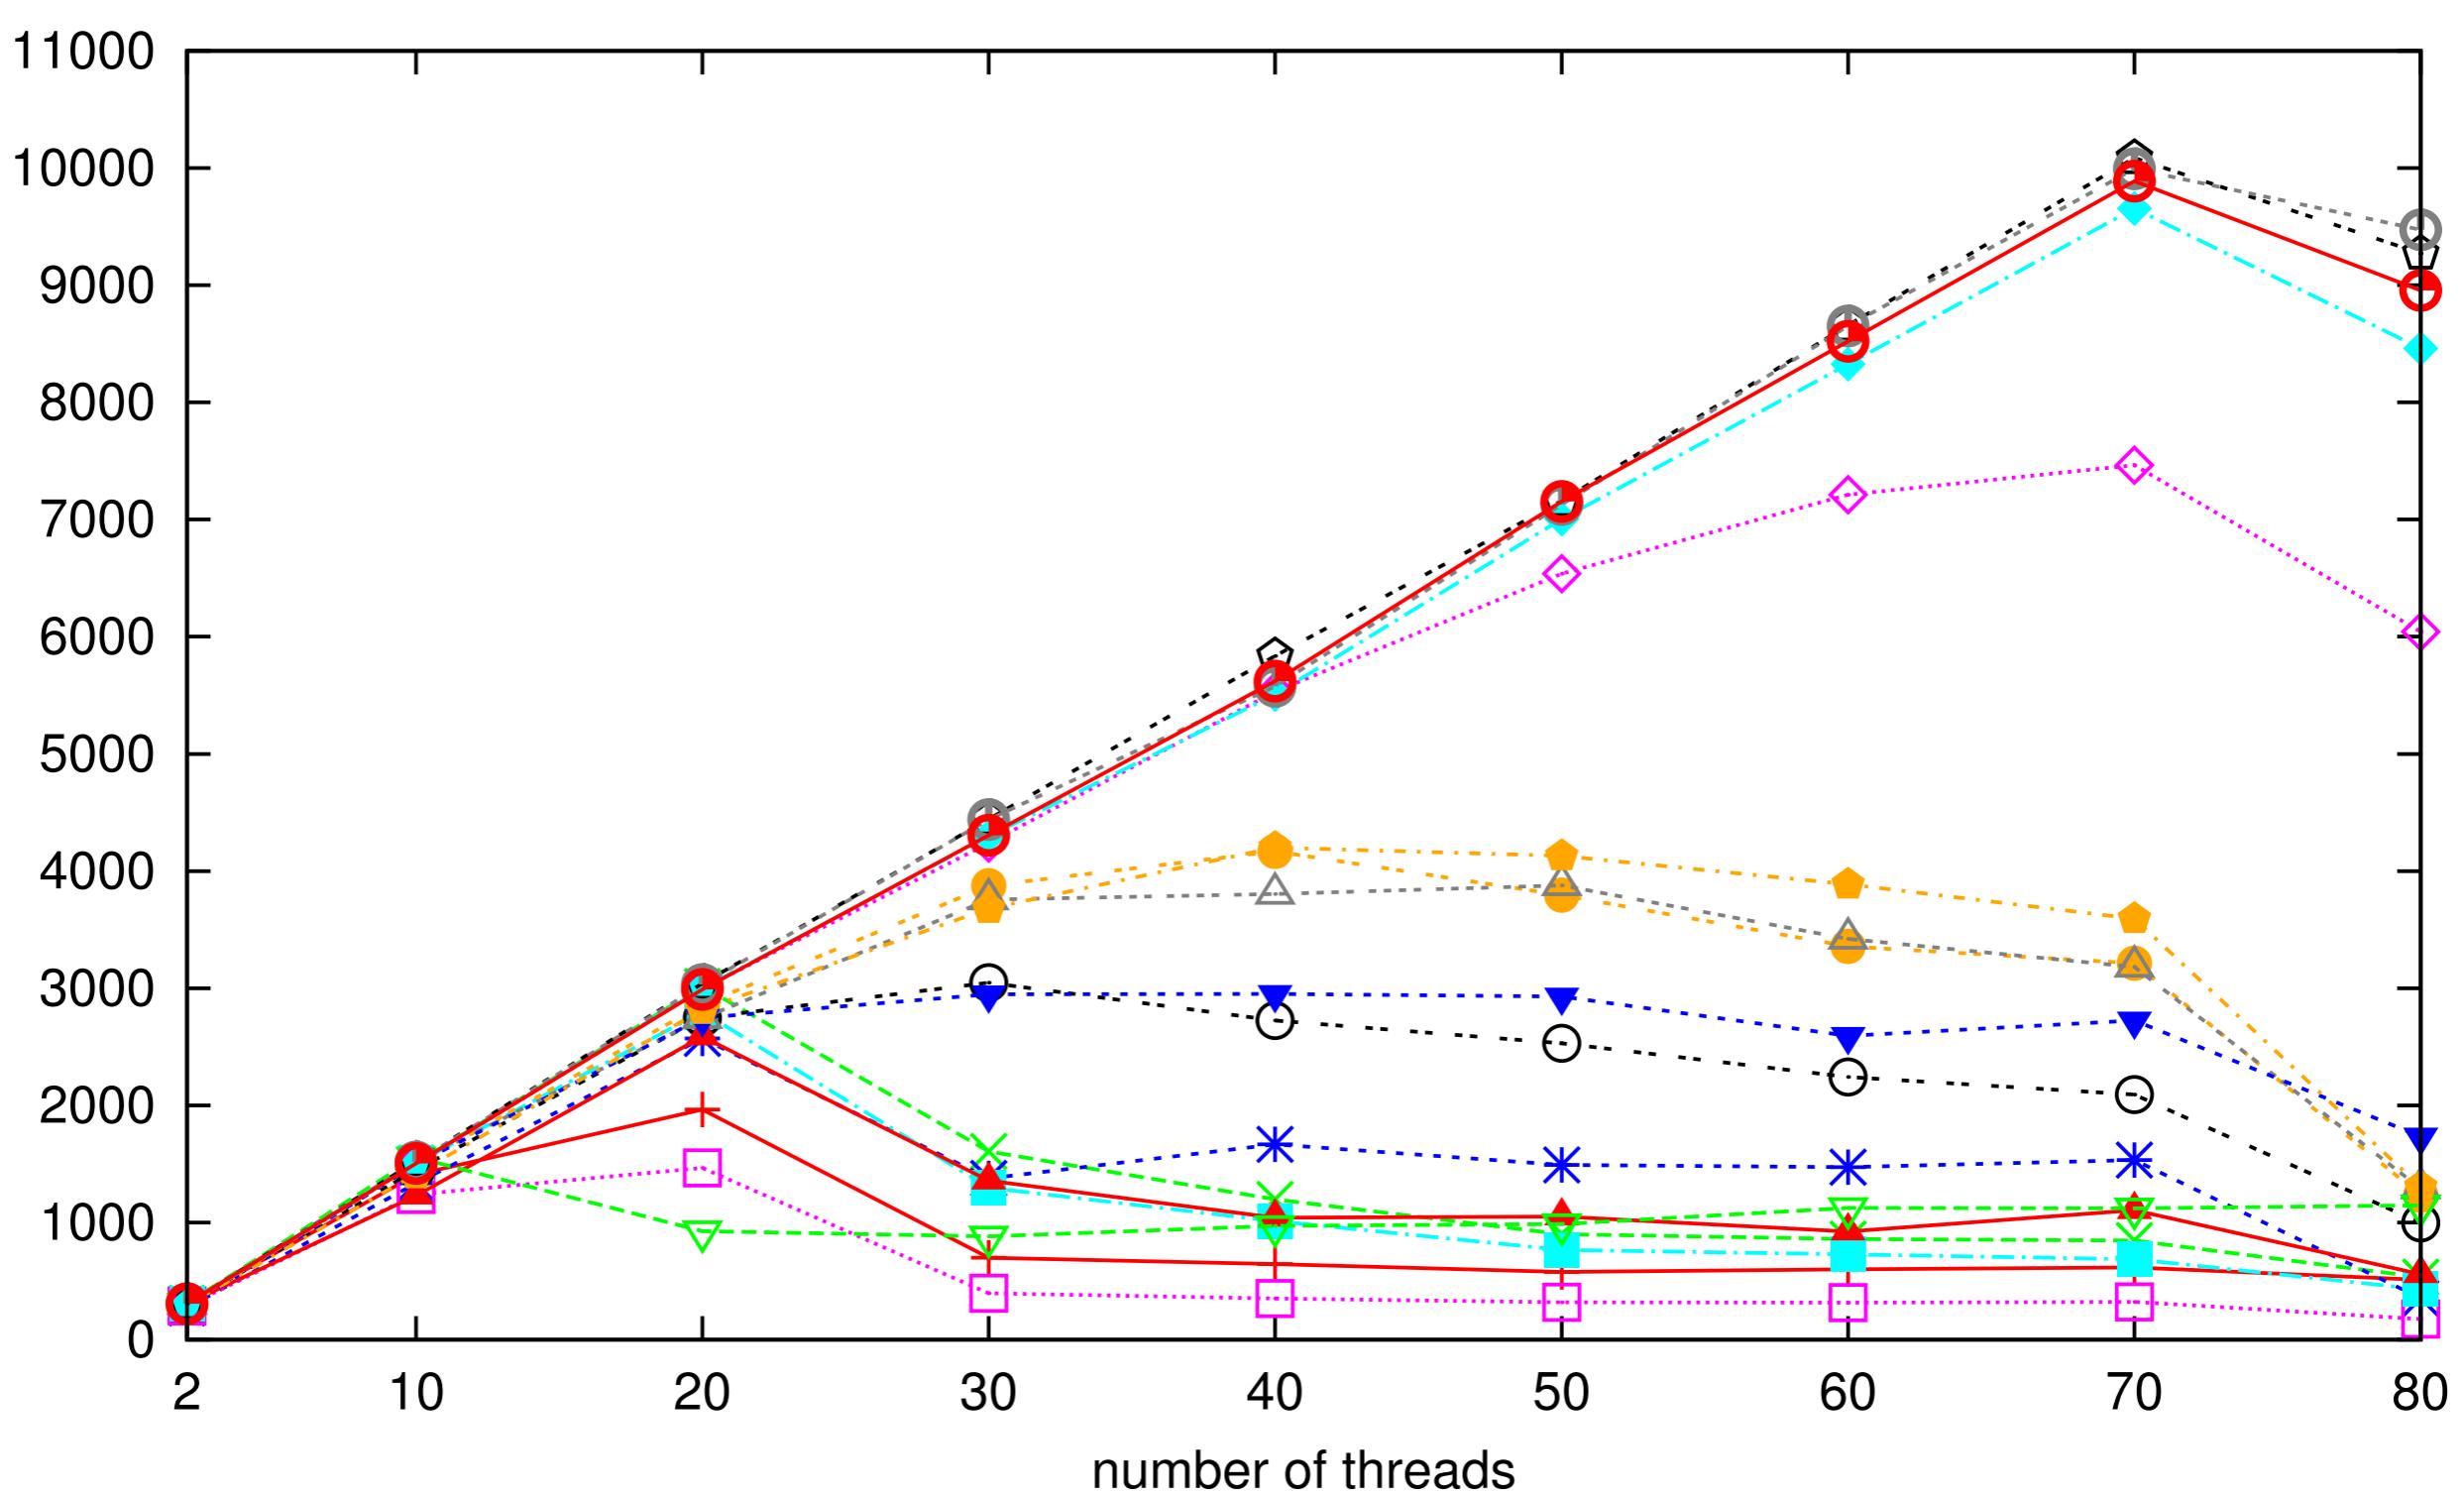
[Sundell et al.'11]

[Afek et al.'11, '10]

[Incze et al.'10]

[Kogan et al.'11]

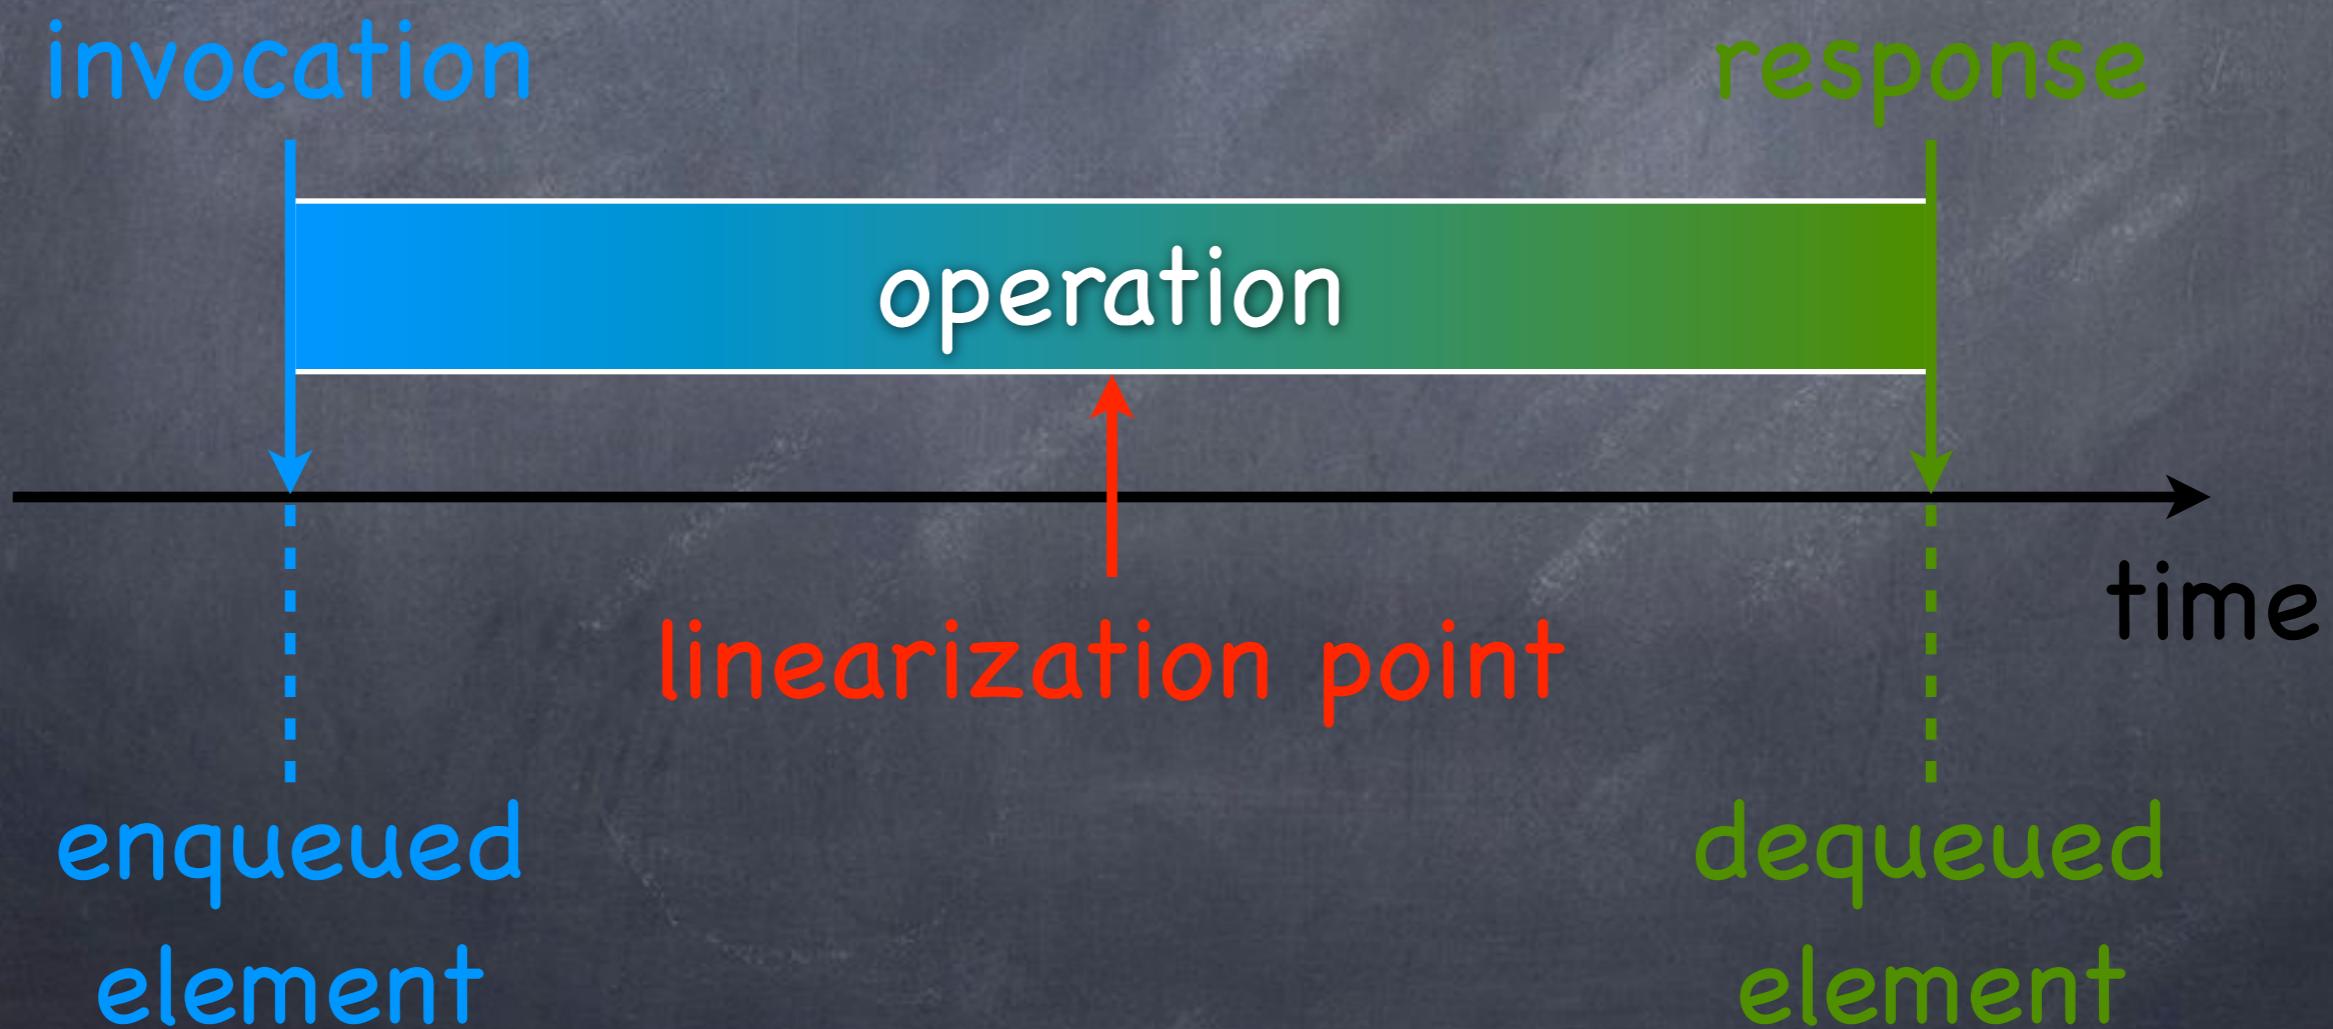
operations per ms (more is better)



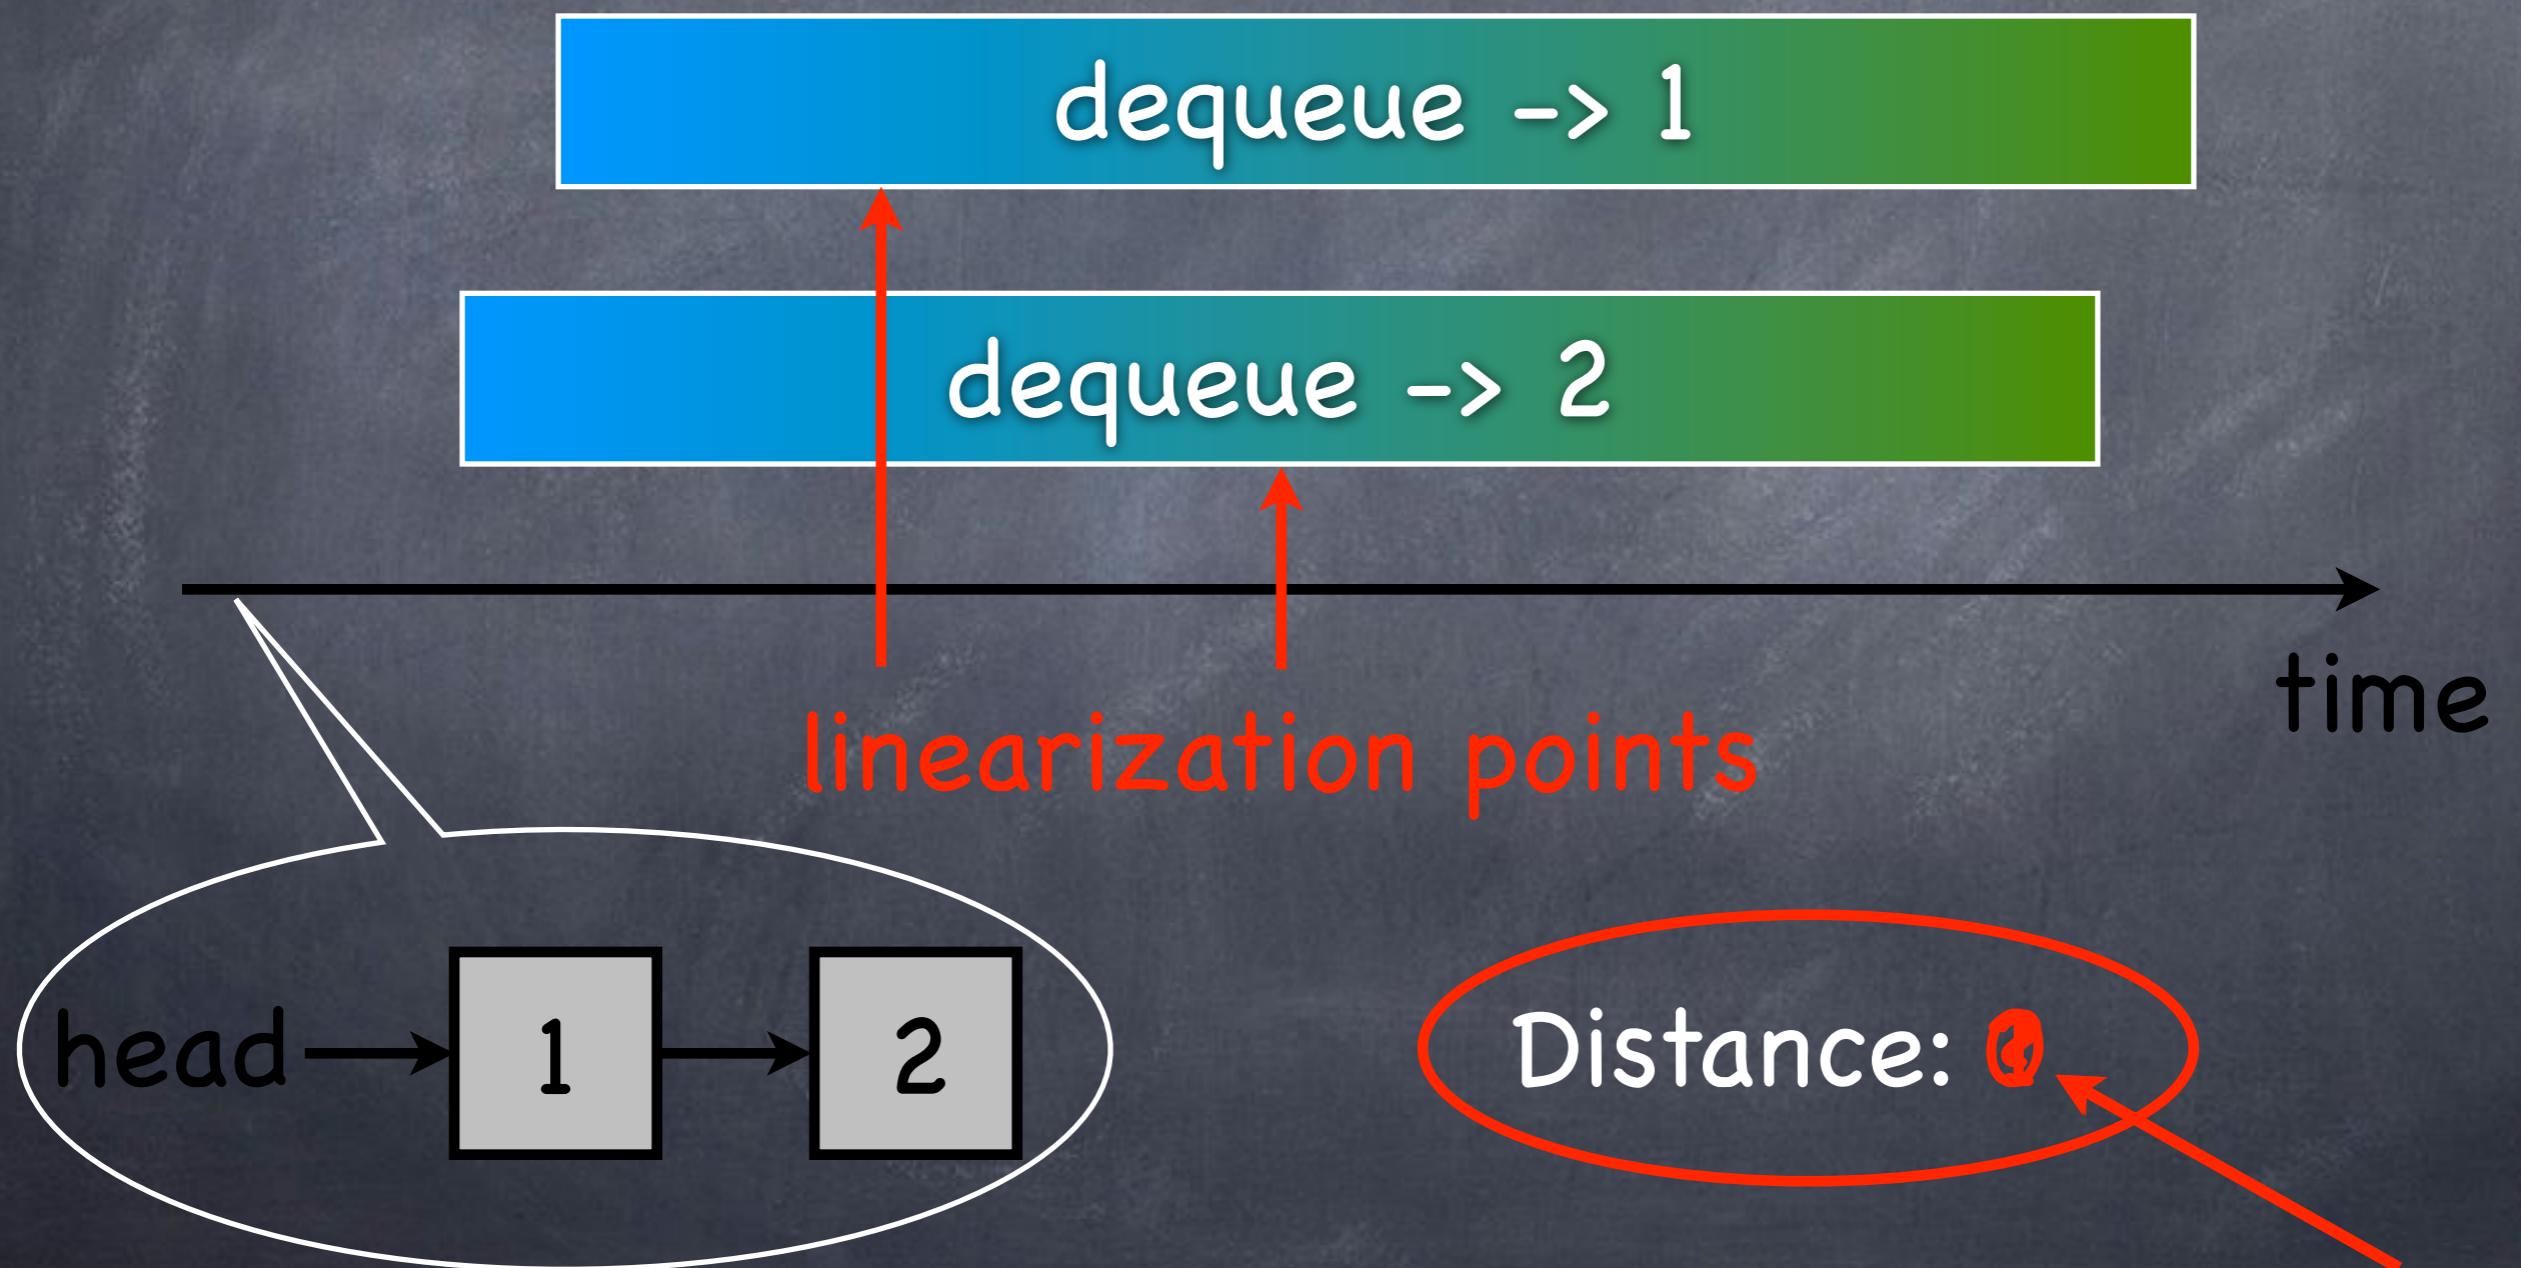
(b) Low contention producer-consumer microbenchmark ($c = 2000$)

(Enhanced) Concurrent History

Sequence of Time-stamped Invocation and Response Events
as well as Time-stamped Linearization Points (Approximative)



Measuring “Out-of-Order Distance”



The Actual-Time Linearization
of a concurrent history

is

the sequence of its operations
ordered by
their linearization points

The Actual-Time Distance
of a concurrent history
is
the average
out-of-order distance
of
its actual-time linearization

Actual-Time Distance
measures
re-ordering
due to
semantical relaxation!

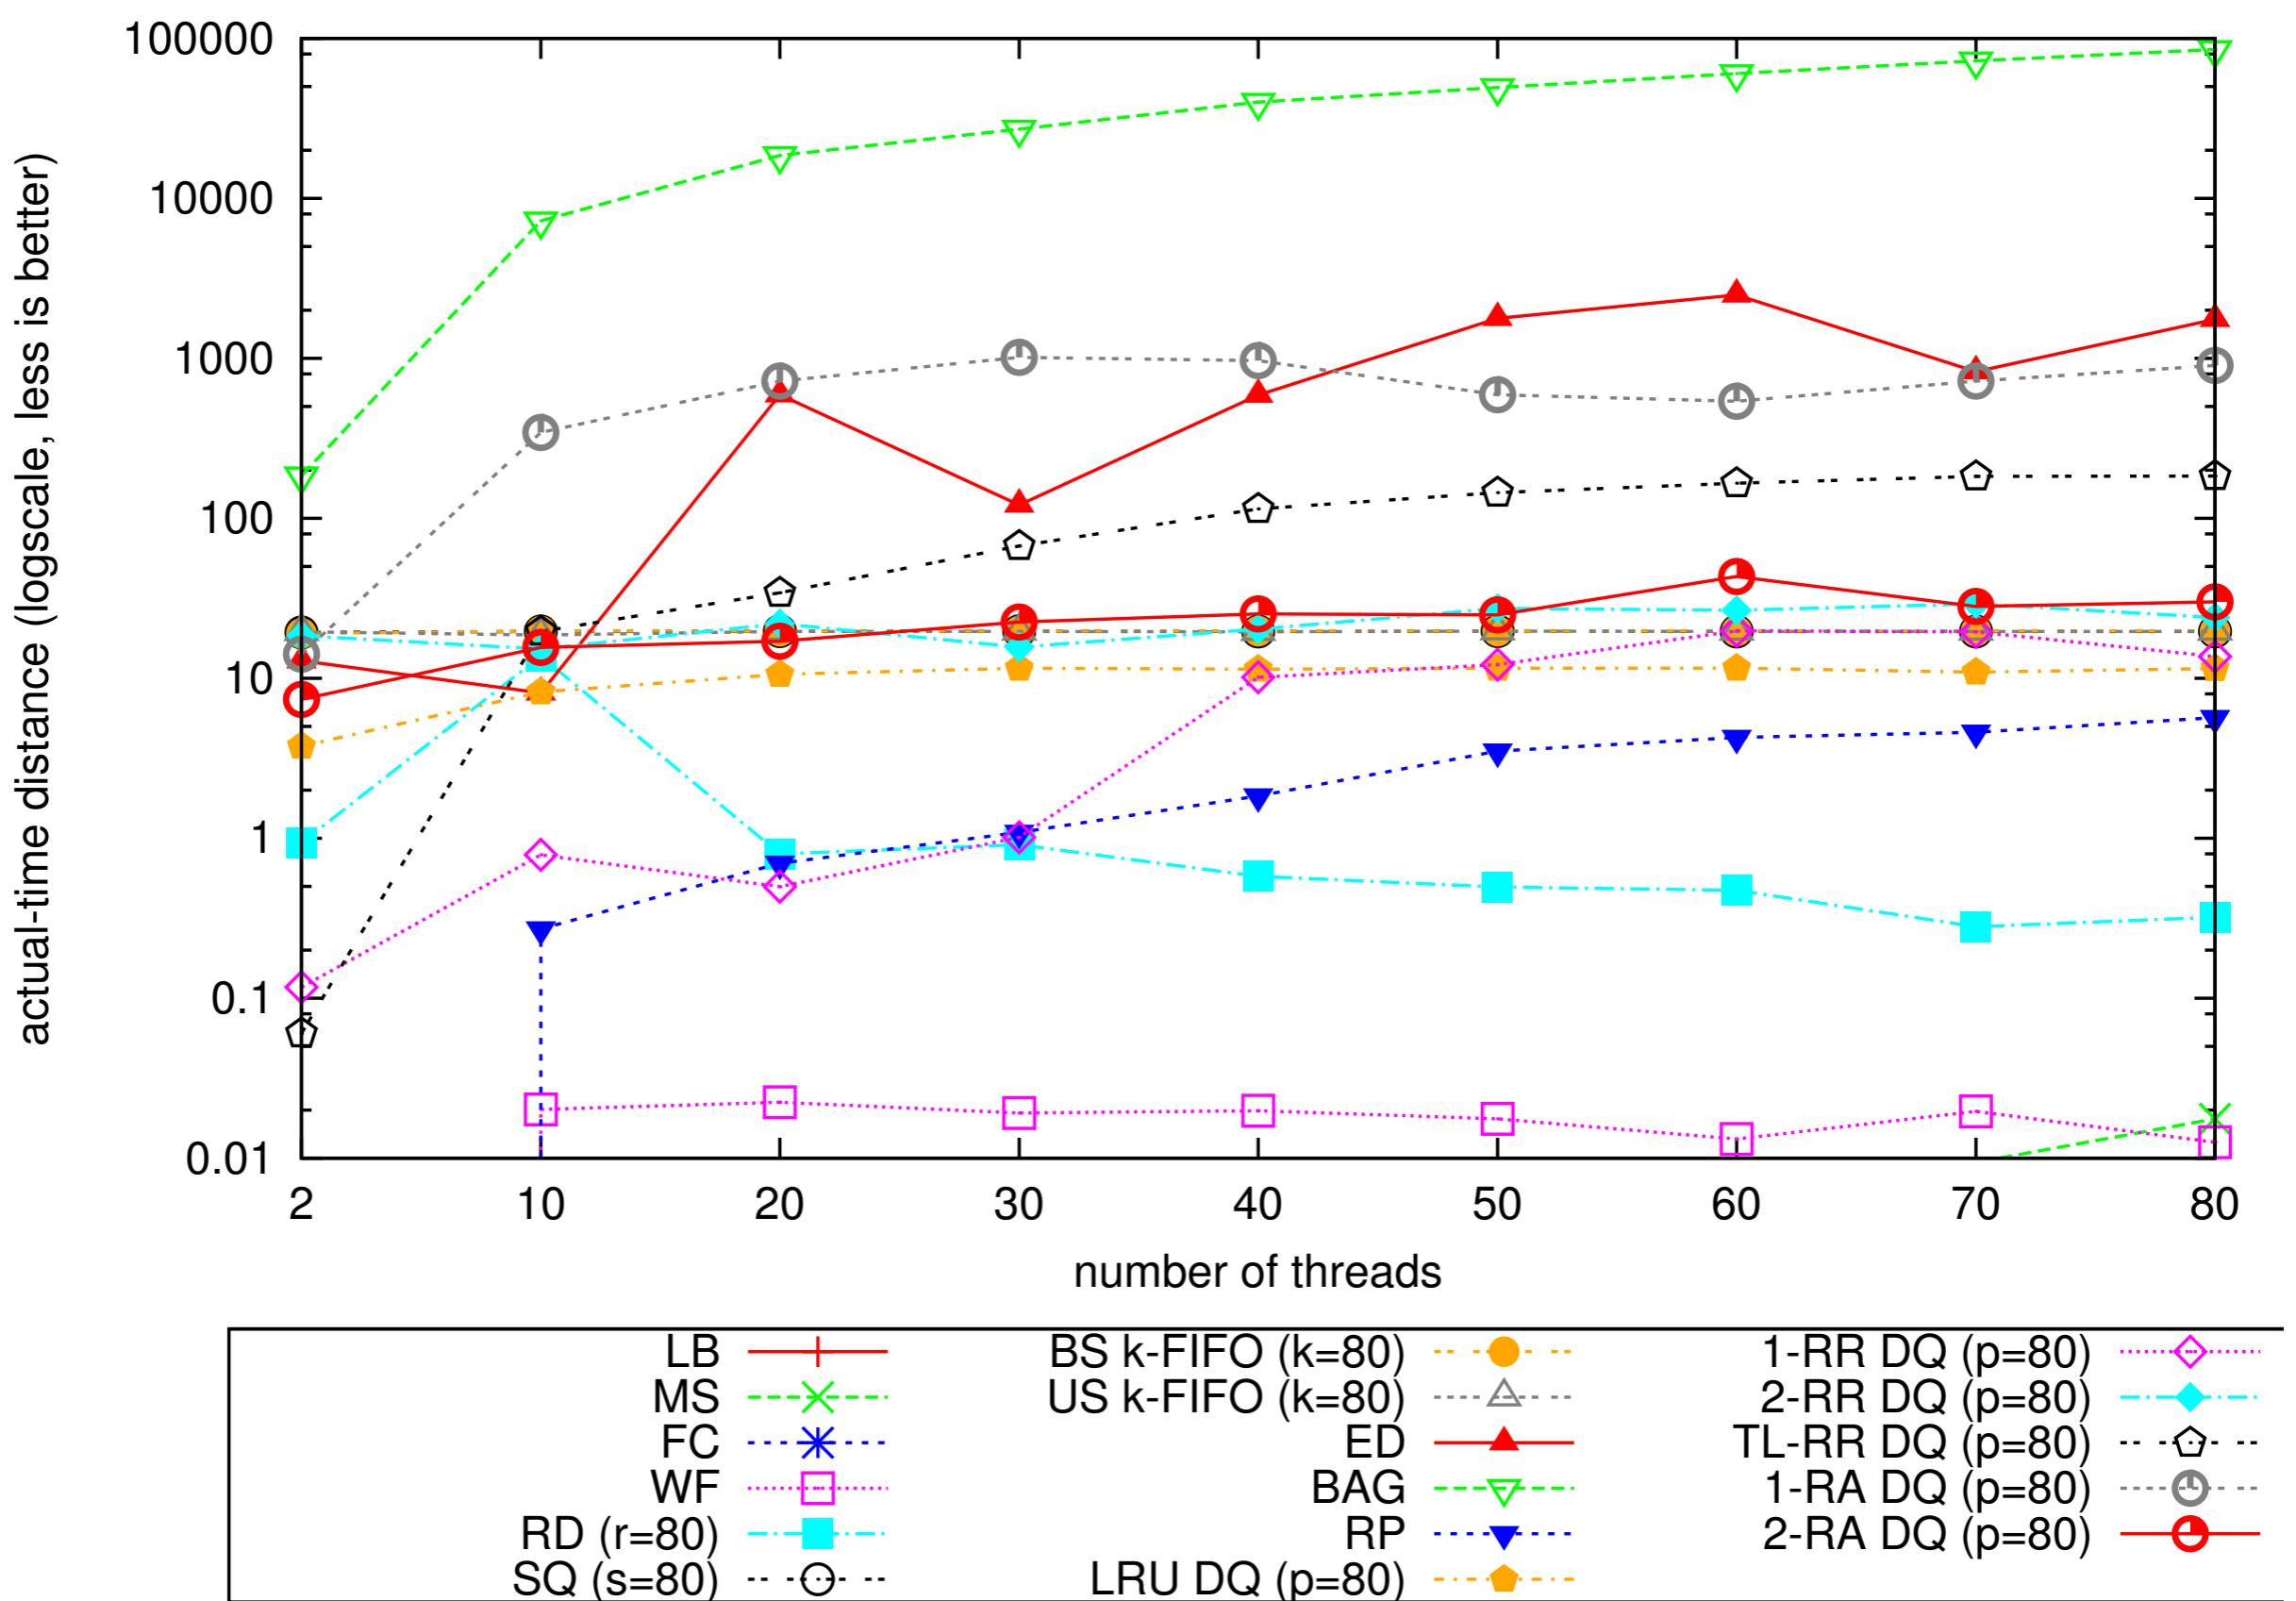
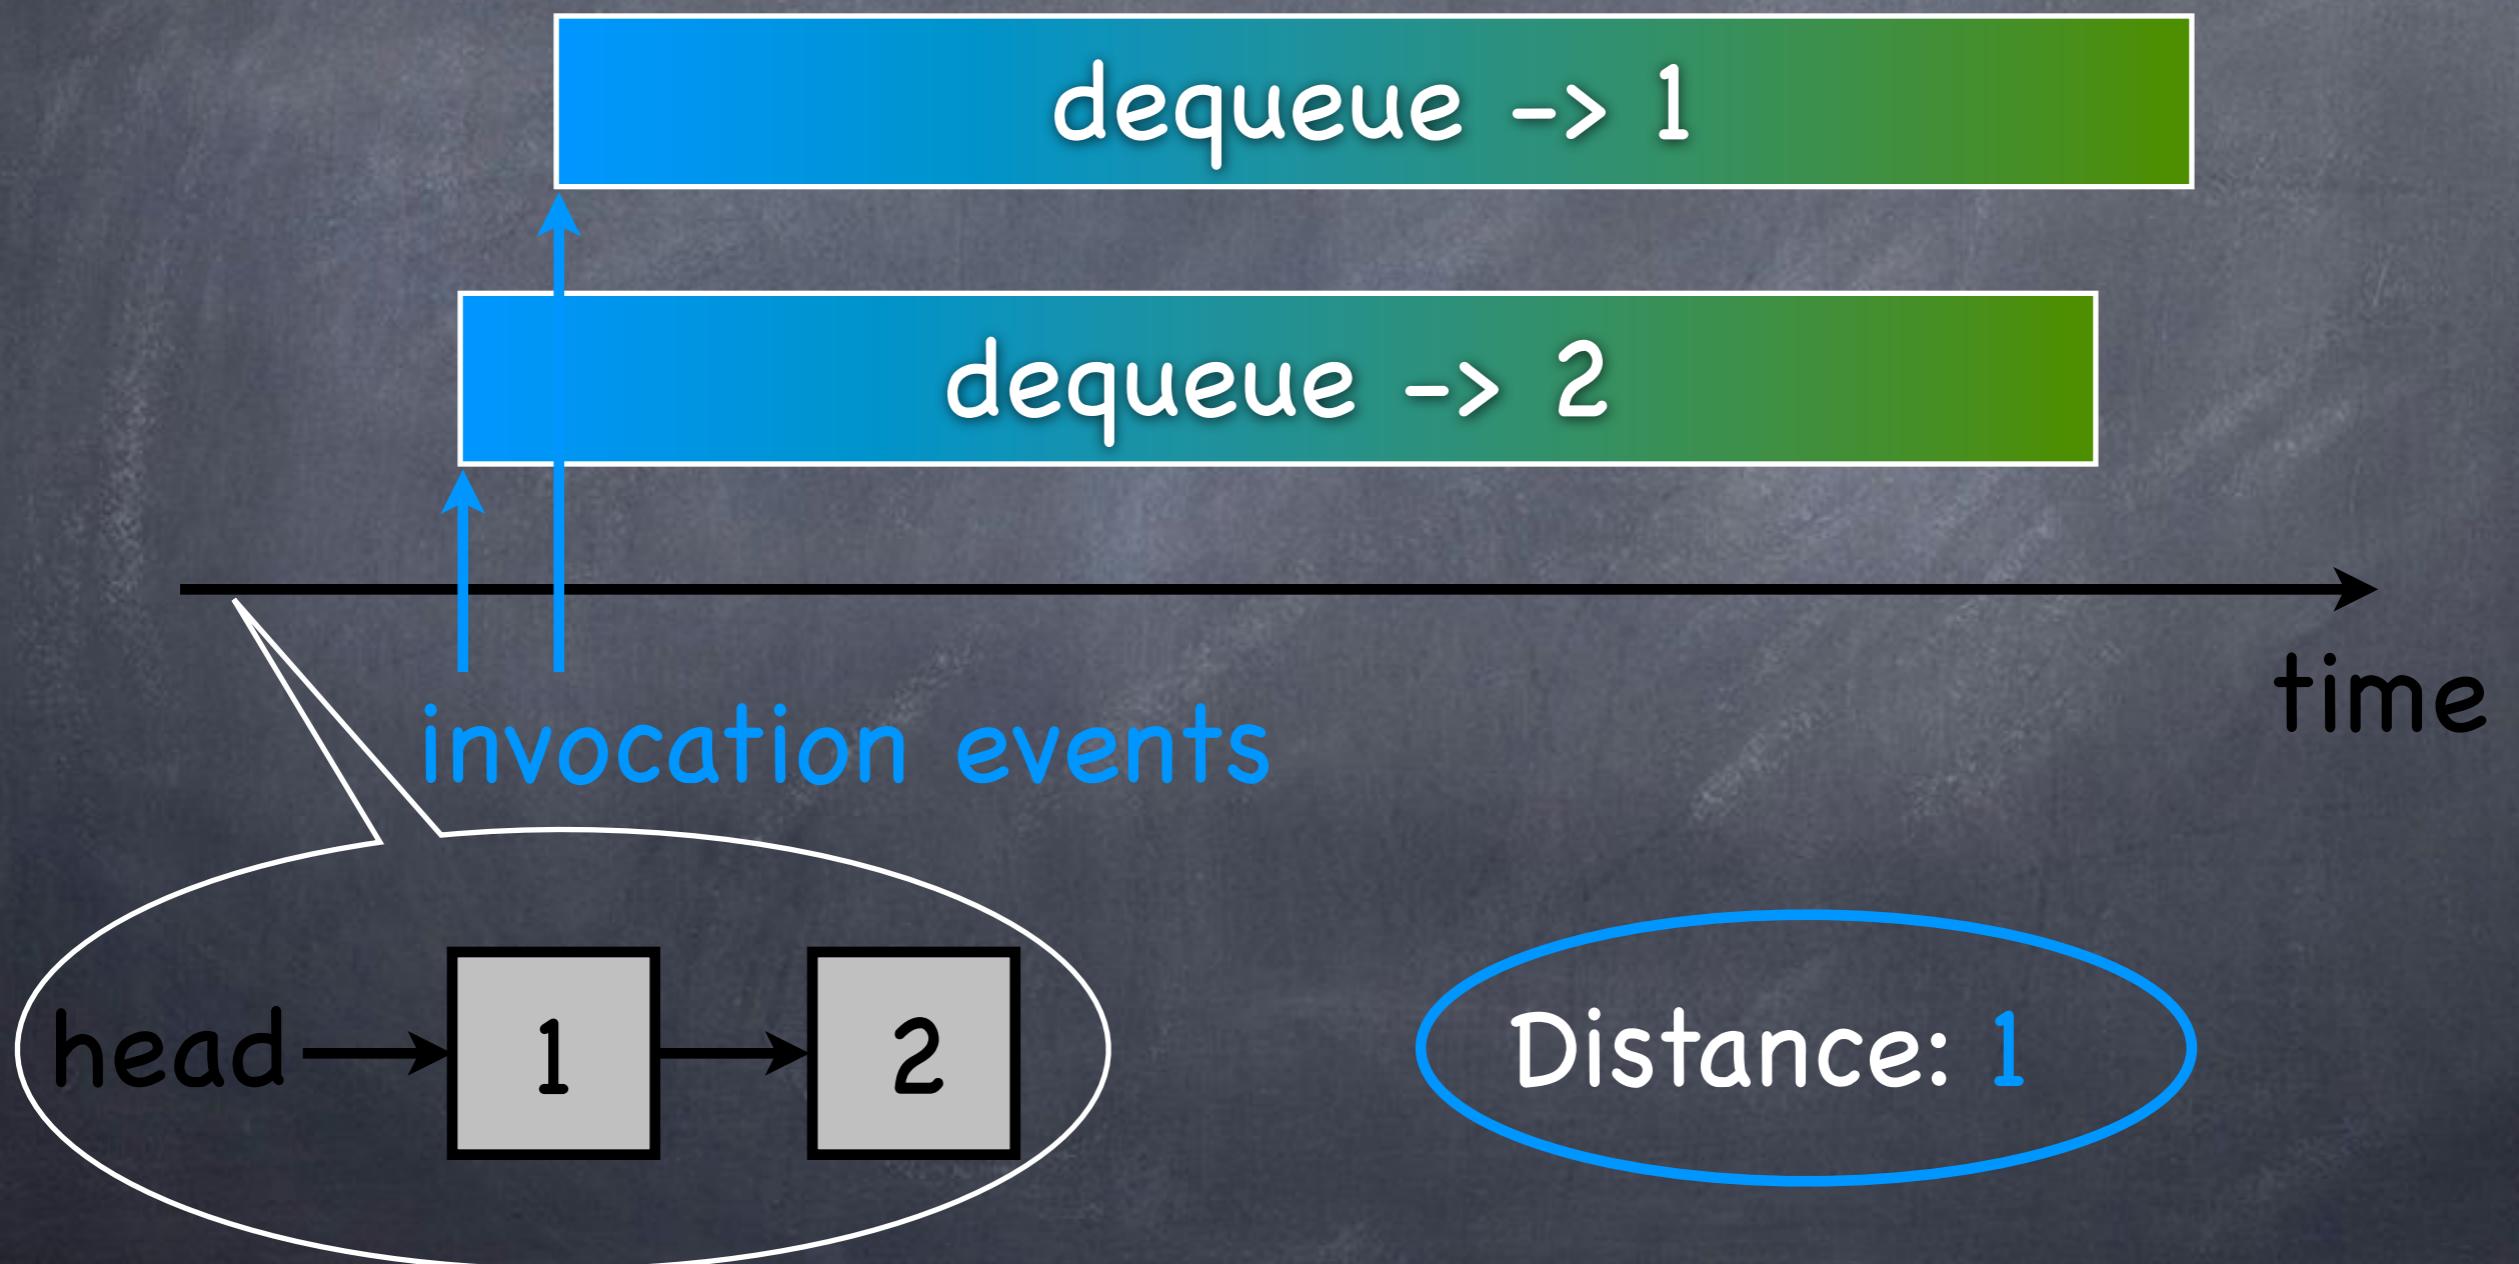


Figure 6: Actual-time distance of the high contention producer-consumer microbenchmark ($c = 250$)

Invocation vs. Linearization



The Zero-Time Linearization
of a concurrent history
is
the sequence of its operations
ordered by
their invocation events

The Zero-Time Distance
of a concurrent history
is
the average
out-of-order distance
of
its zero-time linearization

Zero-Time Distance

measures
re-ordering
due to
semantical relaxation
and
linearizability!

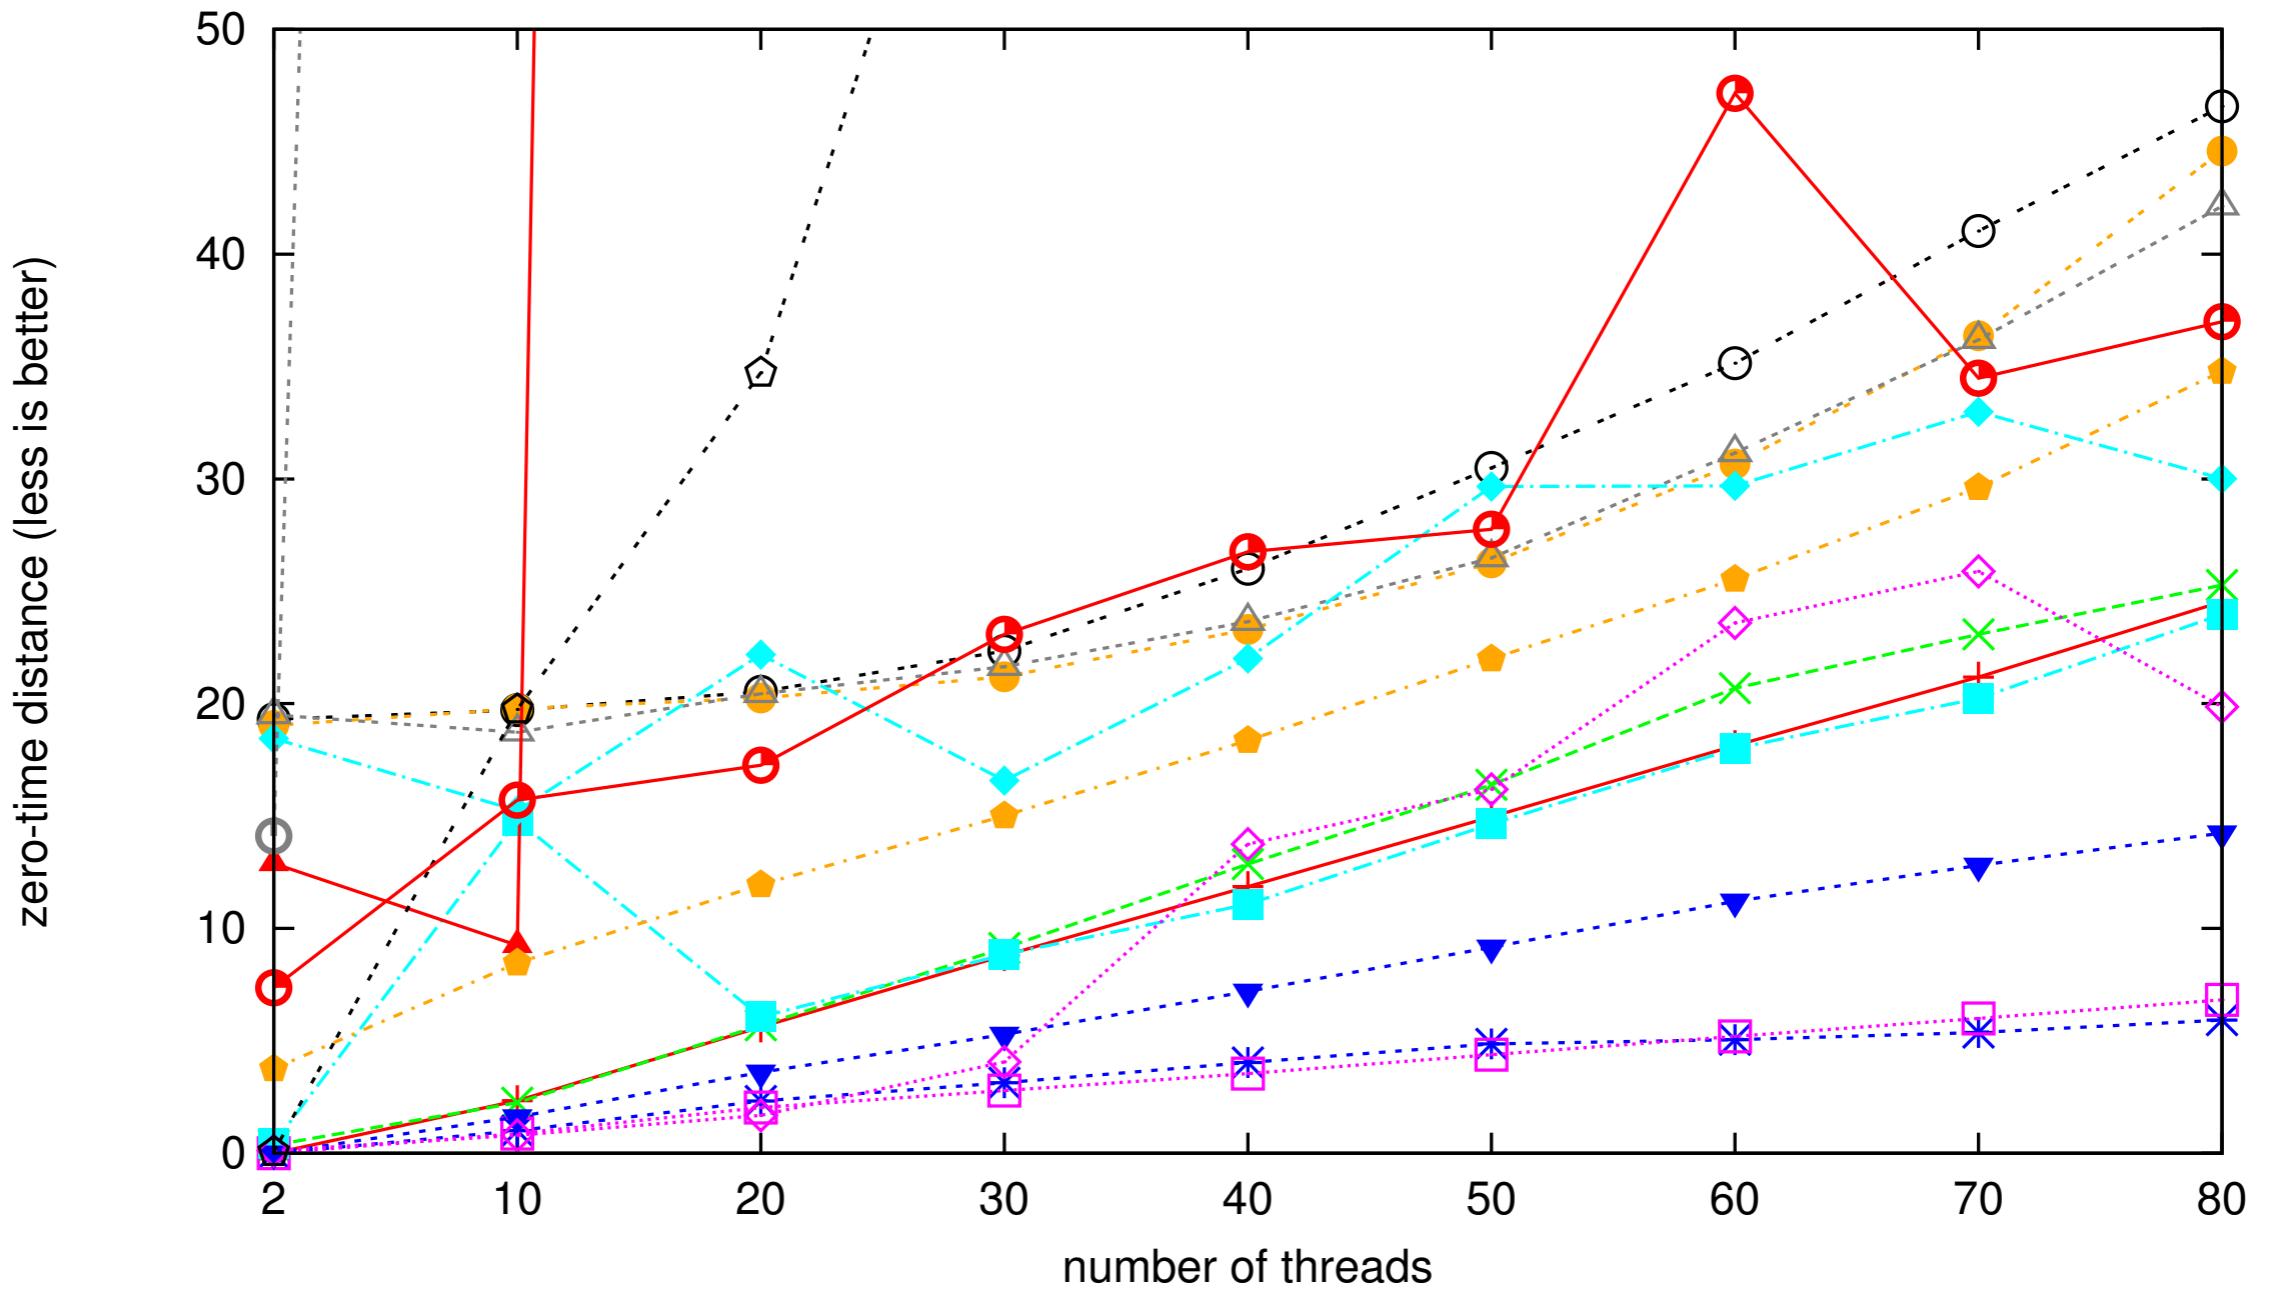


Figure 7: Zero-time distance of the high contention producer-consumer microbenchmark, zoomed-in for better resolution cutting off ED, BAG, TL-RR, and 1-RA ($c = 250$)

Linearization Difference
(difference of zero- and
actual-time distance)

measures
re-ordering
due to
linearizability!

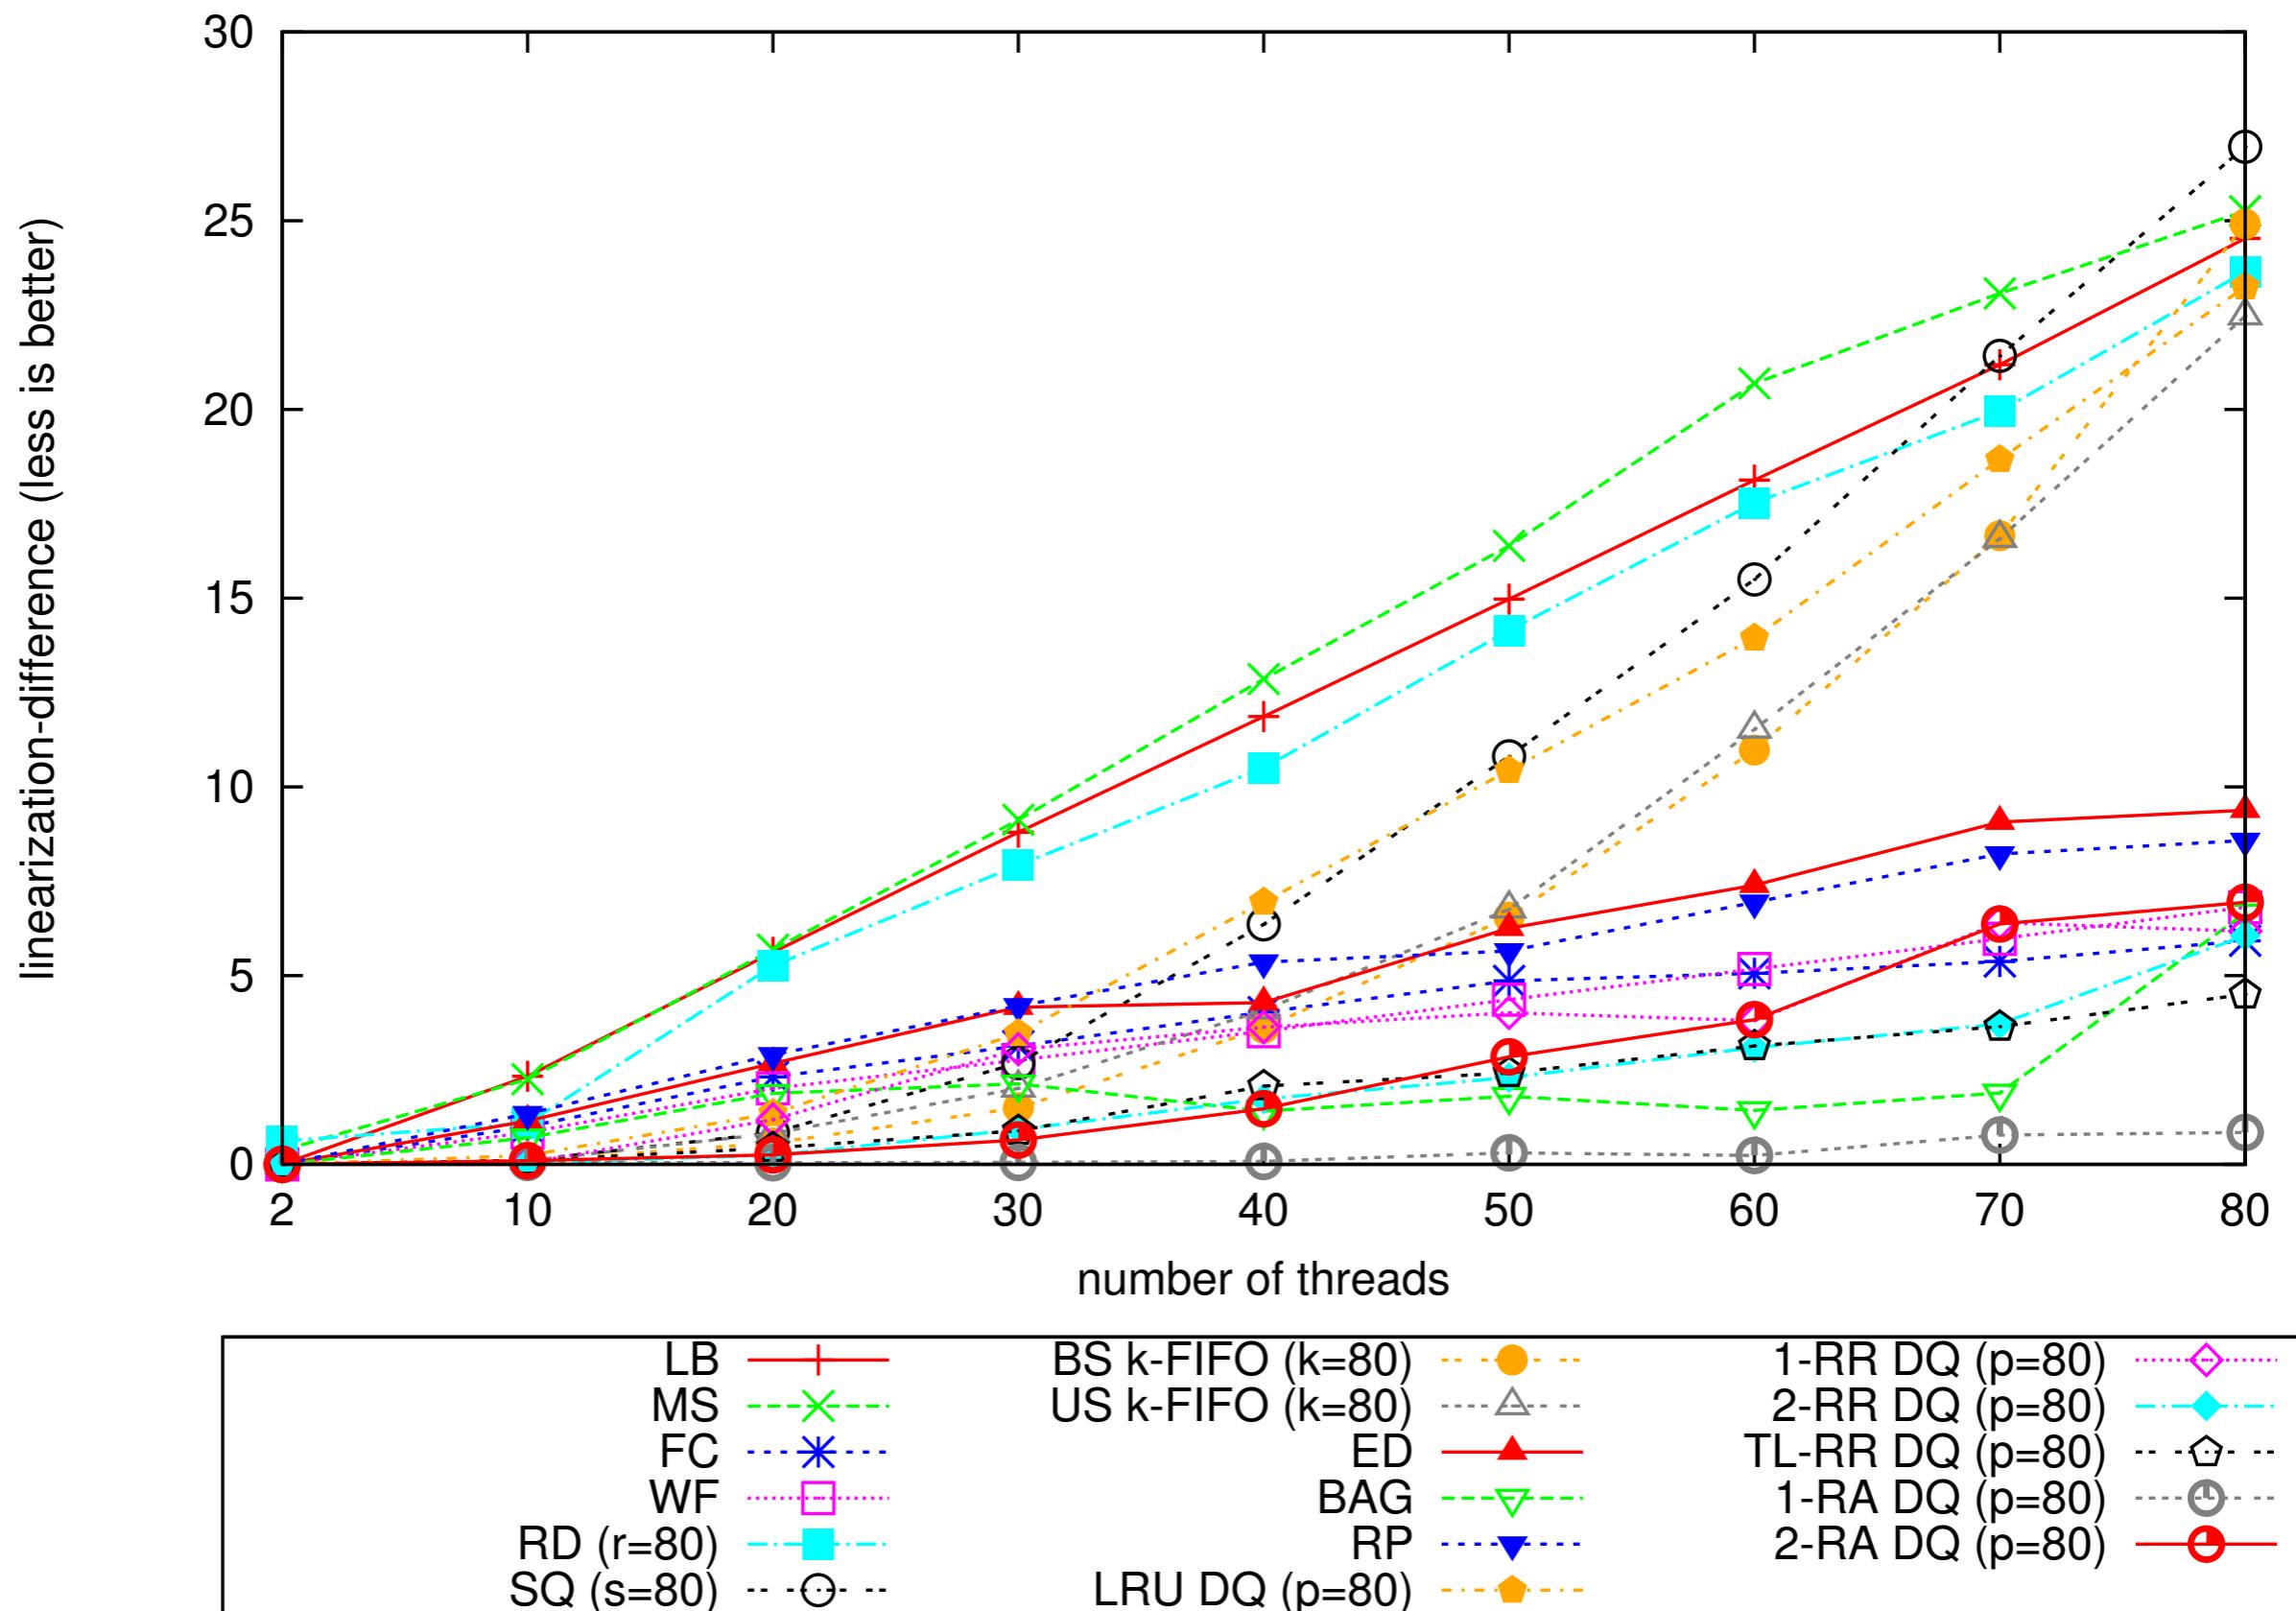


Figure 8: Linearization difference of the high contention producer-consumer microbenchmark ($c = 250$)

Thank you

